

CS-446/646

Synchronization

C. Papachristos

Robotic Workers (RoboWork) Lab
University of Nevada, Reno



Synchronization

Producer-Consumer with Threads

```
void* produce(void *arg) {  
    int i;  
    for(i=0; i<1e7; ++i)  
        ++ balance;  
}
```

```
void* consume(void *arg) {  
    int i;  
    for(i=0; i<1e7; ++i)  
        -- balance;  
}
```

```
int balance = 0;  
int main() {  
    pthread_t t1, t2;  
    pthread_create(&t1, NULL, produce, (void*)1);  
    pthread_create(&t2, NULL, consume, (void*)2);  
    pthread_join(t1, NULL);  
    pthread_join(t2, NULL);  
    printf("all done: balance = %d\n", balance);  
    return 0;  
}
```

```
$ gcc -Wall -pthread -o bank bank.c  
$ ./bank  
all done: balance = 0  
$ ./bank  
all done: balance = 140020  
$ ./bank  
all done: balance = -94304  
$ ./bank  
all done: balance = -191009
```



Synchronization

Producer-Consumer with Threads

Why?

➤ Load – Increment/Decrement – Store

```
$ objdump -d bank
...
08048464 <produce>:
...
// ++ balance
8048473: a1 80 97 04 08    mov 0x8049780,%eax
8048478: 83 c0 01          add $0x1,%eax
804847b: a3 80 97 04 08    mov %eax,0x8049780
...
0804849b <consume>:
...
// -- balance
80484aa: a1 80 97 04 08    mov 0x8049780,%eax
80484af: 83 e8 01          sub $0x1,%eax
80484b2: a3 80 97 04 08    mov %eax,0x8049780
...
```

One possible *Thread Schedule*

Thread 1

```
mov 0x8049780,%eax
add $0x1,%eax
mov %eax,0x8049780
```

Thread 2

```
balance: 0
eax0: 0
eax0: 1
balance: 1
mov 0x8049780,%eax
eax1: 1
sub $0x1,%eax
eax1: 0
mov %eax,0x8049780
balance: 0
```



Synchronization

Producer-Consumer with Threads

Why?

- Load – Increment/Decrement – Store

```
$ objdump -d bank
...
08048464 <produce>:
...
08048473: a1 80 97 04 08    // ++ balance
08048478: 83 c0 01          mov 0x8049780,%eax
0804847b: a3 80 97 04 08    add $0x1,%eax
0804847b: a3 80 97 04 08    mov %eax,0x8049780
...
0804849b <consume>:
...
080484aa: a1 80 97 04 08    // -- balance
080484af: 83 e8 01          mov 0x8049780,%eax
080484af: 83 e8 01          sub $0x1,%eax
080484b2: a3 80 97 04 08    mov %eax,0x8049780
...
```

A more “problematic” *Thread Schedule*

Thread 1	Thread 2
	balance: 0
mov 0x8049780,%eax	
	eax0: 0
add \$0x1,%eax	
	eax0: 1
	mov 0x8049780,%eax
	eax1: 0
mov %eax,0x8049780	
	balance: 1
	sub \$0x1,%eax
	eax1: -1
	mov %eax,0x8049780
	balance: -1

- Interrupt can occur before and after any *Instruction* (but not during it)



Synchronization

Race Condition

Definition: A timing-dependent error involving *Shared* state

Very bad

- “*Non-Deterministic*”
 - Can't know what the output will be, and it is likely to be different across runs
- Hard to detect
 - Too many possible *Schedules*
- Hard to debug
 - “*Heisen-bug*” : debugging changes timing so it can hide the bugs (vs “*Bohr-bug*”)



Synchronization

Avoiding *Race Conditions*

Atomic Operations

- No other *Instructions* can be interleaved
- Entire operation is executed “as a unit” - Guaranteed by Hardware

Possible approach:

- Have a dedicated *Atomic Instruction* for the job:
 - `add $0x1, 0x8049780` ←

```
// ++ balance  
mov 0x8049780,%eax  
add $0x1,%eax  
mov %eax,0x8049780
```

Problem:

- Can't anticipate every possible way we want *Atomicity*
- Increases Hardware complexity, slows down other *Instructions*



Synchronization

Layered Approach to *Synchronization*

- Hardware provides simple low-level *Atomic* Operations
 - Upon which we can build high-level *Synchronization* Primitives
 - Upon which we can implement *Critical Sections* and build correct Multi-Threaded/Multi-Processing programs

Properly synchronized Application

High-level *Synchronization Primitives*

Hardware-provided low-level *Atomic* operations

- Example low-level *Atomic* Operations
 - On Uniprocessor, disable/enable *Interrupts*
 - On x86, *Aligned-Load* and *Aligned-Store* of words
 - Special instructions: Test-Set-Lock/Exchange (TSL, XCHG), Compare-and-Swap (lock CMPXCHG)
- Example high-level *Synchronization* Primitives
 - *Lock, Semaphore, Monitor*



Synchronization

The Problem with *Threads*

x is a global variable initialized with 0

Thread 1	Thread 2
<pre>void foo() { x++; }</pre>	<pre>void bar() { x--; }</pre>

After both *Threads* finish, what is **x** ?

➤ 0, 1, -1

- *Assembly-level Instruction* sequence + *Time-Slice Interrupt* causing *Thread Switching* in the middle
 - Would run into same situation with pre-increment as well



Synchronization

The Problem with *Threads*

Global `int p = 0, ready = 0;`

Process 1

```
p = 1000;  
ready = 1;
```

Process 2

```
while (!ready);  
use(p);
```

What value of **p** is read by *Process 2*?

➤ 0, 1000

➤ Compiler is free to *Reorder* if it can “prove” no side-effects



Synchronization

Synchronization Motivation

Threads cooperate in Multithreaded programs

- To **share** resources, access shared data structures
- To **coordinate** their execution

For correct execution, control of this cooperation is required

- *Thread Scheduling* is non-deterministic (i.e. runtime behavior changes on same program re-runs)
 - *Scheduling* is not under the Program's control
 - *Scheduler* is part of OS
 - *Threads* interleave executions arbitrarily and at different rates
- Multi-Word operations are not *Atomic*
- Compiler (e.g. *Instructions*) and/or Hardware (e.g. *Memory*) *Reordering*



Synchronization

Shared Resources

Initially focus on controlling access to *Shared Resources*

Basic problem

- If two concurrent *Threads* (/Processes) are accessing a shared variable, and that variable is read/modified/written by those *Threads*, then access to the variable must be controlled

We need

- Mechanisms to control access to *Shared Resources*
 - *Locks, Mutexes, Semaphores, Monitors, Condition Variables*, etc.
- Patterns for coordinating accesses to *Shared Resources*
 - *Bounded-Buffer, Producer-Consumer*, etc.



Synchronization

Example: Bank Account Balance

Implement a function to handle withdrawals from a bank account

```
int withdraw (account, amount) {  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    return balance;  
}
```

Problem: Suppose 2 people go to separate ATMs and simultaneously initiate withdrawal

➤ Bank server runs the 2 Threads:

```
int withdraw (account, amount) {  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    return balance;  
}
```

```
int withdraw (account, amount) {  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    return balance;  
}
```



Synchronization

Example: Bank Account Balance

A Bad Schedule

```
balance = get_balance(account);  
balance = balance - amount;
```

```
balance = get_balance(account);  
balance = balance - amount;  
put_balance(account, balance);
```

Thread Context Switching

```
put_balance(account, balance);
```

- Initial balance: **1000**, 2 x Withdrawal amount (each): **100**
- Final balance: **900**



Synchronization

Example: **Bank Account Balance**

Can get a lot more interleaved

➤ *Remember: Case of Producer-Consumer Assembly-level Thread Schedule*

Assumptions:

- We have to assume that the only *Atomic* operations are *Instructions*
 - e.g. reads and writes of Words
 - even for that, the Hardware has to explicitly provide such support
- A *Context Switch* can happen at any time
- A *Thread* can be delayed indefinitely as long as it is not forever
 - no *Real-Time* guarantee



Synchronization

Shared Resources

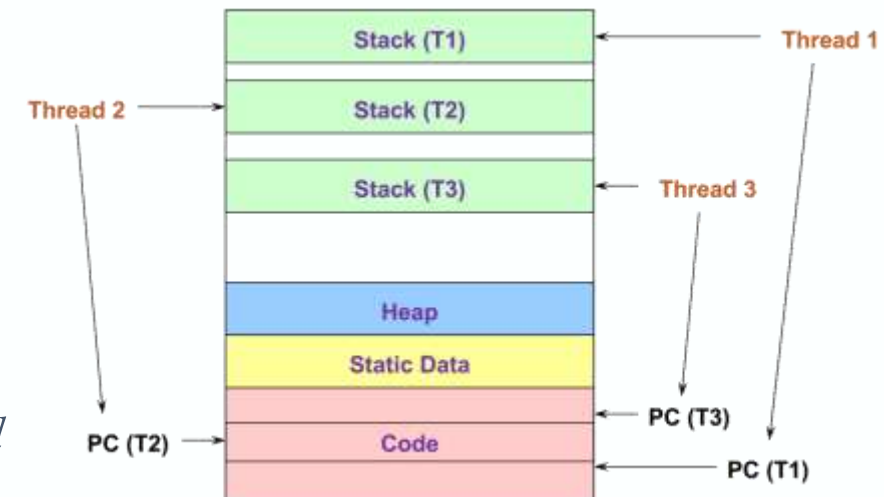
The previously demonstrated problem came from accessing a *Shared Resource* without proper *Synchronization*

➤ *Race Condition*

Controlled-access mechanisms to *Shared Data Structures* (bank account, queues, lists, hash tables, etc.) are required to deal with *Concurrency*, so we can ensure a degree of determinism in program execution.

What is *Shared*:

- Local variables are not shared
 - Refer to data on the *Stack*
 - Each *Thread* has its own *Stack*
 - Potentially dangerous to pass/share/store a pointer to a local variable on the *Stack* of one *Thread* to another
- Global and **static** variables are shared
 - Stored in the *Static Data Segment*, accessible by any *Thread*
- Dynamic and other *Heap* data are shared
 - Allocated from *Heap* with **malloc/free**



Synchronization

Mutual Exclusion

We use *Mutual Exclusion* to *Synchronize* access to *Shared Resources*

- Allows us to build larger *Atomic* blocks

Critical Section

Code that uses *Mutual Exclusion* to *Synchronize* its execution

- Only one *Thread's* execution at a time can enter-or-be in the *Critical Section*
- All other *Threads* are forced to wait on entry
- When a *Thread* leaves a *Critical Section*, another can enter



Synchronization

Critical Section Requirements

Mutual Exclusion (Mutex)

- If one *Thread* is in the *Critical Section*, then no other is

Liveness (Progress)

- If some *Thread T* is not in the *Critical Section*, then *T* cannot prevent some other thread *S* from entering
 - If multiple *Threads* simultaneously request to enter *Critical Section*, one must be allowed to proceed
 - A *Thread's* operations outside the *Critical Section* should not be able to prevent another one to proceed
- A *Thread* in the *Critical Section* will eventually leave it

Bounded Waiting (Starvation-free)

- If some *Thread T* is waiting on the *Critical Section*, then *T* will eventually enter the *Critical Section*

Performance

- The overhead of entering and exiting the *Critical Section* is small with respect to the work being done within it



Synchronization

Critical Section Desired Properties

Safety : Nothing bad should happen (#1 Priority)

- *Mutex*

Liveness : Something useful should be happening

- *Progress, Bounded Waiting*

Performance :

- *Efficiency*: Don't consume too many Resources while waiting
 - Don't *Busy-Wait (Spin-Wait)*. Better to relinquish CPU and let another *Thread* run
- *Fairness*: Don't make one *Thread* wait longer than others.
 - Hard to do efficiently
- *Simplicity*: Should be simple to use
- Properties hold for each run, while *Performance* is quantified by all runs



Synchronization

Critical Section Implementation – using `pthread_mutex_t`

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

Acquire *Mutex* (/Lock) exclusively; wait if not available

- The *Mutex* object referenced by **mutex** shall be **Locked** by calling `pthread_mutex_lock()`. If the *Mutex* is already locked, the calling *Thread* shall **Block** until the **mutex** becomes available.

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

Release exclusive access to *Mutex* (/Lock)

- Shall **Release** the *Mutex* object referenced by **mutex**.

```
pthread_mutex_t l = PTHREAD_MUTEX_INITIALIZER;
```

```
void* produce(void *arg) {  
    int i;  
    for(i=0; i<1e7; ++i)  
        pthread_mutex_lock(&l);  
        ++ balance;  
        pthread_mutex_unlock(&l);  
}
```

```
void* consume(void *arg) {  
    int i;  
    for(i=0; i<1e7; ++i)  
        pthread_mutex_lock(&l);  
        -- balance;  
        pthread_mutex_unlock(&l);  
}
```

} Critical Section



Synchronization

Implementing *Locks* – v1

On a Uniprocessor we can cheat

- Implement *Mutual Exclusion* by disabling/enabling *Interrupts*

```
void lock()                void unlock()
{                          {
    disable_interrupts();   enable_interrupts();
}
```

Good:

- Simple

Bad:

- Both operations are *Privileged, User-Level* program not allowed to use them
- Doesn't work on Multiprocessor
 - On multi-core architectures, enabling/disabling *Interrupts* is on a per-core basis. One *Thread* might be running on a different core, so we would either have to disable *Interrupts* on all cores, or have an architecture-dependent implementation using *Inter-Processor Interrupts* (IPIs).



Synchronization

Implementing *Locks* – v2

Software-based *Lock*

Desired specifications for a Software-based *Lock* algorithm:

Good:

- Shouldn't require much from hardware

Only assumptions:

- *Loads* and *Stores* are *Atomic*
- They execute *In-Order*
 - (vs *Out-of-Order* execution)
- Does not require special hardware Instructions



Synchronization

Implementing *Locks* – v2

Software-based *Lock* – 1st Attempt

```
// 0: lock is available, 1: lock is held by a thread  
int flag = 0;
```

```
void lock() {  
    while (flag == 1); // spin wait  
    flag = 1;  
}
```

```
void unlock() {  
    flag = 0;  
}
```

Idea: Use one *Flag*, *Test* then *Set*; if unavailable *Spin-Wait*

Problem?

- Not *Safe*: Both *Threads* can be in *Critical Section*
 - Both can execute the *Test* before one proceeds to execute the line that does the *Set*
- Not *Efficient*: *Busy-waiting*, particularly bad on Uniprocessor (will address this later)



Synchronization

Implementing *Locks* – v2

Software-based *Lock* – 2nd Attempt

```
// 1: a thread wants to enter critical section, 0: it doesn't  
int flag[2] = {0, 0};
```

```
void lock() {  
    flag[self] = 1; // I need lock  
    while (flag[1-self] == 1); // spin wait  
}
```

```
void unlock() {  
    // not any more  
    flag[self] = 0;  
}
```

Idea: Use per-Thread *Flags*, *Set* then *Test*, to achieve *Mutual Exclusion*

Problem?

- Not *Live*: Can enter a *Deadlock*
 - Both can execute the *Set* before one proceeds to *Test*, therefore both will forever *Spin-Wait*
- Not *Efficient*: *Busy-waiting*, particularly bad on Uniprocessor (will address this later)



Synchronization

Implementing *Locks* – v2

Software-based *Lock* – 3rd Attempt

```
// whose turn is it?  
int turn = 0;
```

```
void lock() {  
    // wait for my turn  
    while (turn == 1 - self); // spin wait  
}
```

```
void unlock() {  
    // I'm done. your turn  
    turn = 1 - self;  
}
```

Idea: Strict *Alternation* to achieve *Mutual Exclusion*

Problem?

- Not *Live*: Depends on *Threads* operations outside *Critical Section*
 - *Thread 1* can go into an infinite loop after its *Critical Section* (after it **unlocks**)
 - *Thread 2* will get to execute once, but then *Thread 1* will never again alternate the **turn** over to it



Synchronization

Implementing *Locks* – v2

Software-based *Lock* – *Peterson's* Algorithm – Final Attempt (combine previous ideas)

```
// whose turn is it?  
int turn = 0;  
// 1: a thread wants to enter critical section, 0: it doesn't  
int flag[2] = {0, 0};
```

```
void lock() {  
    flag[self] = 1; // I need lock  
    turn = 1 - self; // wait my turn (set NOT my turn)  
    while (flag[1-self] == 1 && turn == 1 - self);  
    // spin wait while the other thread has intent  
    // AND it is the other thread's turn  
}
```

```
void unlock() {  
    // not any more  
    flag[self] = 0;  
}
```

- *Safe*
- *Live*: One of the two will have executed the “*Set: Not My Turn*” operation last, before entering the *Spin-Wait* phase, i.e. the other will proceed



Synchronization

Implementing *Locks* – v3

Atomic Operation-based Lock

- Problem with Software-based Lock: Hard to implement for > 2 *Threads*
- Also modern CPUs can perform operations *Out-of-Order* (need *Memory Barrier*)

```
// 0: lock is available, 1: lock is held by a thread
int flag = 0;
```

```
void lock() {
    while( test_and_set( &flag ) );
}
```

```
void unlock() {
    unset( &flag );
}
```

Note:

Approach better thought of as “Set-and-Test-Previous-Value”

- If previously 1 (by another *Thread*), will just *Spin-Wait*
 - (and set 1, but irrelevant)
- If previously 0, will *immediately* set 1 and proceed

Remember: Problem with the *Test-then-Set* approach is it is not *Atomic*

Idea:

- Make *Atomic Test-and-Set*:

```
int test_and_set (int *lock) {
    int old = *lock;
    *lock = 1;
    return old;
}
```

Should *Atomically* return prior value of ***lock** and set ***lock** to new value of 1



Synchronization

Implementing *Locks* – v3

Implementing **test_and_set** on x86

```
long test_and_set(volatile long* lock) {  
    int old;  
    asm("xchgl %0, %1"  
        : "=r" (old), "+m" (*lock) // output  
        : "0" (1) // input  
        : ["cc"], "memory" // ... the compiler does not assume that any values read from memory  
    ); // before an asm remain unchanged after that asm; it reloads them  
    // as needed. ... Using the "memory" clobber effectively forms a  
    // read/write memory barrier for the compiler.  
    return old;  
}
```

Note:

The data that **lock** points to is **volatile**: Disable compiler optimizations (for this object) that can result in a variable being assumed that it does not change outside the scope of the current function (e.g. by an *ISR*, by another *Thread*, etc.), and enforce that it is always read from memory afresh (instead of keeping a cached copy in a temporary *Register*)

➤ Extended Assembly (<https://gcc.gnu.org/onlinedocs/gcc/Extended-Asm.html>)

Atomic Instruction **xchg** of **reg** (**old**), **addr**: (***lock**) *Atomically* swaps them

- Most *Spin-Locks* on x86 are implemented using this *Instruction*
 - e.g. xv6 **spinlock.h**, **spinlock.c**, **x86.h**



Synchronization

Implementing *Locks* – v3

- More modern CPU *Atomic* Instructions unlock more possibilities

Atomic Compare-and-Swap (or *Compare-and-Exchange*):

- Checks whether content of memory location matches a value, and if so, modifies it to a new value
- Can now store *Thread ID* of owning *Thread*, instead of just a true/false variable

```
// 0: lock is available, !0: tid of thread holding lock
```

```
int tid_lock = 0;
```

```
void lock() {  
    while( !compare_and_swap(&tid_lock, 0, getpid()) );  
}
```

```
int compare_and_swap (int *addr, int test, int new) {  
    if (*addr != test)  
        return 0;  
    *addr = new;  
    return 1;  
}
```

Note:

MACRO to get current *Thread ID* in Linux

```
#define getpid()  
((pid_t) syscall(SYS_gettid))
```

Note:

Can use x86 **lock cmpxchg** *Instruction*
(the **lock** prefix ensures CPU exclusive ownership of Cache line – possibly with a *Memory Bus Lock*)



Synchronization

Spin-Waiting vs Blocking

Problem: Waste of CPU cycles

- Worst case scenario: a *Thread* holding a Busy-Wait *Lock* (i.e. is inside the *Critical Section*) gets *Preempted*, while some other *Threads* try to acquire the same *Lock*

On Uniprocessor: Should not use a *Spin-Lock*

- **yield** CPU when *Lock* is not available (need OS support)

On Multiprocessor

- If a *Thread* holding *Lock* gets *Preempted*, the correct action depends on how long before the *Lock* would be released



Synchronization

Spin-Waiting vs Blocking

Problem with the simple *Yield*:

```
void lock() {  
    while( test_and_set( &flag ) )  
        sched_yield();  
}
```

- Uncontrollably results in a lot of *Context Switches*
 - *Thundering Herd*
- *Starvation* due to lack of control over which *Thread* gets the *Lock* becomes possible

Why?

- No control over who gets the *Lock* next
- Need explicit control to ensure which *Thread* should get the *Lock*



Synchronization

Implementing *Locks* – v4

// 0: lock is available, 1: lock is held by a thread

int flag = 0;

```
void lock() {  
    while( test_and_set( &flag ) ) {  
        // add myself (back) to wait queue  
        sched_yield();  
    }  
}
```

```
void unlock() {  
    unset( &flag );  
    if( any_thread_in_wait_queue ) {  
        // wake up one thread from wait queue  
    }  
}
```

Idea: Have a *Wait Queue* with those *Threads* that are actually waiting on this specific *Lock*

- (Re-)Add *Thread* to *Wait Queue* while *Lock* remains unavailable
- In `unlock()`, wake up one *Thread* from *Wait Queue*

Problem 1: Lost wakeup

Because it wastes time performing (re-)enqueueing itself; in that time another *Thread* reaches the *Test-and-Set* and grabs the *Lock*

- Fix: **Need** the *Spin Lock* to be fast
- *Spin-Wait* will still take place, but should not be expected to be active for too long... (*How?*)

Problem 2: Wrong *Thread* gets *Lock*

- Fix: `unlock()` directly transfers *Lock* to waiting *Thread*
- No other *Thread* should be possible to acquire *Lock*... (*How?*)



Synchronization

Implementing *Locks* – v4 : mutex

```
typedef struct __mutex_t {  
    int guard; // simple guard lock to avoid losing wakeups  
    int flag;  // 0: mutex is available, 1: mutex is not available  
    queue_t *queue; // queue of waiting threads  
} mutex_t;
```

```
void lock(mutex_t *m) {  
    //acquire guard lock by spinning  
    while (test_and_set(m->guard));  
    if (m->flag == 0) {  
        m->flag = 1; // mutex acquired  
        unset(m->guard);  
    } else {  
        enqueue(m->queue, self);  
        unset(m->guard);  
        sched_yield();  
    }  
}
```

```
void unlock(mutex_t *m) {  
    //acquire guard lock by spinning  
    while (test_and_set(m->guard));  
    if (empty(m->queue))  
        // release mutex; no one wants mutex  
        m->flag = 0;  
    else  
        // direct transfer mutex to next thread  
        wakeup( dequeue(m->queue) );  
    unset(m->guard);  
}
```



Synchronization

Implementing *Locks* – v4 : mutex

- Now the **m->guard** *Lock* is an internal property of the **mutex_t**, i.e. it only protects its inner *Critical Sections* of **lock()** and **unlock()** (between the *Spin-Lock* on **m->guard**, and the line unsetting it to 0)
 - The *Critical Sections* of **lock()** & **unlock()** is now where the actual marshalling of *Threads* happens, by manipulating the *Mutex* state variables **m->flag** and **m->queue**.

```
void lock(mutex_t *m) {  
    //acquire guard lock by spinning  
    while (test_and_set(m->guard));  
    if (m->flag == 0) {  
        m->flag = 1; // mutex acquired  
        unset(m->guard);  
    } else {  
        enqueue(m->queue, self);  
        unset(m->guard);  
        sched_yield();  
    }  
}
```

```
void unlock(mutex_t *m) {  
    //acquire guard lock by spinning  
    while (test_and_set(m->guard));  
    if (empty(m->queue))  
        // release mutex; no one wants mutex  
        m->flag = 0;  
    else  
        // direct transfer mutex to next thread  
        wakeup( dequeue(m->queue) );  
    unset(m->guard);  
}
```



Synchronization

Reader – Writer Problem

- A *Reader* is a *Thread* that needs to look at the shared data but won't change it
- A *Writer* is a *Thread* that modifies the shared data
 - e.g. making an airline reservation
 - Courtois et al 1971: [Concurrent Control with “Readers” and “Writers”](#)

Problem: With the regular *Lock* approach, there is unnecessary *Synchronization*

- Only one *Writer* should be active at a time
- However, any number of *Readers* can be active simultaneously

Solution:

- Acquire *Lock* for *Read Mode* and *Write Mode*



Synchronization

Reader – Writer Lock

```
rwlock_t lock;
```

```
void* writer(void *arg) {  
    while(true) {  
        write_lock(&lock);  
        ...  
        // write shared data  
        ...  
        write_unlock(&l);  
    }  
}
```

```
void* reader(void *arg) {  
    while(true) {  
        read_lock(&lock);  
        ...  
        // read shared data  
        ...  
        read_unlock(&lock);  
    }  
}
```

read_lock: Acquires *Lock* in *Read (Shared Access) Mode*

- If *Lock* is not acquired or in *Read Mode* → Success
- Otherwise, *Lock* is in *Write Mode* → Wait

write_lock: Acquires *Lock* in *Write (Exclusive Access) Mode*

- If *Lock* is not acquired → Success
- Otherwise → Wait



Synchronization

Implementing *Reader – Writer Lock*

```
struct rwlock_t {  
    int nreader;           // init to 0  
    lock_t guard;         // init to unlocked  
    lock_t datalock;      // init to unlocked  
};
```

```
void write_lock(rwlock_t *l) {  
    lock(&l->datalock);  
}  
  
void write_unlock(rwlock_t *l) {  
    unlock(&l->datalock);  
}
```

```
void read_lock(rwlock_t *l) {  
    lock(&l->guard);  
    ++ nreader;  
    if(nreader == 1) // 1 reader, no more writing  
        lock(&l->datalock);  
    unlock(&l->guard);  
}  
  
void read_unlock(rwlock_t *l) {  
    lock(&l->guard);  
    -- nreader;  
    if(nreader == 0) // 0 readers, can write  
        unlock(&l->datalock);  
    unlock(&l->guard);  
}
```

Problem:

➤ *Writer Starvation* is possible



CS-446/646

Time for Questions !

