

CS-446/646

Semaphores & Monitors

C. Papachristos

Robotic Workers (RoboWork) Lab
University of Nevada, Reno



Semaphores

Semaphore Motivation

Problem with *Lock*:

- Ensures *Mutual Exclusion*, but not execution order

Producer-Consumer problem: Ensuring execution order makes sense

- *Producer*: Creates resources
- *Consumer*: Uses resources
- *Bounded Buffer*: Shared between them
- Execution order: *Producer* should just wait if *Bounded Buffer* is full, *Consumer* should just wait if *Bounded Buffer* is empty
 - e.g. `$ cat entries.txt | sort | uniq | wc`



Semaphores

Semaphore Definition

Abstract data type (i.e. a high-level mechanism) to provide *Synchronization*

- Described by Dijkstra in the “THE (Technische Hogeschool Eindhoven) Operating System” in 1968

A *Synchronization* object that **contains an integer counter** variable

- No operation to access integer counter variable directly
- *Semaphore* safety property: Integer counter value never allowed to go below 0
- Integer counter variable must be initialized to some value:
 - `sem_init (sem_t *s, int pshared, unsigned int value)`
- Operations to manipulate integer counter variable:
- `sem_wait` (or `down()`, `P()` -robiezen): Decrements, *Blocks* until semaphore is *Open*
- `sem_post` (or `up()`, `V()` -erhogen): Increments, allows another *Thread* to enter

```
int sem_wait(sem_t *s) {  
    // 1. wait until value of  
    // semaphore s becomes > 0  
    // 2. decrement value by 1  
}
```

```
int sem_post(sem_t *s) {  
    // 1. increment value of s by 1  
    // 2. if there are 1 or more  
    // threads waiting, wake 1  
}
```



Semaphores

Blocking in Semaphores

Associated with each *Semaphore* is a *Queue* of waiting *Threads*

When **P()** / **sem_wait()** is called by a *Thread*:

- If *Semaphore* is *Open*, *Thread* continues
- If *Semaphore* is *Closed*, *Thread* will *Block* on *Queue*

When **V()** / **sem_post()** *Opens* the *Semaphore*:

- If a *Thread* is waiting on the *Queue*, it is *Unblocked*
- If no *Threads* are waiting on the *Queue*, the **signal is remembered** for the next *Thread*
 - In other words, **V()** has “memory”
 - In contrast to *Condition Vars* (will see these later)
 - This “memory” property is derived from the integer counter value



Semaphores

Semaphore Types

Mutex Semaphore (or Binary Semaphore)

- Represents single access to a resource; **X=1**
- Guarantees *Mutual Exclusion* to a *Critical Section*

Counting Semaphore (or General Semaphore)

- Represents a resource with many units available, or a resource to which we want to limit concurrent access (e.g. reading); **X>1**
 - Is initialized to number of resources available
- Multiple *Threads* can pass the *Semaphore* “wait” test
- Number of *Threads* determined by *Semaphore* “counter”

Note:

No direct access to counter

```
sem_init(&s, 0, X)
...
sem_wait(&s);
// critical section
sem_post(&s);
```

```
int sem_init(sem_t *sem,
             int pshared,
             unsigned int value);
```

Initializes the *Semaphore* at **sem**.

value specifies the initial value for it.

pshared indicates whether this *Semaphore* is to be shared between the *Threads* of a *Process*, or between *Processes* (**sem** can be part of a region of *Shared Memory*).

```
int sem_post(sem_t *sem);
```

Increments (*Unlocks*) the *Semaphore* at **sem**.

```
int sem_wait(sem_t *sem);
```

Decrements (*Locks*) the *Semaphore* at **sem**.

If the *Semaphore* currently has the value zero, then the call *Blocks* until either it becomes possible to perform the decrement or a *Signal* handler *Interrupts* the call.



Semaphores

Semaphore Uses

Mutual Exclusion

- Case of *Binary Semaphore*

```
sem_init(s, 0 or 1, X=1)
```

```
sem_wait(s);  
// critical section  
sem_post(s);
```

```
sem_wait(s);  
// critical section  
sem_post(s);
```

Execution Ordering

- Case of Limiting Concurrent Access → *Counting Semaphore*

```
sem_init(s, 0 or 1, X=0)
```

```
// 1st half  
// of computation  
sem_post(s);
```

```
sem_wait(s);  
// 2nd half  
// of computation
```



Semaphores

Producer-Consumer (Bounded-Buffer) Problem

Bounded Buffer

- size N , Access entry $0 \dots N-1$, then “wraps around” to 0 again

Producer Thread : Writes data to *Bounded Buffer*

Consumer Thread : Reads data from *Bounded Buffer*

Execution ordering constraints:

- *Producer* shouldn't try to produce if *Bounded Buffer* is full
- *Consumer* shouldn't try to consume if *Bounded Buffer* is empty



Semaphores

Producer-Consumer (Bounded-Buffer) Problem

Solution – 1st version

Two *Semaphores*

- `sem_t filled; // # of filled slots`
- `sem_t empty; // # of empty slots`
- *Problem:* Does this also achieve *Mutual Exclusion*?

```
sem_init(&filled, 0, 0);  
sem_init(&empty, 0, N);
```

```
void* producer(void* _arg) {  
    sem_wait(&empty);  
    ... // fill a slot  
    sem_post(&filled);  
}  
  
void* consumer(void* _arg) {  
    sem_wait(&filled);  
    ... // empty a slot  
    sem_post(&empty);  
}
```

Note:
Sequencing
operations



Semaphores

Producer-Consumer (Bounded-Buffer) Problem

Solution – Final version

Three *Semaphores*

- `sem_t filled; // # of filled slots`
- `sem_t empty; // # of empty slots`
- `sem_t mutex; // # mutual exclusion`

Note: Can also use a `pthread_mutex_t`

```
sem_init(&filled, 0, 0);  
sem_init(&empty, 0, N);  
sem_init(&mutex, 0, 1);
```

```
void* producer(void* _arg) {  
    sem_wait(&empty);  
    sem_wait(&mutex);  
    ... // fill a slot  
    sem_post(&mutex);  
    sem_post(&filled);  
}
```

```
void* consumer(void* _arg) {  
    sem_wait(&filled);  
    sem_wait(&mutex);  
    ... // empty a slot  
    sem_post(&mutex);  
    sem_post(&empty);  
}
```

Data Structure
“internal
access”
Critical Section

Note:
Fill / Empty
operations
correspond to
manipulating the
Circular Buffer's
head & tail



Semaphores

Semaphore Summary

- *Semaphores* can be used to solve any of the traditional *Synchronization* problems
- Drawbacks:
 - They are essentially shared global variables
 - Can potentially be accessed anywhere in Program
 - No direct connection between the *Semaphore* and the data being controlled by it
 - Used for both *Critical Sections* (*Mutual Exclusion*) and *Execution Ordering* (*Scheduling*)
 - No control or guarantees for their proper usage
- When used in complex code can lead to bugginess
 - Solution: Leverage *Object-Oriented Programming* to support controlled behaviors



Monitors

Monitors

An *Object-Oriented Language* construct that controls access to shared data

- *Synchronization* code added by compiler, enforced at runtime

A module that encapsulates

- ***Shared Data Structures***
 - ***Procedures*** that operate on the shared data structures
 - ***Synchronization*** between concurrent *Threads* that invoke these procedures
- Guarantees that access of its data through *Threads* is done in legitimate ways only



Monitors

A *Monitor* guarantees *Mutual Exclusion*

- Only one *Thread* can execute **any** *Monitor* Procedure **at a time**
 - The *Thread* is “inside the *Monitor*”
- If a second *Thread* invokes a *Monitor* procedure when a first *Thread* is already executing one, the second *Thread* shall *Block*
 - i.e. the *Monitor* has to have a *Wait Queue*
- If a *Thread* that is “inside a *Monitor*” *Blocks*, then another *Thread* can enter the *Monitor*

Note: A *Monitor Invariant* is a safety property associated with the *Monitor*

- It's an assertion regarding the *Monitored Variables*
- It holds whenever a *Thread* enters or exits the *Monitor*
 - i.e. the assertion holds whenever there is no *Thread* executing “inside the *Monitor*”



Monitors

Monitors

A *Monitor* is like one big *Super-Lock* for a set of operations/methods

➤ It is however a *Language-level* implementation

➤ Compiler automatically inserts the necessary *Synchronization* operations upon entry and exit of

Monitor Procedures

```
monitor account {  
    int balance;  
    public void deposit() {  
        ++balance;  
    }  
    public void withdraw() {  
        --balance;  
    }  
};
```

*Monitor
Procedures*

Example of (part of) the operations inserted at *Compile-Time*.

```
lock(this.m);  
...  
++balance;  
...  
unlock(this.m);
```

```
lock(this.m);  
...  
--balance;  
...  
unlock(this.m);
```

➤ C++ does not have *Monitors*

Note: But check out **synchronized**, C++20 *Synchronized Blocks* (experimental):

https://en.cppreference.com/w/cpp/language/transactional_memory



Monitors

Condition Variables

Remember: A *Monitor* also needs to take care of Wait, Wakeup, Queueing functionalities

- Not just *Locking*
- What if a *Thread* has to wait for something to happen/change, but is already “inside the *Monitor*”?
 - Bad if left to just *Busy-Wait*
 - Worse: No one can now get “inside the *Monitor*” (e.g. not even to take corrective actions)
- Have to be able to let a different *Thread* enter “inside the *Monitor*”

In order to achieve the above, a *Monitor* can use a different *Synchronization* mechanism:

Condition Variables

- A *Condition Variable* is associated with a condition needed for a *Thread* to make progress once it is “inside the *Monitor*”



Monitors

Condition Variables (with respect to *Monitors*)

Operations on *Condition Variables*

wait()

Suspends the calling *Thread* and releases the *Monitor Lock* (when it resumes, it will reacquire the *Lock*)

For **wait()** to be called, the *Thread* has to already be “inside the *Monitor*”

- (Should be) called when the *Condition Predicate* is **false**

signal()

Resumes one *Thread* waiting in **wait()**, if any

- (Should be) called once *Condition Predicate* becomes **true**, and wants to **Wakeup one** waiting *Thread*

broadcast(): Resumes all *Threads* waiting in **wait()**

- (Should be) called once *Condition Predicate* becomes **true**, and wants to **Wakeup all** waiting *Threads*

Note: Condition Variables are not **boolean** objects; they are *associated* with a **boolean** *Condition Predicate*

- **if (cv) then ...** does not make sense
- ✓ **if (num_resources == 0) then wait(cv)** does



Monitors

Condition Variables (with respect to *Monitors*)

Although operations have similar names with *Semaphores*, they are different

- But one can be used to implement the other

Access to the *Monitor* is controlled by a *Lock*

wait(): Blocks the calling *Thread*, and gives up the *Lock*

- To call **wait()**, the *Thread* has to be “inside the Monitor” (hence holds the *Lock*)
 - *Semaphore*’s **sem_wait()** just blocks the *Thread* on the *Queue*

signal(): Causes a waiting *Thread* to Wakeup

- If there is no **wait()**ing *Thread*, the **signal()** is lost
 - Remember: *Semaphore*’s **sem_post()** increases its count, allowing future entry even if no *Thread* is Waiting right now
- I.e. *Semaphores* are “sticky”, *Condition Variables* have no “memory”
 - If no one is Waiting for a **signal()**, it is lost



Monitors

Condition Variables (with respect to *Monitors*)

Producer-Consumer with *Monitors*

```
monitor ProducerConsumer {  
    int nfilled = 0;  
    cond has_empty, has_filled;  
    void produce() {  
        if (nfilled == N)  
            wait (has_empty);  
        ... // fill a slot  
        ++ nfilled;  
        signal (has_filled);  
    }  
    void consume() {  
        if (nfilled == 0)  
            wait (has_filled);  
        ... // empty a slot  
        -- nfilled;  
        signal (has_empty);  
    }  
};
```

A (one) *Monitor* with two *Condition Variables*:

- **has_empty**: Buffer has at least one empty slot
- **has_filled**: Buffer has at least one filled slot

nfilled: Number of filled slots

E.g.:

- If a *Thread* tries to **consume()** and the Buffer is empty, it will be blocked at the *CV*. If another *Thread* tries to **consume()** again, it will also be blocked at the *CV*, etc.
- If a third *Thread* tries to **produce()**, it will pass the other *CV*'s wait, and **signal()** (one of) the first 2 *Threads*

I.e. (each) *Condition Variable* also has to have a *Queue*



Monitors

Condition Variable Signal Semantics

When **signal()** wakes up a waiting *Thread*, which *Thread* to run “inside the *Monitor*”?

- The *Signaling Thread* (/ *Signaler*), or the *Waiting Thread* (/ *Waiter*) ?

Hoare Semantics:

Suspends *Signaler*, and immediately transfers control to a *Waiter*

- The *Condition* that the *Waiter* was anticipating is guaranteed to hold when waiter executes
- Difficult to implement in practice, *Signaler* must restore *Monitor Invariants* before signaling

Remember: Assertions that hold whenever no *Thread* is “inside the *Monitor*”; i.e. implementation needs to remember state because *Thread* hasn’t completed yet

Mesa Semantics

Signal moves a single *Waiter* from the blocked state to a runnable state, then the *Signaler* continues until it “exits the *Monitor*”

- Problem: *Condition Variable’s Predicate* is not necessarily true when *Waiter* gets to run again
 - Return from **wait()** is only a hint that something changed, **always have to recheck Predicate**
- E.g. *Spurious Wakeup* – Fill one single slot and **signal()**, but before a scheduled woken consumer grabs the *Queue Lock* to continue, a different (e.g. fourth) *Thread* enters the *Queue*, grabs the *Lock*, consumes the one filled slot. The woken *Thread* will find the *Predicate* changed once it runs.



Monitors

Condition Variables

Producer-Consumer with Monitors

```
monitor ProducerConsumer {
    int nfilled = 0;
    cond has_empty, has_filled;

    void produce() {
        while (nfilled == N)
            wait (has_empty);
        ... // fill a slot
        ++ nfilled;
        signal (has_filled);
    }

    void consume() {
        while (nfilled == 0)
            wait (has_filled);
        ... // empty a slot
        -- nfilled;
        signal (has_empty);
    }
};
```

Spurious Wakeup – pthread

- **pthread_cond_signal()** is only guaranteed to unblock **at least one Thread**
- Even worse, a *Thread* blocked in **pthread_cond_wait** can return with no **pthread_cond_signal/broadcast()** call

Spurious Wakeup Fix:

- When woken up, a *Thread* **must** recheck the *Predicate* associated to the *Condition Variable* it was waiting on
- Most systems use *Mesa Semantics*
 - e.g. **pthread**



Monitors

Monitor & Condition Variables with **pthread**

Producer-Consumer with *Monitors*

```
class ProducerConsumer {
    int nfull = 0;
    pthread_mutex_t m;
    pthread_cond_t has_empty,
                    has_full;

public:
    void produce() {
        pthread_mutex_lock(&m);
        while (nfull == N)
            pthread_cond_wait(&has_empty,
                              &m);

        ... // fill slot
        ++ nfull;
        pthread_cond_signal(has_full);
        pthread_mutex_unlock(&m);
    }
    ...
};
```

C/C++ don't provide *Monitors*, but we can implement such functionality using **pthread_mutex_t** and **pthread_cond_t**

For the *Producer-Consumer* problem, we need 1 *Mutex* and 2 *Condition Variables*

➤ Manually lock and unlock *Mutex* for *Monitor* procedures

➤ **int pthread_cond_wait(**
 pthread_cond_t *restrict cond,
 pthread_mutex_t *restrict mutex);

Atomically waits on **cond** and releases **mutex**

The function shall *Block* on a *Condition Variable*. It shall be called with *mutex* locked by the calling *Thread* or Undefined Behavior results.

The function **atomically** releases **mutex** and causes the calling *Thread* to *Block* on **cond**... Upon successful return, the **mutex** shall have been locked and shall be owned by the calling *Thread*.



Monitors

Monitor & Condition Variables with **pthread**

Producer-Consumer with *Monitors*

```
class ProducerConsumer {
    int nfull = 0;
    pthread_mutex_t m;
    pthread_cond_t has_empty,
                    has_full;

public:
    void produce() {
        pthread_mutex_lock(&m);
        while (nfull == N)
            pthread_cond_wait(&has_empty,
                              &m);

        ... // fill slot
        ++ nfull;
        pthread_cond_signal(has_full);
        pthread_mutex_unlock(&m);
    }
    ...
};
```

Note: Unlock the *Mutex* after calling
`pthread_cond_signal()`

C/C++ don't provide *Monitors*, but we can implement such functionality using **pthread_mutex_t** and **pthread_cond_t**

For the *Producer-Consumer* problem, we need 1 *Mutex* and 2 *Condition Variables*

➤ Manually lock and unlock *Mutex* for *Monitor* procedures

➤ `int pthread_cond_signal(
pthread_cond_t * cond);`

Atomically waits on **cond** and releases **mutex**

The function shall *Unblock* **at least one** of the *Threads* that are *Blocked* on the specified *Condition Variable* **cond**... may be called by a *Thread* whether or not it currently owns the *Mutex* that *Threads* calling **pthread_cond_wait()** ... have associated with the *Condition Variable*... however, if predictable *Scheduling* behaviour is required, then that *Mutex* is *Locked* by the **pthread_cond_signal()**-calling *Thread*



CS-446/646

Time for Questions !

