

**CS-446/646**

**I/O & Disks**

**C. Papachristos**

**Robotic Workers (RoboWork) Lab  
University of Nevada, Reno**



# I/O & Disks

## *OS Abstractions*

### **Virtualization**

Processes  
Scheduling  
Virtual Memory

### **Concurrency**

Threads  
Synchronization  
Semaphores & Monitors

### **Persistence**

I/O  
Disks  
Filesystems

I/O Management is another major component of OS

- Important aspect of computer operation
- I/O Devices vary greatly
  - Various methods to control them
- New types of Devices pop up constantly



# I/O & Disks

## I/O Devices




Issues to address:

- How should I/O be integrated into systems?
- What are the general mechanisms?
- How can we manage them efficiently?



# I/O & Disks

## Structure of Input/Output (I/O) Device



## I/O Device Interfaces

### ➤ *Port*

Connection point (I/O special Address) for Device

- *Serial Port*

### ➤ *Bus*

Daisy chain for Devices sharing a common set of wires

- *PCI Bus* (Parallel-Interface) common in PCs and servers
- *PCI Express (PCIe) Bus* (high-speed Serial-Interface)
- *Expansion Bus* (connects relatively slow Devices)

### ➤ *Controller*

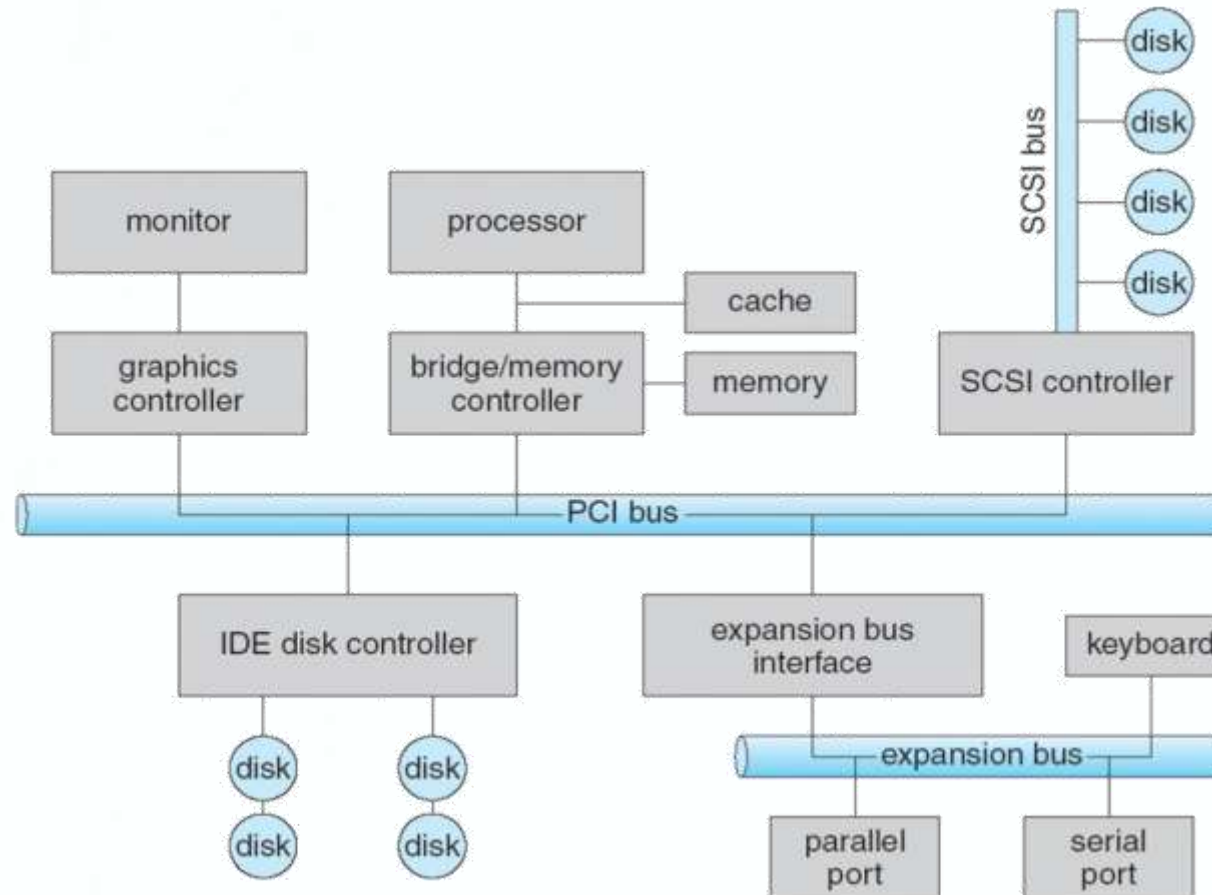
Electronics that operate *Port, Bus, Device*

- Sometimes integrated, sometimes separate circuit board (Host Adapter)
- Contains *Processor, Microcode*, private Memory, *Bus Controller*, etc.
- Some talk to per-Device *Controller* with *Bus Controller, Microcode*, Memory, etc.



# I/O & Disks

## I/O Bus – *Example: Peripheral Component Interconnect (PCI) Bus*





# I/O & Disks

## Device “Standardized” I/O *Port* Mappings on PCs

“I/O *Port*”: Technical term for a specific “special use” Address on the x86’s I/O *Bus*

### ➤ Legacy

- Used by older Hardware that was present on pre-PCI systems
  - e.g. Floppy Drive, Serial Port, Parallel Port
- Modern Architectures utilize “*Memory-Mapped*” I/O (more later) for Device communication
  - do not even have a predefined I/O *Bus*

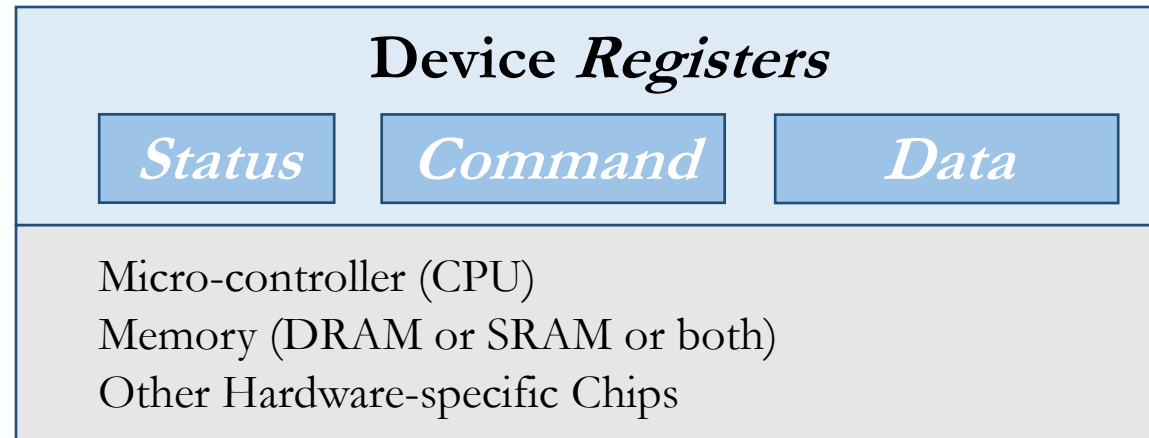
Port range	Summary
0x0000-0x001F	The first legacy <a href="#">DMA controller</a> , often used for transfers to floppies.
0x0020-0x0021	The first <a href="#">Programmable Interrupt Controller</a>
0x0022-0x0023	Access to the Model-Specific Registers of Cyrix processors.
0x0040-0x0047	The <a href="#">PIT</a> (Programmable Interval Timer)
0x0060-0x0064	The " <a href="#">8042</a> " <a href="#">PS/2 Controller</a> or its predecessors, dealing with keyboards and mice.
0x0070-0x0071	The <a href="#">CMOS</a> and <a href="#">RTC</a> registers
0x0080-0x008F	The <a href="#">DMA</a> (Page registers)
0x0092	The location of the fast <a href="#">A20</a> gate register
0x00A0-0x00A1	The second <a href="#">PIC</a>
0x00C0-0x00DF	The second <a href="#">DMA</a> controller, often used for soundblasters
0x00E9	Home of the <a href="#">Port E9 Hack</a> . Used on some emulators to directly send text to the hosts' console.
0x0170-0x0177	The secondary <a href="#">ATA</a> harddisk controller.
0x01F0-0x01F7	The primary <a href="#">ATA</a> harddisk controller.
0x0278-0x027A	Parallel port
0x02F8-0x02FF	Second <a href="#">serial port</a>
0x03B0-0x03DF	The range used for the <a href="#">IBM VGA</a> , its direct predecessors, as well as any modern video card in legacy mode.
0x03F0-0x03F7	<a href="#">Floppy disk controller</a>
0x03F8-0x03FF	First <a href="#">serial port</a>

*Note:* Also see [https://wiki.osdev.org/I/O\\_Ports](https://wiki.osdev.org/I/O_Ports)



## Canonical I/O Device

OS reads/writes  
to these



**Interface**

**Internals**





## Hardware Interface Of Canonical Device

*Registers-based:*

- By reading or writing the three *Registers*, OS controls Device behavior
- *Status* : Read the current operating status of the Device
- *Command* : Write to command the Device to perform a certain task
- *Data* : Write data to the Device, or read data from the Device
- Typical interaction example:

```
while (STATUS == BUSY); //wait until device is not busy  
  
write data to data register  
write command to command register //doing this starts the device and executes the command  
  
while (STATUS == BUSY); //wait until device is done with your request
```



## Device Interaction

How the OS communicates with the Device:

- I/O *Instructions* control Devices
  - **in** and **out** *Instructions* on x86
    - “I/O Mapping”: Special control signal from the CPU to indicate that access is performed to an I/O Port rather than a regular *Memory* location
  - Devices usually have *Registers*
    - Device *Driver* places *Commands*, *Addresses*, and *Data* there, in order to perform read/write
- “Memory-Mapped” I/O
  - Device *Registers* available as if they were *Memory* locations
    - I/O Ports “Memory-Mapped” within the same unified *Address Space* as ordinary *Memory* (but in a special reserved *Address Region*)
  - OS performs **loads** (to read) or **stores** (to write) to the Device, instead of Main *Memory*



## x86 I/O *Instructions*

- Used in conjunction with I/O *Port* Mappings

*Example:* Pintos **threads/io.h**

```
static inline uint8_t inb (uint16_t port) {
    uint8_t data;
    asm volatile ("inb %w1, %b0" : "=a" (data) : "Nd" (port));
    return data;
}

static inline void outb (uint16_t port, uint8_t data) {
    asm volatile ("outb %b0, %w1" : : "a" (data), "Nd" (port));
}

static inline void insw (uint16_t port, void *addr, size_t cnt) {
    asm volatile ("rep insw" : "+D" (addr), "+c" (cnt) : "d" (port) : "memory");
}
```



# I/O & Disks

## *Example: IDE Disk Driver w/ x86 I/O Instructions*

```
void IDE_ReadSector(int disk,
                    int off,
                    void *buf) {
    // Select Drive
    outb( 0x1F6 , disk == 0 ? 0xE0 : 0xF0);
    IDEWait();
    // Read length (1 Sector = 512 B)
    outb( 0x1F2 , 1); // 1 Sector
    outb( 0x1F3 , off); // Logical Block Address low
    outb( 0x1F4 , off >> 8); // Logical Block Address mid
    outb( 0x1F5 , off >> 16); // Logical Block Address high
    outb( 0x1F7 , 0x20); // Read command
    insw( 0x1F0 , buf, 256); // Read 256 words
}

void IDE_Wait() {
    // Discard status 4 times
    inb( 0x1F7 ); inb( 0x1F7 );
    inb( 0x1F7 ); inb( 0x1F7 );
    // Wait for status BUSY flag to clear
    while ((inb( 0x1F7 ) & 0x80) != 0);
}
```

*Remember:*

0x01F0-0x01F7	The primary ATA harddisk controller.
---------------	--------------------------------------



## *Memory-Mapped I/O*

I/O Port Mappings & **in/out** Instructions are slow and clunky

- *Instruction* format restricts what *Registers* you can use
- Only allows 216 different *Port* numbers
- Per-*Port* access control turns out to not be as useful
  - any *Port* access allows you to disable all *Interrupts*

Devices can achieve same effect with *Physical Addresses*, e.g.:

```
volatile int32_t *device_control = (int32_t *) (0xc0100 + PHYS_BASE);  
*device_control = 0x80; // write  
int32_t status = *device_control; // read
```

- Kernel must ensure specific Mapping of these *Physical* to *Virtual Addresses* across entire OS, and ensure they are *non-Cacheable*



# I/O & Disks

## *Polling*

OS waits until the Device is ready by repeatedly reading the *Status Register*

- Positive: Simple and working
- Negative: Wastes CPU time continuously waiting for the Device
  - Switching to another Ready-state *Task* would be better utilization of the CPU

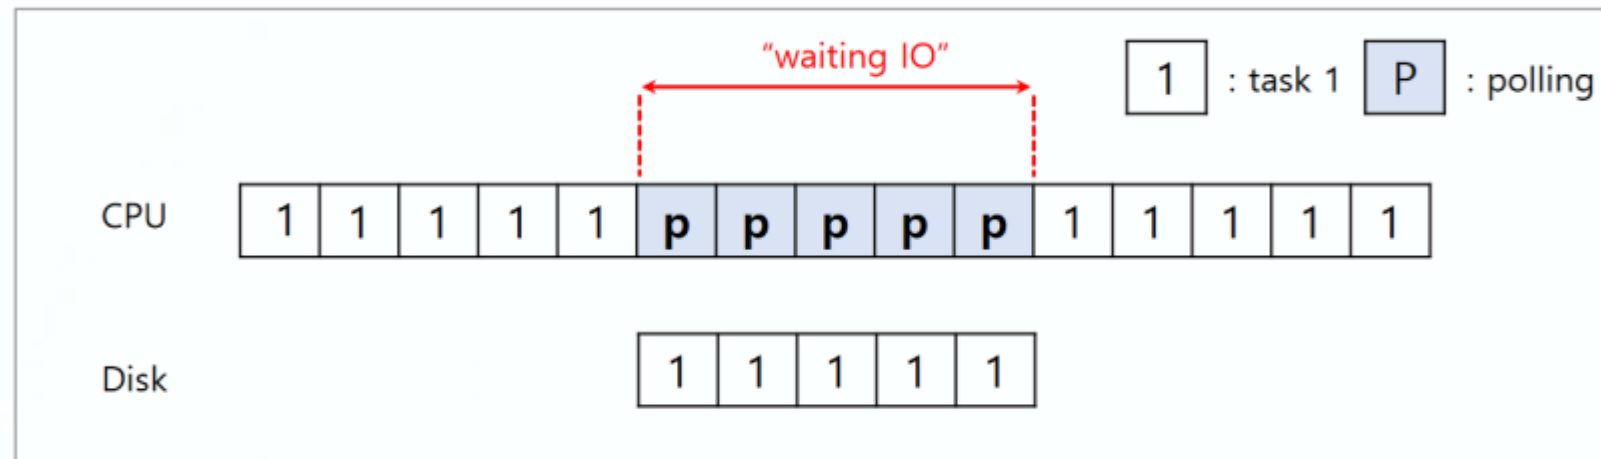


Diagram of CPU utilization when *Polling*





# I/O & Disks

## *Interrupts*

OS puts the I/O-requesting *Process* to Sleep and *Context Switches* to another

When the Device is finished, the *Process* is woken-up via *Interrupt*

- Positive: CPU and the I/O Device are more appropriately utilized
- Negative: ...




Diagram of CPU utilization with *Interrupts*





## *Polling vs Interrupts*

However, “*Interrupts* is not always the best solution”

➤ If Device performs very quickly, *Interrupts* will “slow down” the system

e.g. high *Network Packet* arrival rate

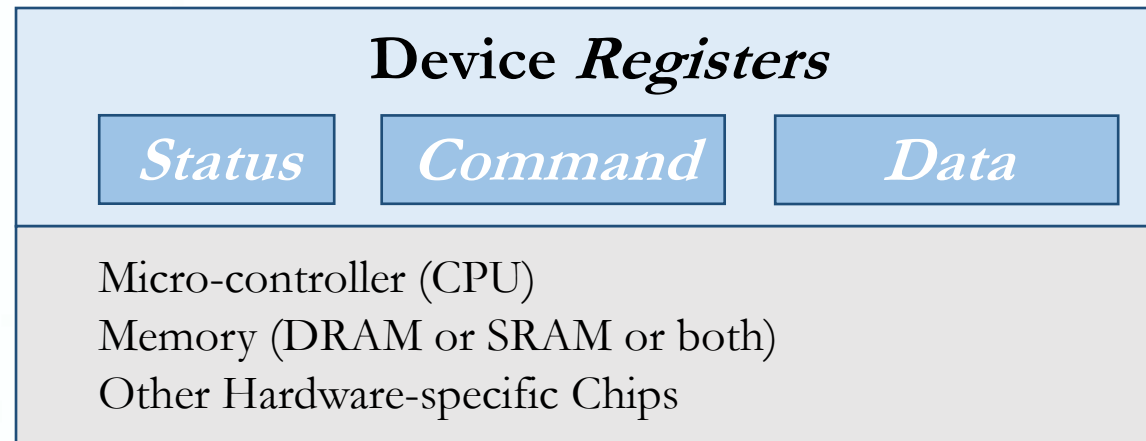
- *Network Packets* can arrive faster than OS can process them
- *Interrupts* are very expensive (*Context Switching*)
- *Interrupt Handlers* have high *Priority*
- Worst case: Spend 100% of time in *Interrupt Handlers* never making any progress – “*Receive Livelock*”
- Best case: Adaptive switching between *Interrupts* and *Polling*

## Rule-of-Thumb

- If Device is fast → *Polling*
- If it is slow → *Interrupts*



## Protocol Variants



- *Status* checking
  - *Polling* –vs– *Interrupts*
- *Data*
  - *Programmed I/O (PIO)* –vs– *Direct Memory Access (DMA)*
- *Control*
  - *Special Command Instructions* –vs– *Memory-Mapped I/O*



## Variety is a Challenge

Problem:

- Many different Devices
- Each has its own Protocol

We want to avoid writing a slightly different OS for each piece of Hardware

Solution: *Abstraction*

- Build a common Interface
- Write a specific Device *Driver* for each Device
- *Drivers* are 70% of Linux source code

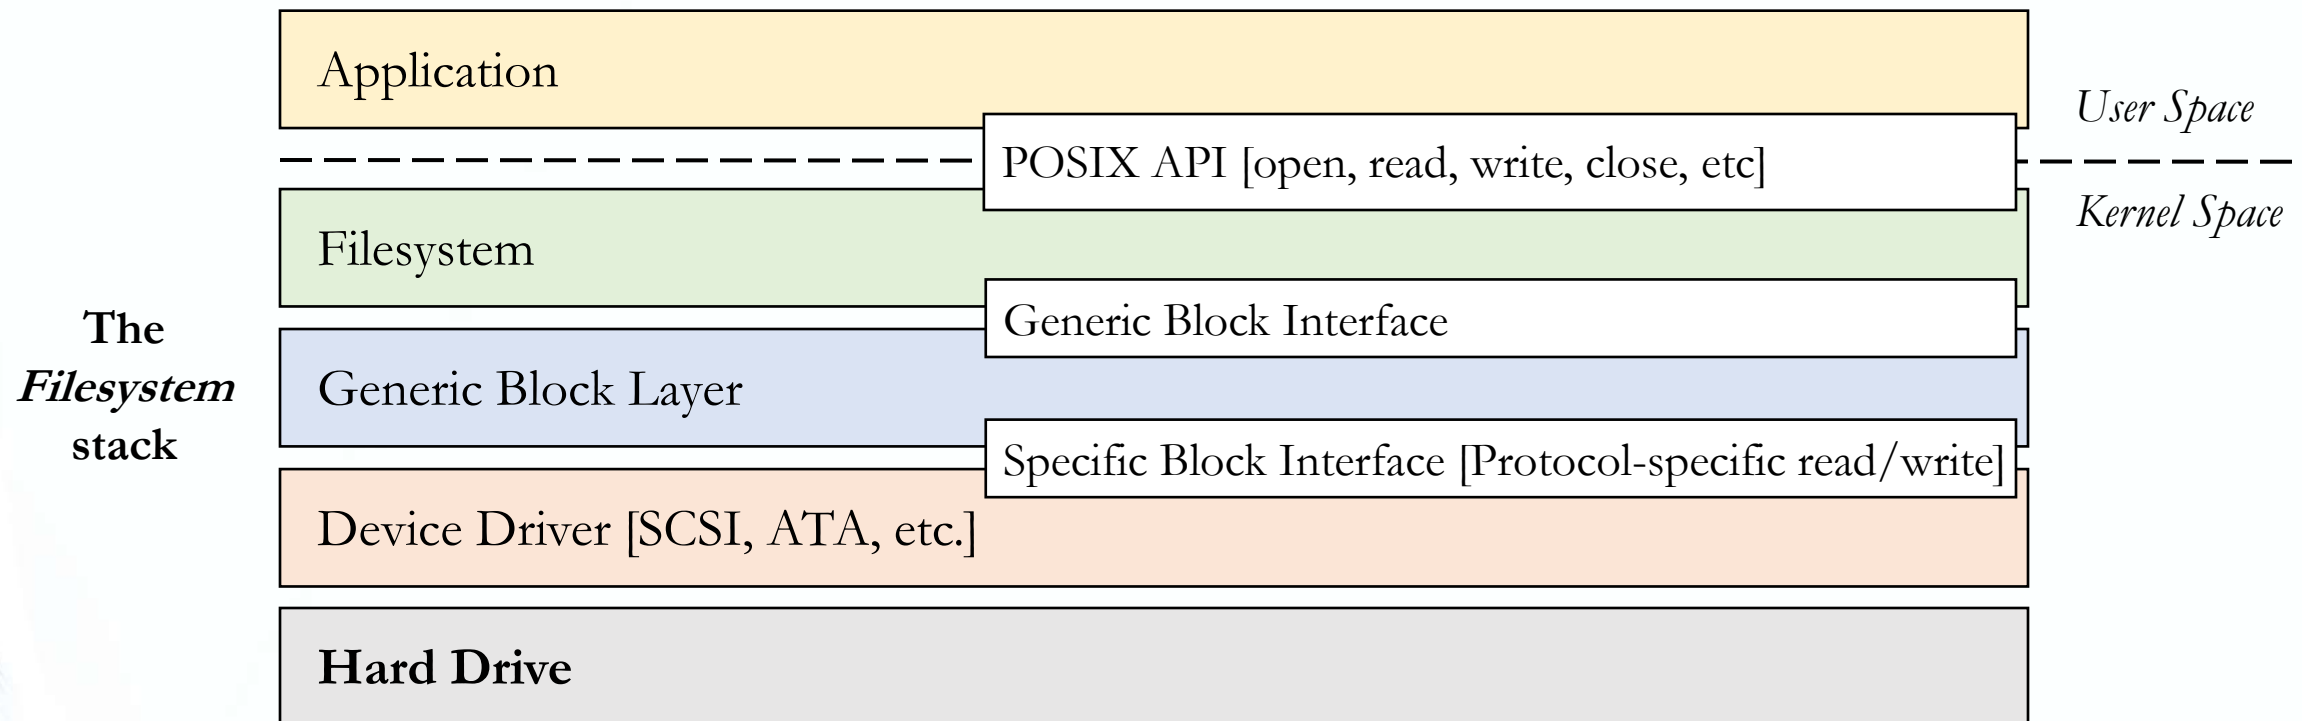


# I/O & Disks

## *Filesystem Abstraction*

*Filesystem* specifics of which Disk class it is using

➤ e.g. it issues *Block* read and write request to the *Generic Block Interface* layer



## Hard Disks – Basic Interface

- Disk *Interface* represents a linear array of *Sectors*
  - *Sector*: Historically 512 Bytes for *Hard Disk Drives* (HDDs)
  - Written atomically (even if there is a power failure)
  - 4 KB in newer “*Advanced Format*” (AF) Disks – HDDs and *Solid State Drives* (SSDs)
    - “*Torn Write*” – In an untimely power loss, only a portion of a larger write may complete
- *Disk Controller* maps the (linear) *Logical Sector Numbers* to *Physical Sectors*
  - *Physical Sectors* identified by *Surface #*, *Track #*, *Sector #* (next slides)
- OS doesn't know *Logical Sector Number* to *Physical Sector Mapping*



## Hard Disks – Basic Geometry

- *Platter* (Aluminum coated with a thin magnetic layer)
  - Disk-shaped
  - Data is stored persistently by inducing magnetic changes to it
  - Each *Platter* has 2 sides, each of which is called a *Surface*





## Hard Disks – Basic Geometry

### ➤ *Spindle*

- *Spindle* is connected to a motor that spins the *Platters* around
- The rate of rotations is measured in *RPM (Revolutions Per Minute)*
  - Typical modern values : 7,200 RPM to 15,000 RPM

### ➤ *Track*

- Concentric circles of *Sectors*
- Data is encoded on each *Surface* in a *Track*
- A single *Surface* contains many thousands of *Tracks*

### ➤ *Cylinder*


- A stack of *Tracks* of fixed radius
- Heads record and sense data along *Cylinders*
- Generally only one *Head* active at a time






# I/O & Disks

## *Cylinders, Tracks, Sectors*



# I/O & Disks

## A simple Hard Disk Drive




A single *Track* + a *Head*

- Disk *Head* – One *Head* per *Surface* of the Drive
  - The process of reading and writing is accomplished by the Disk *Head*
  - Attached to a single Disk arm, which moves across the *Surface*



## Single-track Latency: The Rotational Delay



A single *Track* + a *Head*

- *Rotational Delay*: Time for the desired *Sector* to rotate
  - Example: Full Rotational delay is  $R$  and we start at Sector 6
  - Read sector 0: *Rotational Delay* =  $R/2$
  - Read sector 5: *Rotational Delay* =  $R-1$  (worst case)



# I/O & Disks

## Multiple Tracks: Start a Read




➤ Goal: Read **Sector 12**



# I/O & Disks

**Multiple Tracks:** *Seek to Track (/ Cylinder)*



➤ Goal: Read **Sector 12**


➤ *Seek Time*: Slow (e.g. more than 0.5 – 2 ms)





# I/O & Disks

## Multiple Tracks: Wait for *Rotation*




- Goal: Read **Sector 12**
  - *Rotation Time*: Still slow (depends on mechanical motion)



# I/O & Disks

## Multiple Tracks: *Transfer Data*



- Goal: Read **Sector 12**
  - *Transfer Rate*: Fast (e.g. 125 MB/s)
    - *Transfer Time*: Fast





# I/O & Disks

**Multiple Tracks:** *Transaction Complete*




➤ Goal: Read **Sector 12**



# I/O & Disks

## Disk Latencies

- *Seek* : Move the Disk arm to the correct *Track (/ Cylinder)*
- *Seek Time* : Time to move *Head* to *Track* that contains the desired *Sector*
  - One of the most costly Disk operations
- *Rotational Delay* : Disk rotation until the correct *Sector* is reached by the *Head*
  - Still slow
- *Transfer* : Transfer of bits of information from the desired *Sector*
- *Transfer Time* : Time to perform *Transfer*
  - Fast
- *I/O Time* :  $Seek + Rotation + Transfer$



## *Seek, Rotate, Transfer*

- *Acceleration* → *Coasting* → *Deceleration* → *Settling*
- *Acceleration*: The Disk Arm gets moving
- *Coasting*: The Arm is moving at full speed
- *Deceleration*: The Arm slows down
- *Settling*: The Head is carefully positioned over the correct track
- *Seeks* often take several milliseconds!
  - *Settling* alone can take 0.5 to 2ms
  - Entire *Seek* often takes 4 - 10 ms



# I/O & Disks

## *Seek, **Rotate**, Transfer*

- Depends on *Revolutions Per Minute (RPM)*
  - 7,200 RPM is common, 15,000 RPM is high-end
- At 7,200 RPM, time to complete 1 revolution?
  - $1 \text{ min} / 7,200 \text{ RPM} = 1 \text{ second} / 120 \text{ revolutions} = 8.3 \text{ ms} / \text{revolution}$
- “Average” revolution time
  - $8.3 \text{ ms} / 2 = 4.15 \text{ ms}$



# I/O & Disks

## *Seek, Rotate, Transfer*

- The final phase of I/O
  - Data is either read from or written to the *Surface*
- Fast – Depends on RPM and *Sector* density
  - 100+ MB/s is typical for maximum transfer rate
- *Example*: Time to transfer 512 B:
  - $512 \text{ B} * (1 \text{ sec} / 100 \text{ MB}) = 5 \mu\text{s}$



# I/O & Disks

## *Workload*

- *Seeks* are slow
- *Rotations* are slow
- *Transfers* are fast

What kind of *Workload* is fastest for Disks?

- *Sequential* :  
Access *Sectors* in order (*Transfer*-dominated)
- *Random* :  
Access *Sectors* arbitrarily (*Seek*-&-*Rotation*-dominated)




# I/O & Disks

## *Sector Mapping*

- Mapping of the (linear) *Logical Sectors* to *Physical Sectors*

### *Logical Sector 0*

- The first *Sector* of the first (outermost) *Track* of the first *Surface*
- *Logical Sector* Address incremented within *Track*, then *Tracks* within *Cylinder*, then across *Cylinders*, from outermost to innermost
- “*Track Skew*”





## *Sector Mapping*

### ➤ *Default Mapping*

#### Advantages

- Simple to implement
- *Default Mapping* reduces *Seek Time* for *Sequential Access*

#### Limitations

- *Filesystem* can't infer mapping
- Reverse-engineering of mapping in OS is difficult
  - Number of *Sectors* per *Track* changes (i.e. with radius)
  - Disk Hardware can silently remap *Bad Sectors*



## *Disk Cache*


- Separate internal *Memory* (8MB - 32MB) that is used as Hardware Cache
- “*Read-Ahead*”: Acts as *Track Buffer*
  - Read contents of entire *Track* into *Memory* during *Rotational Delay*
- Write-caching with volatile *Memory*
  - “*Write-Back*” or “*Immediate Reporting*”:
    - Claim written to *Disk* when not actually done yet
      - Faster, but data could be lost on power failure
  - “*Write-through*”:
    - *Ack* after data actually written to *Platter*



# I/O & Disks

## *Disk Scheduling*

- *Disk Scheduler* decides which I/O request to schedule next
- Goal: Minimize *Positioning Time*
  - Performed by both OS and Disk itself (Why?)
- Schedule requests in order received ( FCFS )
  - Advantage: *Fairness*
  - Disadvantage: High *Seek* cost (+ *Rotation*)
- Handle nearest *Cylinder* next ( SSTF )
  - Advantage: Reduces arm movement (*Seek Time*)
  - Disadvantage: Unfair, can *Starve* some requests
- One-direction *Sweeping* of Disk ( SCAN / C-SCAN )
- If request comes for a *Block* already serviced in this *Sweep*, queue it for next *Sweep*



## *Disk Scheduling* – FCFS

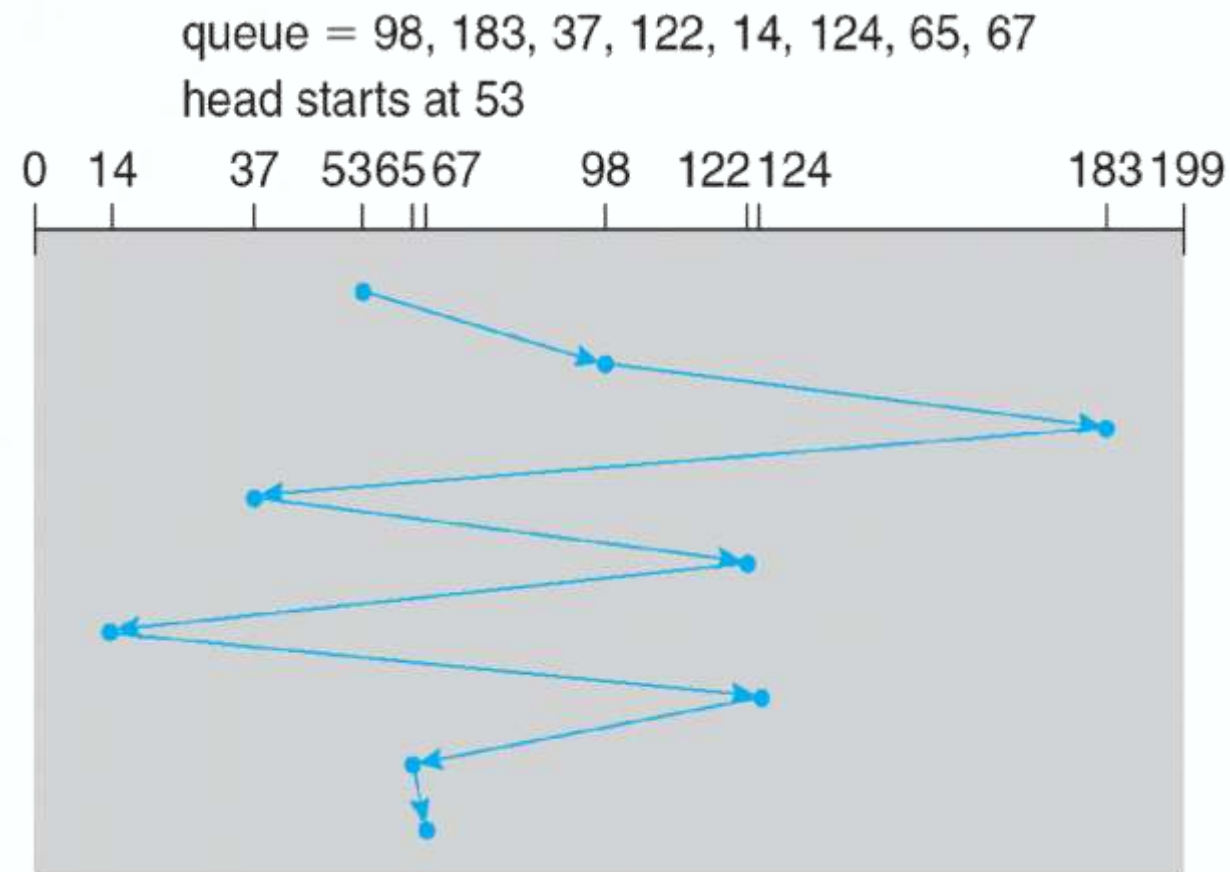
*First Come First Served* (FCFS )

- Process Disk requests in the order they are received
- Advantages
  - Easy to implement
  - Good *Fairness*
- Disadvantages
  - Cannot exploit request *Locality*
  - Increases *Average Latency*, decreasing *Throughput*



## *Disk Scheduling – FCFS*

➤ *Example:*



## ***Disk Scheduling – SSTF (/SPTF)***

*Shortest Seek-Time First ( SSTF ) –or– Shortest Positioning Time First ( SPTF )*


- Order the queue of I/O requests by *Track*
- Pick requests on the nearest *Track* to complete first
- Advantages
  - Exploits *Locality* in Disk requests
  - Higher *Throughput*
- Disadvantages
  - *Starvation*
  - Can't always know which request will be the fastest





## *Disk Scheduling – SSTF (/SPTF)*

➤ *Example:*



## *“Elevator” Scheduling – SCAN (/C-SCAN)*


Sweep across Disk, servicing all requests passed

- Like SSTF, but next *Seek* must be in same direction
- Switch directions only if no further requests in same direction
- Advantages
  - Takes advantage of *Locality*
  - Bounded waiting
- Disadvantages
  - Cylinders in the middle get better service
  - Might miss some *Locality* which SSTF could exploit
- *Circular-SCAN* ( C-SCAN ): Only sweep in one direction
  - Very commonly used algorithm in Unix



## *Disk Scheduling – C-SCAN*

➤ *Example:*



## Disk Technology Trends

- Data → More dense
  - More bits per square inch
  - Disk *Head* closer to *Surface*
  - Create smaller Disks with same *Capacity*
- Disk geometry → Smaller
  - Spin faster → Increase *Bandwidth*, reduce *Rotational Delay*
  - Faster *Seek*
  - More lightweight
- Disk price → Cheaper
  - Density improving more than speed (mechanical limitations)



## New *Mass Storage* Technologies

- New Solid-State Memory-based *Mass Storage* technologies avoid *Seek Time* and Rotational Delay
  - NAND Flash
  - Battery-backed DRAM (NVRAM)

## Disadvantages

- Price: More expensive than same *Capacity* Disk
- Reliability: More likely to lose data
- Open research question:  
How to effectively use *Flash* in Commercial Storage systems



## *Flash Memory*

Today, we increasingly use *Flash Memory*

- Completely Solid-State (no moving parts)
  - Remembers Data by storing charge
  - Lower power consumption and heat
  - No mechanical *Seek* Times to worry about
- Limited # of overwrites possible
  - *Blocks* wear out after 10,000 (MLC) – 100,000 (SLC) erases
  - Requires *Flash Translation Layer (FTL)* to provide wear leveling, so repeated writes to same *Logical Blocks* don't wear out *Physical Blocks*
  - FTL can seriously impact performance
- Limited durability
  - Charge wears out over time
  - Turning off Device for a year may even lead to loss of data





## *Redundant Array of Independent Disks (RAID)*

Motivation:

- Performance
  - Disks are slow compared to CPU
  - Disk speed improves slowly compared to CPU
- Reliability
  - In single-Disk systems, one Disk failure leads to data loss
- Cost
  - A single fast & reliable Disk is expensive



## *Redundant Array of Independent Disks (RAID)*

Idea:

- Use redundancy to improve Performance and Reliability
  - Redundant array of cheap Disks as one storage unit
  - Fast: Simultaneous read and write of Disks in the array
  - Reliable: Use “*XOR-Magic*” Parity:
    - To detect Errors
    - To rebuild missing data in case of any 1 Disk Failure
- RAID can have different redundancy levels, achieving different Performance and Reliability criteria
  - Seven different RAID levels (RAID 0-6)



## *Redundant Array of Independent Disks (RAID)*

Evaluating RAID:

- Cost-wise
  - Storage utilization:  $\text{Data Capacity} / \text{Total Capacity}$
- Reliability-wise
  - Tolerance against Disk failures
- Performance-wise
  - Perform (large) *Sequential* read, write, read-modify-write
  - Perform (small) *Random* read, write, read-modify-write
  - Measure speedup over a single Disk



## *Redundant Array of Independent Disks (RAID)*

Evaluating RAID:

- Computing Cost:
  - $G$  = Number of Data Disks in a RAID group
  - $C$  = Number of Check/Parity Disks in a RAID group
  - $\text{Cost} = C/(G+C)$



## *Redundant Array of Independent Disks (RAID)*

### Evaluating RAID:

#### ➤ Computing Reliability:

- $N$  = Total number of Disks
- $G$  = Number of Data Disks in a RAID group
- $C$  = Number of Check/Parity Disks in a RAID group
- $MTTF(\text{disk})$  = Mean time to failure for a single Disk
  - Estimated as  $MTTF(\text{in years}) = 1 / AFR$  (Annual Failure Rate (in percentage))
  - Ex: 114 years (1M hours) =  $1 / 0.88\%$
  - Source: "Disk failures in the real world: What does an MTTF of 1,000,000 hours mean to you?", FAST'07
- $MTTR(\text{disk})$  = Mean time to repair for a failed Disk

#### Compute:

- $MTTF(\text{group})$  = Mean time to two failed Disks before first gets repaired in one group
- $MTTF(\text{raid})$  = Mean time to failure over entire array
- $MTTF(\text{raid}) = MTTF(\text{group}) / \text{Num. groups}$



## *Redundant Array of Independent Disks (RAID)*

### Evaluating RAID:

#### ➤ Computing Reliability:

- Assume single-error tolerance in one group
  - If another error comes before repair, group fails
- $MTTF(\text{group}) = MTTF(1 \text{ disk}) / \text{Prob}[\text{Another failure within } MTTR]$ 
  - If  $\text{Prob} \approx 1$ ,  $MTTF(\text{group})$  same as  $MTTF(1 \text{ disk})$  – no benefit of RAID
  - If  $\text{Prob} \approx 0$ ,  $MTTF(\text{group})$  approaches  $\infty$  – good
- $MTTF(1 \text{ disk}) = MTTF(\text{disk}) / (D+C)$
- $MTTF(\text{another disk}) = MTTF(\text{disk}) / (D+C-1)$
- $\text{Prob}[\text{Another failure within } MTTR] = MTTR / (MTTF(\text{disk}) / (D+C-1))$
- $MTTF(\text{group}) = MTTF(1 \text{ disk}) / \text{Prob}[\text{Another failure within } MTTR]$   
 $= (MTTF(\text{disk}))^2 / ((D+C) * (D+C-1) * MTTR)$
- Num groups  $G = N / (D+C)$
- $MTTF(\text{raid}) = MTTF(\text{group}) / G = MTTF(\text{group}) / (N / (D+C))$
- Thus:  $MTTF(\text{raid}) = (MTTF(\text{disk}))^2 / (N * (D+C-1) * MTTR)$





## RAID 0: Non-Redundant Striping


- Structure:
- A set of Data Sectors striped across (Data) Disks
  - No Parity Disks

Advantages:

- Good performance – with N Disks, roughly N-times speedup

Disadvantages:

- Poor Reliability – one Disk failure → Data loss
- $MTTF(raid) = MTTF(disk) / N$



## RAID 0 Performance

Large read of 100 blocks:

- One Disk:  $100 * t$ ,
- Raid0:  $100/N * t * S$
- S: Slowdown. Need to wait for slowest Disk to complete before returning

Performance:

- Large read:  $N/S$
- Large write:  $N/S$
- Large R-M-W:  $N/S$
- Small read:  $N$
- Small write:  $N$
- Small R-M-W:  $N$

## RAID 1: Mirroring


Structure: ➤ Keep a Mirrored (Shadow) copy of each Data Sector

Advantages:

- Good Reliability: 1 Disk failure OK
- Good read Performance

Disadvantages:

- High cost: Each Data Disk requires one Parity Disk



## RAID 1 Performance

- $\text{Cost} = C/(D+C) = 1/(1+1) = 50\%$
- $\text{MTTF}(\text{raid}) = \text{MTTF}(\text{disk})^2 / (N * \text{MTTR})$

### Performance

- Large read:  $N/S$
- Large write:  $N/2S$
- Large R-M-W:  $2N/3S$ 
  - $X$  sectors,  $2X$  events ( $X$  reads,  $X$  writes)
  - Speedup (w.r.t. to 1 Disk) =  $2X / (X/(N/S) + X/(N/2S)) = 2N/3S$
- Small read:  $N$  (no  $S$  here since only two Disks)
- Small write:  $N/2$
- Small R-M-W:  $2N/3$



## RAID 2: Memory-Style Error-Correcting Parity


- Structure:
- A Data Sector striped across Data Disks
  - Compute Error-Correcting Parity and store in separate Parity Disks

Advantages:

- Good Reliability with higher Storage Utilization than Mirroring

Disadvantages:

- Unnecessary cost: 1 Parity Disk can already detect Failure (coming up next)
- Poor *Random* Performance



## RAID 3: Bit-Interleaved Parity


- Structure:
- A Data Sector striped across Data Disks
  - Single Parity Disk (XOR of each stripe of a Data Sector)

Advantages:

- Same Reliability with 1 Disk Failure as RAID 2 (since Disk Controller can determine which Disk Failed)
- Higher Storage Utilization

Disadvantages:

- Poor *Random* read Performance
- Poor *Random* write and read-modify-write Performance (bottleneck: Parity Disk updating)





## RAID 4: Block-Interleaved Parity


Structure: ➤ A set of Data Sectors (Parity Group) striped across Data Disks

Advantages:

- Same Reliability as RAID 3
- Good *Random* read Performance

Disadvantages:

- Poor *Random* write and read-modify-write Performance (bottleneck: Parity Disk updating)



## RAID 4 Performance

- One Parity Disk (XOR of Data Sectors)
  - Write Data Disk + Parity Disk
  - To update Parity, don't have to read all Disk Sectors
  - $\text{Parity} = \text{oldParity} \text{ XOR (changed bits)} = \text{oldParity} \text{ XOR (newData XOR oldData)}$
- Number of groups:  $G = N / (D + 1)$  (=number of Parity Check Disks)

### Performance

- Large read:  $(N - G) / S$
- Large write:  $(N - G) / S$
- Large R-M-W:  $(N - G) / S$
- Small read:  $N - G$
- Small write:  $\frac{1}{2} * G$  (for each Block, need a read and a write to Parity Disk)
  - RAID:  $X$  sectors:  $X / ((X/1) + (X/1)) = \frac{1}{2}$
- Small R-M-W:  $1 * G$ 
  - RAID:  $X$  sectors:  $2X / ((X/1) + (X/1)) = 1$




## RAID 5: Block-Interleaved Distributed Parity

Structure: ➤ Parity Sectors distributed across all Disks

Advantages:

- Same Reliability as RAID 3 & 4 (can tolerate 1 Disk Failure)
- Good *Small* write and read-modify-write Performance



## RAID 5 Performance

- Same as RAID 4 except no single Parity Disk
  - Good small write and read-modify-write Performance

### Performance

- Large read:  $(N-G)/S$
- Large write:  $(N-G)/S$
- Large R-M-W:  $(N-G)/S$
- Small read:  $N$
- Small write:  $N/4$ 
  - One disk:  $X \text{ sectors} * t$
  - Raid 5:  $(X \text{ (read original)} + X \text{ (read parity)} + X \text{ (write original)} + X \text{ (write parity)}) / N * t$
  - Raid5 can do  $4X$  over all  $N$  Disks
- Small R-M-W:  $N/2$ 
  - Same as small write, except read-original is not wasted




## RAID 6: P+Q Redundancy


Structure: ➤ Same as RAID 5 except using two Parity Sectors per Parity Group

Advantages:

➤ Can tolerate 2 Disk Failures



## RAID Levels



(a) RAID 0: non-redundant striping.



(b) RAID 1: mirrored disks.



(c) RAID 2: memory-style error-correcting codes.



(d) RAID 3: bit-interleaved parity.



(e) RAID 4: block-interleaved parity.



(f) RAID 5: block-interleaved distributed parity.



(g) RAID 6: P + Q redundancy.





**CS-446/646**

Time for Questions !

