

CS-446/646

Fast File System

C. Papachristos

Robotic Workers (RoboWork) Lab
University of Nevada, Reno



Fast File System

Original Unix *Filesystem*

- From Bell Labs by Ken Thompson

Simple and elegant:

Unix Disk Layout



- Components

- *Data Blocks*
- *inodes (inode Table)*
- *Freelist*
- *Superblock*

- Specifies number of *Blocks* in *Filesystem*, counts Max # of *Files*, has Pointer to *Head of Freelist*

Problem: slow

- Only gets 2% of Disk maximum (20Kb/s) even for *Sequential* Disk Transfers



Fast File System

Original Unix *Filesystem*

- Why so slow?
- Problem 1: *Blocks* too small (512 Bytes)
 - *inode Table* too large
 - Requires more *Indirect Blocks*
 - *Transfer Rate* low (get one *Block* at time)
- Problem 2: Unorganized *Freelist*
 - Consecutive *File Blocks* not close together
 - Pay *Seek* cost even for *Sequential Access*
 - *Aging*: Becomes *Fragmented* over time
- Problem 3: Poor *Locality*
 - *inode Table* far from *Data Blocks*
 - *inodes* for *Directories* not close together
 - Poor performance doing enumeration (e.g. `ls`, `grep foo *.c`)



Fast File System

Unix *Fast File System* (FFS)

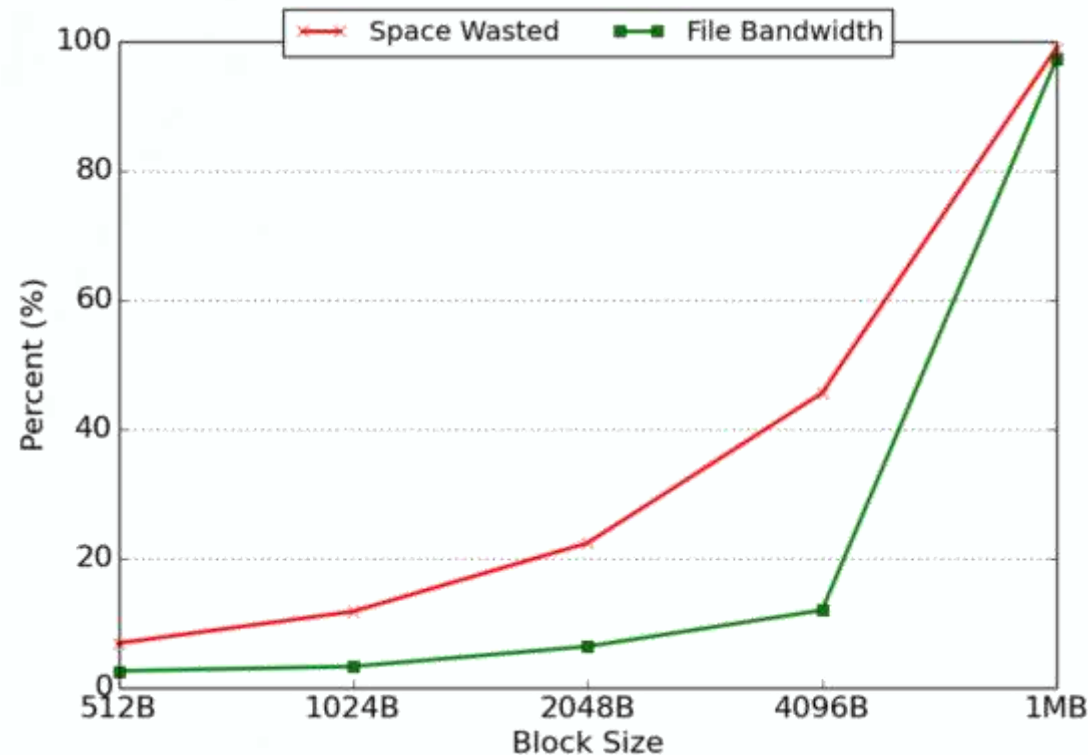
- Designed by a Berkeley research group for the BSD UNIX
 - Seminal *Filesystems* paper to read: “*A Fast File System for UNIX*”, McKusick et al.
- Approach:
 - Measure state of the art *Filesystems*
 - Identify and understand the fundamental problems
 - The original *Filesystem* treats Disks like *Random-Access Memory* !
 - Build a better *Filesystem*
- Idea: Design *Filesystem* structures and Allocation Policies to be “Disk-aware”
- Next: Performance problems and how FFS fixes them



Fast File System

Problem 1: Blocks Too Small

Measurement:




- Larger *Block* increases Bandwidth, but how to deal with *Wastage* (/ *Internal Fragmentation*)?
 - Use idea from `malloc()`: Split unused portion



Fast File System

FFS Solution: *Fragments*

- BSD FFS:
 - Has large *Block Size* (4096B or 8192B)
 - Allow large *Block* to be chopped into smaller ones called “*Fragments*”
 - Ensure *Fragments* only used either for a) small *Files* or b) Ends of *Files*



- *Fragment Size* specified at the time that the *Filesystem* is created
- Limit number of *Fragments* per *Block* to 2, 4, or 8


Advantages:

- High *Transfer Speed* for larger *Files*
- Low *Wasted Space* for a) small *Files* or b) Ends of *Files*



Fast File System


Fragments Example



Fast File System

Fragments Example


```
write(fd1, "A"); // append A to first file
```



Fast File System

Fragments Example

```
write(fd1, "A"); // append A to first file  
write(fd1, "A");
```



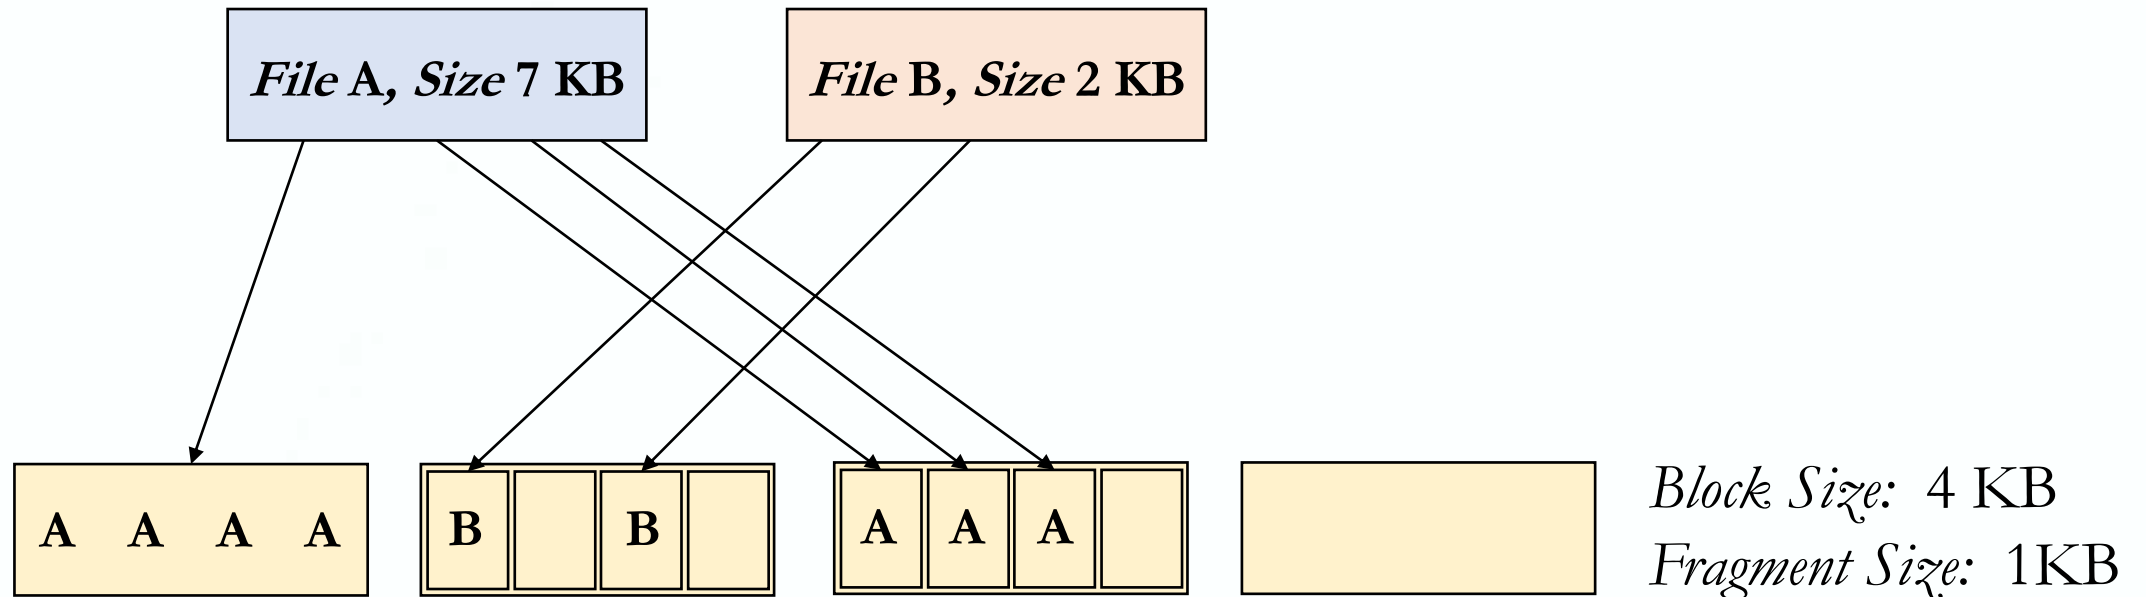
➤ But, not allowed to use *Fragments* across multiple *Blocks*...



Fast File System

Fragments Example

```
write(fd1, "A"); // append A to first file  
write(fd1, "A");
```



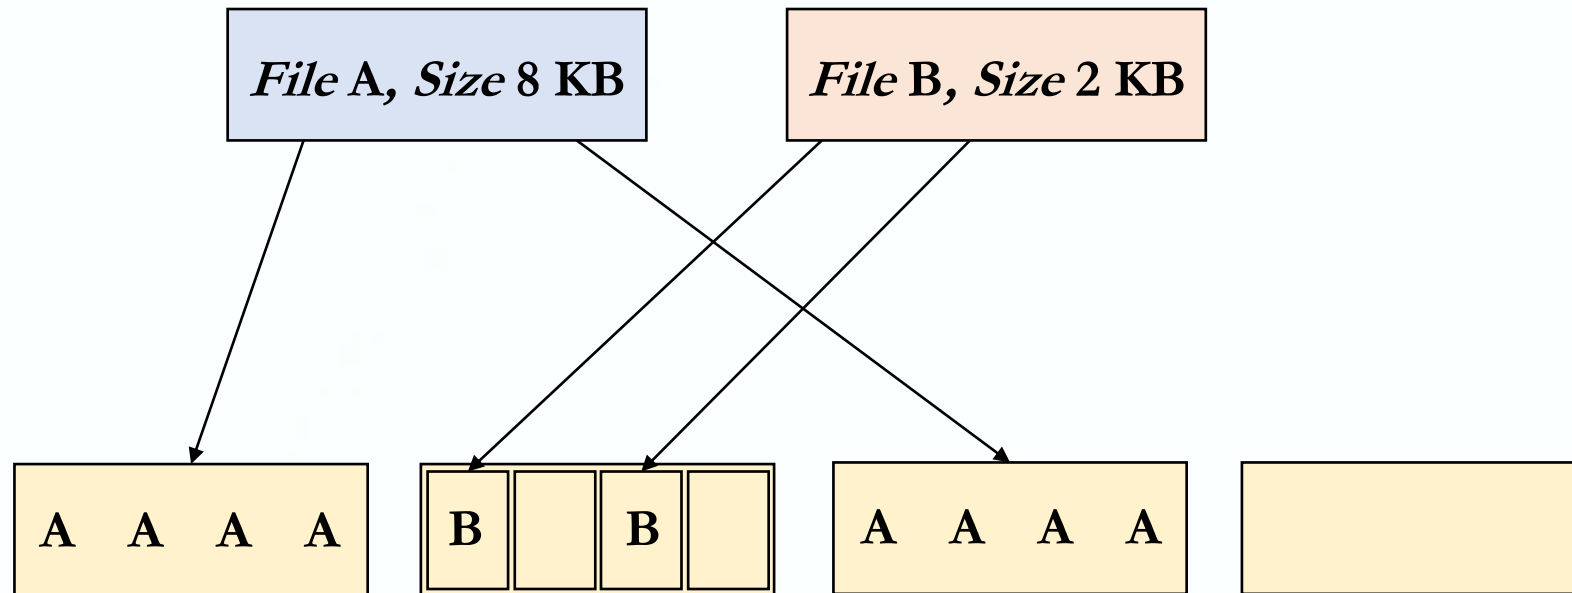
- ... so, copy old *Fragments* to new *Block*
- Any new Data will use remaining *Fragments*



Fast File System

Fragments Example

```
write(fd1, "A"); // append A to first file  
write(fd1, "A");  
write(fd1, "A");
```




Block Size: 4 KB
Fragment Size: 1KB




Fast File System

Problem 2: Unorganized *Freelist*

- Leads to random-like allocation of *Sequential File Blocks* over time



Initial Performance good



Gets worse over time

Measurement:

- New *Filesystem*: 17.5% of Disk Bandwidth
- Few weeks old: 3% of Disk Bandwidth



Fast File System

Solutions for Unorganized *Freelist*

- Periodic Disk *Defragmentation*
 - Cons: Locks-up Disk Bandwidth during operation
- Keep adjacent *Free Blocks* together on *Freelist*
 - Cons: Costly to maintain
- FFS Solution: *Bitmap* of *Free Blocks*
 - Each bit indicates whether *Block* is *Free*
 - e.g. 1010101111111000001111111000101100
 - Easier to find Contiguous *Free Blocks*
 - Small, so usually keep entire thing in *Memory*
 - Time to find *Free Blocks* increases if fewer *Free Blocks* are available
 - Consideration:
Also handle *Block Fragments*

Bits in map	XXXX	XX00	00XX	0000
Fragment numbers	0-3	4-7	8-11	12-15
Block numbers	0	1	2	3



Fast File System

FFS Solution: *Bitmap of Free Blocks*

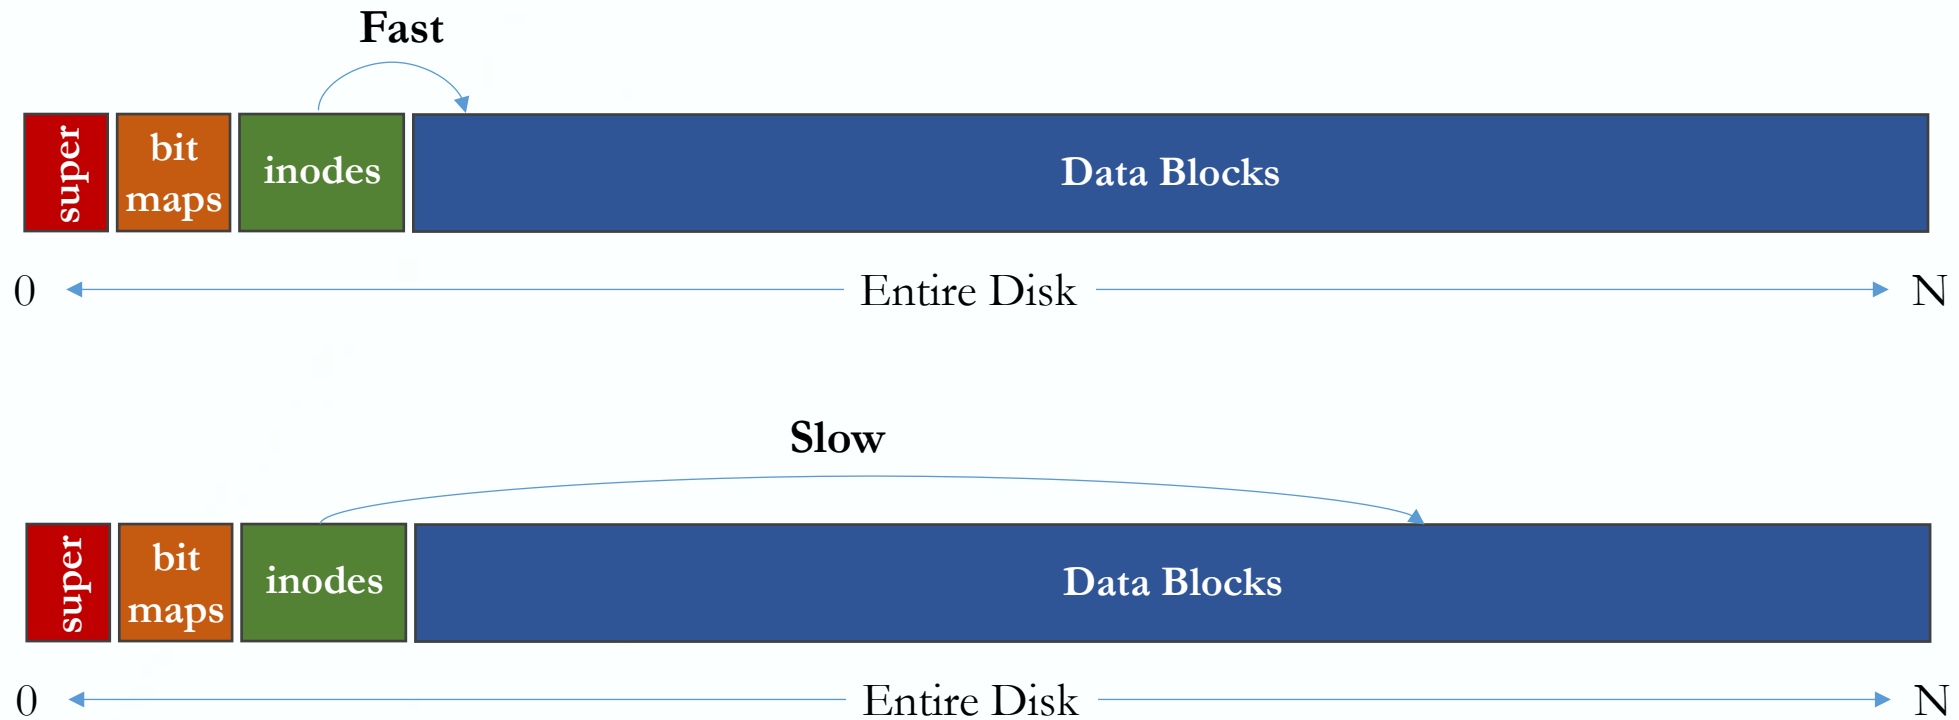
- Usually keep entire *Bitmap* in Memory:
 - 4 GiB disk / 4 KiB *Blocks* → *Bitmap* Size: 1 Mi *Entries* → 1 Mbit = 125 KB
- Allocate *Block* close to *Block x*?
 - Check for *Blocks* near `bmap[x/32]` (e.g. assuming `int32_t bmap[125*1024]`)
 - If Disk almost empty, will likely find one near
 - As Disk becomes full, search becomes more expensive and less effective
- Trade Space for Time (*Search Time, File Access Time*)
 - Instead of *Freelist* (effectively just a Pointer to *Head*), use *Bitmap of Free Blocks*



Fast File System

Problem 3: Poor *Locality*


- Desired to keep *inode* close to *Data Block*



Fast File System


Problem 3: Poor *Locality*

- Desired to keep *inode* close to *Data Block*
- Example bad Layout and Access Sequence:



Fast File System


Remember. Cylinders, Tracks, & Sectors



Fast File System

FFS Solution: *Cylinder Group*

- Group sets of consecutive *Cylinders* into “*Cylinder Groups*”



- Key Concepts:
- Can access any *Block* in a *Cylinder* without performing a *Seek*
 - Next fastest place is adjacent *Cylinder*
 - Try to put everything related in same *Cylinder Group*
 - Try to put everything unrelated in different *Groups*




Fast File System

Clustering in FFS


➤ Try to put *Sequential Blocks* in adjacent *Sectors*

- Assumption: If Accessing one *Block*, probably will need to Access next *Block* as well



➤ Try to keep *inode* in same *Cylinder* as *File Data*:

- Assumption: If accessing an *inode*, most likely will also access the *File's Data* too



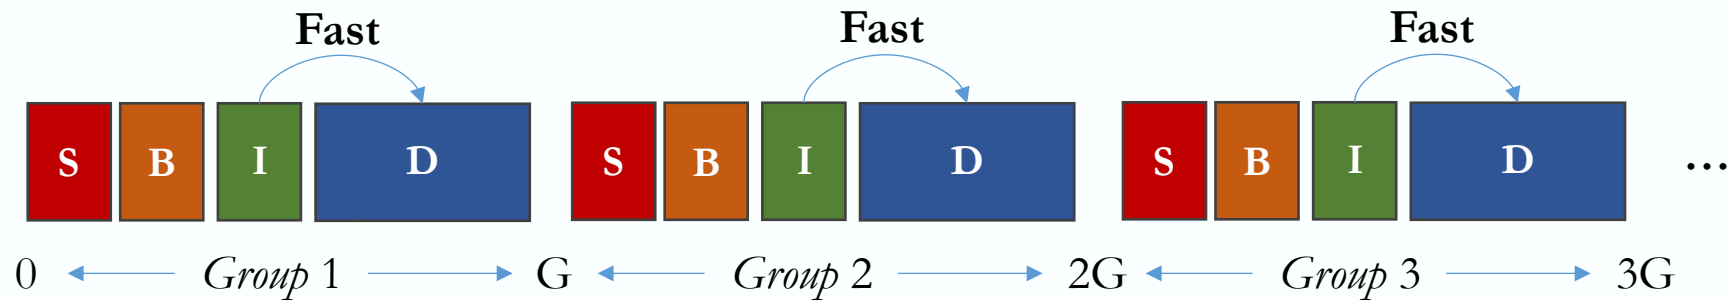
➤ Try to keep all *inodes* in a *Directory* in same *Cylinder Group*

- Assumption: If accessing one Name, frequently need to access many locally, e.g. `ls -l`)



Fast File System

Resulting FFS Disk Layout:



Try to keep *inodes* close to *Data Blocks*

- Use *Groups* across Disk
- Strategy: Allocate *inodes* and *Data Blocks* in same *Cylinder Group*
- Each *Cylinder Group* basically a mini-Unix *Filesystem*
- Utility of multiple *Superblocks*:
 - *Reliability*: Superblock contains general *Filesystem* description, necessary to mount it
 - Each *Group's S* is a copy the FFS *Superblock*



Fast File System

FFS Results

- Performance improvements:
 - Able to get 20-40% of Disk Bandwidth for large *Files*
 - 10-20x of original Unix *Filesystem*
 - Stable over *Filesystem* lifetime
 - Better small *File* performance
 - *Locality*: Small *Files* take up *Fragments* of same *Block*
- Other enhancements
 - Long *Filenames*
 - Parameterization
 - Maintains *Free Space* reserve (10%) that only admin can allocate *Blocks* from



CS-446/646

Time for Questions !

