

**CS-446/646**

Security

**C. Papachristos**

Robotic Workers (RoboWork) Lab  
University of Nevada, Reno



# Mandatory Access Control

## *DAC* vs *MAC*

Most people are familiar with *Discretionary Access Control (DAC)*

- Unix *Permission bits* are an example
  - e.g. might set *File* **private** such that only *Group* **friends** can read it:
  - **-rw-r--- 1 dm friends 1254 Feb 11 20:22 private**
- Anyone with access to information can further propagate that information at their discretion:
  - **\$ mail sigint@enemy.gov < private**

*Mandatory Access Control (MAC)* can restrict Propagation

- *Note:* MAC here not be confused with *Message Authentication Codes* or *Medium Access Control*
- Security administrator may allow you to read but not disclose *File*



# Mandatory Access Control


## **MAC Motivation**

- Prevent users from disclosing sensitive information (whether accidentally or maliciously)
  - e.g. Classified information requires such protection
- Prevent Software from surreptitiously leaking data
  - Seemingly innocuous Software may steal secrets in the background
  - Such a program is known as a “*Trojan Horse*”
- Case study: *Symantec AntiVirus 10*
  - Contained a “*Remote Exploit*” (attacker could run arbitrary code)
  - Inherently required access to all of a user’s *Files* to scan them
  - Can an OS protect private (e.g. Classified) *File* contents under such circumstances?



# Mandatory Access Control

*Example: Anti-Virus (AV) Software*

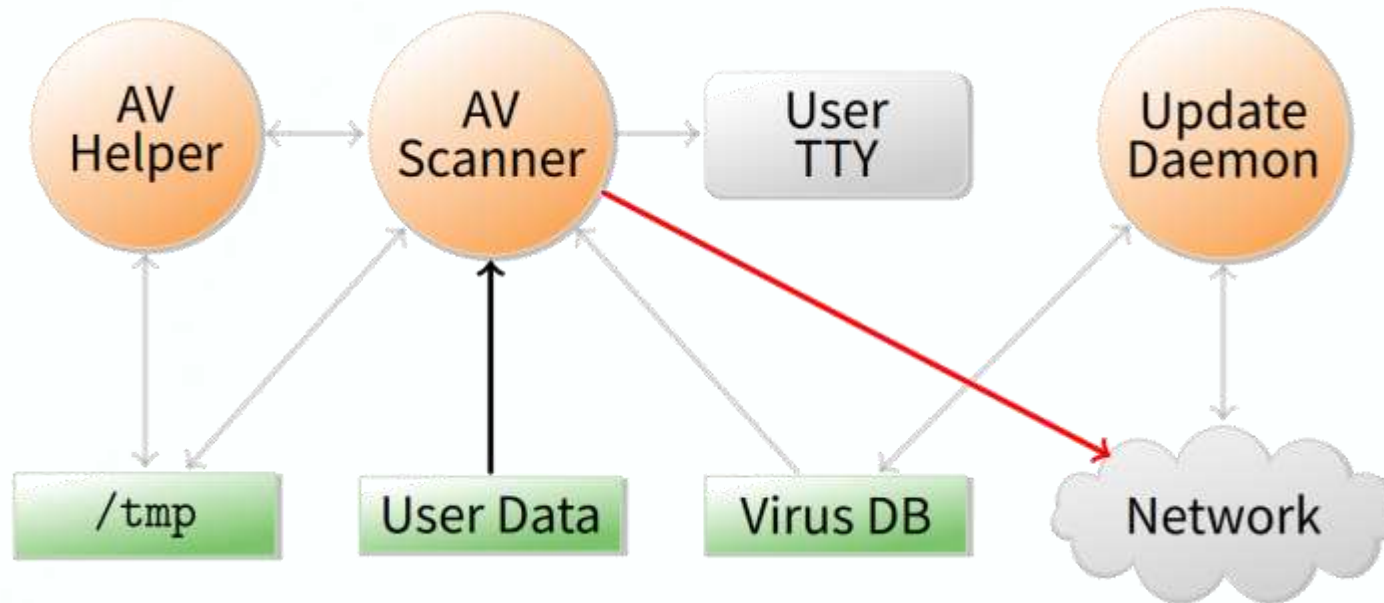


- *Scanner* : Checks for Virus Signatures
- *Update Daemon* : Downloads new Virus Signatures
- How can OS enforce *Security* without blindly trusting AV Software?
  - Must not leak contents of your *Files* to Network
  - Must not tamper with contents of your *Files*



# Mandatory Access Control

*Example: Anti-Virus (AV) Software*



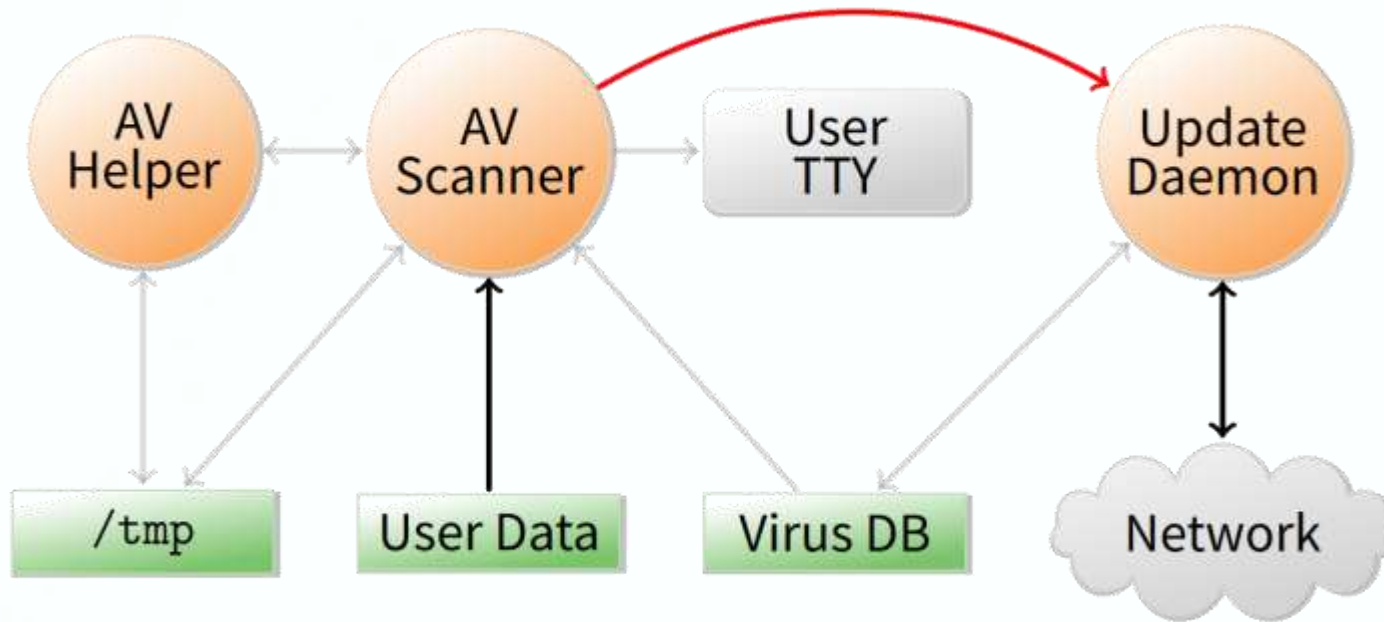
- *Scanner* can write your private data to Network
- Prevent *Scanner* from invoking any *System Call* that might send a Network message?





# Mandatory Access Control

*Example: Anti-Virus (AV) Software*

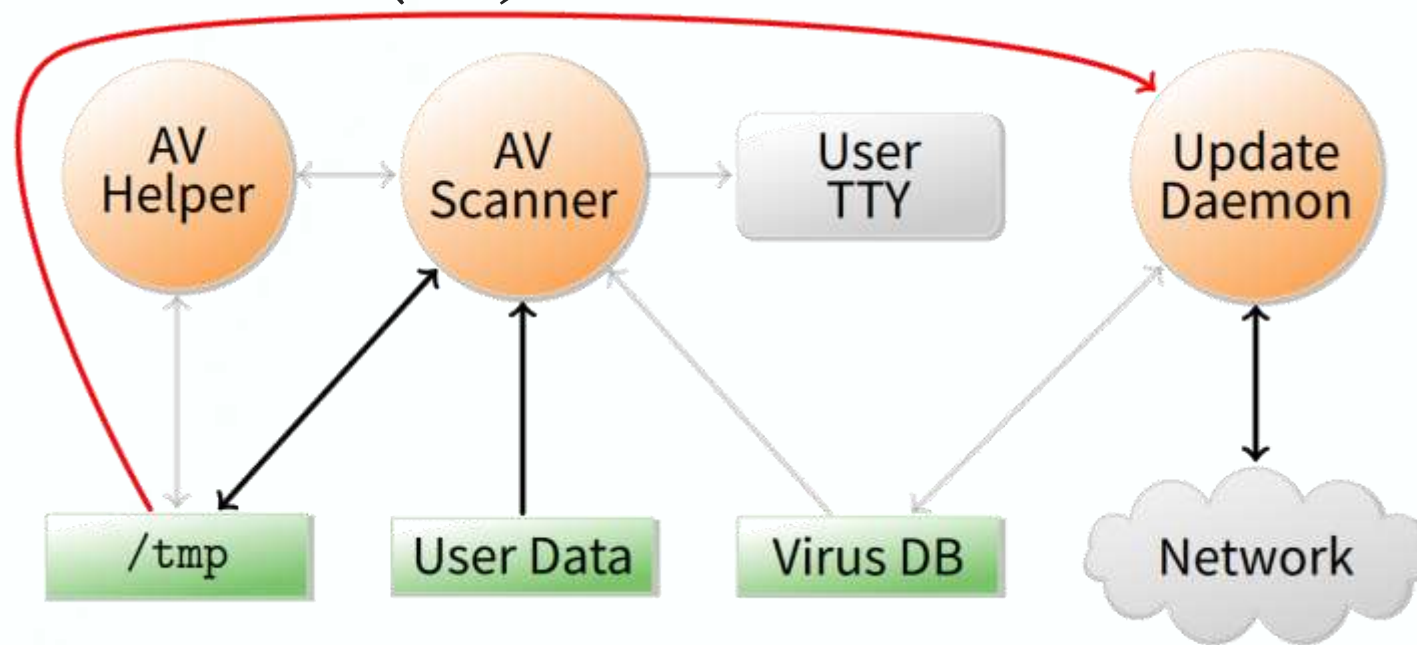


- *Scanner* can send private data to *Update Daemon*
- *Update Daemon* sends data over Network
  - Can cleverly disguise secrets in order/timing of update requests
- Block *Inter-Process Communication & Shared Memory*–related *System Calls* in *Scanner*?



# Mandatory Access Control

*Example: Anti-Virus (AV) Software*




- *Scanner* can write data to world-readable *File* in **/tmp**
- *Update Daemon* later reads and discloses *File*
- Prevent *Update Daemon* from using **/tmp**?



# Mandatory Access Control

*Example: Anti-Virus (AV) Software*



- *Scanner* can acquire read *Locks* on Virus Signature Database
  - **Encode** private User data by *Locking* various ranges of *File*
- *Update Daemon* **decodes** data by detecting *Locks*
  - Then disclose private User data over the Network
- Have *Trusted* Software copy Virus DB for *Scanner*?





# Mandatory Access Control

## The list goes on...

- *Scanner* can call **setproctitle** with User data
  - *Update Daemon* can then extract User data by running **ps**
- *Scanner* can **bind** particular TCP or UDP *Port* numbers
  - Sends no Network traffic, but used *Port* numbers detectable by *Update Daemon*
- *Scanner* can relay data through another *Process*
  - Call **ptrace** to take over *Process*, then write to Network
  - Use **sendmail**, **httpd**, or **portmap** to reveal data
- Disclose data by modulating free *Disk* space
- Can be certain we've covered all possible communication channels?
  - Not without a more systematic approach to the problem



# Labels and Lattices

## *Secrecy / Confidentiality*

- Achieve Controlled Access to Classified information

Typical way to think of *Security* in this context:

- A subject at a given *Security Level* should not be able to read information at higher *Security levels*, **and**
- A subject at a given *Security Level* should not be able to write information (leak it) at lower *Security levels*
- I.e. a User can create content only at or above their own *Security Level* (e.g. a Secret-level researcher can create Secret-level or Top-Secret-level *Files* but may not create Public-level *Files*). Conversely, a User can view content only at or below their own *Security Level* (e.g. a Secret-level researcher can view Public-level or Secret-level *Files*, but may not view Top-Secret-level *Files*).
- Transfer of information from a high-Secrecy document to a lower-Secrecy document may happen via *Trusted* subjects only



# Labels and Lattices

## Bell-La Padula model [[BL](#)]

- View the system as Subjects accessing Objects
  - *Access Control* : Take Requests as input and output Decisions
- Four modes of Access are possible:
  - *execute* : No Observation or Alteration
  - *read* : Observation
  - *append* : Alteration
  - *write* : Both Observation and Alteration
- An *Access Matrix*  $M$  encodes permissible *Access* modes
  - Subjects are rows, Objects are columns
- The current *Access Set* :  $b$  is  $(Subj, Obj, Attr) : (S, O, A)$  triples
  - Encodes Accesses in progress (e.g. open *Files*)
  - At a minimum,  $(S, O, A) \in b$  requires  $A$  permitted by cell  $M_{S,O}$



# Labels and Lattices

## *Security Levels*

- A *Security Level* (or *Label*)  $L_i$  is a pair  $\langle c, s \rangle$  where:
  - $c = \text{Classification}$  – e.g. 1 = Unclassified, 2 = Secret, 3 = Top-Secret
  - $s = \text{Category-Set}$  – e.g. Nuclear, Crypto
- $\langle c_1, s_1 \rangle$  “*Dominates*”  $\langle c_2, s_2 \rangle$  iff  $c_1 \geq c_2$  and  $s_1 \supseteq s_2$ 
  - $L_1$  *Dominates*  $L_2$  relation is sometimes written  $L_1 \propto L_2$  or  $L_1 \supseteq L_2$
  - *Labels* then form a *Lattice* (partial order with “least upper-bound” & “greatest lower-bound”)
- Inverse of “*Dominates*” relation is : “*Can Flow To*”, written  $\sqsubseteq$ 
  - $L_1 \sqsubseteq L_2$  :  $L_1$  *Can Flow To*  $L_2$  (means that  $L_2$  *Dominates*  $L_1$ )
- Subjects and Objects are assigned *Security Levels*
  - $\text{Level}(S), \text{Level}(O)$  : *Security Level* of Subject/Object
  - $\text{CurrentLevel}(S)$  : Subject may operate at a lower *Security Level*
  - $\text{Level}(S)$  *Bounds*  $\text{CurrentLevel}(S)$  :  $\text{CurrentLevel}(S) \sqsubseteq \text{Level}(s)$
  - Since  $\text{Level}(S)$  is max, sometimes called  $S$ ’s “*Clearance*”



# Labels and Lattices

## *Security Properties*

➤ Two Access Control *Properties* with respect to *Security Labels*:

### 1. *Simple Security* or *SS-Property* (DAC)

- For any  $(S, O, A) \in b$ , if  $A$  includes *observation*, then  $Level(O) \sqsubseteq Level(S)$  (i.e.  $Level(S)$  must *Dominate*  $Level(O)$ )
  - e.g. an Unclassified User cannot *read* a Top-Secret File – “No Read-Up”

### 2. *Star Security* or *\*-Property* (MAC)

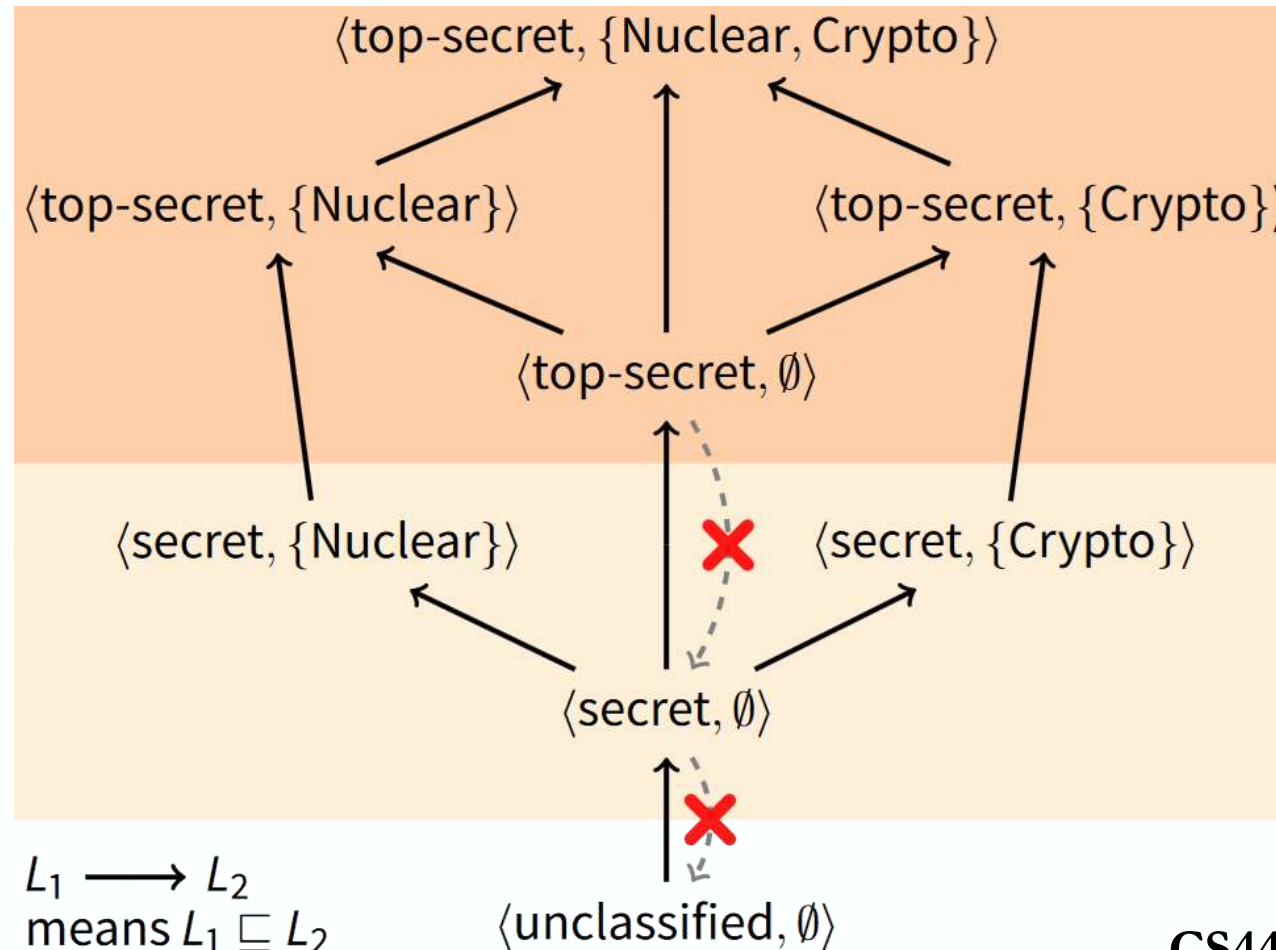
- If any *Subject* *observes*  $O_1$  and *modifies*  $O_2$ , then  $Level(O_1) \sqsubseteq Level(O_2)$  (i.e.  $Level(O_2)$  *Dominates*  $Level(O_1)$ )
  - e.g. no *Subject* can *read* a Top-Secret File, then *write* a Secret File – “No Write-Down”
- More precisely, given an  $(S, O, A) \in b$ :
  - if  $A = r$  then  $Level(O) \sqsubseteq CurrentLevel(S)$  – “No Read-Up”
  - if  $A = a$  then  $CurrentLevel(S) \sqsubseteq Level(O)$  – “No Write-Down”
  - if  $A = w$  then  $CurrentLevel(S) = Level(O)$





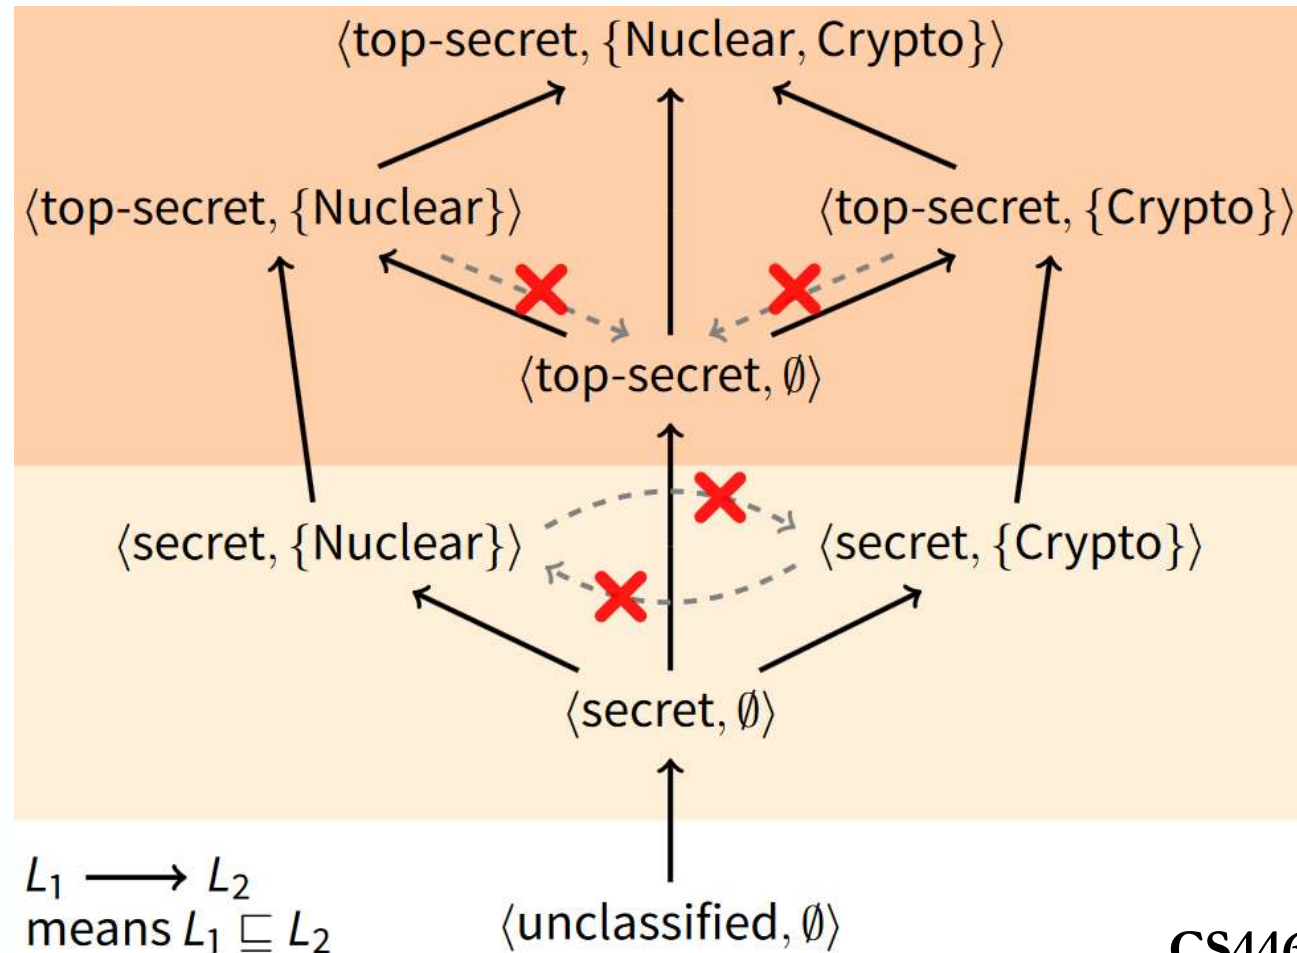
# Labels and Lattices

*Labels form a Lattice* [Denning]



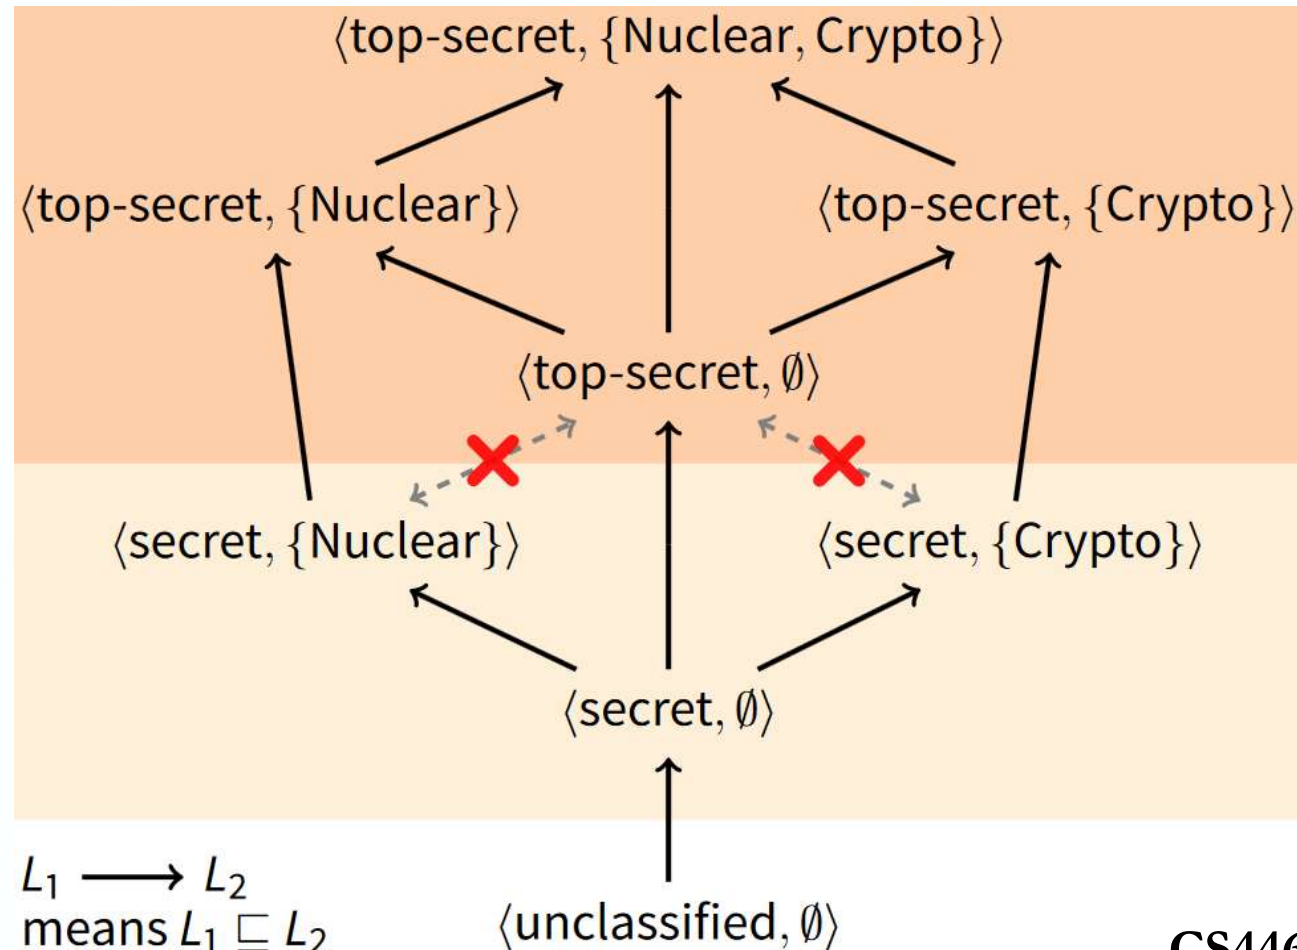
# Labels and Lattices

*Labels form a Lattice* [Denning]




# Labels and Lattices

*Labels form a Lattice* [Denning]



# Labels and Lattices

$\sqsubseteq$  is *Transitive*




- *Transitivity* makes it easier to reason about *Security*
- Example: *Label* User data (*File*) so it cannot *Flow-To* Internet:  $L_U \not\sqsubseteq L_{net}$ 
  - Policy holds regardless of what other Software does



# Labels and Lattices

$\sqsubseteq$  is *Transitive*




- *Transitivity* makes it easier to reason about *Security*
- Example: *Label* User data (*File*) so it cannot *Flow-To* Internet:  $L_U \not\sqsubseteq L_{net}$ 
  - Policy holds regardless of what other Software does
- Suppose untrustworthy Software *reads* *File*





# Labels and Lattices

$\sqsubseteq$  is *Transitive*




- *Transitivity* makes it easier to reason about *Security*
- Example: *Label* User data (*File*) so it cannot *Flow-To* Internet:  $L_U \not\sqsubseteq L_{net}$ 
  - Policy holds regardless of what other Software does
- Suppose untrustworthy Software *reads* File
  - Process that is *Labeled*  $L_{bug}$  *reads* File, so must have  $L_U \sqsubseteq L_{bug}$
  - If  $L_U \sqsubseteq L_{bug}$  and  $L_U \not\sqsubseteq L_{net}$ , it follows that  $L_{bug} \not\sqsubseteq L_{net}$



# Labels and Lattices

$\sqsubseteq$  is *Transitive*



- *Transitivity* makes it easier to reason about *Security*
- Example: *Label* User data (*File*) so it cannot *Flow-To* Internet:  $L_U \not\sqsubseteq L_{\text{net}}$ 
  - Policy holds regardless of what other Software does
- Conversely, a *Process* that can *write* to the Network cannot *read* the *File*



# Labels and Lattices

## Naïve MAC Implementation

- Take an ordinary Unix system
- Put *Labels* on all *Files* and *Directories* to track *Security Levels*
- Each User  $U$  assigned a *Security Clearance Level*( $U$ ) on login
- Determine *Current*(*Security*)*Level* dynamically
  - When  $U$  logs in, start with lowest *CurentLevel*
  - Increase *CurentLevel* as higher-level *Files* are *observed*
    - (sometimes called a “*Floating Label*” system)
  - If  $U$ ’s *Level* does not *Dominate* *CurentLevel*, kill program
    - e.g. kill Program that *writes* to *File* if the Program’s *CurentLevel* can’t *Flow-To* the *File*’s *Label*
- Is this *Secure*?



# Labels and Lattices

## No: *Covert Channels*

- Operating System rife with *Covert Storage Channels*
  - *Low CurrentLevel* Process executes another Program
  - New Program reads sensitive *File*, gets *High CurrentLevel*
  - *High* Program can then exploit *Covert Channels* to pass data to *Low CurrentLevel Process*
  - Example: *High* Program inherits *File Descriptor*
    - Can pass 4-Byte information chunks to *Low* Program through the ***File offset*** !
- Other *Covert Storage Channels*:
  - **exit value, signals, *File* locks, terminal escape codes, ...**
- If we eliminate all *Covert Storage Channels*, is system *Secure*?



# Labels and Lattices

## Still No: *Timing Channels*

### ➤ *Example:* **CPU Utilization**

- To send a 0 bit, use 100% of CPU in busy-loop
- To send a 1 bit, sleep and relinquish CPU
- Repeat to transfer more bits

### ➤ *Example:* **Resource Exhaustion**

- *High* Program allocates all Physical Memory if bit is 1
- If *Low* Program slows down due to *Paging*, knows less Memory available

### ➤ Other *Examples*:

- **Disk head position, Processor Cache/TLB pollution, ...**





# Labels and Lattices

## Reducing *Covert Channels*

- Observation: *Covert Channels* come from *Sharing*
  - If you have no *Shared Resources*, no *Covert Channels* exist
  - Extreme example: Just use two computers (common in DoD)
- Problem: *Sharing* is needed
  - e.g. *read* Unclassified data when preparing (*writing*) Classified
- In general, can only hope to bound Bandwidth of *Covert Channels*
- One approach: **Strict *Partitioning* of *Resources***
  - Strictly *Partition* and *Schedule Resources* between *Levels*
  - Occasionally reapportion *Resources* based on usage [[Browne](#)]
  - Do so infrequently to bound leaked information
  - Approach still not so good if many *Security Levels* possible



# Labels and Lattices

## *Declassification*

- Sometimes need to prepare Unclassified report from Classified data
- *Declassification* happens outside of traditional Access Control Model
  - Present *File* to security officer for downgrade
- Job of *Declassification* often not trivial
  - e.g., Microsoft Word saves a lot of Undo information
    - This might be all the secret stuff you cut from document
  - Another bad mistake:
    - Redact PDF using black censor bars over or under text, leaving text selectable
      - e.g. [[Cluley](#)]



# Labels and Lattices

## *Integrity*

- Achieve Controlled Access to Classified information

Typical way to think of *Security* in this context:

- A Subject should not be able to corrupt Data in an *Integrity Level* ranked higher than themselves
- A Subject should not be able to become corrupted by Data from a lower *Integrity Level*

Preservation of data *Integrity* – Goals:

- Prevent Data modification by unauthorized parties
- Prevent unauthorized Data modification by authorized parties
- Maintain Internal and External Consistency (i.e. Data reflects the real world)



# Labels and Lattices

## *Biba Integrity Model* [[Biba](#)]

- Problem: How to protect *Integrity*
  - Suppose text editor gets Trojan-ed, subtly modifies *Files*
  - Might mess up critical operations even without leaking anything
- Observation: *Integrity* is the converse of *Secrecy*
  - In *Secrecy*, want to avoid writing to *Lower-Secrecy Files* – Prevent *Leaking*
  - In *Integrity*, want to avoid writing *Higher-Integrity Files* – Prevent *Corruption*

Use *Integrity* hierarchy parallel to *Secrecy* one

- Now *Label* (/ *Security Level*) is a  $\langle c, i, s \rangle$  triple, where  $i = \text{Integrity}$
- $\langle c_1, i_1, s_1 \rangle \sqsubseteq \langle c_2, i_2, s_2 \rangle$  iff  $c_1 \leq c_2$  and  $i_1 \geq i_2$  and  $s_1 \subseteq s_2$ 
  - Only *Trusted* Users can operate at *Higher Integrity*  
(which is visually lower in the *Lattice* – Opposite of *Secrecy*)
  - If you *read* less authentic data, your *CurrentIntegrityLevel* gets lowered (placing you higher in the *Lattice*), and you can no longer *write Higher-Integrity Files*



## ***LOMAC*** [[Fraser](#)]

- MAC not widely accepted outside military

*Low water Mark Access Control (LOMAC)*'s goal:

- Make MAC more palatable
- Concentrates on *Integrity*
  - More important goal for many settings
    - e.g. don't want Viruses tampering with all your *Files*
  - Also don't have to worry as much about *Covert Channels*
- Provides reasonable defaults (minimally obtrusive)
- Has actually had impact
  - Originally available for Linux (2.2)
  - Now ships with [FreeBSD](#)
  - Windows introduced similar [Mandatory Integrity Control \(MIC\)](#)





## ***LOMAC*** Overview

Subjects are “*Jobs*” (essentially *Processes* – possibly groups of *Processes*; more later)

- Each Subject labeled with an *Integrity* number (e.g. 1, 2)
- Higher numbers mean higher *Integrity* (so –unfortunately–  $2 \sqsubseteq 1$  using earlier notation)
- Subjects can be *Reclassified* (change *CurrentIntegrityLevel*) on *observation* of *Low-Integrity* data

Objects (*Files*, *Pipes*, etc.) also labeled with *Integrity Level*

- Object *Integrity Level* is fixed and cannot change


In the context of Security:

- *Low-Integrity* Subjects cannot write to *High-Integrity* Objects
- New Objects acquire *Integrity Level* of their creator



## LOMAC Defaults

*Note: Can-Flow-To is downward;  
opposite of earlier diagram*



- Two *Integrity Levels*: 1 and 2
- Level 2 (*High-Integrity*) contains:
  - FreeBSD/Linux *Files* intact from distro, static Web Server config
  - The console, *Trusted* terminals, *Trusted* Network
- Level 1 (*Low-Integrity*) contains:
  - NICs connected to Internet, *Untrusted* terminals, etc.
- Idea: Suppose Worm compromises your Web Server
  - Worm comes from external Network → Level 1
  - Won't be able to muck with System *Files* or Web Server config




## The *Self-Revocation* Problem

- Want to integrate with Unix unobtrusively
- Problem: Application expectations
  - Kernel performs Access Checks usually at *File open* time (i.e. not *read* / *write*)
    - Remember: *Access Control Lists & Capabilities (File Descriptors)*
  - Legacy Applications don't pre-declare they will *observe Low-Integrity* data
  - An Application can “*taint*” itself unexpectedly, *Revoking* its own Permission to access an Object it created




## *Self-Revocation* Example

- User has *High-Integrity* (Level 2) Shell
- Runs: **ps | grep user**
  - *Pipe* created before **ps** reads *Low-Integrity* data
  - **ps** becomes “*tainted*”, can no longer *write* to **grep**



## *Self-Revocation* Example


- User has *High-Integrity* (Level 2) Shell
- Runs: **ps | grep user**
  - *Pipe* created before **ps** reads *Low-Integrity* data
  - **ps** becomes “*tainted*”, can no longer *write* to **grep**






## *Self-Revocation* Example

- User has *High-Integrity* (Level 2) Shell
- Runs: **ps | grep user**
  - *Pipe* created before **ps** reads *Low-Integrity* data
  - **ps** becomes “*tainted*”, can no longer *write* to **grep**



## *Self-Revocation* Example

- User has *High-Integrity* (Level 2) Shell
- Runs: **ps | grep user**
  - *Pipe* created before **ps** reads *Low-Integrity* data
  - **ps** becomes “*tainted*”, can no longer *write* to **grep**



## Solution

- Don't consider Pipes to be real Objects
- Join multiple *Processes* together in a “*Job*”
  - Pipe ties *Processes* together in *Job*
  - Any *Processes* tied to *Job* when they *read* or *write* to Pipe
  - So will lower *Integrity* of both **ps** and **grep**
- Similar idea applies to *Shared Memory* and *Inter-Process Communication*
- Summary: LOMAC applies MAC to non-military systems
  - But doesn't allow military-style Security policies
    - i.e. with *Secrecy*, various *Categories*, etc.



**CS-446/646**

Time for Questions !

