

CS 447/647

DNS & DHCP

DHCP Overview

What is DHCP?

How does DHCP work?


What are some common DHCP problems?

What is DHCP

- Dynamic Host Configuration Protocol (RFC 2131)
- Automatically assigns IP addresses to devices on a network
- Saves time that would be spent manually assigning IP addresses
- Prevents typos that cause network problems

DHCP Session

- Client broadcasts DHCPDISCOVER message
- DHCP server responds with a DHCPOFFER message containing an available IP address and network settings
- Client sends a DHCPREQUEST to accept the DHCP offer
- DHCP server confirms with a DHCPACK



DHCP Key Components

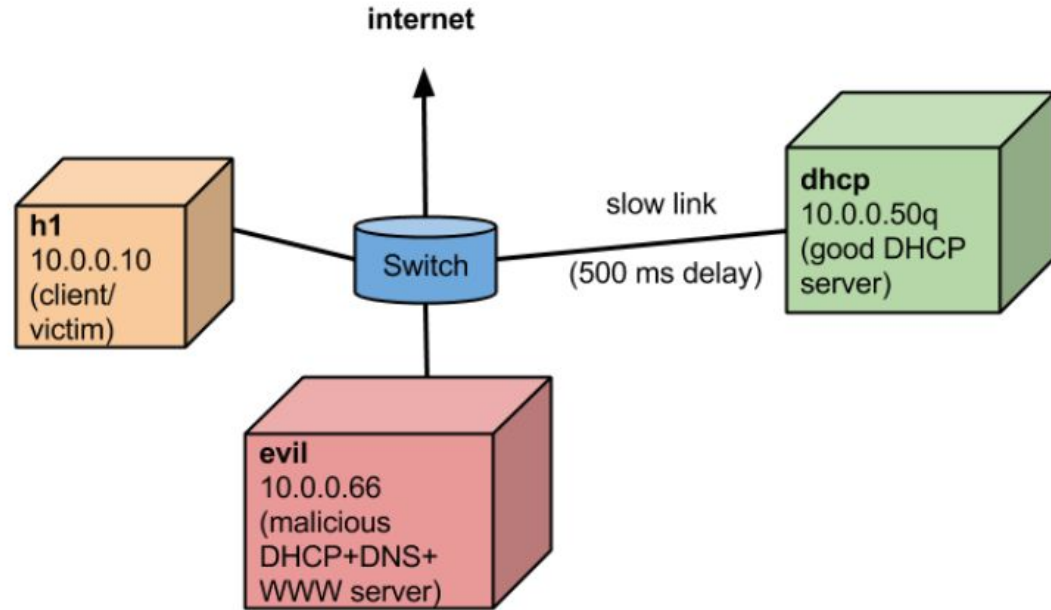
- DHCP Server
 - Assigns IP addresses
 - Manages network configurations
- DHCP Client
 - Requests and receives IP configuration from the DHCP server
- DHCP Relay Agent
 - Forwards DHCP messages across networks

Configuring a DHCP Server

- Defining IP Address Pools
 - Specify IP ranges available for dynamic assignment
 - Example: 192.168.0.100-192.168.0.199
- Setting lease durations
 - Default lease time
 - Max lease time
- Static IP reservations
 - Critical devices can have a static IP reservation
- Default route
- DNS domain
- Name servers

DHCP Security

- Rogue DHCP servers
 - Malicious DHCP server intercepts DHCP requests
- DHCP starvation attacks
 - Client spams DHCP requests with spoofed MAC addresses to exhaust all available IP addresses
- DHCP snooping
 - Security feature on network switches that blocks unauthorized DHCP servers



DNS Overview

What is DNS?

How is DNS managed?

What are some common DNS resource records?

What is the SOA record?

What command line tools help you query DNS?

How to setup and manage dnsmasq and bind.

Domain Name System

- Maps a hostname to an IP
 - google.com -> 172.217.6.46
- Essential for the global Internet
- Used for:
 - Mail
 - Service Discovery
 - Authentication
 - SSL
 - WWW

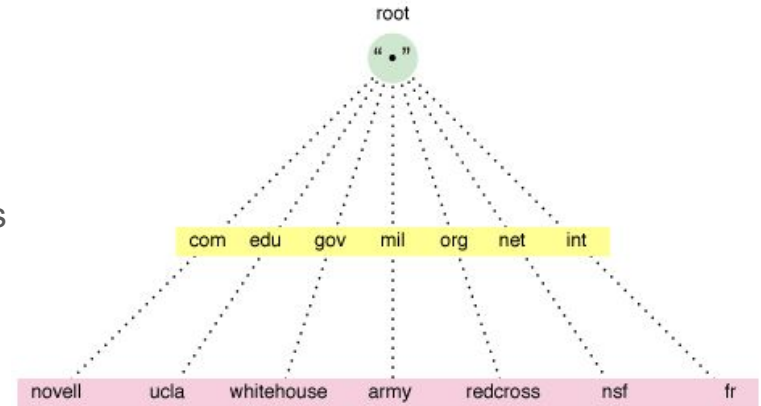


<https://xkcd.com/1361/>



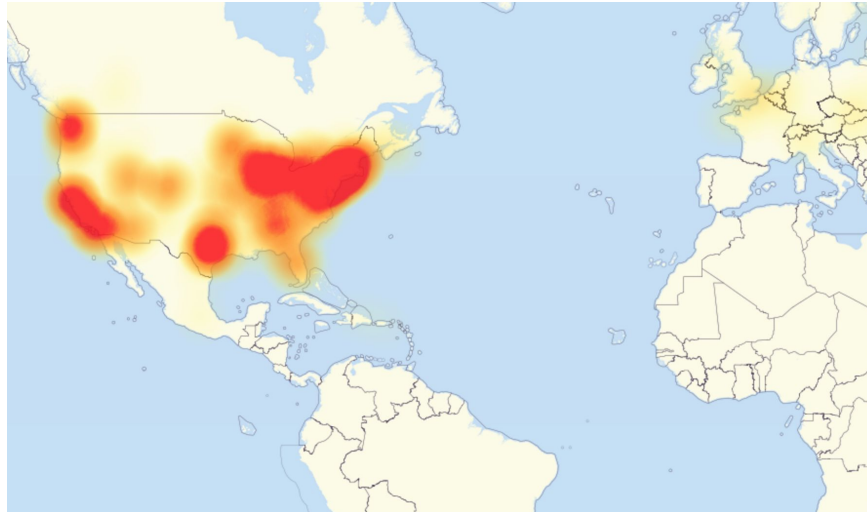
DNS Architecture

- Distributed Database
 - . is managed by a group of nonprofits and companies
 - ICANN operates one
 - *US Army*
 - *US DoD*
 - *NASA*
 - Verisign
 - 10/13 are in the US
- Each site maintains its own database
 - Company
 - University
 - Individual (zachest.com)
- Globally administered
 - IP network portion
 - Domain
- Local administrators must prevent duplicates



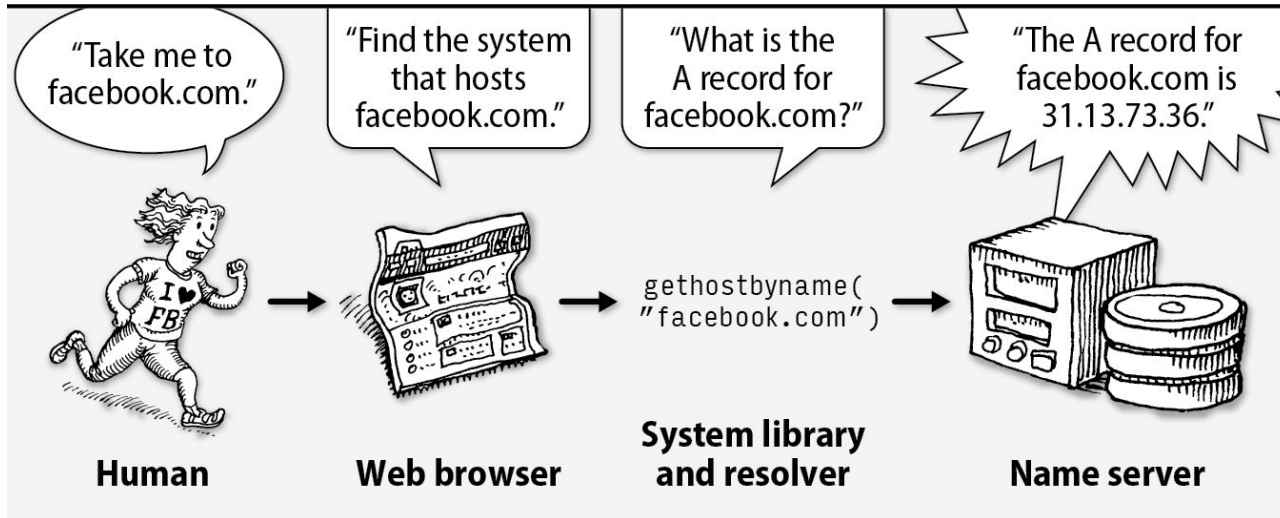
DNS Architecture

- DNS disappears, the Internet vanishes
- Mirai 2016 Distributed Denial of Service attack
 - DynDNS
 - Spotify, Twitter, Github, PayPal



DNS Architecture

- Query, two parts
 - Name: google.com
 - Record type: A
- Response
 - Resource Records



dig eecs.mit.edu A

Use Unix "dig" utility to look up IP address ("A") for hostname eecs.mit.edu via DNS

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.                IN      A

;; ANSWER SECTION:
eecs.mit.edu.                21600   IN      A      18.62.1.6

;; AUTHORITY SECTION:
mit.edu.                     11088   IN      NS      BITSY.mit.edu.
mit.edu.                     11088   IN      NS      W20NS.mit.edu.
mit.edu.                     11088   IN      NS      STRAWB.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.              126738  IN      A      18.71.0.151
BITSY.mit.edu.               166408  IN      A      18.72.0.3
W20NS.mit.edu.               126738  IN      A      18.70.0.160
```

dig eecs.mit.edu A

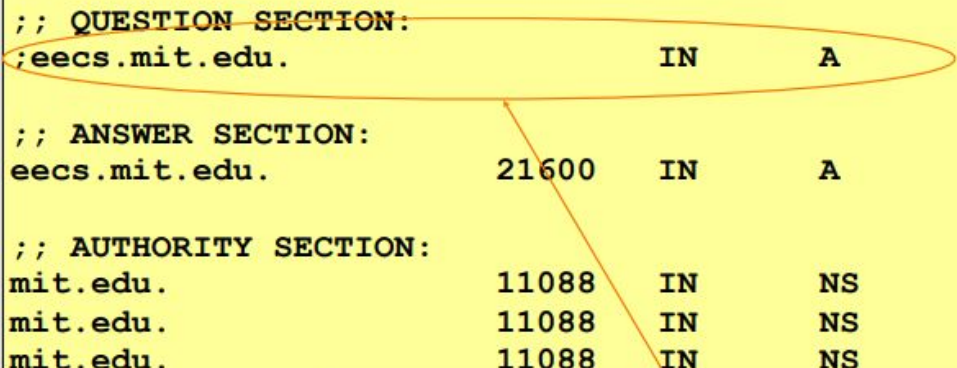
```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.                IN      A

;; ANSWER SECTION:
eecs.mit.edu.                21600   IN      A      18.62.1.6

;; AUTHORITY SECTION:
mit.edu.                     11088   IN      NS      BITSY.mit.edu.
mit.edu.                     11088   IN      NS      W20NS.mit.edu.
mit.edu.                     11088   IN      NS      STRAWB.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.              126738  IN      A      18.71.0.151
BITSY.mit.edu.               166408  IN      A      18.72.0.3
W20NS.mit.edu.               126738  IN      A      18.70.0.160
```



The question we asked the server

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.                IN      A

;; ANSWER SECTION:
eecs.mit.edu.                 21600   A

;; AUTHORITY SECTION:
mit.edu.                      11088   IN      NS      BITSY.mit.edu.
mit.edu.                      11088   IN      NS      W20NS.mit.edu.
mit.edu.                      11088   IN      NS      STRAWB.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.               126738  IN      A        18.71.0.151
BITSY.mit.edu.                166408  IN      A        18.72.0.3
W20NS.mit.edu.                126738  IN      A        18.70.0.160
```

A 16-bit **transaction identifier** that enables the DNS client (dig, in this case) to match up the reply with its original request

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
```

```
;; global options: +cmd
```

```
;; Got answer:
```

```
;; ->>HEADER<<- opcode:
```

```
;; flags: qr rd ra; QUE
```

"Answer" tells us the IP address associated with eecs.mit.edu is 18.62.1.6 and we can cache the result for 21,600 seconds

ONAL: 3

```
;; QUESTION SECTION:
```

```
;eecs.mit.edu.
```

IN

A

```
;; ANSWER SECTION:
```

```
eecs.mit.edu.
```

21600

IN

A

18.62.1.6

```
;; AUTHORITY SECTION:
```

```
mit.edu.
```

11088

IN

NS

BITSY.mit.edu.

```
mit.edu.
```

11088

IN

NS

W20NS.mit.edu.

```
mit.edu.
```

11088

IN

NS

STRAWB.mit.edu.

```
;; ADDITIONAL SECTION:
```

```
STRAWB.mit.edu.
```

126738

IN

A

18.71.0.151

```
BITSY.mit.edu.
```

166408

IN

A

18.72.0.3

```
W20NS.mit.edu.
```

126738

IN

A

18.70.0.160

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.                IN      A

;; ANSWER SECTION:
eecs.mit.edu.                21600   IN      A      18.62.1.6

;; AUTHORITY SECTION:
mit.edu.                     11088   IN      NS      BITSY.mit.edu.
mit.edu.                     11088   IN      NS      BITSY.mit.edu.
mit.edu.                     11088   IN      NS      BITSY.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.              166408  IN      A      18.72.0.3
BITSY.mit.edu.               126738  IN      A      18.70.0.160
W20NS.mit.edu.
```

In general, a single Resource Record (RR) like this includes, left-to-right, a DNS name, a time-to-live, a family (IN for our purposes - ignore), a type (A here), and an associated value

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
```

```
;; global options: +cm
```

```
;; Got answer:
```

```
;; ->>HEADER<<- opcode
```

```
;; flags: qr rd ra; QU
```

```
;; QUESTION SECTION:
```

```
;eecs.mit.edu.
```

```
;; ANSWER SECTION:
```

```
eecs.mit.edu.
```

```
;; AUTHORITY SECTION:
```

```
mit.edu.
```

```
mit.edu.
```

```
mit.edu.
```

```
;; ADDITIONAL SECTION:
```

```
STRAWB.mit.edu.
```

```
BITSY.mit.edu.
```

```
W20NS.mit.edu.
```

“Authority” tells us the name servers responsible for the answer. Each RR gives the **hostname** of a different name server (“NS”) for names in mit.edu. We should cache each record for 11,088 seconds.

If the **“Answer”** had been empty, then the resolver's next step would be to send the original query to one of these name servers.

```
21600 IN A 18.02.1.0
```

11088	IN	NS
11088	IN	NS
11088	IN	NS

BITSY.mit.edu.
W20NS.mit.edu.
STRAWB.mit.edu.

126738	IN	A	18.71.0.151
166408	IN	A	18.72.0.3
126738	IN	A	18.70.0.160

dig eecs.mit.edu A

```
; ; <<>> DiG 9.6.0-APPLE-P2 <<>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.

;; ANSWER SECTION
eecs.mit.edu.

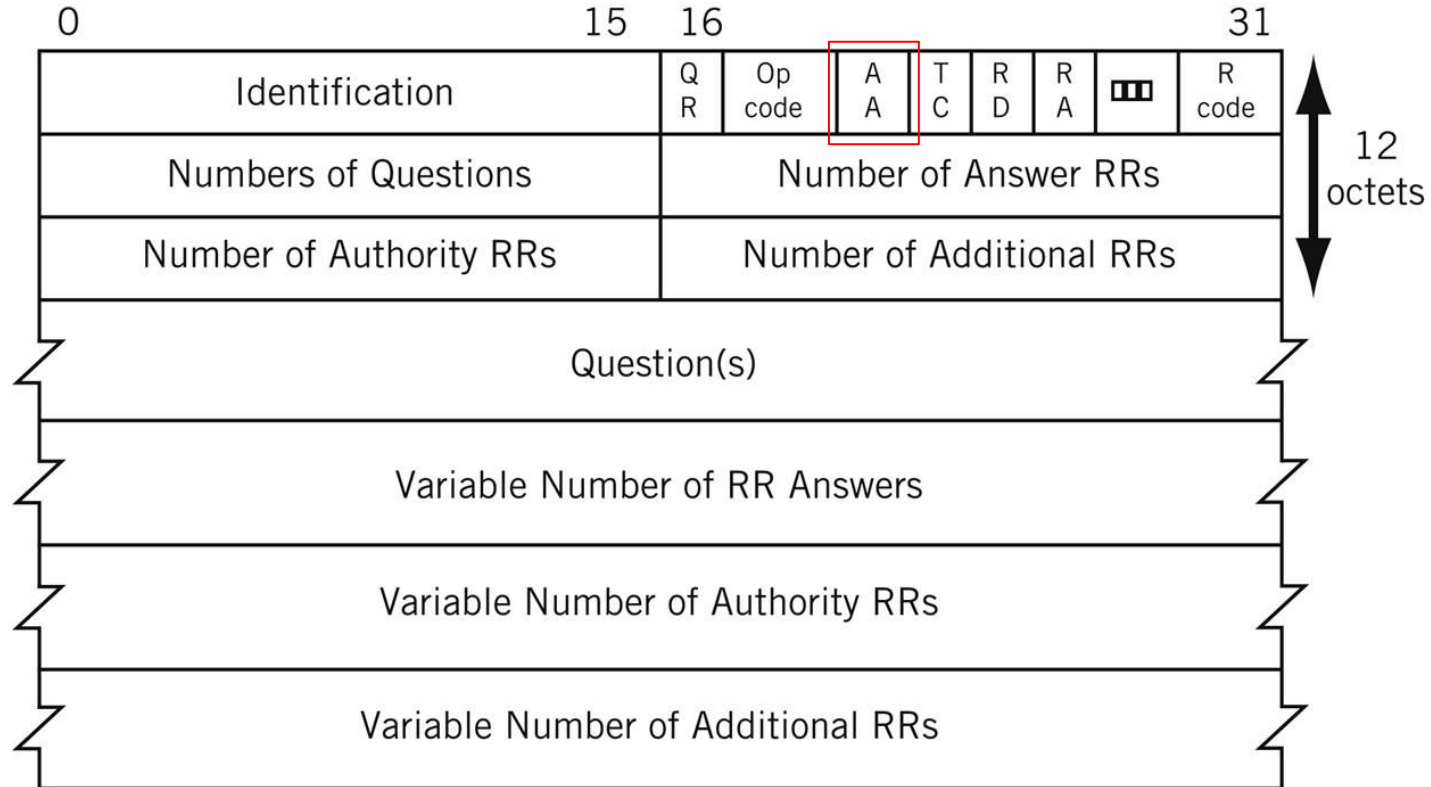
;; AUTHORITY SECTION
mit.edu.                11088      IN         NS         BITSY.mit.edu.
mit.edu.                11088      IN         NS         W20NS.mit.edu.
mit.edu.                11088      IN         NS         STRAWB.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.         126738     IN         A           18.71.0.151
BITSY.mit.edu.          166408     IN         A           18.72.0.3
W20NS.mit.edu.          126738     IN         A           18.70.0.160
```

"Additional" provides extra information to save us from making separate lookups for it, or helps with bootstrapping.

Here, it tells us the IP addresses for the hostnames of the name servers. We add these to our cache.

DNS Request and Response (Same Format)



DNS service providers

- DNS use to be a core sysadmin responsibility
- Products automate DNS
 - Microsoft Active Directory
 - Bluecat
 - Amazon Route 53
- Open Source DIY
 - BIND
 - Unbound
 - dnsmasq
- Still need to understand core concepts.
 - Troubleshooting



DNS for lookups

- Static
 - /etc/hosts
- Stub Resolver
 - /etc/resolv.conf

```
search domainname ...  
nameserver ipaddr
```

```
search atrust.com booklab.atrust.com  
nameserver 63.173.189.1           ; ns1  
nameserver 174.129.219.225        ; ns2
```

nsswitch

- Stored in /etc/nsswitch.conf
- Order of DNS services
 - Left to right

hosts: dns [!UNAVAIL=return] files

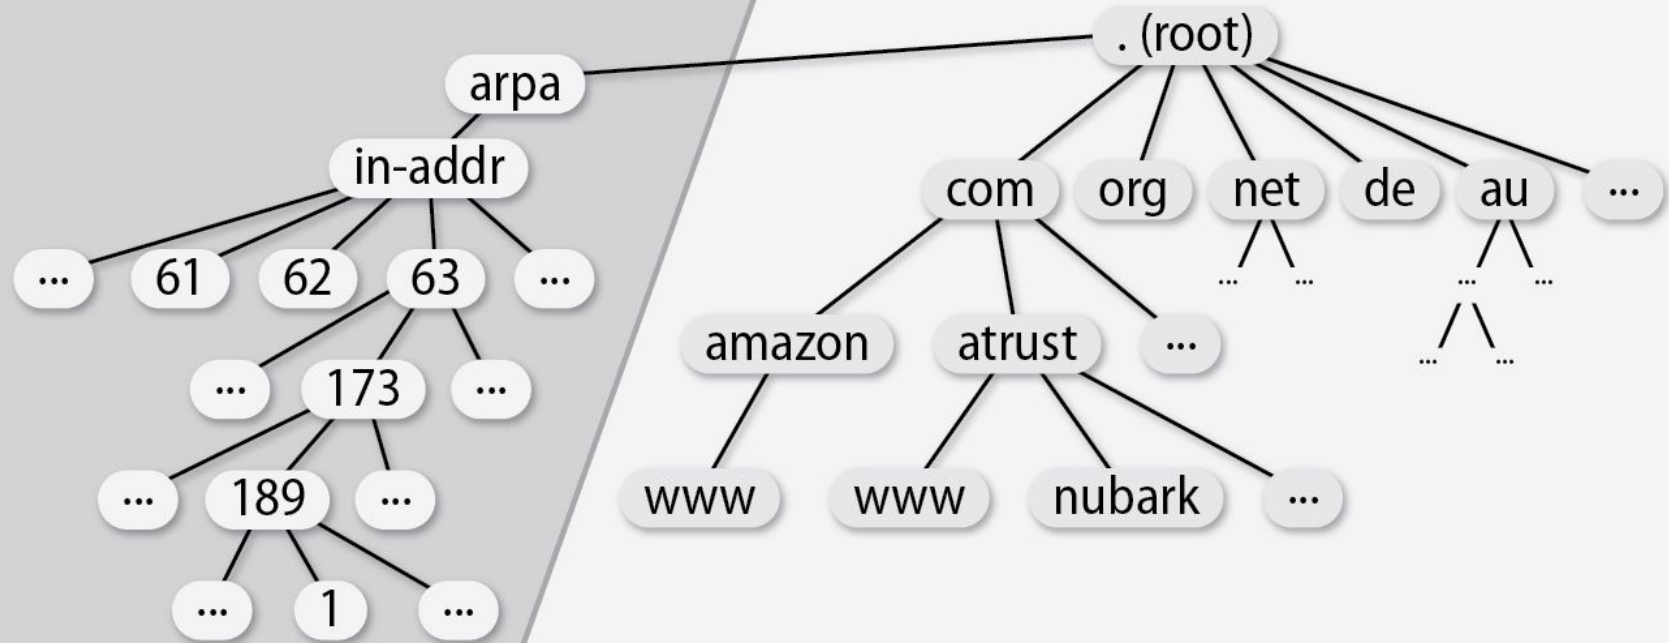
```
files mdns4_minimal [NOTFOUND=return] dns myhostname
```

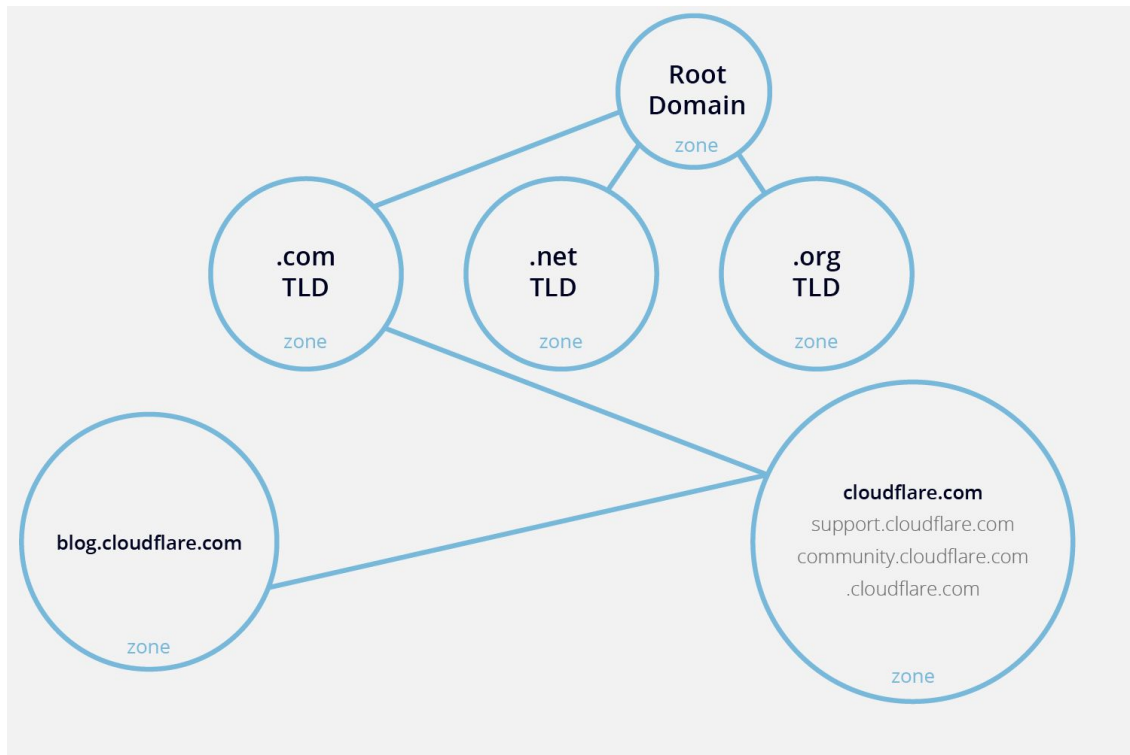

The DNS namespace

- Maintains forward mappings
 - A Record
- Maintains reverse mappings
 - PTR Record - Pointer
 - Inverted IP address
 - 63.173.189.1
 - 1.189.173.63.in-addr.arpa.

Reverse zones

Forward zones





The DNS namespace

- Two top-level domains
 - Country code domains
 - .us
 - .ru
 - .io
 - British Indian Ocean Territory
 - Uncertain of the future of .io
 - Generic top-level domains
 - .com
 - .org
 - .edu
 - .party
 - .dad
 - .lol
 - .ninja

ccTLDs Quiz

What Country owns these ccTLDs?

- .tv.
- .co.
- .ws.
- .ly.
- .re.

New TLDs

- .website
- .press
- .rocks
- .support
- .email
- .pics
- .lgbt
- .red
- .blue
- .wtf

The DNS namespace

- Second level-domains
 - linux.ninja
- Apply at a top-level domain registrar
 - Costs money



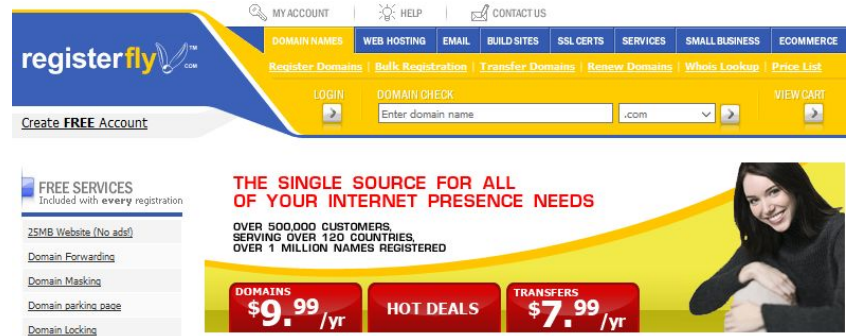
gnulinux.ninja

\$19/year



RegisterFly

- Hosting and Domain Registrar
- Managed over 2,000,000 domains
- Split into RegisterFly.net and RegisterFly.com
 - Former business partners then estranged lovers
- Fraud, lawsuits, counter lawsuits, appeals, then total collapse
 - Liposuction and a \$6,000 chihuahua
- Complete neglect of customers
- ICANN stripped the company of its registrar status



The DNS namespace

- Subdomains
 - .com = Top-level domain
 - bigcompany.com = Second-level domain
 - west.bigcompany.com = subdomain
- Two servers
 - Can work with one.

How DNS works

- Answers queries
 - hostnames and IP addresses
- Forwards requests
- Caches answers
- Synchronizes with other local servers

DNS Servers

- Authoritative
 - Responsible for a zone
 - IE: engr.unr.edu
 - Different Types
 - Primary (master) - Data is on disk
 - Secondary - Data is from master
 - Stub - Copy of zone with a subset of resource records
 - Distribution - Authoritative but not listed aka Stealth
 - Caching
 - Forwarding - Forwards requests, large cache
 - Recursive - Handles referrals
 - Nonrecursive - Sends referrals
 - `dig +recurse @8.8.8.8 www.google.hk`
 - `dig +norecurse @8.8.8.8 www.google.hk`

Resource Records

- Each server is responsible for its own zone
- Text files
- Record for each host
 - A - IPv4
 - AAAA - IPv6
- Load-balancing
 - round-robin

www	IN	A	192.168.0.1
	IN	A	192.168.0.2
	IN	A	192.168.0.3

DNS Database

- Set of files
 - Zone Files
- \$TTL - Time to live
 - Must be first line
- Resource Record
 - [name] [ttl] [class] type data
- Four Groups of Records
 - Zone Infrastructure Record
 - Basic Records
 - Security Records
 - Optional Records

\$TTL 86400

\$TTL 24h

\$TTL 1d

Record Types

	Type	Name	Function
Zone	SOA	Start Of Authority	Defines a DNS zone
	NS	Name Server	Identifies servers, delegates subdomains
Basics	A	IPv4 Address	Name-to-address translation
	AAAA	IPv6 Address	Name-to-IPv6-address translation
	PTR	Pointer	Address-to-name translation
	MX	Mail Exchanger	Controls email routing
Security	DS	Delegation Signer	Hash of signed child zone's key-signing key
	DNSKEY	Public Key	Public key for a DNS name
	NSEC	Next Secure	Used with DNSSEC for negative answers
	NSEC3	Next Secure v3	Used with DNSSEC for negative answers
	RRSIG	Signature	Signed, authenticated resource record set
Optional	CNAME	Canonical Name	Nicknames or aliases for a host
	SRV	Service	Gives locations of a well-known service
	TXT	Text	Comments or untyped information

Start of Authority Record (SOA)

- Each zone has 1 SOA record
 - Name: atrust.com
 - Class: IN (Internet Record)
 - Type: SOA
 - Server: ns1.atrust.com.
 - Email: hostmaster.atrust.com. (hostmaster@atrust.com)

```
; Start of authority record for atrust.com
atrust.com.      IN  SOA      ns1.atrust.com. hostmaster.atrust.com. (
    2017110200    ; Serial number
    10800         ; Refresh   (3 hours)
    1200          ; Retry    (20 minutes)
    36000000      ; Expire   (40+ days)
    3600 )        ; Minimum   (1 hour)
```

Name Server Records (NS)

- Identify the servers that are authoritative for a zone
- Format:
 - zone [ttl] [IN] NS hostname

```
; NS Records  
      IN      NS      ns1.ecc.engr.unr.edu.
```


Address Records (A) and (AAAA)

- Heart of the DNS database
- Format:
 - hostname [ttl] [IN] A ip_address
- “.” denotes fully qualified name
 - No “.” means the default domain is added.
- AAAA for IPv6

```
; NS A Records
ns1      IN      A      134.197.20.131
```

Pointer Records (PTR)

- Maps an IP* to a hostname
- Format:
 - `ip_address [ttl] [IN] PTR hostname`
 - `*130.195.20.172.in-addr.arpa`

```
130      IN      PTR      ecc-a-01.ecc.engr.unr.edu.
```

Mail Exchanger Records (MX)

- Used for routing mail
- Format:
 - name [ttl] [IN] MX preference hostname
 - dig MX cse.unr.edu @134.197.5.1

```
;; ANSWER SECTION:  
cse.unr.edu.          300      IN       MX       10 cse-unr-edu.mail.protection.outlook.com.
```

```
somehost              IN MX     10 mailserver.atrust.com.  
                      IN MX     20 mail-relay3.atrust.com.
```

john@somehost.atrust.com.

Service Records (SRV)

- Specifies the location of services within a domain
- Format:
 - *service.proto.name [ttl] [IN] SRV pri weight port target*
- Kerberos!

```
$ORIGIN foobar.com.  
_kerberos          TXT      "FOOBAR.COM"  
kerberos           CNAME    daisy  
kerberos-1         CNAME    use-the-force-luke  
kerberos-2         CNAME    bunny-rabbit  
_kerberos._udp     SRV      0 0 88 daisy  
                  SRV      0 0 88 use-the-force-luke  
                  SRV      0 0 88 bunny-rabbit  
_kerberos-master._udp SRV      0 0 88 daisy  
_kerberos-adm._tcp  SRV      0 0 749 daisy  
_kpasswd._udp       SRV      0 0 464 daisy
```

Text Records (TXT)

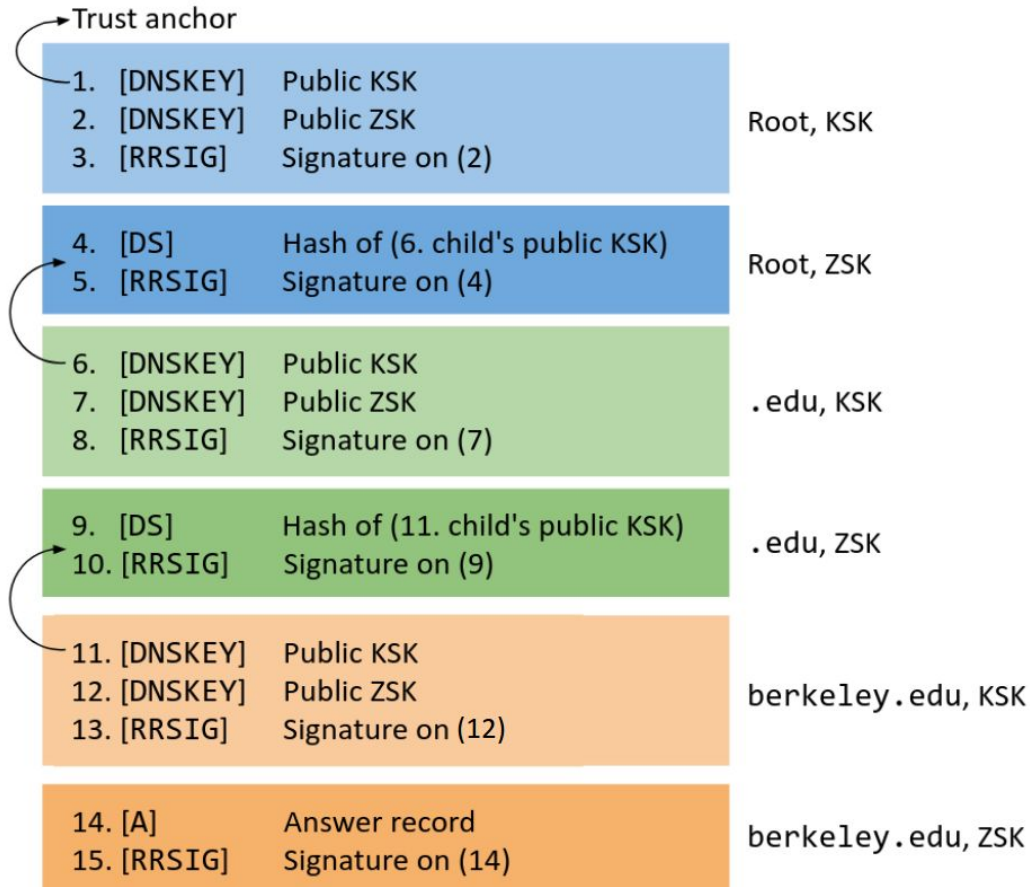
- Adds arbitrary text to a DNS record
- Format:
 - `name [ttl] [IN] TXT info ...`
- Verification
 - Google Apps

 _kerberos	TXT	CSE.UNR.EDU
 201804._domainkey	TXT	v=DKIM1; h=sha256; k=rsa; s=email; p=MII
 _kpasswd._tcp	SRV	[0][100][464]ipa1.cse.unr.edu
 _kerberos._tcp	SRV	[0][100][88]ipa1.cse.unr.edu
 _kpasswd._udp	SRV	[0][100][464]ipa1.cse.unr.edu
 _ldap._tcp	SRV	[0][100][389]ipa1.cse.unr.edu
 _kerberos-master._tcp	SRV	[0][100][88]ipa1.cse.unr.edu
 _kerberos-master._udp	SRV	[0][100][88]ipa1.cse.unr.edu
 _kerberos._udp	SRV	[0][100][88]ipa1.cse.unr.edu
 (Same as Zone)	MX	[10]cse-unr-edu.mail.protection.outlook.com

Certificate Authority Authorization Record (CAA)

- Specifies Certificate Authorities(CA) allowed to issue certificates for the domain
- <https://support.dnssimple.com/articles/caa-record/#what-is-a-caa-record>

 (Same as Zone)	CAA	0 issuewild "letsencrypt.org"
 (Same as Zone)	CAA	0 issue "letsencrypt.org"
 (Same as Zone)	CAA	0 issue "comodoca.com"



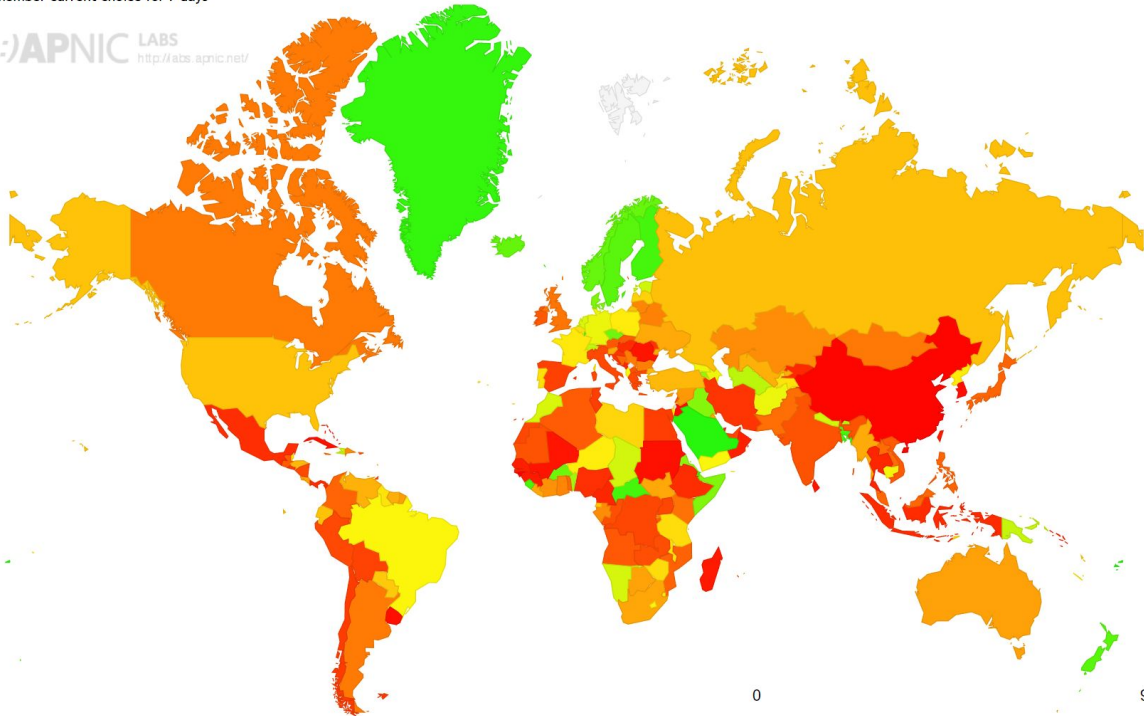
DNSSEC Adoption

DNSSEC Validation Rate by country (%)

[Click here for a zoomable map](#)

☐ Remember current choice for 7 days

(::)APNIC LABS
<http://labs.apnic.net/>



The BIND software

- Berkeley Internet Name Domain system
 - bind9 latest stable version
 - bind10 under development
- Reference implementation for DNS
- Components
 - Named - DNS Server
 - Resolver libraries
 - Command line utilities: dig, nslookup, and host
- Configuration
 - `/etc/bind/named.conf`

Basic Bind configuration

/etc/bind/named.conf.local

```
zone "ecc.engr.unr.edu" {  
    type master;  
    file "/etc/bind/zones/db.ecc.engr.unr.edu"; # zone file path  
};  
  
zone "195.20.172.in-addr.arpa" {  
    type master;  
    file "/etc/bind/zones/db.172.20.195";  
};
```

Basic Bind configuration

/etc/bind/zones/db.ecc.engr.unr.edu

```
;
; BIND data file for local loopback interface
;
$TTL      3600
@          IN      SOA      ns1.ecc.engr.unr.edu. admin.ecc.engr.unr.edu. (
                                3                ; Serial
                                604800           ; Refresh
                                86400            ; Retry
                                2419200          ; Expire
                                604800 )         ; Negative Cache TTL
;

; NS Records
          IN      NS       ns1.ecc.engr.unr.edu.

; NS A Records
ns1       IN      A        134.197.20.131

$INCLUDE  "/etc/bind/zones/imm/labs" ;
```

rndc

Command	Function
dumpdb	Dumps the DNS database to named_dump.db
flush [<i>view</i>]	Flushes all caches or those for a specified <i>view</i>
flushname <i>name</i> [<i>view</i>]	Flushes the specified <i>name</i> from the server's cache
freeze zone [<i>class</i> [<i>view</i>]] ^a	Suspends updates to a dynamic <i>zone</i>
thaw zone [<i>class</i> [<i>view</i>]] ^a	Resumes updates to a dynamic <i>zone</i>
halt	Halts named without writing pending updates
querylog	Toggles tracing of incoming queries
notify zone [<i>class</i> [<i>view</i>]] ^a	Resends notification messages for <i>zone</i>
notrace	Turns off debugging
reconfig	Reloads the config file and loads any new zones
recurring	Dumps queries currently recursing, named.recurring
refresh zone [<i>class</i> [<i>view</i>]] ^a	Schedules maintenance for a <i>zone</i>
reload	Reloads named.conf and zone files
reload zone [<i>class</i> [<i>view</i>]] ^a	Reloads only the specified <i>zone</i> or <i>view</i>
retransfer zone [<i>class</i> [<i>view</i>]] ^a	Recopies the data for <i>zone</i> from the master server
stats	Dumps statistics to named.stats
status	Displays the current status of the running named
stop	Saves pending updates and then stops named
trace	Increments the debug level by 1
trace level	Changes the debug level to the value <i>level</i>
validation newstate	Enables/disables DNSSEC validation on the fly

a. The *class* argument here is the same as for resource records, typically IN for Internet.

dnspython

```
pip install dnspython
```

```
answers = dns.resolver.query(record, rtype)
```

```
for rdata in answers:
```

```
    print('Host: {0}, Preference: {1}\n'.format(rdata, rdata))
```

dnsmasq

```
apt install dnsmasq
```

```
$EDITOR /etc/dnsmasq.conf
```

```
systemctl disable systemd-resolved #Disable the systemd DNS
```

```
systemctl start dnsmasq #Enable our dnsmasq setup
```

dnsmasq configuration

```
interface=lo
interface=br0-qemu

bind-interfaces

server=134.197.5.1
server=134.197.6.1
server=/ncr/192.168.2.1
server=/ecc.engr.unr.edu/134.197.20.131

#DHCP Configuration
dhcp-range=br0-qemu,192.168.200.20,192.168.200.150,255.255.255.0,12h

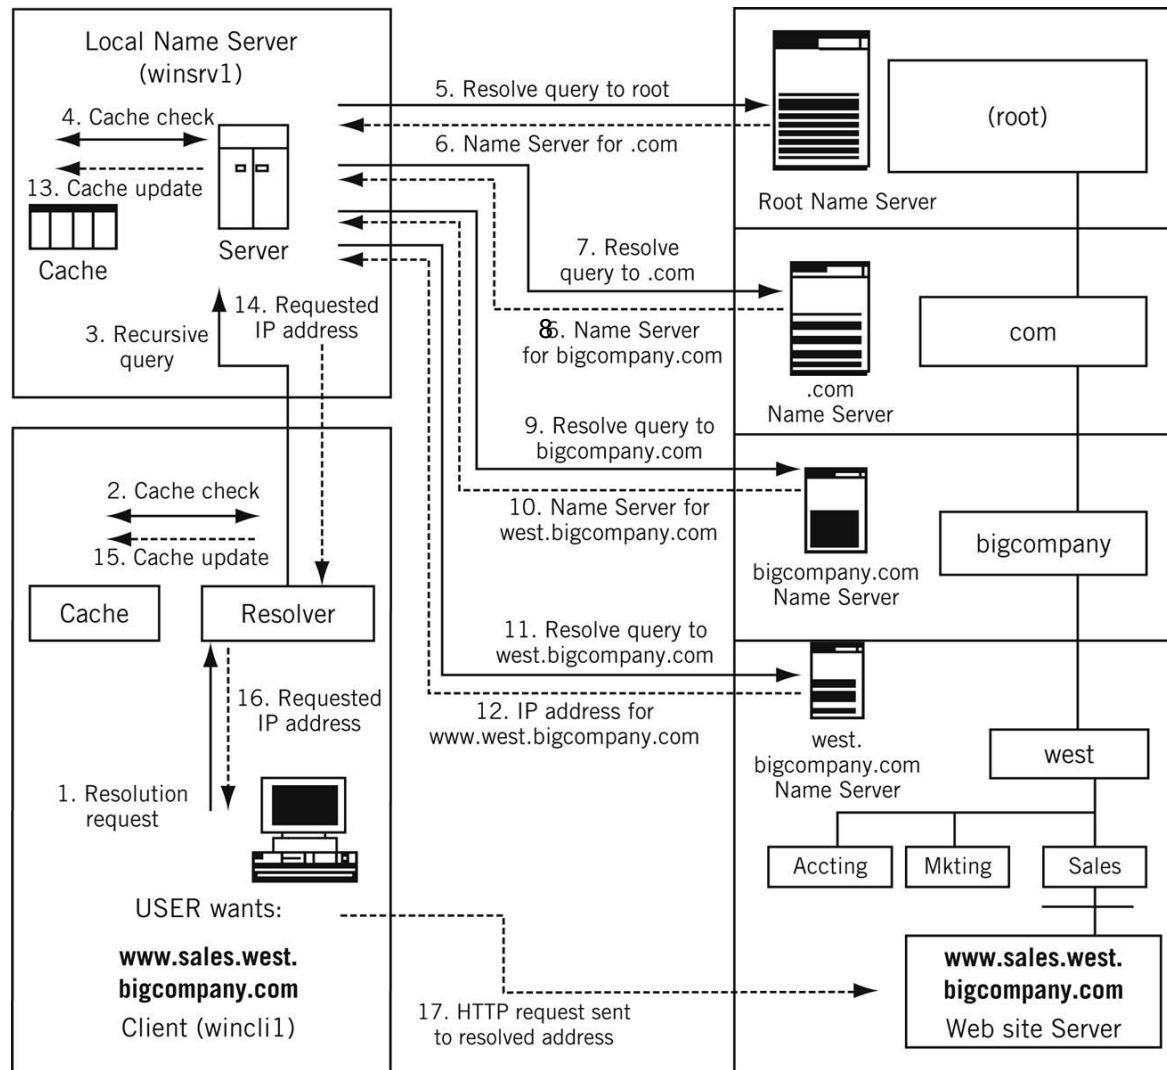
# Set default gateway
#dhcp-option=3,0.0.0.0

# Set DNS servers to announce
#dhcp-option=6,0.0.0.0

#dhcp-host=aa:bb:cc:dd:ee:ff,192.168.47.1
#dhcp-host=aa:bb:cc:ff:dd:ee,192.168.47.2

#####
#Local DNS configuration#
#####
local=/lan/
domain=lan

#Creates entries for /etc/hosts
expand-hosts
```



Additional Reading

The Illustrated Network

<https://learning.oreilly.com/library/view/the-illustrated-network/9780128110287/xhtml/chp023.xhtml>

Network #2: DNS

https://inst.eecs.berkeley.edu/~cs161/fa16/slides/network2_dns.key.pdf

<https://dnsviz.net/>