

# CS 447/647

Shells

*Ken Thompson has an automobile which he helped design. Unlike most automobiles, it has neither speedometer, nor gas gauge, nor any of the other numerous idiot lights which plague the modern driver. Rather, if the driver makes a mistake, a giant “?” lights up in the center of the dashboard. “The experienced driver,” says Thompson, “will usually know what’s wrong.”*

—Anonymous

# Overview

What is a shell?

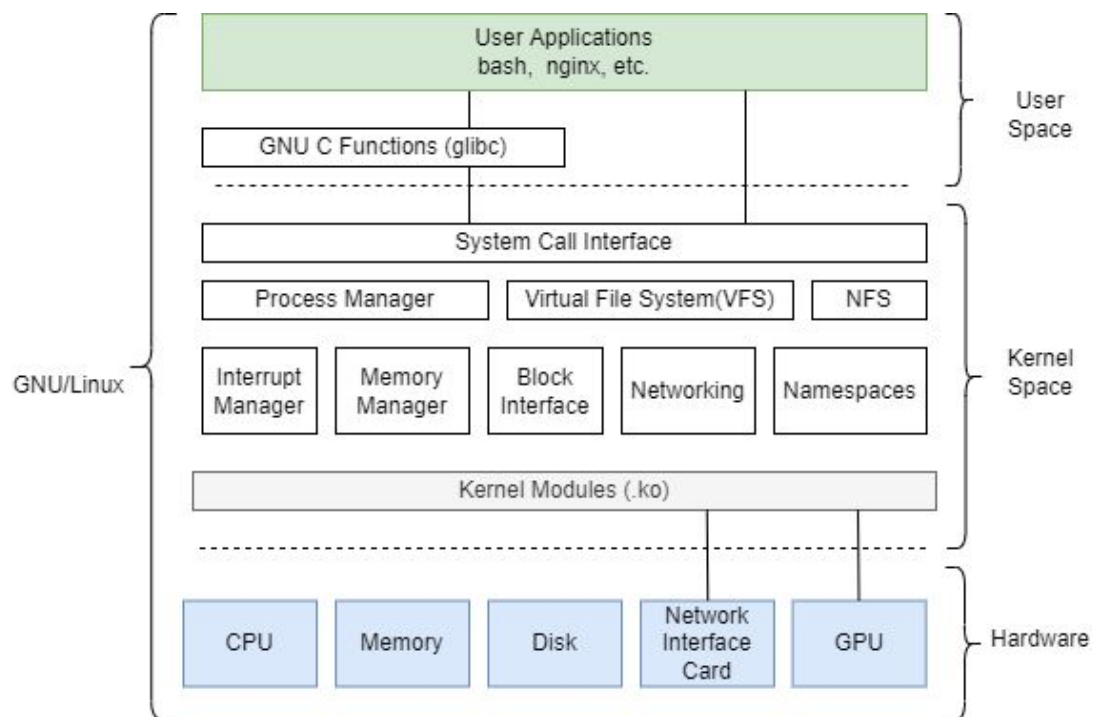
What are commands?

What are ELF binaries?

How are binaries located?

How are shared libraries located?

How do we build a dynamic user space?



# What is a shell?

- A command interpreter

- fork()
- exec()
- wait()

```
158 // Read and run input commands.
159 while(getcmd(buf, sizeof(buf)) >= 0){
160     if(buf[0] == 'c' && buf[1] == 'd' && buf[2] == ' '){
161         // Chdir must be called by the parent, not the child.
162         buf[strlen(buf)-1] = 0; // chop \n
163         if(chdir(buf+3) < 0)
164             printf(2, "cannot cd %s\n", buf+3);
165         continue;
166     }
167     if(fork1() == 0)
168         runcmd(parsecmd(buf));
169     wait();
170 }
171 exit();
```

cd

run

<https://github.com/mit-pdos/xv6-public/blob/master/sh.c> Line 158

# What does it interpret?

- Commands typed by a user or read from a file.
  - Symbols: `<>|&$`()`
  - Keywords: `cd, if, while, for, exit`
- Searches for the commands on the system.
  - PATH environment variable
- What does it execute?
  - language features
    - `|`
    - `for, if, while, function`
  - ELF binaries\*

```
typedef struct elf64_hdr {  
    unsigned char e_ident[EI_NIDENT]; /* ELF "magic number" */  
    Elf64_Half e_type;                /* ELF Type, commonly: executable, shared library */  
    Elf64_Half e_machine;             /* Architecture, commonly: AMD x86-64 */  
    Elf64_Word e_version;             /* ELF version, CURRENT or INVALID */  
    ...  
};
```



# GNU coreutils (ELF binaries)

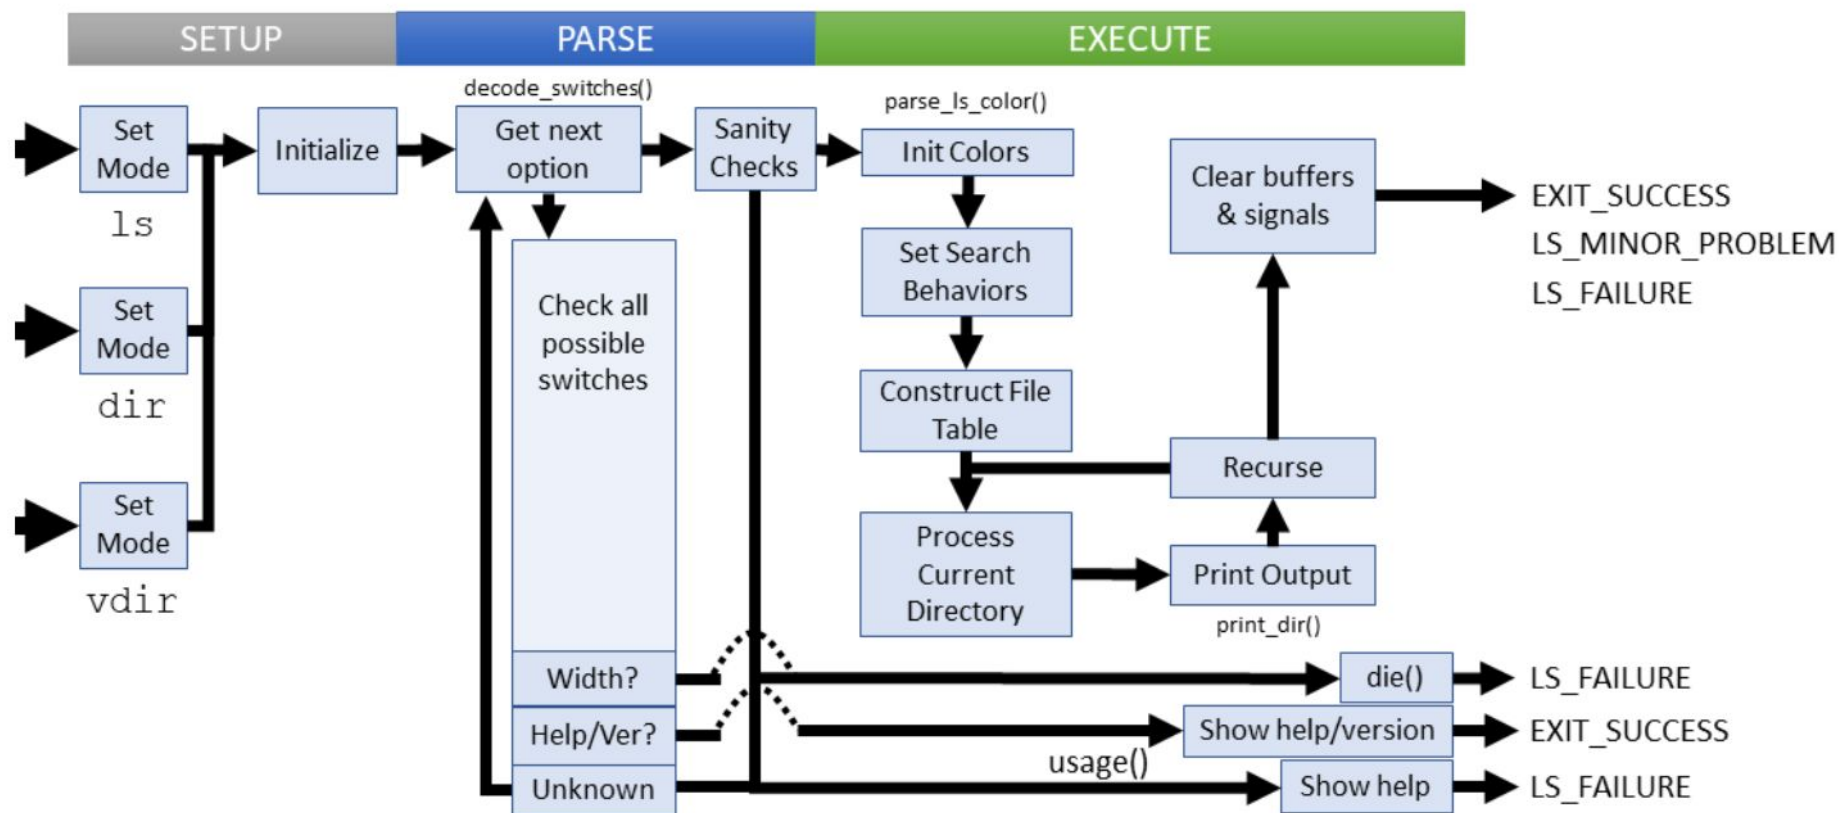
Core utilities which are expected to exist on every operating system

**head**  
**tail**  
**wc**  
md5sum  
**sort**  
shuf  
**uniq**  
comm  
split  
**tr**  
basename  
mktemp  
uptime

**ls**  
**rm**  
**cp**  
**mv**  
**ln**  
**chgrp**  
**chown**  
**chmod**  
**echo**  
**printf**  
**tee**  
false



ls(1)




# Why do we use a shell?

- Managing Services and Hardware
  - Application Layer Protocols
    - DHCP, DNS
- Engineering and Research
  - Scripting
  - Glueing
- Debugging
- Modify runtime parameters



# Why are shells still popular?

- Lightweight
- Accessibility
  - Blanket exception for CLI applications
  - Receive Federal money, must be accessible
- Ease of use\*
- Reproducibility
- Automation
- Change Control
  - git



---

# The UNIX- HATERS Handbook

---



*"Two of the most famous products of Berkeley are LSD and Unix.  
I don't think that is a coincidence."*

*Edited by Simson Garfinkel,  
Daniel Weise,  
and Steven Strassmann*

*Illustrations by John Klossner*

Million Dollar Settlements of Closed Captioning Website Accessibility Lawsuits Highlight Need for Dual Approach

MIT - \$1,000,000 in attorney fees

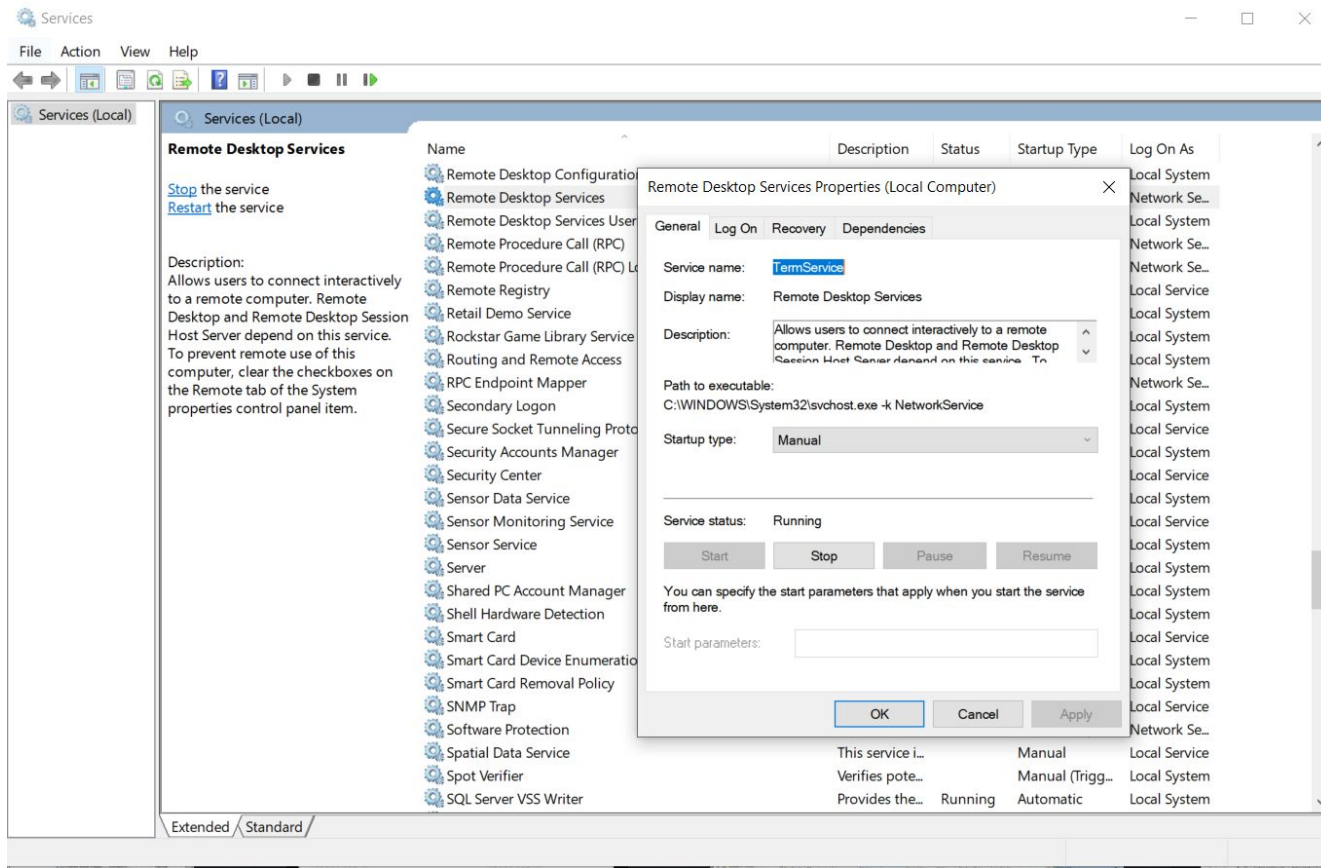


OHIO  
UNIVERSITY



**United States Supreme Court Denies  
Petition from Domino's Pizza**

# Example - Windows Service Management



# Example - Linux

```
systemctl restart bind9  
systemctl status bind9
```

```
# systemctl status bind9  
● named.service - BIND Domain Name Server  
   Loaded: loaded (/lib/systemd/system/named.service; enabled; vendor preset: enabled)  
   Active: active (running) since Thu 2022-02-10 17:44:20 PST; 1 day 18h ago  
     Docs: man:named(8)  
  Main PID: 10800 (named)  
    Tasks: 8 (limit: 9493)  
   Memory: 56.7M  
      CPU: 23.505s  
   CGroup: /system.slice/named.service  
           └─10800 /usr/sbin/named -f -u bind  
  
Feb 10 17:44:20 cs447-s22-newellz2-server named[10800]: automatic empty zone: view internal: B.E.F.IP6.ARPA  
Feb 10 17:44:20 cs447-s22-newellz2-server named[10800]: automatic empty zone: view internal: 8.B.D.0.1.0.0.2.IP6.ARPA  
Feb 10 17:44:20 cs447-s22-newellz2-server named[10800]: automatic empty zone: view internal: EMPTY.AS112.ARPA  
Feb 10 17:44:20 cs447-s22-newellz2-server named[10800]: automatic empty zone: view internal: HOME.ARPA  
Feb 10 17:44:20 cs447-s22-newellz2-server named[10800]: obtaining root key for view external from '/etc/bind/bind.keys'  
Feb 10 17:44:20 cs447-s22-newellz2-server named[10800]: set up managed keys zone for view external, file 'external.mkey'  
Feb 10 17:44:20 cs447-s22-newellz2-server named[10800]: configuring command channel from '/etc/bind/rndc.key'  
Feb 10 17:44:20 cs447-s22-newellz2-server named[10800]: command channel listening on 127.0.0.1#953  
Feb 10 17:44:20 cs447-s22-newellz2-server named[10800]: configuring command channel from '/etc/bind/rndc.key'
```

# bash configuration and dotfiles


- `/etc/bash.bashrc` - System-wide `.bashrc` file for interactive `bash(1)` shells.
- `/etc/profile` - Profile file for `sh(1)` and `bash(1)`
- `/etc/profile.d` - Collections of profile scripts. Often used for `$PATH` modifications
  - modules!
- `~/.bashrc` - Interactive shell configuration file
- `~/.bash_profile` - Executed for login shells



Best way to learn bash?

Use it!

# Getting help - man(1)



“Classic Unix documentation is written to be **telegraphic** but complete... The style assumes an active reader, one who is **able to deduce obvious unsaid consequences** of what is said, and who has the self-confidence to trust those deductions. **Read every word carefully**, because you will seldom be told anything twice.”

Eric Raymond

Pathname	Contents
<b>/bin</b>	Core operating system commands
<b>/boot</b>	Boot loader, kernel, and files needed by the kernel
<b>/compat</b>	On FreeBSD, files and libraries for Linux binary compatibility
<b>/dev</b>	Device entries for disks, printers, pseudo-terminals, etc.
<b>/etc</b>	Critical startup and configuration files
<b>/home</b>	Default home directories for users
<b>/lib</b>	Libraries, shared libraries, and commands used by <b>/bin</b> and <b>/sbin</b>
<b>/media</b>	Mount points for filesystems on removable media
<b>/mnt</b>	Temporary mount points, mounts for removable media
<b>/opt</b>	Optional software packages (rarely used, for compatibility)
<b>/proc</b>	Information about all running processes
<b>/root</b>	Home directory of the superuser (sometimes just /)
<b>/run</b>	Rendezvous points for running programs (PIDs, sockets, etc.)
<b>/sbin</b>	Core operating system commands <sup>a</sup>
<b>/srv</b>	Files held for distribution through web or other servers
<b>/sys</b>	A plethora of different kernel interfaces (Linux)
<b>/tmp</b>	Temporary files that may disappear between reboots
<b>/usr</b>	Hierarchy of secondary files and commands
<b>/usr/bin</b>	Most commands and executable files
<b>/usr/include</b>	Header files for compiling C programs
<b>/usr/lib</b>	Libraries; also, support files for standard programs
<b>/usr/local</b>	Local software or configuration data; mirrors <b>/usr</b>
<b>/usr/sbin</b>	Less essential commands for administration and repair
<b>/usr/share</b>	Items that might be common to multiple systems
→ <b>/usr/share/man</b>	On-line manual pages
<b>/usr/src</b>	Source code for nonlocal software (not widely used)
<b>/usr/tmp</b>	More temporary space (preserved between reboots)
<b>/var</b>	System-specific data and a few configuration files
<b>/var/adm</b>	Varies: logs, setup records, strange administrative bits
<b>/var/log</b>	System log files
<b>/var/run</b>	Same function as <b>/run</b> ; now often a symlink
<b>/var/spool</b>	Spooling (that is, storage) directories for printers, mail, etc.
<b>/var/tmp</b>	More temporary space (preserved between reboots)

a. The distinguishing characteristic of **/sbin** was originally that its contents were statically linked and so had fewer dependencies on other parts of the system. These days, all binaries are dynamically linked and there is no real difference between **/bin** and **/sbin**.

# manpage sections

Section	Contents
1	User-level commands and applications
2	System calls and kernel error codes
3	Library calls
4	Device drivers and network protocols
5	Standard file formats
6	Games and demonstrations
7	Miscellaneous files and documents
8	System administration commands
9	Obscure kernel specs and interfaces

# man

```
STAT(1)                                User Commands                                STAT(1)

NAME
    stat - display file or file system status

SYNOPSIS
    stat [OPTION]... FILE...

DESCRIPTION
    Display file or file system status.

    Mandatory arguments to long options are mandatory for short options
    too.

    -L, --dereference
        Follow links

    -f, --file-system
        display file system status instead of file status

    -c --format=FORMAT
        use the specified FORMAT instead of the default; output a new-
        line after each use of FORMAT

    --printf=FORMAT
        like --format, but interpret backslash escapes, and do not out-
        put a mandatory trailing newline; if you want a newline, include
        \n in FORMAT

    -t, --terse
        print the information in terse form

    --append-exe
        append .exe if cygwin magic was needed

    --help display this help and exit

    --version
        output version information and exit

    The valid format sequences for files (without --file-system):

    %a    access rights in octal (note '#' and '0' printf flags)
    %A    access rights in human readable form
    %b    number of blocks allocated (see %B)
    %B    the size in bytes of each block reported by %b
    %C    SELinux security context string
    %d    device number in decimal
    %D    device number in hex
    %f    raw mode in hex
```

# How are man pages rendered?

groff

- System for typesetting documents
  - Similar to TeX..
  - troff in 1971.
  - groff in 1990
- More like a compiler
- Text input files with embedded formatting

man 7 groff\_man

# man

```
apt install -y man manpages manpages-dev info groff
```

```
mandb #Regenerate manpages from roff source.  
      #Config in /etc/manpath.config
```

```
man -k ext4    # Keyword search for string “ext4”
```

```
man -K ext4    # Page through manpages that contain ext4
```

```
man -a intro   # Page through the intro manual
```

```
man ls         # manpage for ls
```

```
# Why do we need any of this when Google exists?
```