

CS447/647

Storage
LVM and RAID

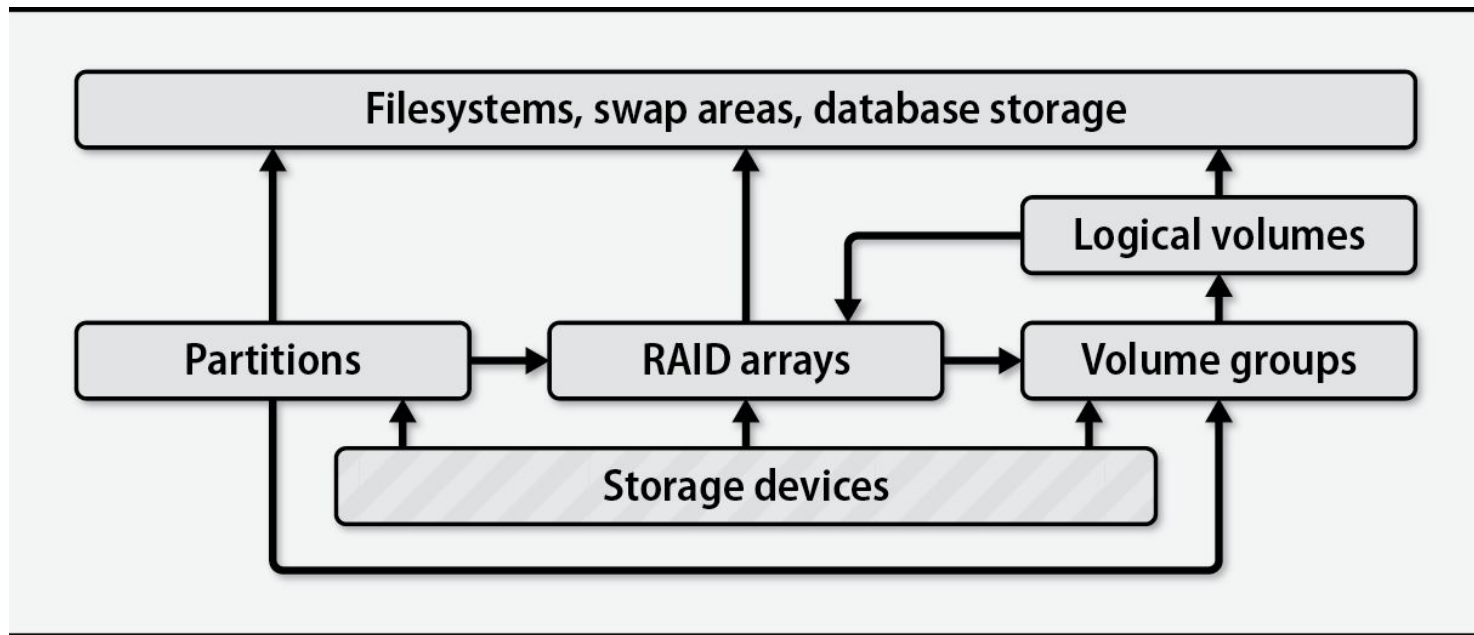
	Cloud Provider 1	Cloud Provider 2	Cloud Provider 3	WhiteBox
Cost per GB	\$ 0.02	\$ 0.02	\$ 0.02	\$ 0.01
Ingress Cost	\$ -	\$ -	\$ -	
Egress Cost	\$ 0.08	\$ 0.09	\$ 0.09	\$ -
Monthly Cost	\$ 2,048.00	\$ 2,048.00	\$ 2,048.00	\$ -
20T Egress	\$ 1,638.40	\$ 1,843.20	\$ 1,843.20	\$ -
1st Year Cost	\$ 24,576.00	\$ 24,576.00	\$ 24,576.00	\$ 5,688.00
5 Years Cost	\$ 124,518.40	\$ 124,723.20	\$ 124,723.20	\$ 5,688.00

References

Nemeth, Evi, et al. *UNIX and Linux System Administration Handbook*. Addison-Wesley, 2018.

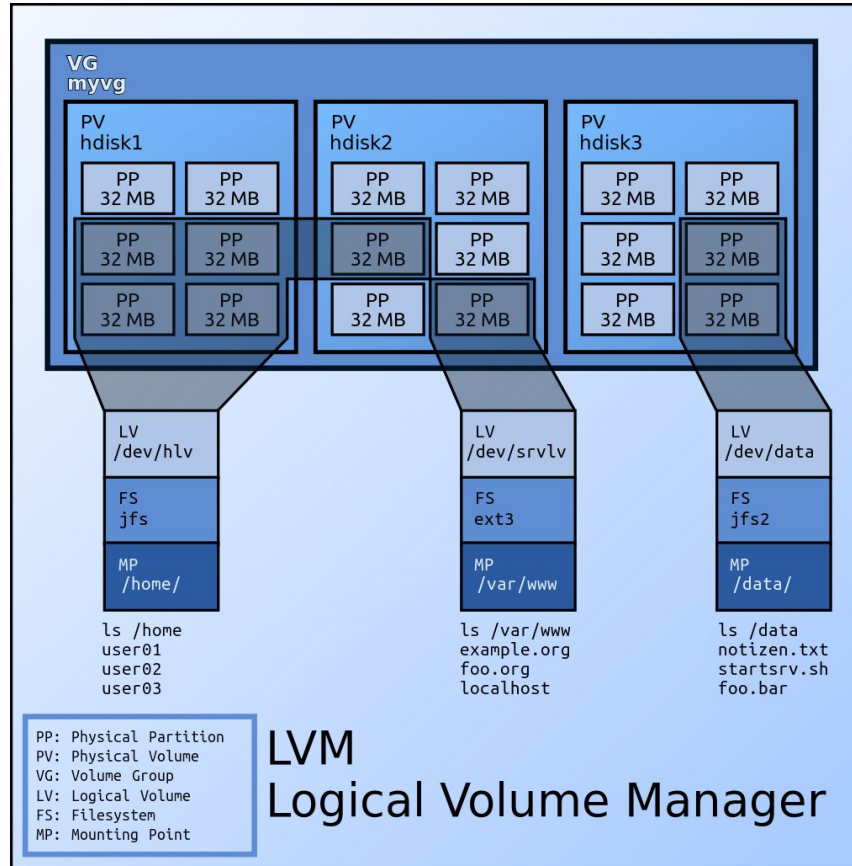
Remzi H., et al. *Operating Systems: Three Easy Pieces*, Arpaci-Dusseau Books, August, 2018 (Version 1.00)

<https://pages.cs.wisc.edu/~remzi/OSTEP/>



LVM - Logical Volume Management

- Provides tools to create virtual block devices from physical devices
- Virtual devices are easier to manage than physical devices
- Three requirements
 - Device-mapper kernel module
 - Userspace device-mapper
 - Userspace lvm2 tools
- Three components
 - Physical Volume
 - Volume Group
 - Logical Volume



LVM - Userspace Basics

```
truncate -s 5G disk1.img
```

```
truncate -s 5G disk2.img
```

```
losetup -d /dev/loop0
```

```
losetup --find --show disk1.img
```

```
losetup --find --show disk2.img
```

```
parted -s /dev/loop0 'mklabel gpt mkpart lvpart1 1M 1G'
```

```
parted -s /dev/loop1 'mklabel gpt mkpart lvpart1 1M 1G'
```

```
pvcreate -v /dev/loop0p1
```

```
pvcreate -v /dev/loop1p1
```

```
pvdisplay
```

LVM - Userspace Basics

```
vgcreate vg0 /dev/loop0p1
vgextend vg0 /dev/loop1p1
vgdisplay
lvcreate -L 1.5G -n lv0 vg0
lvdisplay
parted -s /dev/mapper/vg0-lv0 'print'
lvresize -L -.5G vg0/lv0 #Shrink
lvresize -L +.5G vg0/lv0 #Grow
```


LVM - Userspace Physical Volume Moves

```
truncate -s 5G disk3.img  
losetup --find --show disk3.img  
parted -s /dev/loop2 'mklabel gpt mkpart lvpart1 1M 1G'  
pvcreate -v /dev/loop2p1  
pvmove /dev/loop0p1 /dev/loop2p1  
pvdisplay -m
```

LVM - Userspace Snapshots

```
pvcreate /dev/loop0
pvcreate /dev/loop1
vgcreate vg0 /dev/loop0 /dev/loop1
lvcreate -L 1G -n lv0 vg0
mkfs.ext /dev/mapper/vg0-lv0
mount /dev/mapper/vg0-lv0 /mnt
touch /mnt/file
lvcreate -L1G -s -n lv0-snap vg0/lv0 #Create a snapshot named
lv0-snap
umount /mnt
mount /dev/mapper/vg0-lv0--snap /mnt
umount /mnt && lvremove vg0/lv0-snap
```

RAID

- We often want disks to be
 - faster
 - larger
 - more reliable
- Redundant Array of Inexpensive Disks (RAID)
 - Developed in the late 1980's by the CS department at Berkeley
 - Technique to make multiple disks to appear as a single disk
 - More storage, better performance and reliability
 - Complex
 - Multiple Disks
 - RAM
 - Processors

RAID

- Advantages
 - Performance
 - Capacity
 - Reliability
 - Redundancy - Tolerate the loss of a disk
- Transparency - Easing Deployment
 - New functionality
 - Demands no changes to the rest of the system
 - RAID is a perfect example
 - Looks like one big disk
 - Solved the deployment problem

Interface and RAID Internals

- Filesystem sees one big disk
- When a logical IO request is made, a RAID must:
 - Calculate which disk
 - Issue physical IO request
 - Mirroring results in 2 physical writes for 1 logical
- SATA, SCSI or NVME
 - NVME is software RAID only
 - mdraid
 - Hardware RAID for SATA and SCSI

RAID0 - Striping

- Upper Bound of Performance and Capacity
- “Perfect reliability”
 - One disk fails the whole array fails
- Excellent Performance
 - All disks are utilized
 - Often parallel
- Best Capacity
 - All Disks combined

Disk 0	Disk 1	Disk 2	Disk 3
0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

Figure 38.1: **RAID-0: Simple Striping**

RAID Mapping Problem

- How does the RAID map logical blocks to physical disks?
 - Logical Block A
 - $\text{Disk} = A \% \text{number_of_disk}$
 - $\text{Offset} = A / \text{number_of_disks}$
- So, a write to $A = 14$ with a 4 disk RAID
 - Disk: $14 \% 4 = 2$
 - Offset: $14 / 4 = 3$

Chunk Size



- Affects performance
 - 64Kb, 512Kb (common)
- Group blocks together on a single disk
- Small chunks means files are striped across many disks
- Large chunks reduce intra-file parallelism
- Art more than a science

RAID0 - Performance

- Sequential Workload - Large continuous chunks
- Random Workload - Small requests for random disk locations (blocks)
 - Databases
- In general Sequential > Random
- Number of Disks * Random Rate
- Number of Disks * Sequential Rate
- Full Bandwidth

How do we benchmark?

- `dd(1)` - Basic Sequential Read\Write
- `hdparm(1)` - Basic buffered reads test
- `fio(1)` - synthetic benchmarks, 'real world' workloads



RAID1 - Mirroring

- Copy of each block on a different disk
- Each logical write is two physical writes
 - Slowest of the two
 - Happen in parallel
- Sequential Performance: $(N/2) * \text{Sequential Rate}$
- Random Performance: $(N/2) * \text{Random Rate}$

Disk 0	Disk 1	Disk 2	Disk 3
0	0	1	1
2	2	3	3
4	4	5	5
6	6	7	7

Figure 38.3: **Simple RAID-1: Mirroring**

RAID4 - Parity

- Each stripe has a parity block
- Parity calculated using XOR
- Can lose 1 disk
 - Replacement has to be rebuilt
- Performance
 - $(N - 1) * \text{Rate}$
 - Random write does not improve when you add disks

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	1	2	3	P0
4	5	6	7	P1
8	9	10	11	P2
12	13	14	15	P3

Figure 38.4: RAID-4 With Parity

C0	C1	C2	C3	P
0	0	1	1	$\text{XOR}(0,0,1,1) = 0$
0	1	0	0	$\text{XOR}(0,1,0,0) = 1$

Disk and Flash Drive Rebuild Times

RAID	Capacity TB	Capacity GB	Capacity MB	Seq Write Speed MB/sec	Rebuild Time Minimum secs	Minutes	Hours
Disk	0.72	72	72,000	80	900	15	0.25
	1	1,000	1,000,000	115	8,696	145	2.42
	4	4,000	4,000,000	115	34,783	580	9.66
SSD FlashMax III	2.2	2,200	2,200,000	1,400	1,571	26	0.44
Intel D3600	2	2,000	2,000,000	1,500	1,333	22	0.37
Micron 9100	3.2	3,200	3,200,000	2,000	1,600	27	0.44
Intel DC P3608	4	4,000	4,000,000	3,000	1,333	22	0.37

https://www.theregister.com/2016/05/13/disak_versus_ssd_raid_rebuild_times/

RAID5 - Rotating Parity

- Operates identically to RAID4
- Random read performance slightly better
- Random Write: $(N/4) * R$

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	1	2	3	P0
5	6	7	P1	4
10	11	P2	8	9
15	P3	12	13	14
P4	16	17	18	19

Figure 38.7: RAID-5 With Rotated Parity

	RAID-0	RAID-1	RAID-4	RAID-5
Capacity	$N \cdot B$	$(N \cdot B)/2$	$(N - 1) \cdot B$	$(N - 1) \cdot B$
Reliability	0	1 (for sure) $\frac{N}{2}$ (if lucky)	1	1
Throughput				
Sequential Read	$N \cdot S$	$(N/2) \cdot S^1$	$(N - 1) \cdot S$	$(N - 1) \cdot S$
Sequential Write	$N \cdot S$	$(N/2) \cdot S^1$	$(N - 1) \cdot S$	$(N - 1) \cdot S$
Random Read	$N \cdot R$	$N \cdot R$	$(N - 1) \cdot R$	$N \cdot R$
Random Write	$N \cdot R$	$(N/2) \cdot R$	$\frac{1}{2} \cdot R$	$\frac{N}{4} R$
Latency				
Read	T	T	T	T
Write	T	T	$2T$	$2T$

Figure 38.8: RAID Capacity, Reliability, and Performance

mdraid - Linux Software Raid

- RAID devices are virtual devices created from two or more block devices
- Many devices to one virtual device
- RAID Levels offer performance and redundancy
- Levels
 - **LINEAR** - concatenates devices in a single device. Like LVM Volume Group
 - **RAID0** (striping) - No redundancy, performance
 - **RAID1** (mirroring) - Mirror disks
 - **RAID4** - RAID0 plus a parity disk
 - **RAID5** - RAID4 with parity spread across disks, lose 1 disk
 - **RAID6** - RAID5 with two parity segments, lose two disks
 - **RAID10** - striped mirroring
 - MULTIPATH - Not a RAID. Multiple paths to same storage device. iSCSI
 - FAULTY - provides a layer over a true device that can be used to inject faults
 - CONTAINER - Set of devices

.config - Linux/x86 4.15.18 Kernel Configuration

> Device Drivers > Multiple devices driver support (RAID and LVM)

Multiple devices driver support (RAID and LVM)

Arrow keys navigate the menu. <Enter> selects submenus ---> (or empty submenus ----). Highlighted letters are hotkeys. Pressing <Y> includes, <N> excludes, <M> modularizes features. Press <Esc><Esc> to exit, <?> for Help, </> for Search. Legend: [*] built-in [] excluded <M> module < > module

--- Multiple devices driver support (RAID and LVM)

[*] RAID support

[*] Autodetect RAID arrays during kernel boot

<M> Linear (append) mode

{M} RAID-0 (striping) mode

{M} RAID-1 (mirroring) mode

{M} RAID-10 (mirrored striping) mode

{M} RAID-4/RAID-5/RAID-6 mode

<M> Multipath I/O support

<M> Faulty test module for MD

<M> Cluster Support for MD

<M> Block device as cache

[] Bcache debugging

[] Debug closures

<*> Device mapper support

[] request-based DM: use blk-mq I/O path by default

[] Device mapper debugging support

[] Block manager locking

<M> Crypt target support

<M> Snapshot target

+(<+>)

<Select>

< Exit >

< Help >

< Save >

< Load >

mdadm

```
mdadm -C /dev/md0 -l stripe -n 2 /dev/loop0 /dev/loop1 #RAID0
```

```
mdadm -C /dev/md0 -l raid1 -n 2 /dev/loop0 /dev/loop1 #RAID1
```

#RAID5

```
mdadm -C /dev/md0 -l raid5 -n 3 /dev/loop0 /dev/loop1 /dev/loop2 -x 1 /dev/loop3
```

```
mdadm --detail /dev/md0
```

```
cat /proc/mdstat
```

#Persist across reboots

```
mdadm --detail --scan --verbose > /etc/mdadm.conf
```

mdadm

```
mdadm -C /dev/md0 -l raid1 -n 2 /dev/loop0 /dev/loop1 #RAID1
```

#Fail a disk

```
mdadm --fail /dev/md0 /dev/loop0
```

```
mdadm --remove /dev/md0 /dev/loop0
```

```
mdadm --add /dev/md0 /dev/loop0
```

mdadm

#Fake Corruption


```
dd if=/dev/zero of=/dev/loop0 bs=1M count=128
```

```
echo check > /sys/block/md0/md/sync_action
```

```
cat /proc/mdstat
```

drbd

- Distributed Replicated Block Device
- High-Availability Storage
- RAID1 over the network
- Active/Passive Setup



MySQL High Availability

drbd

```
# begin resource drbd0
```

```
resource drbd0 {  
    protocol C;  
    startup {degr-wfc-timeout 120;}  
    disk {on-io-error detach;}  
    net {}  
    syncer {  
        rate 100m;  
        al-extents 257;  
    }  
    on server1 {  
        device /dev/drbd0;  
        disk /dev/sdb;  
        address 192.168.1.230:7788;  
        meta-disk internal;  
    }  
}
```

```
on server2 {  
    device /dev/drbd0;  
    disk /dev/sdb;  
    address 192.168.1.231:7788;  
    meta-disk internal;  
}
```

DRBD

```
apt install -y drbd-utils
mkdir -p /srv/drbd && cd /srv/drbd
truncate -s 1G res0.img
losetup --show --find res0.img
#Create /etc/drbd.d/drbd0.res
drbdadm create-md drbd0
drbdadm up drbd0
cat /proc/drbd
drbdadm primary --force drbd0
```


Why does this matter?

Complex container and virtual machine managers use LVM, mdadm and drbd extensively. They are heavily scripted.

