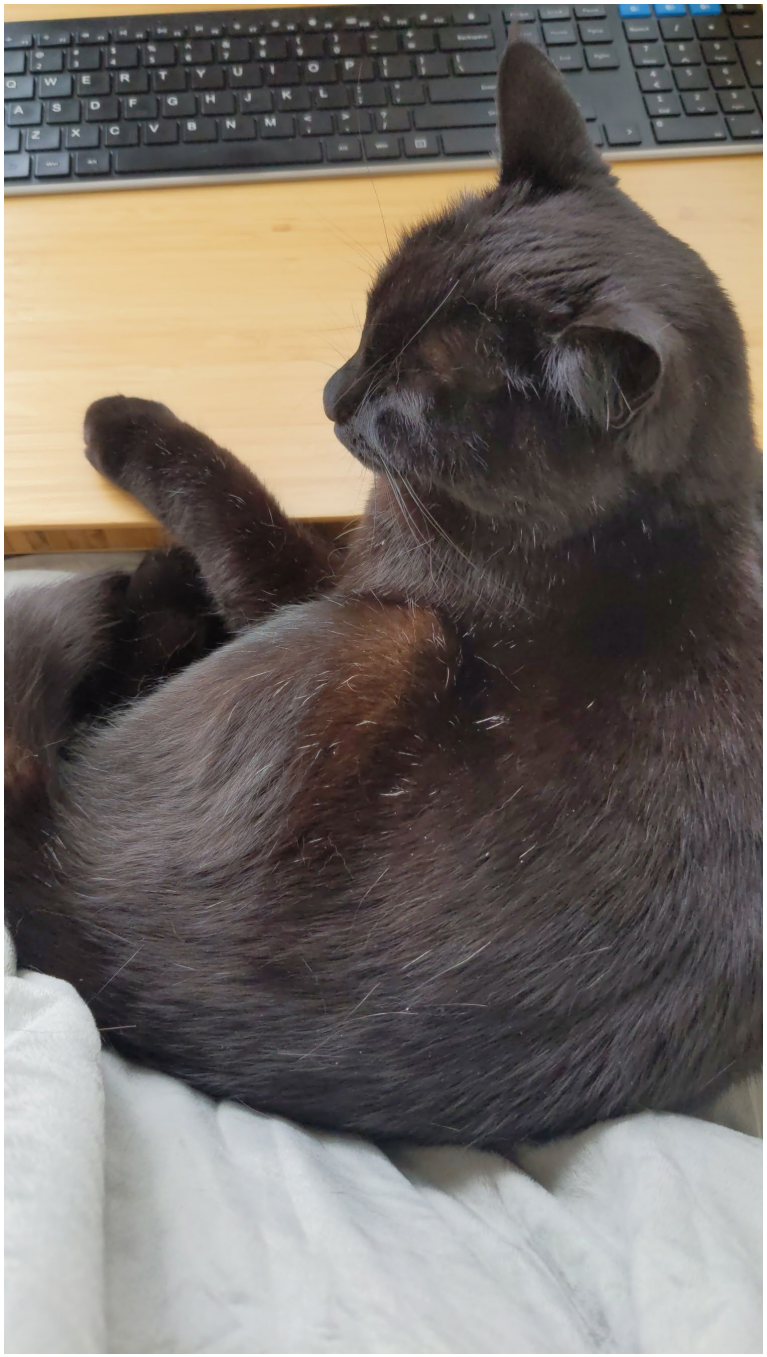

Storage

ERIN KEITH

Goals

1. Review General Storage Methods
2. Introduce Database Storage Methods

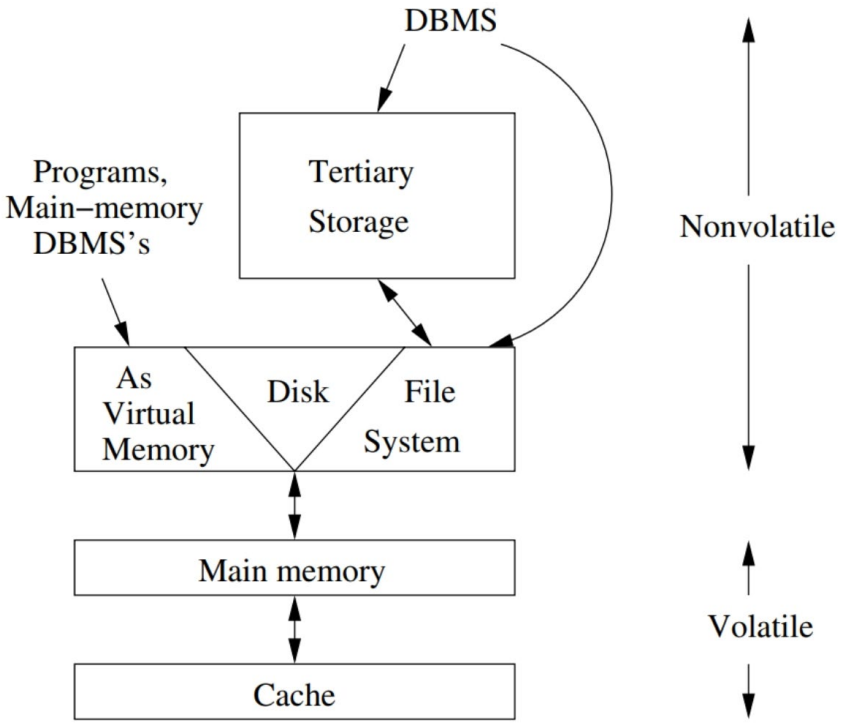

Figure 13.1: The memory hierarchy

Memory

Database systems always involve secondary storage

- Considerations
  - retrieval times
  - volatility

The Database Environment

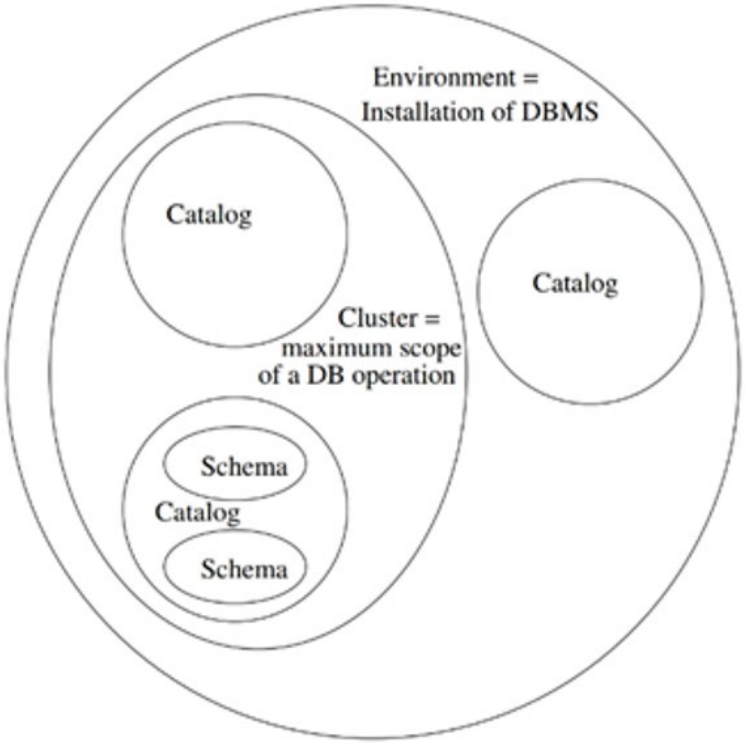

Figure 9.2: Organization of database elements within the environment

Vocabulary Review

record

- represents a tuple or object

tuple

- a row with linked data about a certain entity (it can be any object)

row

- an entry in a table

Database Storage

record

- storage element for an entry in a table
- consecutive bytes in some disk block

Database Storage

record

- The simplest sort of record consists of fixed-length fields, one for each attribute of the represented tuple.
  - ***field***
    - column, property, attribute
    - size becomes important when we're creating tables

Database Storage

header

- Often, the record begins with a header, a fixed-length region where information about the record itself is kept.
  - a pointer to the schema
  - length of the record (for skipping over)
  - timestamps (for last modified/read)
  - pointers to the fields (instead of the schema)

Example

```
CREATE TABLE MovieStar (
    name      CHAR(30),
    address   VARCHAR(255),
    gender    CHAR(1),
    birthdate DATE
);
```

Database Storage

blocks

- Records representing tuples of a relation are stored in blocks of the disk and moved into main memory (along with their entire block) when we need to access or update them.

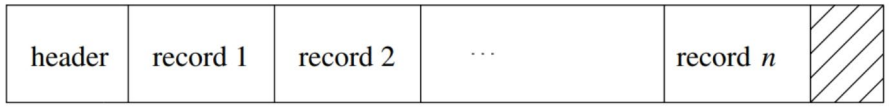

Figure 13.17: A typical block holding records

In addition to the records, there is a *block header* holding information such as:

1. Links to one or more other blocks that are part of a network of blocks such as those that will be described in Chapter 14 for creating indexes to the tuples of a relation.
2. Information about the role played by this block in such a network.
3. Information about which relation the tuples of this block belong to.
4. A "directory" giving the offset of each record in the block.
5. Timestamp(s) indicating the time of the block's last modification and/or access.

Client-Server Systems

Commonly, a database system consists of a server process that provides data from secondary storage to one or more client processes that are applications using the data.

- The server and client processes may be on one machine, or the server and the various clients can be distributed over many machines.
- There is a client virtual memory to server physical memory mapping, but it's not too important.

BLOBs

Binary Large OBjects

- GIF
- JPEG
- MPEG
- stored on a sequence of blocks (hopefully sequential)
- client may receive pieces

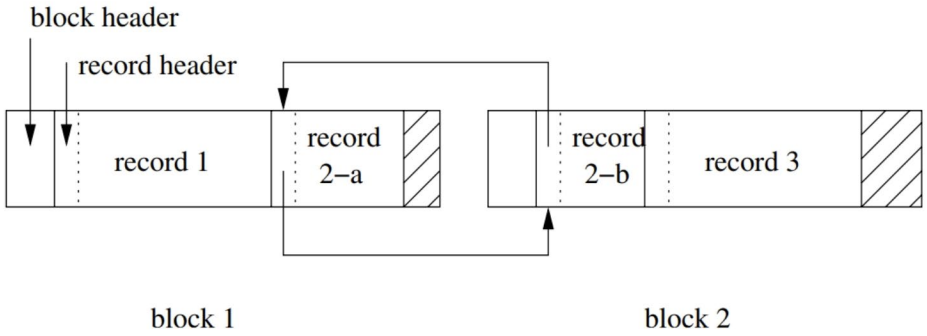

Figure 13.27: Storing spanned records across blocks

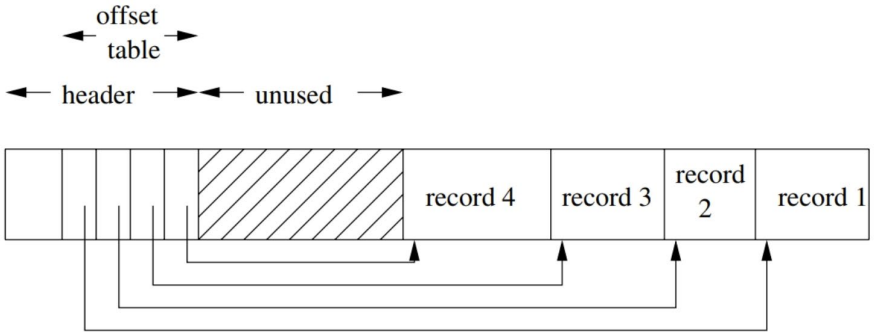

Figure 13.28: An offset table lets us slide records within a block to make room for new records

Operations

Insertion

- adding one or more records to a table
- keeping records ordered can cause complications

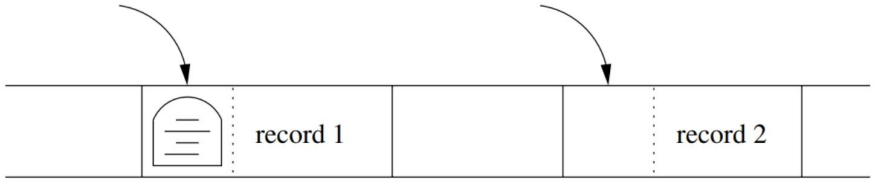

Figure 13.30: Record 1 can be replaced, but the tombstone remains; record 2 has no tombstone and can be seen when we follow a pointer to it

Operations

Deletion

- removing one or more records from a table
- can reclaim or "slide" around
- can replace with a "tombstone"

Operations

Update

- changing one or more records in a table
- same problems if it is a variable length record

Retrieval

- finding one or more records in a table

Indexes

- used behind the scenes to connect tables
- optimizes data retrieval
- decreases performance of database modifications

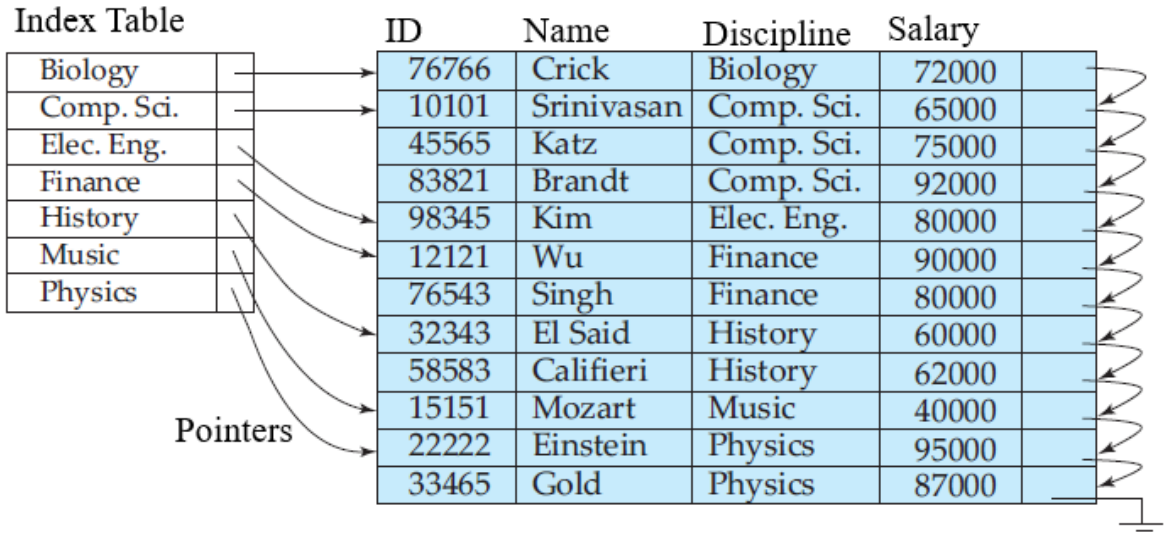

Indexes

When relations are very large, it becomes expensive to scan all the tuples of a relation to find those (perhaps very few) tuples that match a given condition.

- An index on an attribute **`A`** of a relation is a data structure that makes it efficient to find those tuples that have a fixed value for attribute **`A`**.

Next Class

Module:

Week 2: Background, Ch 2

Topic:

The Relational Model