



Design

ERIN KEITH

Goals

1. Background
2. Relations in SQL
3. Introduction to PostgreSQL

SQL

There is a current standard for SQL, called SQL-99. Most commercial database management systems implement something similar, but not identical to, the standard.

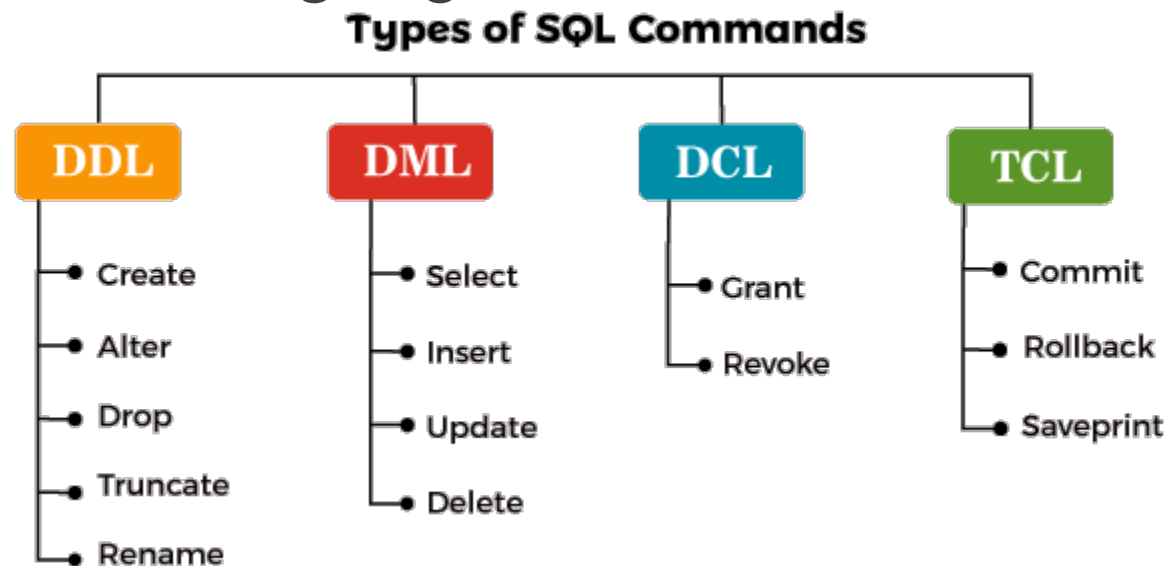
SQL Commands

DDL – Data Definition

DML – Data Manipulation

DCL – Data Control Language

TCL – Transaction Control Language



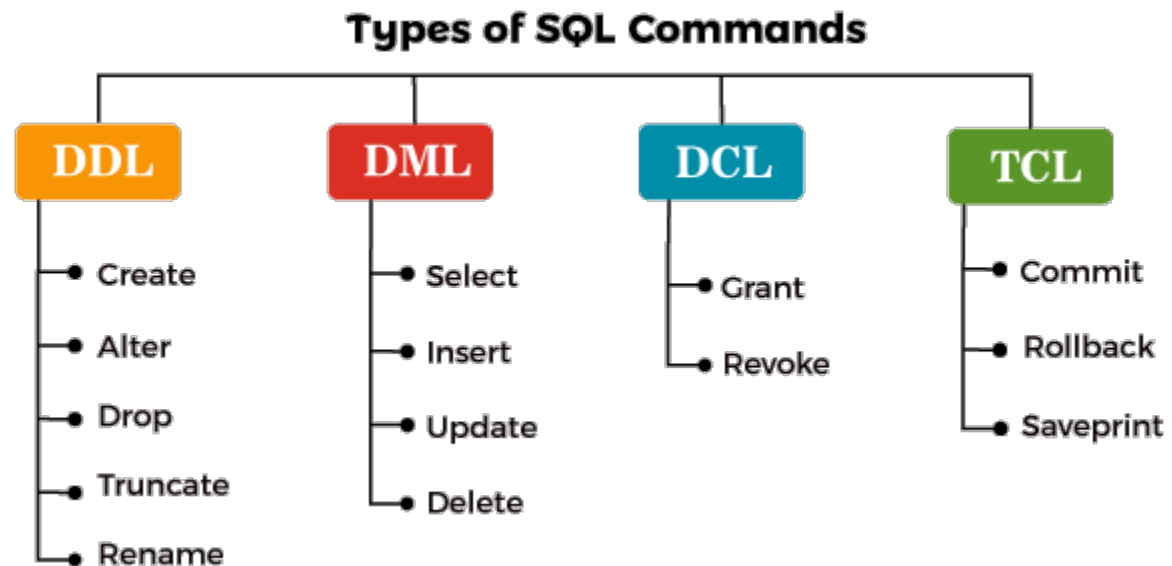
SQL Commands

DDL – tables

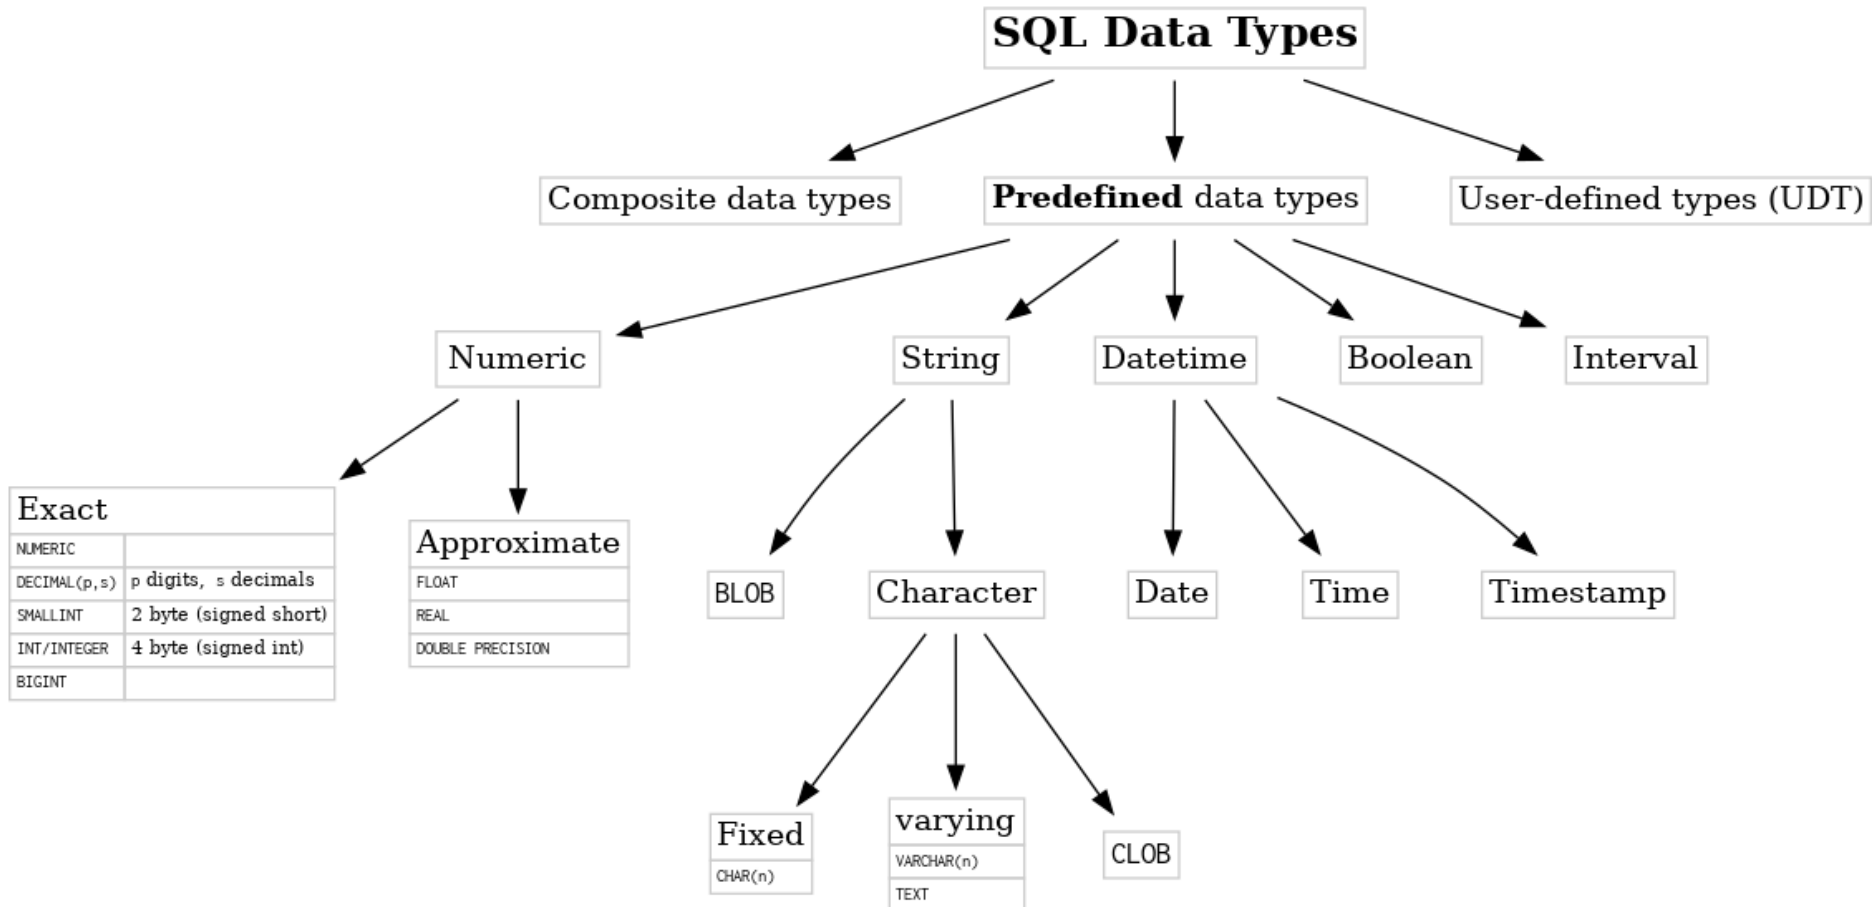
DML – queries

DCL – users

TCL – transactions



SQL Data Types



Create Table



SCHEMA

```
Movies(  
    title:string,  
    year:integer,  
    length:integer,  
    genre:string,  
    studioName:string,  
    producerC#:integer )  
)
```

SQL

```
CREATE TABLE Movies (  
    title        CHAR(100),  
    year         INT,  
    length       INT,  
    genre        CHAR(10),  
    studioName   CHAR(30),  
    producerC#   INT
```



QUESTION 2

SCHEMA

```
MovieStar(  
    name:string,  
    address:string,  
    gender:char,  
    birthdate:date  
)
```

SQL

```
CREATE TABLE MovieStar (  
    name        CHAR(30),  
    address     VARCHAR(255),  
    gender      CHAR(1),  
    birthdate   DATE  
);
```


Keys

- cannot be **NULL**
- the set of keys must be *unique*
- otherwise the insert operation will result in an error

```
CREATE TABLE MovieStar (  
    name CHAR(30) PRIMARY KEY,  
    address VARCHAR(255),  
    gender CHAR(1),  
    birthdate DATE  
);
```

OR

```
CREATE TABLE MovieStar (  
    name CHAR(30),  
    address VARCHAR(255),  
    gender CHAR(1),  
    birthdate DATE,  
    PRIMARY KEY (name)  
);
```

Keys

- cannot be **NULL**
- the set of keys must be *unique*
 - otherwise the insert operation will result in an error

```
CREATE TABLE Movies (  
    title          CHAR(100),  
    year           INT,  
    length         INT,  
    genre          CHAR(10),  
    studioName     CHAR(30),  
    producerC#     INT,  
    PRIMARY KEY (title, year)  
);
```

Delete Table

```
DROP TABLE <TableName>;
```

Alter Table

DROP or **ADD** an attribute on a table

ALTER TABLE <TableName> ADD <Attribute>

```
ALTER TABLE MovieStar ADD phone CHAR(16);
```

```
ALTER TABLE MovieStar DROP birthdate;
```

Default Values

```
birthdate DATE DEFAULT DATE '0000-00-00'
```

Database Servers

There are often many services involved in your system
We'll talk about how to put it all together soon, but for now we'll focus on


Database Servers. These processes run the DBMS and perform queries and modifications at the request of the application servers.

The **D**atab**B**ase **M**anagement **S**ystem is a service that we can **start** or **stop** running.

When the service is running, we can send it queries or commands and receive the results.

PostgreSQL

- <https://www.postgresql.org/docs/current/tutorial-arch.html>



- The database server can be running locally!
- PostgreSQL port: 5432

PostgreSQL

- Clients

CLI – typing SQL commands in the terminal

GUI – like an IDE for the environment

Applications – code programmatically connects and interacts with the database

PostgreSQL

- Download and install:
<https://www.postgresql.org/download/>

Packages and Installers

Select your operating system family:

Linux



macOS



Windows



BSD



Solaris



- Admin users: <https://www.postgresql.org/docs/current/notation.html>
- You may have to add the bin dir to the **PATH** env var
- You may have to add the absolute path to the binary folder to the **PATH** environment variable

PostgreSQL

- You can login as the admin user

```
> psql -U postgres
```

- You can create a user which matches your system username

```
C: Users\Work OneDrive\Documents\UNR\CS457> createuser -U postgres -P -s -e Work
```

- Then you can create a database and access it

```
> createdb Movies  
> psql Movies
```

PostgreSQL Activity



PostgreSQL Activity

CS 457/657
Fall 2023 – PostgreSQL Activity

Student Names:



ALL QUESTIONS ARE REFERRING TO THE PostgreSQL DATABASE SYSTEM.

How do you create a new database?

Terminal: `createdb dbName`
OR
Invoke DBMS: `CREATE DATABASE dbName`
*name should be on same line

What are possible issues you might run into while attempting to create a new database?

Don't have permission
Incorrect PATH env var
Incorrect installation
Insufficient disk space
Database exists
Server is not running (or wrong port, etc.)

How do you delete a database?

`dropdb dbName`
OR
`DROP DATABASE dbName`

How do you access a database?

`psql dbName`

more on the back →

How do you create a table?

```
CREATE TABLE tblName(  
  name VARCHAR(80), --a descriptive comment  
  ...  
);
```

How do you indicate a primary key?

```
CREATE TABLE tblName(  
  name VARCHAR(80) PRIMARY KEY, --a descriptive comment  
  ...  
);
```

How do you indicate a foreign key?

```
CREATE TABLE tblName(  
  name VARCHAR(80) PRIMARY KEY, --a descriptive comment  
  partNum INT REFERENCES tbl(col),  
  ...  
  --FOREIGN KEY (partNum) REFERENCES tbl(col)  
);
```

How do you add rows to a table?

```
INSERT INTO tblName VALUES('Erin', 134243);  
*multiple rows???
```

How do you query a table?

```
SELECT * FROM tblName
```

How do you join two tables?

```
SELECT * FROM a b WHERE a.id = b.id;  
SELECT * FROM t1 JOIN t2 ON t1.partNum = t2.partNum
```

PostgreSQL Movies

CS 457/657

Fall 2023 – PostgreSQL Movies

Example 2.21: Consider the two relations from our running movie database:

```
Movies(title, year, length, genre, studioName, producerC#)
MovieExec(name, address, cert#, netWorth)
```

How do you create the new database?

```
createdb Movies
```

How do you access the database?

```
sql Movies
```

How do you create the MovieExec table?

```
CREATE TABLE MovieExec(
  name VARCHAR(80),
  address VARCHAR(255),
  cert_num INT PRIMARY KEY,
  net_worth INT
);
```

How do you create the Movie table?

```
CREATE TABLE Movie(
  title VARCHAR(80),
  year INT,
  length INT,
  genre VARCHAR(80),
  studio_name VARCHAR(80),
  producer_num INT REFERENCES MovieExec (cert_num),
  PRIMARY KEY (title, year)
);
```

How do you get the name and net worth for producers of movies made since the turn of the century?

SQL Query Practice

Example 2.21: Consider the two relations from our running movie database:

`Movies(title, year, length, genre, studioName, producerC#)`

`MovieExec(name, address, cert#, netWorth)`

How do you get the name and net worth for producers of movies made since the turn of the century?

PostgreSQL

1.

```
Movies=# CREATE TABLE MovieExec(name VARCHAR(80), address VARCHAR(255), cert_num INT PRIMARY KEY, net_worth INT);  
CREATE TABLE
```

2.

```
Movies=# CREATE TABLE Movie(title VARCHAR(80), year INT, length INT, genre VARCHAR(80), studio_name VARCHAR(80), producer_num INT REFERENCES MovieExec (cert_num), PRIMARY KEY (title, year));  
CREATE TABLE
```

- You can show the tables!

```
Movies=# \dt  
          List of relations  
 Schema |   Name   | Type  | Owner  
-----+-----+-----+-----  
 public | movie    | table | Work  
 public | movieexec | table | Work  
(2 rows)
```

PostgreSQL

- You can show the columns!

```
Movies=# \d movie;
```

Column	Type	Collation	Nullable	Default
title	character varying(80)		not null	
year	integer		not null	
length	integer			
genre	character varying(80)			
studio_name	character varying(80)			
producer_num	integer			

Indexes:
"movie_pkey" PRIMARY KEY, btree (title, year)

Foreign-key constraints:
"movie_producer_num_fkey" FOREIGN KEY (producer_num) REFERENCES movieexec(cert_num)

```
Movies=# \d movieexec
```

Column	Type	Collation	Nullable	Default
name	character varying(80)			
address	character varying(255)			
cert_num	integer		not null	
net_worth	integer			

Indexes:
"movieexec_pkey" PRIMARY KEY, btree (cert_num)

Referenced by:
TABLE "movie" CONSTRAINT "movie_producer_num_fkey" FOREIGN KEY (producer_num) REFERENCES movieexec(cert_num)

PostgreSQL

1.

```
Movies=# INSERT INTO
Movies=# MovieExec(name, address, cert_num, net_worth)
Movies=# VALUES
Movies=#   ('Gary Kurtz', 'London, England', 12345, 7000000),
Movies=#   ('Mark Johnson', 'Washington DC', 67890, 17500000000),
Movies=#   ('Lorne Michaels', 'New York, NY', 99999, 5000000000);
ERROR:  integer out of range
```

◦ Uh oh!

PostgreSQL

1.

```
Movies=# ALTER TABLE MovieExec ALTER COLUMN net_worth TYPE BIGINT;  
ALTER TABLE
```

2.

```
Movies=# \d movieexec
```

Table "public.movieexec"				
Column	Type	Collation	Nullable	Default
name	character varying(80)			
address	character varying(255)			
cert_num	integer		not null	
net_worth	bigint			

```
Indexes:
```

```
    "movieexec_pkey" PRIMARY KEY, btree (cert_num)
```

```
Referenced by:
```

```
    TABLE "movie" CONSTRAINT "movie_producer_num_fkey" FOREIGN KEY (producer_num) REFERENCES movieexec(cert_num)
```

PostgreSQL

1.

```
Movies=# INSERT INTO
Movies=# MovieExec(name, address, cert_num, net_worth)
Movies=# VALUES
Movies=# ('Gary Kurtz', 'London, England', 12345, 7000000),
Movies=# ('Mark Johnson', 'Washington DC', 67890, 17500000000),
Movies=# ('Lorne Michaels', 'New York, NY', 99999, 5000000000);
INSERT 0 3
```

2.

```
Movies=# SELECT * FROM MovieExec;
   name   | address      | cert_num | net_worth
-----+-----+-----+-----
 Gary Kurtz | London, England |    12345 |    7000000
Mark Johnson | Washington DC  |    67890 | 17500000000
 Lorne Michaels | New York, NY   |   99999 | 5000000000
(3 rows)
```

PostgreSQL

1.

```
Movies=# INSERT INTO
Movies=# Movie(title, year, length, genre, studio_name, producer_num)
Movies=# VALUES
Movies=#      ('Star Wars', 1977, 124, 'SciFi', 'Fox', 12345),
Movies=#      ('Galaxy Quest', 1999, 104, 'comedy', 'DreamWorks', 67890),
Movies=#      ('Wayne''s World', 1992, 95, 'comedy', 'Paramount', 99999);
INSERT 0 3
```

- If we had tried to insert into Movie first

```
Movies=# INSERT INTO
Movies=# Movie(title, year, length, genre, studio_name, producer_num)
Movies=# VALUES
Movies=#      ('Star Wars', 1977, 124, 'SciFi', 'Fox', 12345),
Movies=#      ('Galaxy Quest', 1999, 104, 'comedy', 'DreamWorks', 67890),
Movies=#      ('Wayne''s World', 1992, 95, 'comedy', 'Paramount', 99999);
ERROR:  insert or update on table "movie" violates foreign key constraint "movie_producer_num_fkey"
```

PostgreSQL

- How do you get the name and net worth for producers of movies made since the turn of the century?

1. Join

```
Movies=# SELECT * FROM Movie JOIN MovieExec on producer_num = cert_num;
```

title	year	length	genre	studio_name	producer_num	name	address	cert_num	net_worth
Star Wars	1977	124	SciFi	Fox	12345	Gary Kurtz	London, England	12345	7000000
Galaxy Quest	1999	104	comedy	DreamWorks	67890	Mark Johnson	Washington DC	67890	17500000000
Wayne's World	1992	95	comedy	Paramount	99999	Lorne Michaels	New York, NY	99999	500000000

(3 rows)

- Add data so we can test our query

```
Movies=# INSERT INTO MovieExec VALUES('Paul Webster', 'United Kingdom', 25327, 5000000);
INSERT 0 1
```

```
Movies=# INSERT INTO Movie VALUES('Pride and Prejudice', 2005, 127, 'period', 'Universal', 25327);
INSERT 0 1
```

PostgreSQL

- How do you get the name and net worth for producers of movies made since the turn of the century?

1. Join (combines the table)

```
Movies=# SELECT * FROM Movie JOIN MovieExec on producer_num = cert_num;
```

title	year	length	genre	studio_name	producer_num	name	address	cert_num	net_worth
Star Wars	1977	124	SciFi	Fox	12345	Gary Kurtz	London, England	12345	7000000
Galaxy Quest	1999	104	comedy	DreamWorks	67890	Mark Johnson	Washington DC	67890	17500000000
Wayne's World	1992	95	comedy	Paramount	99999	Lorne Michaels	New York, NY	99999	500000000
Pride and Prejudice	2005	127	period	Universal	25327	Paul Webster	United Kingdom	25327	5000000

(4 rows)

2. Selection (refines the query)

```
Movies=# SELECT * FROM Movie JOIN MovieExec on producer_num = cert_num WHERE year >= 2000;
```

title	year	length	genre	studio_name	producer_num	name	address	cert_num	net_worth
Pride and Prejudice	2005	127	period	Universal	25327	Paul Webster	United Kingdom	25327	5000000

(1 row)

PostgreSQL

- How do you get the name and net worth for producers of movies made since the turn of the century?

2. Selection (refines the query)

```
Movies=# SELECT * FROM Movie JOIN MovieExec on producer_num = cert_num WHERE year >= 2000;
```

title	year	length	genre	studio_name	producer_num	name	address	cert_num	net_worth
Pride and Prejudice	2005	127	period	Universal	25327	Paul Webster	United Kingdom	25327	5000000

(1 row)

3. Projection (isolate the columns)...

```
Movies=# SELECT name, net_worth FROM Movie JOIN MovieExec ON producer_num = cert_num WHERE year >= 2000;
```

name	net_worth
Paul Webster	5000000

(1 row)

PostgreSQL

- How do you get the name and net worth for producers of movies made since the turn of the century?
- Join + Selection + Projection

```
Movies=# SELECT name, net_worth FROM Movie JOIN MovieExec ON producer_num = cert_num WHERE year >= 2000;
 name      | net_worth 
-----+-----
 Paul Webster | 50000000
(1 row)
```

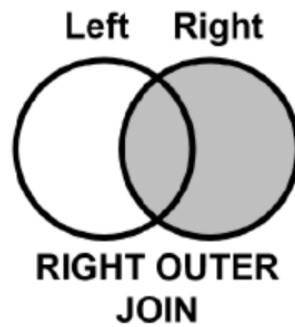
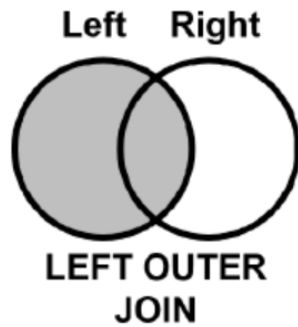
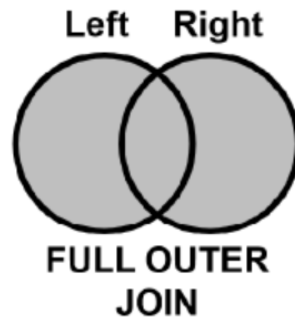
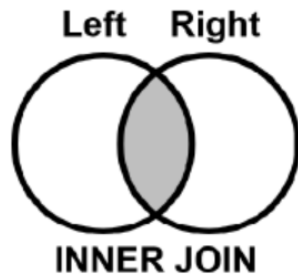
```
SELECT
    name,
    net_worth
FROM Movie JOIN MovieExec ON producer_num = cert_num
WHERE year >= 2000;
```


PostgreSQL

- How do you get the name and net worth for producers of movies made since the turn of the century?
- Join + Selection + Projection (Aliasing)

```
Movies=# SELECT me.name, me.net_worth FROM Movie AS m JOIN MovieExec AS me ON m.producer_num = me.cert_num WHERE m.year >= 2000;
 name      | net_worth 
-----+-----
 Paul Webster | 50000000
(1 row)
```

```
SELECT
    me.name,
    me.net_worth
FROM Movie AS m JOIN MovieExec AS me ON m.producer_num = me.cert_num
WHERE m.year >= 2000;
```






TABLE 1

TABLE 2

RESULT OF INNER JOIN




TABLE 1

TABLE 2

RESULT OF OUTER JOIN




TABLE 1

TABLE 2

RESULT OF LEFT JOIN




TABLE 1

TABLE 2

RESULT OF RIGHT JOIN

?


Example 2.21: Consider the two relations from our running movie database:

`Movies(title, year, length, genre, studioName, producerC#)`

`MovieExec(name, address, cert#, netWorth)`

QUESTION 1

- Which joins produce the same results for these tables?



```
Movies(title, year, length, genre, studioName, producerC#)
StarsIn(movieTitle, movieYear, starName)
MovieExec(name, address, cert#, netWorth)
```

Joins

ON StarsIn.starName = MovieExec.name

- **JOIN (INNER JOIN)**
 - Only results where they match
 - Values from each table in each tuple
- **OUTER JOIN**
 - Every tuple from each table
 - If actor doesn't produce, no values in MovieExec cols
 - If producer doesn't act, no values in StarsIn cols
- **LEFT OUTER JOIN**
 - Every tuple from StarsIn table
 - If actor doesn't produce, no values in MovieExec cols
 - If producer doesn't act, no values in StarsIn cols
- **RIGHT OUTER JOIN**
 - Every tuple from MovieExec table
 - If producer doesn't act, no values in StarsIn cols

Query Activity

SQL Query Practice

Example 2.21: Consider the two relations from our running movie database:

`Movies(title, year, length, genre, studioName, producerC#)`

`MovieExec(name, address, cert#, netWorth)`

How do you get the title and year for the movie produced by the producer with the highest net worth?

PostgreSQL

1. `SELECT * FROM Movie JOIN MovieExec ON producer_num = cert_num;`
2. `SELECT * FROM Movie JOIN MovieExec ON producer_num = cert_num
ORDER BY net_worth;`
3. `SELECT * FROM Movie JOIN MovieExec ON producer_num = cert_num
ORDER BY net_worth DESC;`
4. `SELECT * FROM Movie JOIN MovieExec ON producer_num = cert_num
ORDER BY net_worth DESC LIMIT 1;`
5. **`SELECT title, year FROM Movie JOIN MovieExec ON producer_num =
cert_num ORDER BY net_worth DESC LIMIT 1;`**

Subqueries

1. can return a single constant to be compared with another value in a **WHERE** clause
2. can return relations that can be used in various ways in **WHERE** clauses
3. can appear in **FROM** clauses, followed by a tuple variable that represents the tuples in the result of the subquery

Subqueries

1. can return a single constant to be compared with another value in a **WHERE** clause

```
1)  SELECT name
2)  FROM MovieExec
3)  WHERE cert# =
4)      (SELECT producerC#
5)      FROM Movies
6)      WHERE title = 'Star Wars'
      );
```

```
SELECT name
FROM Movie JOIN MovieExec ON producer_num = cert_num
WHERE title = 'Star Wars';
```

Subqueries

2. can return relations that can be used in various ways in **WHERE** clauses

Finding the producers of Harrison Ford's movies

```
Movies(title, year, length, genre, studioName, producerC#)  
StarsIn(movieTitle, movieYear, starName)  
MovieExec(name, address, cert#, netWorth)
```

Subqueries

2. can return relations that can be used in various ways in **WHERE** clauses

```
1)  SELECT name
2)  FROM MovieExec
3)  WHERE cert# IN
4)      (SELECT producerC#
5)      FROM Movies
6)      WHERE (title, year) IN
7)          (SELECT movieTitle, movieYear
8)          FROM StarsIn
9)          WHERE starName = 'Harrison Ford'
      )
    );
```

Subqueries

2. can return relations that can be used in various ways in **WHERE** clauses

```
SELECT name
FROM MovieExec JOIN Movie ON cert_num = producer_num
JOIN StarsIn ON title = movie_title AND year =
movie_year
WHERE star_name = 'Harrison Ford';
```

Subqueries

3. can appear in **FROM** clauses, followed by a tuple variable that represents the tuples in the result of the subquery

```
1)  SELECT name
2)  FROM MovieExec, (SELECT producerC#
3)                        FROM Movies, StarsIn
4)                        WHERE title = movieTitle AND
5)                        year = movieYear AND
6)                        starName = 'Harrison Ford'
7)                        ) Prod
8)  WHERE cert# = Prod.producerC#;
```

Figure 6.11: Finding the producers of Ford's movies using a subquery in the FROM clause

Practice

Let's revisit our schema for a College Ranking System.

Next Class

Module:

Week 8: Midterm

Topic:

Midterm Review

