

# Analysis of Algorithms

## CS 477/677


---

Instructor: Monica Nicolescu

Lecture 13

# Midterm Results - CS 477

---




Min = 23

Max = 110

Average = 78.2

# Midterm Results - CS 677

---



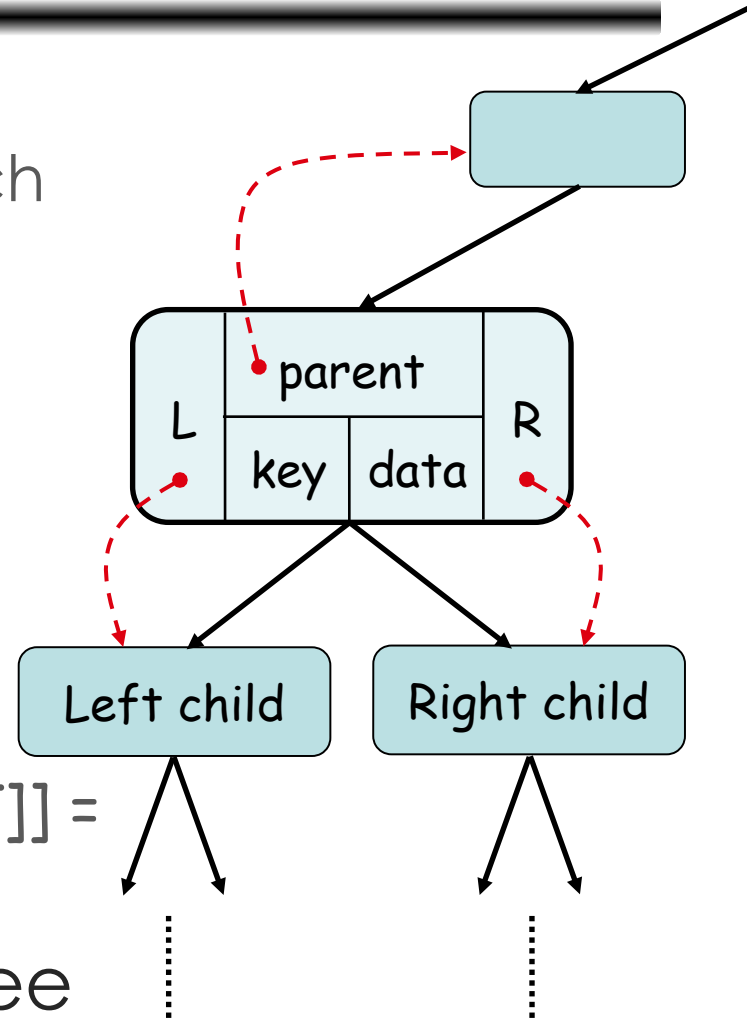
Min = 57

Max = 111

Average = 89.26

# Binary Search Trees


- Tree representation:
  - A linked data structure in which each node is an object
- Node representation:
  - **Key** field
  - Satellite data
  - **Left**: pointer to left child
  - **Right**: pointer to right child
  - **p**: pointer to parent ( $p[\text{root}[T]] = \text{NIL}$ )
- Satisfies the binary search tree property



# Binary Search Tree Example

---

- Binary search tree property:
  - If  $y$  is in left subtree of  $x$ ,  
then  $\text{key}[y] \leq \text{key}[x]$
  - If  $y$  is in right subtree of  $x$ ,  
then  $\text{key}[y] \geq \text{key}[x]$



# Binary Search Trees

---

- Support many dynamic set operations
  - SEARCH, MINIMUM, MAXIMUM, PREDECESSOR, SUCCESSOR, INSERT, DELETE
- Running time of basic operations on binary search trees
  - On average:  $\Theta(\lg n)$ 
    - The expected height of the tree is  $\lg n$
  - In the worst case:  $\Theta(n)$ 
    - The tree is a linear chain of  $n$  nodes

# Red-Black Trees

---

- “Balanced” binary trees guarantee an  $O(\lg n)$  running time on the basic dynamic-set operations
- Red-black tree
  - Binary tree with an additional attribute for its nodes: **color** which can be **red** or **black**
  - Constrains the way nodes can be colored on any path from the root to a leaf
    - Ensures that no path is more than twice as long as another  $\Rightarrow$  the tree is balanced
  - The nodes inherit all the other attributes from the binary-search trees: **key**, **left**, **right**, **p**

# Red-Black Trees Properties


---

1. Every **node** is either **red** or **black**
2. The **root** is **black**
3. Every **leaf (NIL)** is **black**
4. If a node is red, then both its children are black
  - No two red nodes in a row on a simple path from the root to a leaf
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes




# Example: RED-BLACK TREE

---



- For convenience we use a sentinel **NIL[T]** to represent all the **NIL** nodes at the leafs
  - **NIL[T]** has the same fields as an ordinary node
  - **Color[NIL[T]] = BLACK**
  - The other fields may be set to arbitrary values

# Black-Height of a Node



- **Height of a node:** the number of edges in a longest path to a leaf
- **Black-height** of a node  $x$ :  $bh(x)$  is the number of black nodes (including NIL) on a path from  $x$  to leaf, not counting  $x$


# Properties of Red-Black Trees

- **Claim**

- Any node with height  $h$  has black-height  $\geq h/2$

- **Proof**

- By property 4, there are at most  $h/2$  **red** nodes on the path from the node to a leaf
- Hence at least  $h/2$  are **black**



Property 4: if a node is **red** then both its children are black

# Properties of Red-Black Trees

---

## Claim

The subtree rooted at any node  $x$  contains at least  $2^{bh(x)} - 1$  internal nodes


**Proof:** By induction on height of  $x$

**Basis:**  $height[x] = 0 \Rightarrow$

$x$  is a leaf ( $NIL[T]$ )  $\Rightarrow$

$bh(x) = 0 \Rightarrow$

# of internal nodes:  $2^0 - 1 = 0$




# Properties of Red-Black Trees

---


## Inductive step:

- Let  $\text{height}(x) = h$  and  $\text{bh}(x) = b$
- Any child  $y$  of  $x$  has:
  - $\text{bh}(y) = b$  (if the child is **red**), or
  - $\text{bh}(y) = b - 1$  (if the child is **black**)



# Properties of Red-Black Trees

- Want to prove:
  - The subtree rooted at any node  $x$  contains at least  $2^{bh(x)} - 1$  internal nodes
- Assume true for children of  $x$ :
  - Their subtrees contain at least  $2^{bh(x) - 1} - 1$  internal nodes
- The subtree rooted at  $x$  contains at least:  
$$(2^{bh(x) - 1} - 1) + (2^{bh(x) - 1} - 1) + 1 =$$
$$2 \cdot (2^{bh(x) - 1} - 1) + 1 =$$
$$2^{bh(x)} - 1 \text{ internal nodes}$$




# Properties of Red-Black Trees

**Lemma:** A red-black tree with  $n$  internal nodes has height at most  $2\lg(n + 1)$ .

**Proof:**

$$\begin{array}{lcl} n & \geq 2^b - 1 & \geq 2^{h/2} - 1 \\ \text{number of internal nodes} & & \text{since } b \geq h/2 \end{array}$$



- Add 1 to all sides and then take logs:

$$\begin{aligned} n + 1 &\geq 2^b \geq 2^{h/2} \\ \lg(n + 1) &\geq h/2 \Rightarrow \\ h &\leq 2 \lg(n + 1) \end{aligned}$$

# Operations on Red-Black Trees

---


- The non-modifying binary-search tree operations **MINIMUM**, **MAXIMUM**, **SUCCESSOR**, **PREDECESSOR**, and **SEARCH** run in  $O(h)$  time
  - They take  $O(\lg n)$  time on red-black trees
- What about TREE-INSERT and TREE-DELETE?
  - They will still run in  $O(\lg n)$
  - We have to guarantee that the modified tree will still be a red-black tree




# INSERT

INSERT: what color to make the new node?

- Red? Let's insert 35!
  - Property 4: if a node is red, then both its children are black
- Black? Let's insert 14!
  - Property 5: all paths from a node to its leaves contain the same number of black nodes




# DELETE



DELETE: what color was the node that was removed? **Red?**

1. Every **node** is either **red** or **black** OK!
2. The **root** is **black** OK!
3. Every **leaf (NIL)** is **black** OK!
4. If a node is red, then both its children are black  
OK! Does not change any black heights  
OK! Does not create two red nodes in a row
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes

# DELETE



DELETE: what color was the node that was removed? **Black?**

1. Every node is either **red** or **black** OK!
2. The root is **black** **Not OK!** If removing the root and the child that replaces it is **red**
3. Every leaf (NIL) is **black** OK!
4. If a node is red, then both its children are black **Not OK!** Could create two red nodes in a row
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes **Not OK!** Could change the black heights of some nodes


# Rotations

---

- Operations for restructuring the tree after insert and delete operations on red-black trees
- Rotations take a red-black tree and a node within the tree and:
  - Together with some node re-coloring they help restore the red-black tree property
  - Change some of the pointer structure
  - Do not change the binary-search tree property
- Two types of rotations:
  - Left & right rotations

# Left Rotations



- Assumption for a left rotation on a node  $x$ :
  - The right child of  $x$  ( $y$ ) is not NIL



- Idea:
  - Pivots around the link from  $x$  to  $y$
  - Makes  $y$  the new root of the subtree
  - $x$  becomes  $y$ 's left child
  - $y$ 's left child becomes  $x$ 's right child


# Example: LEFT-ROTATE

---




# LEFT-ROTATE( $T, x$ )

1.  $y \leftarrow \text{right}[x]$       ► Set  $y$
2.  $\text{right}[x] \leftarrow \text{left}[y]$       ►  $y$ 's left subtree becomes  $x$ 's right subtree
3. **if**  $\text{left}[y] \neq \text{NIL}$
4.     **then**  $p[\text{left}[y]] \leftarrow x$       ► Set the parent relation from  $\text{left}[y]$  to  $x$
5.  $p[y] \leftarrow p[x]$       ► The parent of  $x$  becomes the parent of  $y$
6. **if**  $p[x] = \text{NIL}$
7.     **then**  $\text{root}[T] \leftarrow y$
8.     **else if**  $x = \text{left}[p[x]]$
9.         **then**  $\text{left}[p[x]] \leftarrow y$
10.        **else**  $\text{right}[p[x]] \leftarrow y$
11.  $\text{left}[y] \leftarrow x$       ► Put  $x$  on  $y$ 's left
12.  $p[x] \leftarrow y$       ►  $y$  becomes  $x$ 's parent




LEFT-ROTATE( $T, x$ )  
 .....>>>



# Right Rotations

- Assumption for a right rotation on a node  $x$ :
  - The left child of  $y$  ( $x$ ) is not NIL



- Idea:
  - Pivots around the link from  $y$  to  $x$
  - Makes  $x$  the new root of the subtree
  - $y$  becomes  $x$ 's right child
  - $x$ 's right child becomes  $y$ 's left child




# Insertion

---


- Goal:
  - Insert a new node  $z$  into a red-black tree
- Idea:
  - Insert node  $z$  into the tree as for an ordinary binary search tree
  - Color the node **red**
  - Restore the red-black tree properties
    - Use an auxiliary procedure RB-INSERT-FIXUP

# RB-INSERT( $T, z$ )



1.  $y \leftarrow \text{NIL}$
  2.  $x \leftarrow \text{root}[T]$
  3. **while**  $x \neq \text{NIL}$
  4.     **do**  $y \leftarrow x$
  5.         **if**  $\text{key}[z] < \text{key}[x]$
  6.             **then**  $x \leftarrow \text{left}[x]$
  7.             **else**  $x \leftarrow$   
               $\text{right}[x]$
  8.  $p[z] \leftarrow y$
- Initialize nodes  $x$  and  $y$
  - Throughout the algorithm  $y$  points to the parent of  $x$
  - Go down the tree until reaching a leaf
  - At that point  $y$  is the parent of the node to be inserted
  - Sets the parent of  $z$  to be  $y$

# RB-INSERT( $T, z$ )



9. if  $y = \text{NIL}$

10. then  $\text{root}[T] \leftarrow z$

The tree was empty:  
set the new node to be the root

11. else if  $\text{key}[z] < \text{key}[y]$

12. then  $\text{left}[y] \leftarrow z$

13. else  $\text{right}[y] \leftarrow z$

Otherwise, set  $z$  to be the left  
or right child of  $y$ , depending  
on whether the inserted node  
is smaller or larger than  $y$ 's key

14.  $\text{left}[z] \leftarrow \text{NIL}$

15.  $\text{right}[z] \leftarrow \text{NIL}$

16.  $\text{color}[z] \leftarrow \text{RED}$

Set the fields of the newly added node

17.  $\text{RB-INSERT-FIXUP}(T, z)$

Fix any inconsistencies that could have  
been introduced by adding this new red  
node


# RB Properties Affected by Insert

1. Every **node** is either **red** or **black** OK!
2. The **root** is **black** If  $z$  is the root  $\Rightarrow$  **not OK**
3. Every **leaf** (NIL) is **black** OK!
4. If a node is red, then both its children are black

If  $p(z)$  is red  $\Rightarrow$  **not OK**  
 $z$  and  $p(z)$  are both red

OK!

5. For each node, all paths from the node to descendant leaves contain the same number of black nodes






# RB-INSERT-FIXUP – Case 1

$z$ 's “uncle” ( $y$ ) is **red**

**Idea:** ( $z$  is a right child)

- $p[p[z]]$  ( $z$ 's grandparent) must be black:  $z$  and  $p[z]$  are both red
- Color  $p[z]$  **black**
- Color  $y$  **black**
- Color  $p[p[z]]$  **red**
  - Push the **red** node up the tree
- Make  $z = p[p[z]]$



# RB-INSERT-FIXUP – Case 1




$z$ 's "uncle" ( $y$ ) is **red**

**Idea:** ( $z$  is a left child)

- $p[p[z]]$  ( $z$ 's grandparent) must be black:  $z$  and  $p[z]$  are both red

- $\text{color}[p[z]] \leftarrow \text{black}$
- $\text{color}[y] \leftarrow \text{black}$
- $\text{color}[p[p[z]]] \leftarrow \text{red}$
- $z = p[p[z]]$  Case 1

– Push the **red** node up the tree




# RB-INSERT-FIXUP – Case 3

Case 3:

- $z$ 's “uncle” ( $y$ ) is **black**
- $z$  is a left child

Idea:

- $\text{color}[p[z]] \leftarrow \text{black}$
- $\text{color}[p[p[z]]] \leftarrow \text{red}$
- $\text{RIGHT-ROTATE}(T, p[p[z]])$  Case3
- No longer have 2 reds in a row
- $p[z]$  is now black



# RB-INSERT-FIXUP – Case 2



Case 2:

- $z$ 's “uncle” ( $y$ ) is **black**
- $z$  is a right child

Idea:

- $z \leftarrow p[z]$
- $\text{LEFT-ROTATE}(\mathcal{T}, z)$  Case2

$\Rightarrow$  now  $z$  is a left child, and both  $z$  and  $p[z]$  are red  $\Rightarrow$  case 3






# RB-INSERT-FIXUP( $T, z$ )



1. **while**  $\text{color}[p[z]] = \text{RED}$  ← The while loop repeats only when case1 is executed:  $O(\lg n)$  times
2.     **do if**  $p[z] = \text{left}[p[p[z]]]$  }
3.     **then**  $y \leftarrow \text{right}[p[p[z]]]$  } Set the value of x's "uncle"
4.         **if**  $\text{color}[y] = \text{RED}$
5.             **then Case1**
6.             **else if**  $z = \text{right}[p[z]]$
7.                 **then Case2**
8.                 **Case3**
9.         **else** (same as **then** clause with "right" and "left" exchanged)
10.  $\text{color}[\text{root}[T]] \leftarrow \text{BLACK}$  ← We just inserted the root, or the red node reached the root

# Example



Insert 4





Case 1



Case 2



Case 3



# Analysis of RB-INSERT

---

- Inserting the new element into the tree  
 $O(\lg n)$
- RB-INSERT-FIXUP
  - The while loop repeats only if CASE 1 is executed
  - The number of times the while loop can be executed is  $O(\lg n)$
- Total running time of RB-INSERT:  $O(\lg n)$


# Red-Black Trees - Summary

---

- Operations on red-black trees:
  - SEARCH  $O(h)$
  - PREDECESSOR  $O(h)$
  - SUCCESSOR  $O(h)$
  - MINIMUM  $O(h)$
  - MAXIMUM  $O(h)$
  - INSERT  $O(h)$
  - DELETE  $O(h)$
- Red-black trees guarantee that the height of the tree will be  $O(\lg n)$

# Readings

---



- For this lecture
  - Sections 6.3, 6,5
  - Chapter 13
- Coming next
  - Chapter 17