


Analysis of Algorithms

CS 477/677

Instructor: Monica Nicolescu

Lecture 21

Midterm 2 Results – CS 477




min = 28

max = 110

average = 79.81

Midterm 2 Results – CS 677



min = 40

max = 118

average = 87.56

Midterm 2 - Feedback

- Heap elements do not have pointers to children
- In RBTs nodes are inserted as red
- OS_Select finds the i -th order statistic not element with $key = i$
- When working with trees, always check if the tree root is NULL first, before accessing $T.left$ or $T.right$
- Cannot add pointers
- The height of a node is the maximum of the left/right subtree heights + 1, not their sum
- Recursive algorithms:
 - Pay attention to the return value for the base case
 - Make sure the return value is always the same type
 - Use the return values

Huffman Codes

- Idea:
 - Use the frequencies of occurrence of characters to build a optimal way of representing each character

	a	b	c	d	e	f
Frequency (thousands)	45	13	12	16	9	5

Constructing a Huffman Code

- Let's build a greedy algorithm that constructs an optimal prefix code (called a **Huffman code**)
- Assume that:
 - \mathcal{C} is a set of n characters
 - Each character has a frequency $f(c)$
- Idea:





f: 5	e: 9	c: 12	b: 13	d: 16	a: 45
------	------	-------	-------	-------	-------

 - The tree T is built in a bottom up manner
 - Start with a set of $|\mathcal{C}| = n$ leaves
 - At each step, merge the two least frequent objects: the frequency of the new node = sum of two frequencies
 - Use a min-priority queue Q , keyed on f to identify the two least frequent objects

Example

f: 5 e: 9 c: 12 b: 13 d: 16 a: 45

c: 12 b: 13 d: 16 a: 45



Building a Huffman Code

Alg.: HUFFMAN(\mathcal{C})

Running time: $O(n \lg n)$

1. $n \leftarrow |\mathcal{C}|$
 2. $Q \leftarrow \mathcal{C}$ $\longleftarrow O(n)$
 3. **for** $i \leftarrow 1$ **to** $n - 1$
 4. **do** allocate a new node z
 5. $\text{left}[z] \leftarrow x \leftarrow \text{EXTRACT-MIN}(Q)$
 6. $\text{right}[z] \leftarrow y \leftarrow \text{EXTRACT-MIN}(Q)$
 7. $f[z] \leftarrow f[x] + f[y]$
 8. INSERT (Q, z)
 9. **return** EXTRACT-MIN(Q)
- } $O(n \lg n)$

Greedy Choice Property


Let \mathcal{C} be an alphabet in which each character $c \in \mathcal{C}$ has frequency $f[c]$. Let x and y be two characters in \mathcal{C} having the lowest frequencies.

Then, there exists an optimal prefix code for \mathcal{C} in which the codewords for x and y have the same (maximum) length and differ only in the last bit.

Proof of the Greedy Choice


- Idea:
 - Consider a tree T representing an arbitrary optimal prefix code
 - Modify T to make a tree representing another optimal prefix code in which x and y will appear as sibling leaves of maximum depth
- ⇒ The codes of x and y will have the same length and differ only in the last bit

Proof of the Greedy Choice (cont.)



- **a**, **b** – two characters, sibling leaves of max. depth in T
- Assume: $f[a] \leq f[b]$ and $f[x] \leq f[y]$
- $f[x]$ and $f[y]$ are the two lowest leaf frequencies, in order
 $\Rightarrow f[x] \leq f[a]$ and $f[y] \leq f[b]$
- Exchange the positions of **a** and **x** (T') and of **b** and **y** (T'')

Proof of the Greedy Choice (cont.)




$$\begin{aligned}
 B(T) - B(T') &= \sum_{c \in C} f(c)d_T(c) - \sum_{c \in C} f(c)d_{T'}(c) \\
 &= f[x]d_T(x) + f[a]d_T(a) - f[x]d_{T'}(x) - f[a]d_{T'}(a) \\
 &= f[x]d_T(x) + f[a]d_T(a) - f[x]d_T(a) - f[a]d_T(x) \\
 &= \underbrace{(f[a] - f[x])}_{\geq 0} \underbrace{(d_T(a) - d_T(x))}_{\geq 0}
 \end{aligned}$$

x is a minimum frequency leaf a is a leaf of maximum depth

≥ 0

Proof of the Greedy Choice (cont.)



$$B(T) - B(T') \geq 0$$

Similarly, exchanging **y** and **b** does not increase the cost:

$$B(T') - B(T'') \geq 0$$

$\Rightarrow B(T'') \leq B(T)$. Also, since **T** is optimal, $B(T) \leq B(T'')$


Therefore, $B(T) = B(T'') \Rightarrow T''$ is an optimal tree, in which **x** and **y** are sibling leaves of maximum depth

Discussion

- Greedy choice property:
 - Building an optimal tree by mergers can begin with the greedy choice: merging the two characters with the lowest frequencies
 - The cost of each merger is the sum of frequencies of the two items being merged
 - Of all possible mergers, HUFFMAN chooses the one that incurs the least cost


Interval Partitioning

- Lecture j starts at s_j and finishes at f_j
- Goal: find minimum number of classrooms to schedule all lectures so that no two occur at the same time in the same room
 - Ex: this schedule uses 4 classrooms to schedule 10 lectures



Interval Partitioning

- Lecture j starts at s_j and finishes at f_j
- Goal: find minimum number of classrooms to schedule all lectures so that no two occur at the same time in the same room
 - Ex: this schedule uses only 3




Interval Partitioning: Lower Bound on Optimal Solution

- The **depth** of a set of open intervals is the maximum number that contain any given time
- Key observation:
 - The number of classrooms needed \geq depth
- Ex: Depth of schedule below = 3 \Rightarrow schedule below is optimal

↑
a, b, c all contain 9:30

- Does there always exist a schedule equal to depth of intervals?




Greedy Strategy

- Consider lectures in increasing order of start time: assign lecture to any compatible classroom
 - Labels set $\{1, 2, 3, \dots, d\}$, where d is the depth of the set of intervals
 - Overlapping intervals are given different labels
 - Assign a label that has not been assigned to any previous interval that overlaps it

Greedy Algorithm

1. Sort intervals by start times, such that $s_1 \leq s_2 \leq \dots \leq s_n$
(let I_1, I_2, \dots, I_n denote the intervals in this order)
2. **for** $j = 1$ **to** n
3. Exclude from set $\{1, 2, \dots, d\}$ the labels of preceding and overlapping intervals I_i from consideration for I_j
4. **if** there is any label from $\{1, 2, \dots, d\}$ that was not excluded
 assign that label to I_j
5. **else**
6. leave I_j unlabeled

Example



Claim

- Every interval will be assigned a label
 - For interval I_j , assume there are t intervals earlier in the sorted order that overlap it
 - We have $t + 1$ intervals that pass over a common point on the timeline
 - $t + 1 \leq d$, thus $t \leq d - 1$
 - At least one of the d labels is not excluded by this set of t intervals, which we can assign to I_j

Claim

- No two overlapping intervals are assigned the same label
 - Consider I and I' that overlap, and I precedes I' in the sorted order
 - When I' is considered, the label for I is excluded from consideration
 - Thus, the algorithm will assign a different label to I


Greedy Choice Property

- The greedy algorithm schedules every interval on a resource, using a number of resources equal to the depth of the set of intervals. This is the optimal number of resources needed.
- Proof:
 - Follows from previous claims
- Structural proof
 - Discover a simple “structural” bound asserting that every possible solution must have a certain value
 - Then show that your algorithm always achieves this bound

- Example:

24

Readings



- For this lecture
 - Chapter 15
- Coming next
 - Chapter 15