

Analysis of Algorithms

CS 477/677

Instructor: Monica Nicolescu

Lecture 22


Scheduling to Minimizing Lateness

- Single resource processes one job at a time
- Job j requires t_j units of processing time, is due at time d_j
- If j starts at time s_j , it finishes at time $f_j = s_j + t_j$
- Lateness: $\ell_j = \max \{ 0, f_j - d_j \}$

- Goal: schedule all jobs to minimize **maximum** lateness $L = \max \ell_j$:

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15

- Example:



Greedy Algorithms

- Greedy strategy: consider jobs in some order
 - **[Shortest processing time first]** Consider jobs in ascending order of processing time t_j

counterexample

	1	2
t_j	1	10
d_j	100	10

Choosing t_1 first: $l_2 = 1$

Choosing t_2 first: $l_2 = l_1 = 0$

- **[Smallest slack]** Consider jobs in ascending order of slack $d_j - t_j$

counterexample

	1	2
t_j	1	10
d_j	2	10

Choosing t_2 first: $l_1 = 9$

Choosing t_1 first: $l_1 = 0$ and $l_2 = 1$

Greedy Algorithm

- Greedy choice: earliest deadline first

Sort n jobs by deadline so that $d_1 < d_2 < \dots < d_n$

$t = 0$


for $j = 1$ to n

 Assign job j to interval $[t, t + t_j]$

$s_j = t, f_j = t + t_j$


$t = t + t_j$

output intervals $[s_j, f_j]$



Minimizing Lateness: No Idle Time

- Observation: The greedy schedule has no idle time
- Observation: There exists an optimal schedule with no **idle time**




Minimizing Lateness: Inversions

- An **inversion** in schedule S is a pair of jobs i and j such that: $d_i < d_j$ but j scheduled before i



- Observation: greedy schedule has no inversions

Greedy Choice Property



- Optimal solution: $d_i < d_j$ but j scheduled before i
- Greedy solution: i scheduled before j
 - Job i finishes sooner, no increase in latency

$$\text{Lateness}(\text{Job } j)_{\text{GREEDY}} = f_i - d_j$$

\leq

➔ No increase in latency

$$\text{Lateness}(\text{Job } i)_{\text{OPT}} = f_i - d_i$$

Greedy Analysis Strategies

- Exchange argument
 - Gradually transform any solution to the one found by the greedy algorithm without hurting its quality
- Structural
 - Discover a simple “structural” bound asserting that every possible solution must have a certain value, then show that your algorithm always achieves this bound
- Greedy algorithm stays ahead
 - Show that after each step of the greedy algorithm, its solution is at least as good as any other algorithm's

Coin Changing

- Given currency denominations: 1, 5, 10, 25, 100, devise a method to pay amount to customer using fewest number of coins

- Ex: 34¢



- Ex: \$2.89



Greedy Algorithm

- Greedy strategy: at each iteration, add coin of the largest value that does not take us past the amount to be paid

Sort coins denominations by value: $c_1 < c_2 < \dots < c_n$.

↙ coins selected

```
S = {}  
while (x > 0) {  
    let k be largest integer such that  $c_k \leq x$   
    if (k = 0)  
        return "no solution found"  
    x = x -  $c_k$   
    S = S U {k}  
}  
return S
```

Greedy Choice Property

- Algorithm is optimal for U.S. coinage: 1, 5, 10, 25, 100

$$\text{Change} = D * 100 + Q * 25 + D * 10 + N * 5 + P$$

- Consider optimal way to change $c_k \leq x < c_{k+1}$: greedy takes coin k
- We claim that any optimal solution must also take coin k
- If not, it needs enough coins of type c_1, \dots, c_{k-1} to add up to x
- Problem reduces to coin-changing $x - c_k$ cents, which, by induction, is optimally solved by greedy algorithm

Greedy Choice Property

- Algorithm is optimal for U.S. coinage: 1, 5, 10, 25, 100

$$\text{Change} = DI * 100 + Q * 25 + D * 10 + N * 5 + P$$

- Optimal solution: $DI \quad Q \quad D \quad N \quad P$
- Greedy solution: $DI' \quad Q' \quad D' \quad N' \quad P'$

1. $\text{Value} < 5$

- Both optimal and greedy use the same # of coins

2. $10 (D) > \text{Value} > 5 (N)$

- Greedy uses one N and then pennies after that
- If OPT does not use N, then it should use pennies for the entire amount \Rightarrow could replace 5 P for 1 N

Greedy Choice Property

$$\text{Change} = DI * 100 + Q * 25 + D * 10 + N * 5 + P$$

- Optimal solution: $DI \quad Q \quad D \quad N \quad P$
- Greedy solution: $DI' \quad Q' \quad D' \quad N' \quad P'$

3. $25 (Q) > \text{Value} > 10 (D)$

- Greedy uses dimes (D's)
- If OPT does not use D's, it needs to use either 2 coins (2 N), or 6 coins (1 N and 5 P) or 10 coins (10 P) to cover 10 cents
- Could replace those with 1 D for a better solution

Greedy Choice Property

$$\text{Change} = DI * 100 + Q * 25 + D * 10 + N * 5 + P$$

- Optimal solution: $DI \quad Q \quad D \quad N \quad P$
- Greedy solution: $DI' \quad Q' \quad D' \quad N' \quad P'$

4. $100 (DI) > \text{Value} > 25 (Q)$

- Greedy picks at least one quarter (Q), OPT does not
- If OPT has no Ds: take all the Ns and Ps and replace 25 cents into one quarter (Q)
- If OPT has 2 or fewer dimes: it uses at least 3 coins to cover one quarter, so we can replace 25 cents with 1 Q
- If OPT has 3 or more dimes (e.g., 40 cents: with 4 Ds): take the first 3 Ds and replace them with 1 Q and 1 N


Coin-Changing

US Postal Denominations

- Observation: greedy algorithm is sub-optimal for US postal denominations:
 - \$.01, .02, .03, .04, .05, .10, .20, .32, .40, .44, .50, .64, .65, .75, .79, .80, .85, .98
 - \$1, \$1.05, \$2, \$4.95, \$5, \$5.15, \$18.30, \$18.95
- Counterexample: 160¢
 - Greedy: 105, 50, 5
 - Optimal: 80, 80

Selecting Breakpoints

- Road trip from Princeton to Palo Alto along fixed route
- Refueling stations at certain points along the way (red marks)
- Fuel capacity = C
- Goal:
 - makes as few refueling stops as possible
- Greedy strategy:
 - go as far as you can before refueling



Greedy Algorithm

Sort breakpoints so that: $0 = b_0 < b_1 < b_2 < \dots < b_n = L$


$S = \{0\}$ \longleftarrow breakpoints selected
 $x = 0$ \longleftarrow current location

```
while (x < b_n)
    let p be largest integer such that b_p ≤ x + C
    if (b_p = x)
        return "no solution"
    x = b_p
    S = S ∪ {p}
return S
```

- Implementation: $O(n \log n)$
 - Use binary search to select each breakpoint p

Greedy Choice Property

- Let $0 = g_0 < g_1 < \dots < g_p = L$ denote set of breakpoints chosen by the greedy
- Let $0 = f_0 < f_1 < \dots < f_q = L$ denote set of breakpoints in an optimal solution with $f_0 = g_0, f_1 = g_1, \dots, f_r = g_r$
- Note: $g_{r+1} > f_{r+1}$ by greedy choice of algorithm



Problem – Buying Licenses

- Your company needs to buy licenses for n pieces of software
- Licenses can be bought only one per month
- Each license currently sells for \$100, but becomes more expensive each month
 - The price increases by a factor $r_j > 1$ each month
 - License j will cost $100 * r_j^t$ if bought t months from now
 - $r_i < r_j$ for license $i < j$
- In which order should the company buy the licenses, to minimize the amount of money spent?


Solution

- Greedy choice:
 - Buy licenses in decreasing order of rate r_j
 - $r_1 > r_2 > r_3 \dots$
- Proof of greedy choice property
 - Optimal solution: r_i r_j $r_i < r_j$
 - Greedy solution: r_j r_i
 - Cost by optimal solution: $100 * r_i^t + 100 * r_j^{t+1}$
 - Cost by greedy solution: $100 * r_j^t + 100 * r_i^{t+1}$
 - $CG - CO = 100 * (r_j^t + r_i^{t+1} - r_i^t - r_j^{t+1}) < 0$
 - $r_i^{t+1} - r_i^t < r_j^{t+1} - r_j^t$
 - $r_i^t(r_i - 1) < r_j^t(r_j - 1)$


OK! (because $r_i < r_j$)

Graphs


- Applications that involve not only a set of items, but also the connections between them




Maps




Schedules



Computer networks



Hypertext




Circuits

Graphs - Background


Graphs = a set of nodes (vertices) with edges (links) between them.

Notations:


- $G = (V, E)$ - graph
- V = set of vertices (size of $V = n$)
- E = set of edges (size of $E = m$)



Directed
graph




Undirected
graph



Acyclic
graph


Other Types of Graphs

- A graph is **connected** if there is a path between every two vertices




Connected Not connected

- A **bipartite graph** is an undirected graph $G = (V, E)$ in which $V = V_1 + V_2$ and there are edges only between vertices in V_1 and V_2




Graph Representation

- **Adjacency list representation** of $G = (V, E)$
 - An array of n lists, one for each vertex in V
 - Each list $\text{Adj}[u]$ contains all the vertices v such that there is an edge between u and v
 - $\text{Adj}[u]$ contains the vertices adjacent to u (in arbitrary order)
 - Can be used for both directed and undirected graphs




Undirected graph




Properties of Adjacency List Representation

- Sum of the lengths of all the adjacency lists
 - Directed graph: size of E (m)
 - Edge (u, v) appears only once in u 's list
 - Undirected graph: $2 \times$ size of E ($2E$)
 - u and v appear in each other's adjacency lists: edge (u, v) appears twice




Directed graph




Undirected graph

Properties of Adjacency List Representation

- Memory required
 - $\Theta(m+n)$
- Preferred when
 - the graph is sparse: $m \ll n^2$
- Disadvantage
 - no quick way to determine whether there is an edge between node u and v
- Time to list all vertices adjacent to u :
 - $\Theta(\text{degree}(u))$
- Time to determine if (u, v) exists:
 - $O(\text{degree}(u))$




Undirected graph



Directed graph

Graph Representation

- **Adjacency matrix representation** of $G = (V, E)$
 - Assume vertices are numbered $1, 2, \dots, n$
 - The representation consists of a matrix $A_{n \times n}$
 - $a_{ij} = \begin{cases} 1 & \text{if } (i, j) \text{ belongs to } E \\ 0 & \text{otherwise} \end{cases}$



Undirected graph

	1	2	3	4	5
1	0	1	0	0	1
2	1	0	1	1	1
3	0	1	0	1	0
4	0	1	1	0	1
5	1	1	0	1	0

For undirected graphs matrix A is symmetric:

$$a_{ij} = a_{ji}$$

$$A = A^T$$

Properties of Adjacency Matrix Representation

- Memory required
 - $\Theta(n^2)$, independent on the number of edges in G
- Preferred when
 - The graph is dense: m is close to n^2
 - We need to quickly determine if there is an edge between two vertices
- Time to list all vertices adjacent to u :
 - $\Theta(n)$
- Time to determine if (u, v) belongs to E :
 - $\Theta(1)$

Weighted Graphs

- **Weighted graphs** = graphs for which each edge has an associated weight $w(u, v)$
 $w: E \rightarrow \mathbb{R}$, weight function
- Storing the weights of a graph
 - Adjacency list:
 - Store $w(u, v)$ along with vertex v in u 's adjacency list
 - Adjacency matrix:
 - Store $w(u, v)$ at location (u, v) in the matrix

Searching in a Graph

- **Graph searching** = systematically follow the edges of the graph so as to visit the vertices of the graph
- Two basic graph searching algorithms:
 - Breadth-first search
 - Depth-first search
- The difference between them is in the order in which they explore the unvisited edges of the graph
- Graph algorithms are typically elaborations of the basic graph-searching algorithms

Breadth-First Search (BFS)

- **Input:**

- A graph $G = (V, E)$ (directed or undirected)
- A **source** vertex s from V


- **Goal:**

- Explore the edges of G to “discover” every vertex reachable from s , **taking the ones closest to s first**

- **Output:**

- $d[v]$ = distance (smallest # of edges) from s to v , for all v from V
- A “breadth-first tree” rooted at s that contains all reachable vertices

Readings



- For this lecture
 - Chapter 15
- Coming next
 - Chapter 20