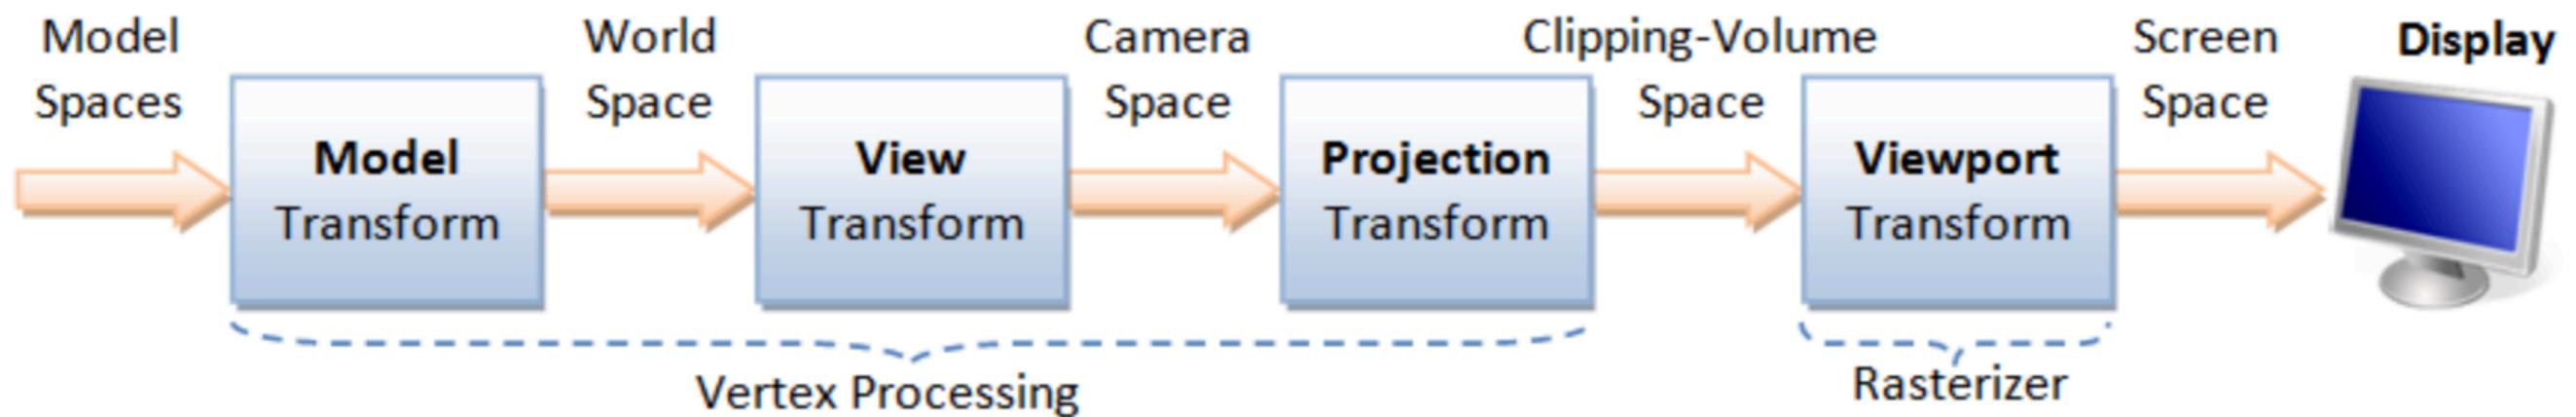


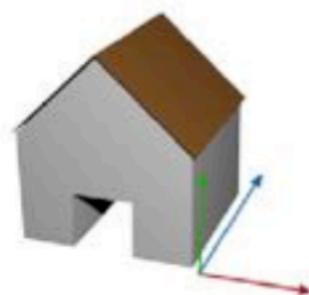


CS/PSY 484: Visual Rendering

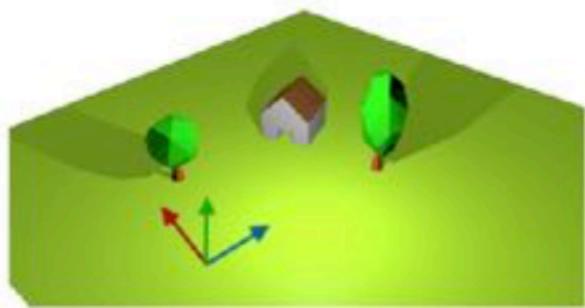


remember this from chapter 4?

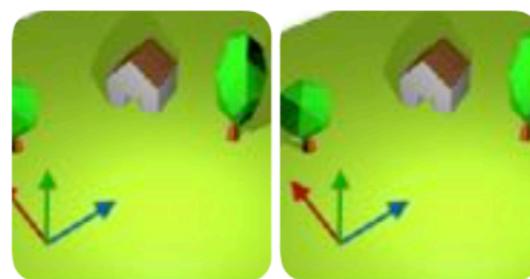
Model-World-View Projection



model space

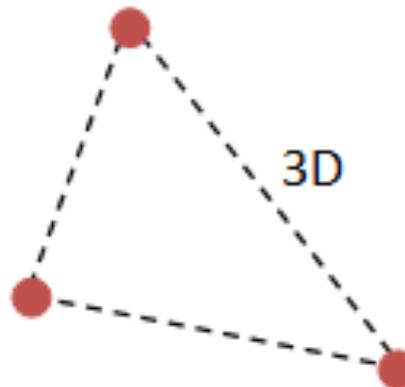


World space

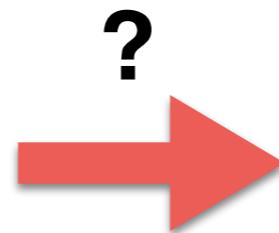


Stereo view space

how is every pixel colored?



A *primitive* is formed by one or more *vertices*. Vertices are not grid-aligned



?



All primitives are merged to produce 2D *pixels* on the display

rendering process

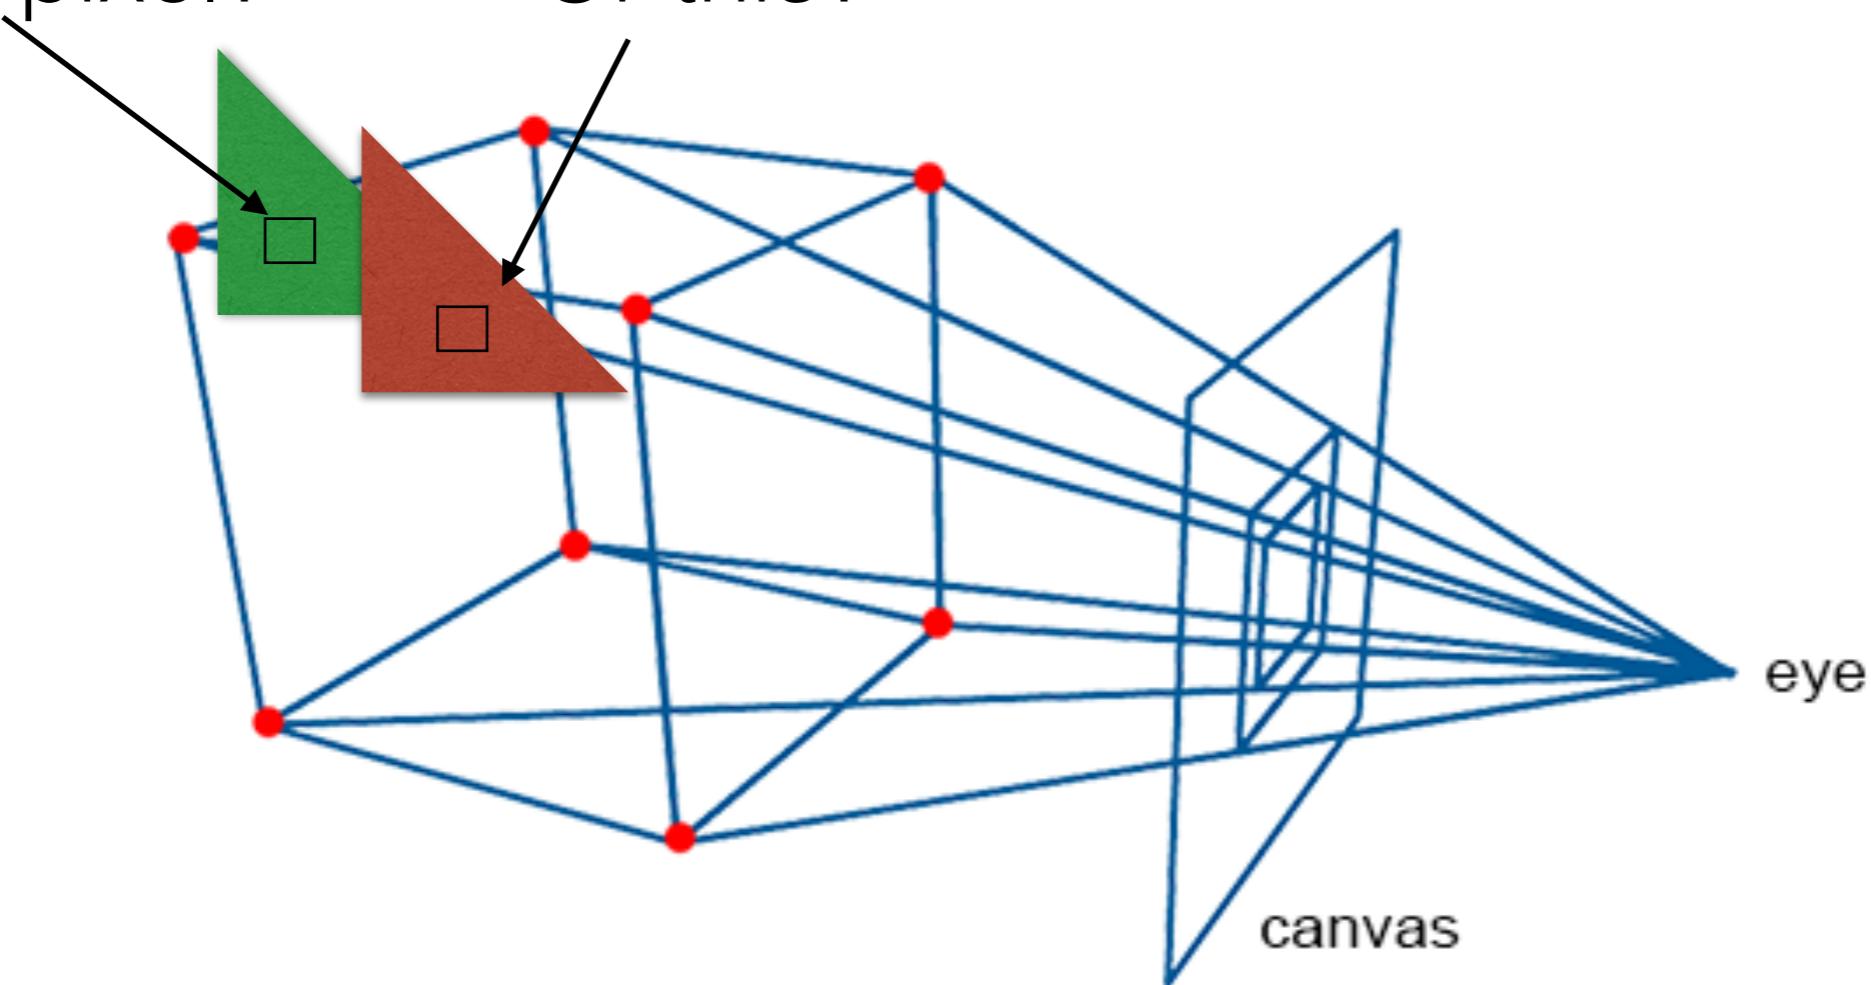
- Visibility problem (pixel visible)
- Shading (if so what color?)

} combined

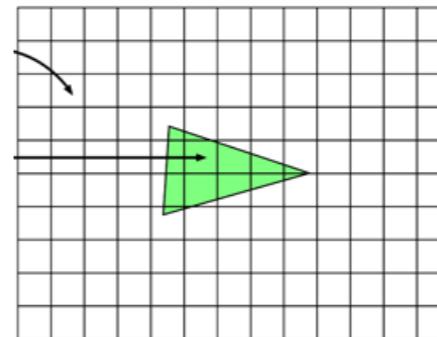
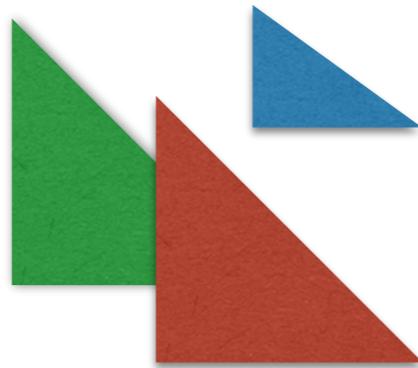
visibility problem

See this pixel?

Or this?



Rendering process



Two distinct rendering methods:

- Object order (iterate over triangles) complex ~ 1M triangles to
- Image order (iterate over pixels) 1080p ~ 2M pixels to consider
- Complicated for VR due to high refresh rates >100 Hz.

Raytracing

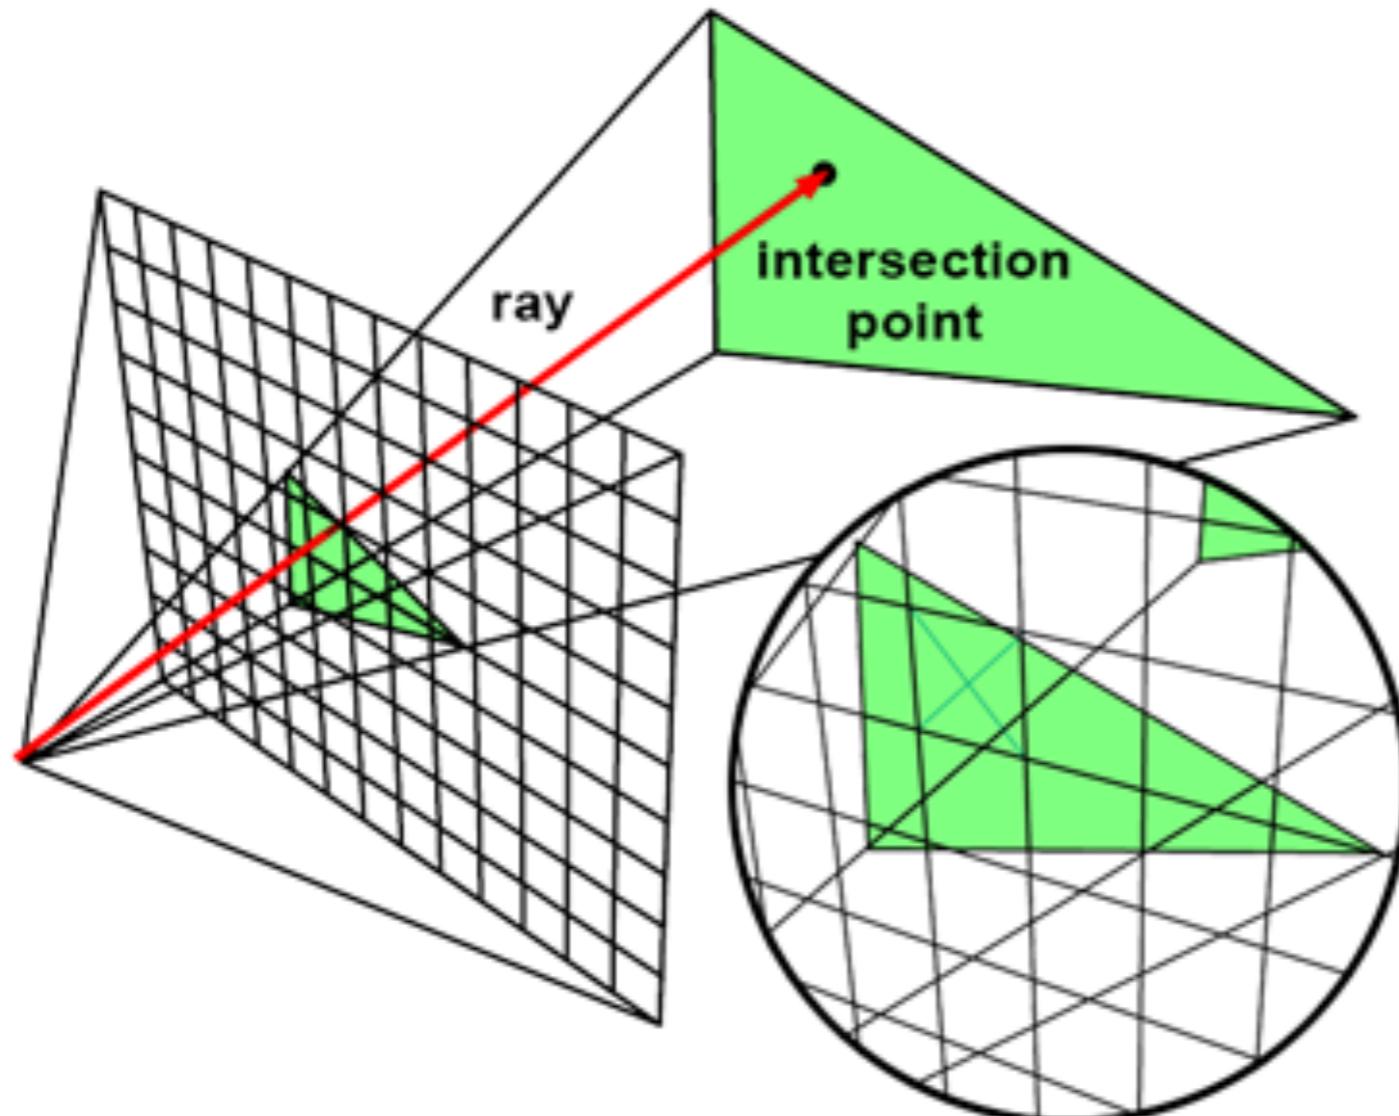
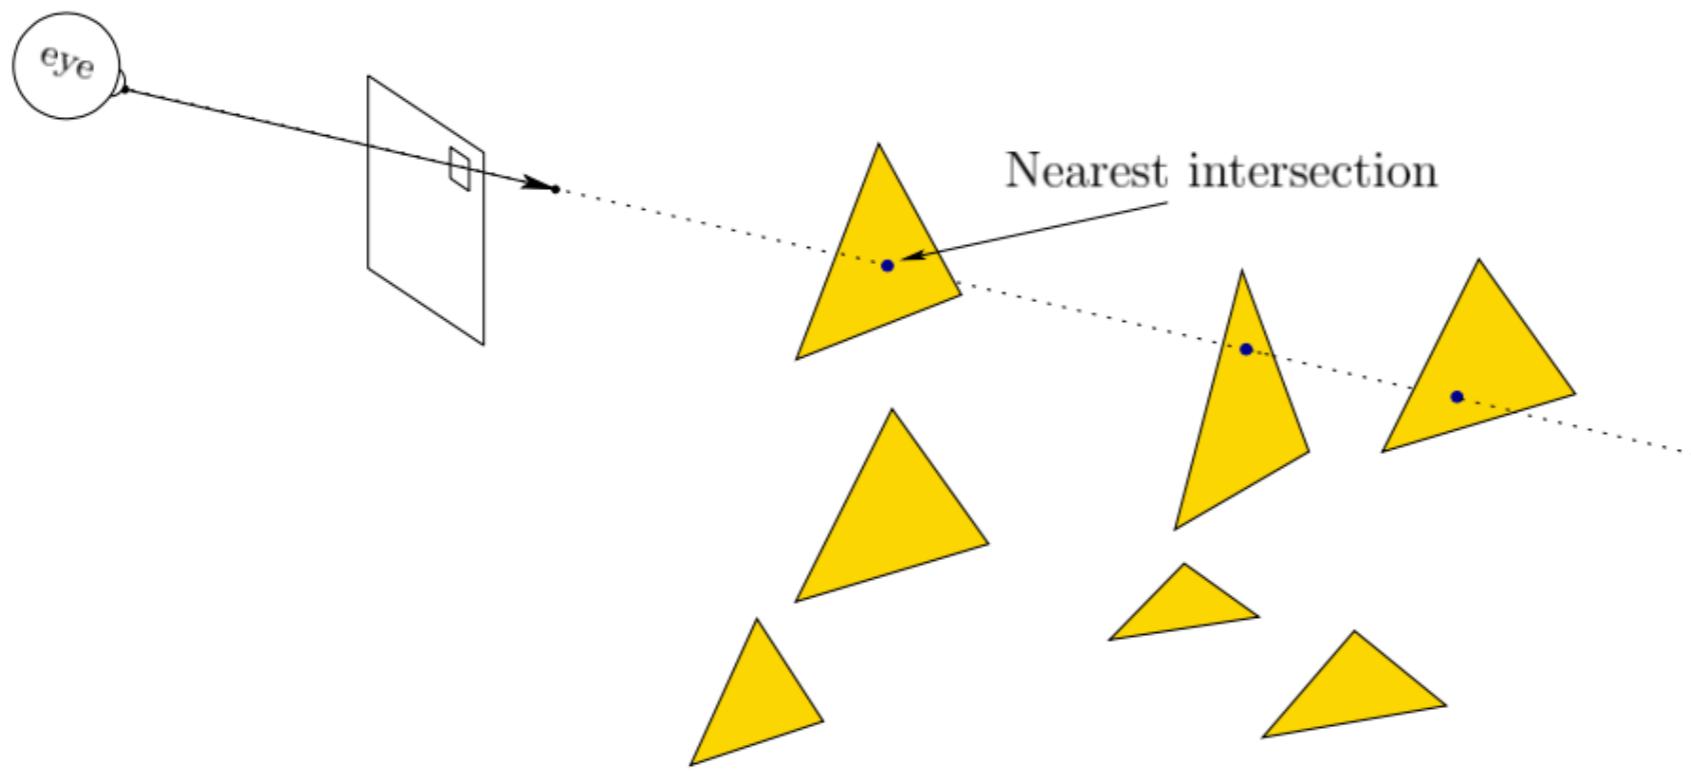


image order rendering: loops over pixels, then over triangles

Two steps: 1) ray casting and 2) shading

Ray casting

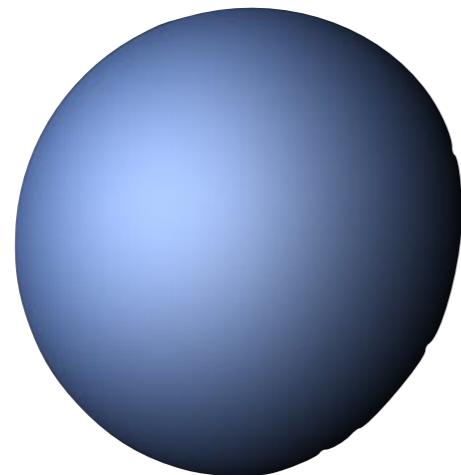


Using analytical geometry: take a pixel, cast ray from the eye to find the closest triangle.

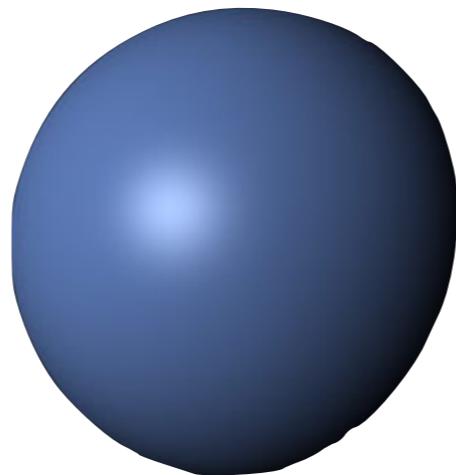
Arrange triangles into a 3D data structure

Shading

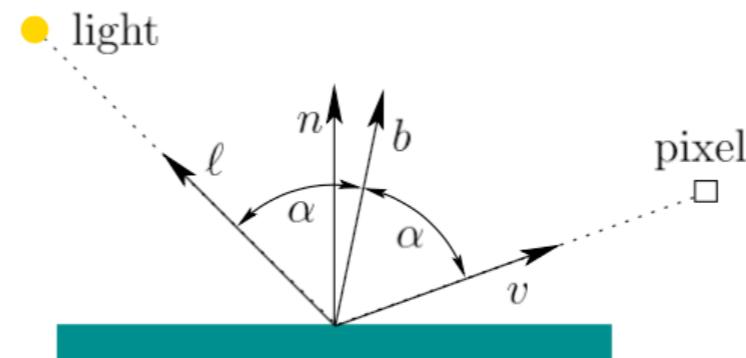
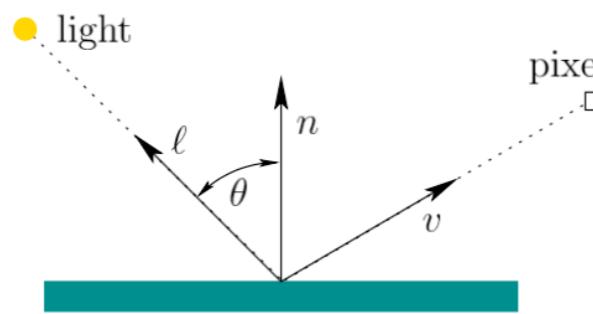
What color should each pixel have?



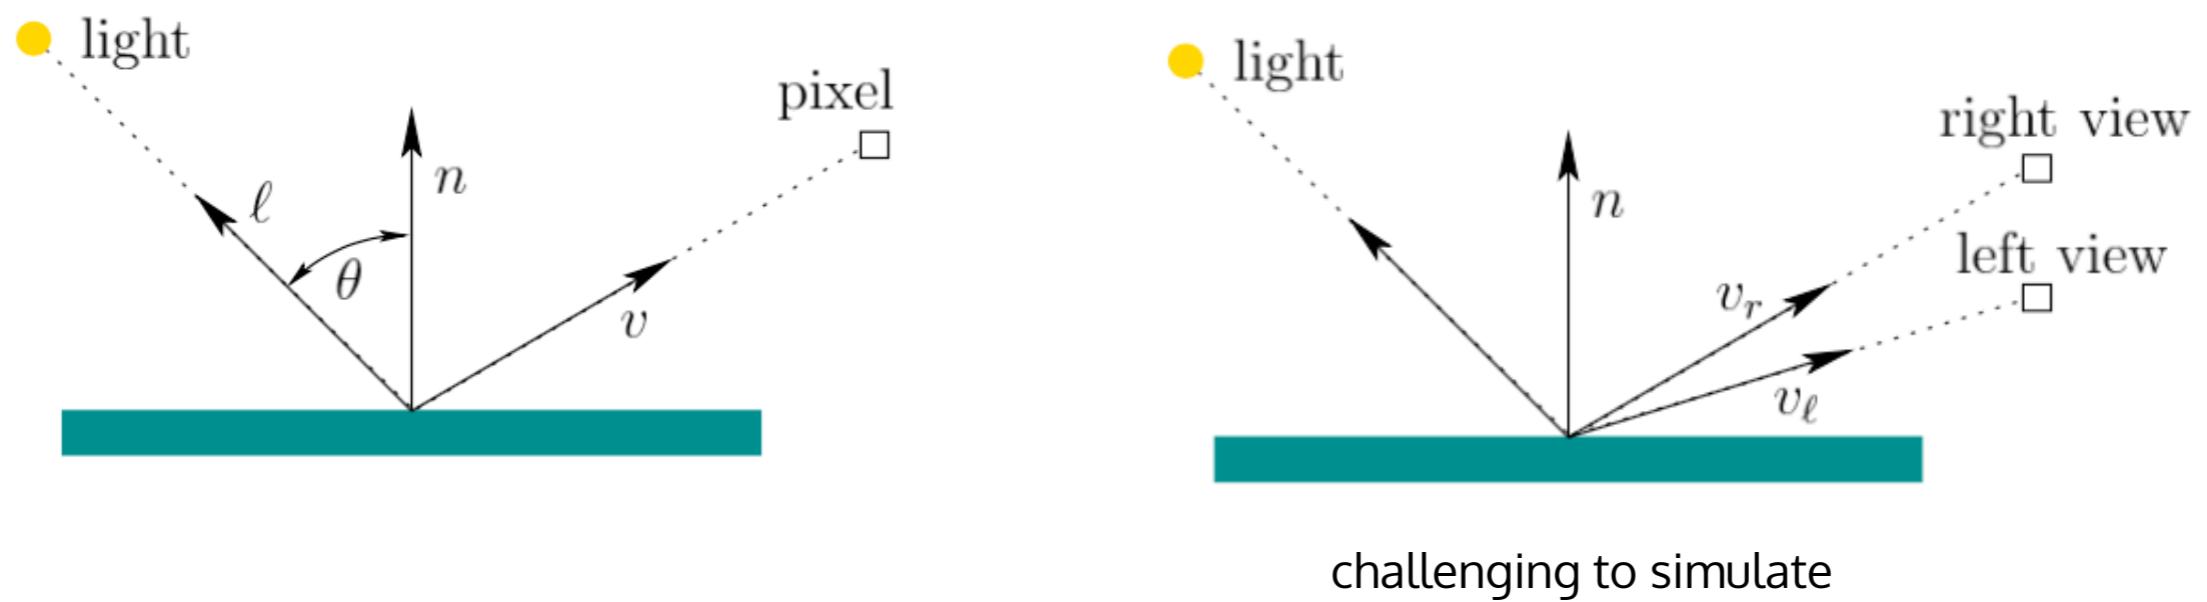
Lambert



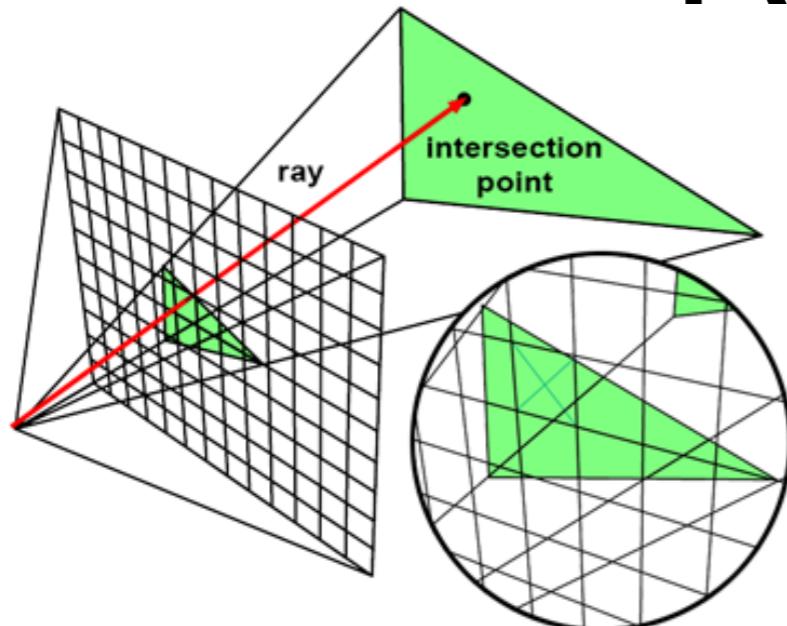
Blinn-Phong



Shading for VR



Raytracing



Doesn't scale.
1080p image @90hz
Requires 180M raycasts per second

Real time Raytracing for
Games demonstrated
You need a pricey GPU

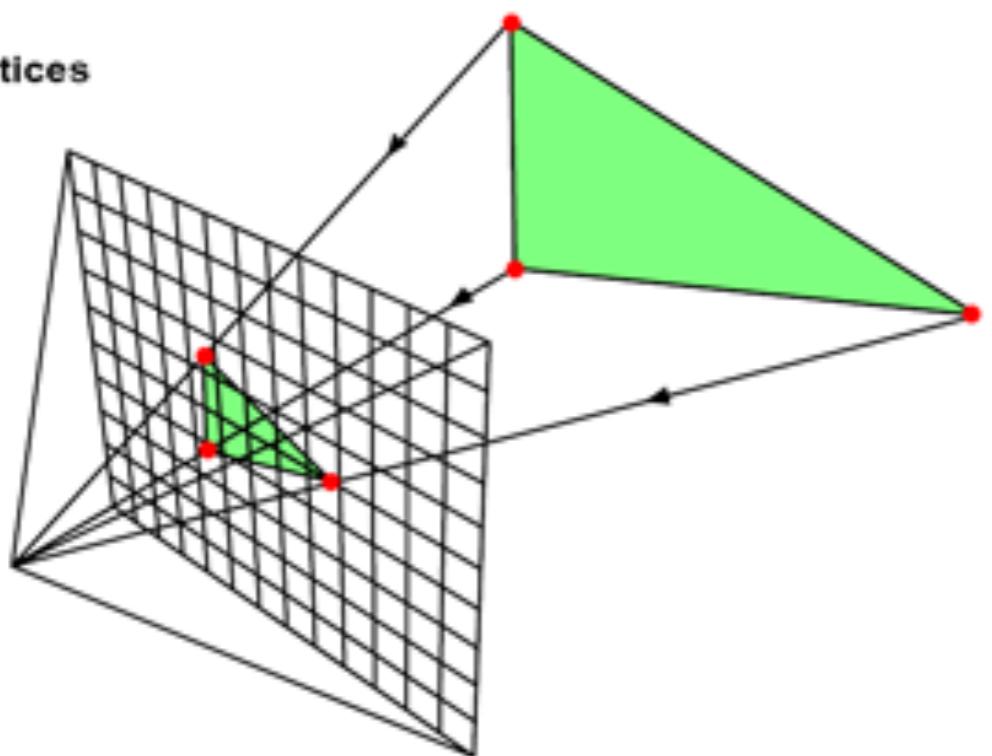
<https://youtu.be/J3ue35ago3Y>



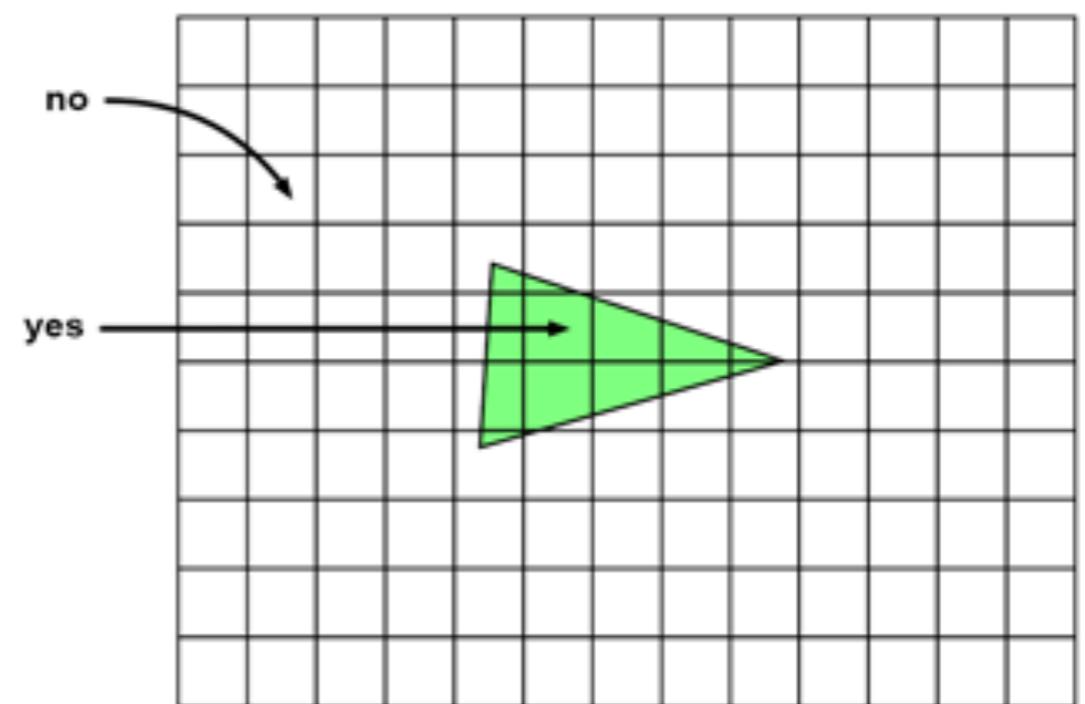
EVENT COVERAGE

Rasterization

1) Project vertices



2) Loop over pixels. Does the pixel lie in the triangle?



Object order rendering: loops over triangles, then over pixels
But how do we know which triangle is visible?
(visibility problem)

painters algorithm

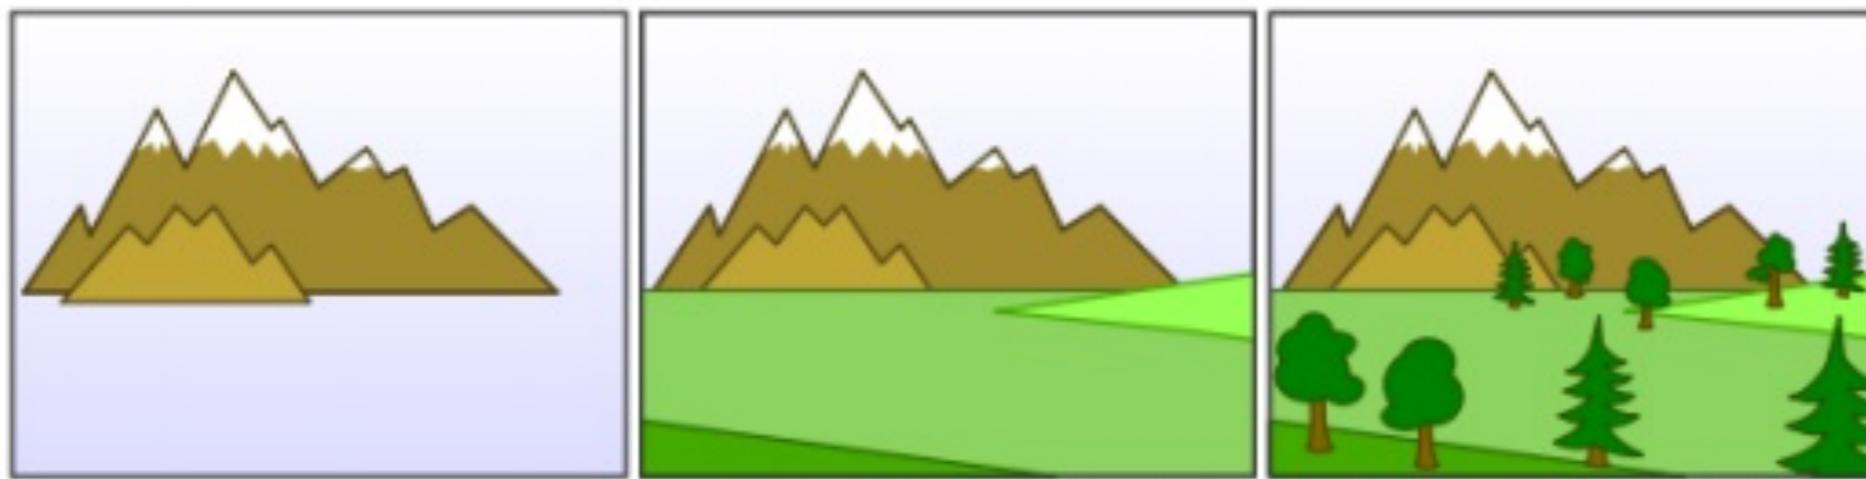
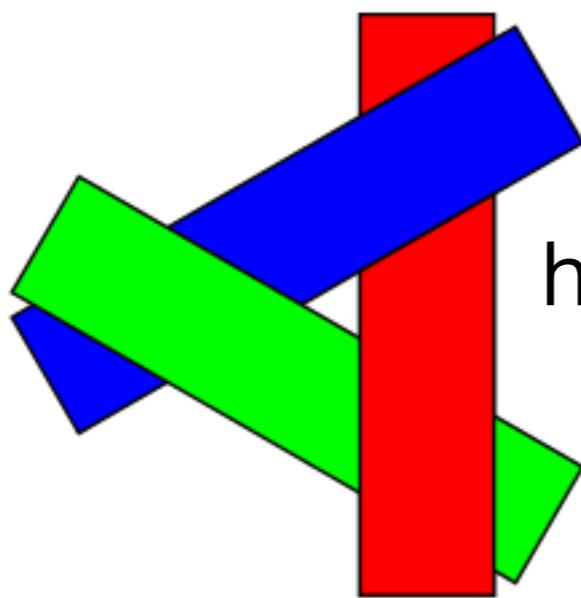


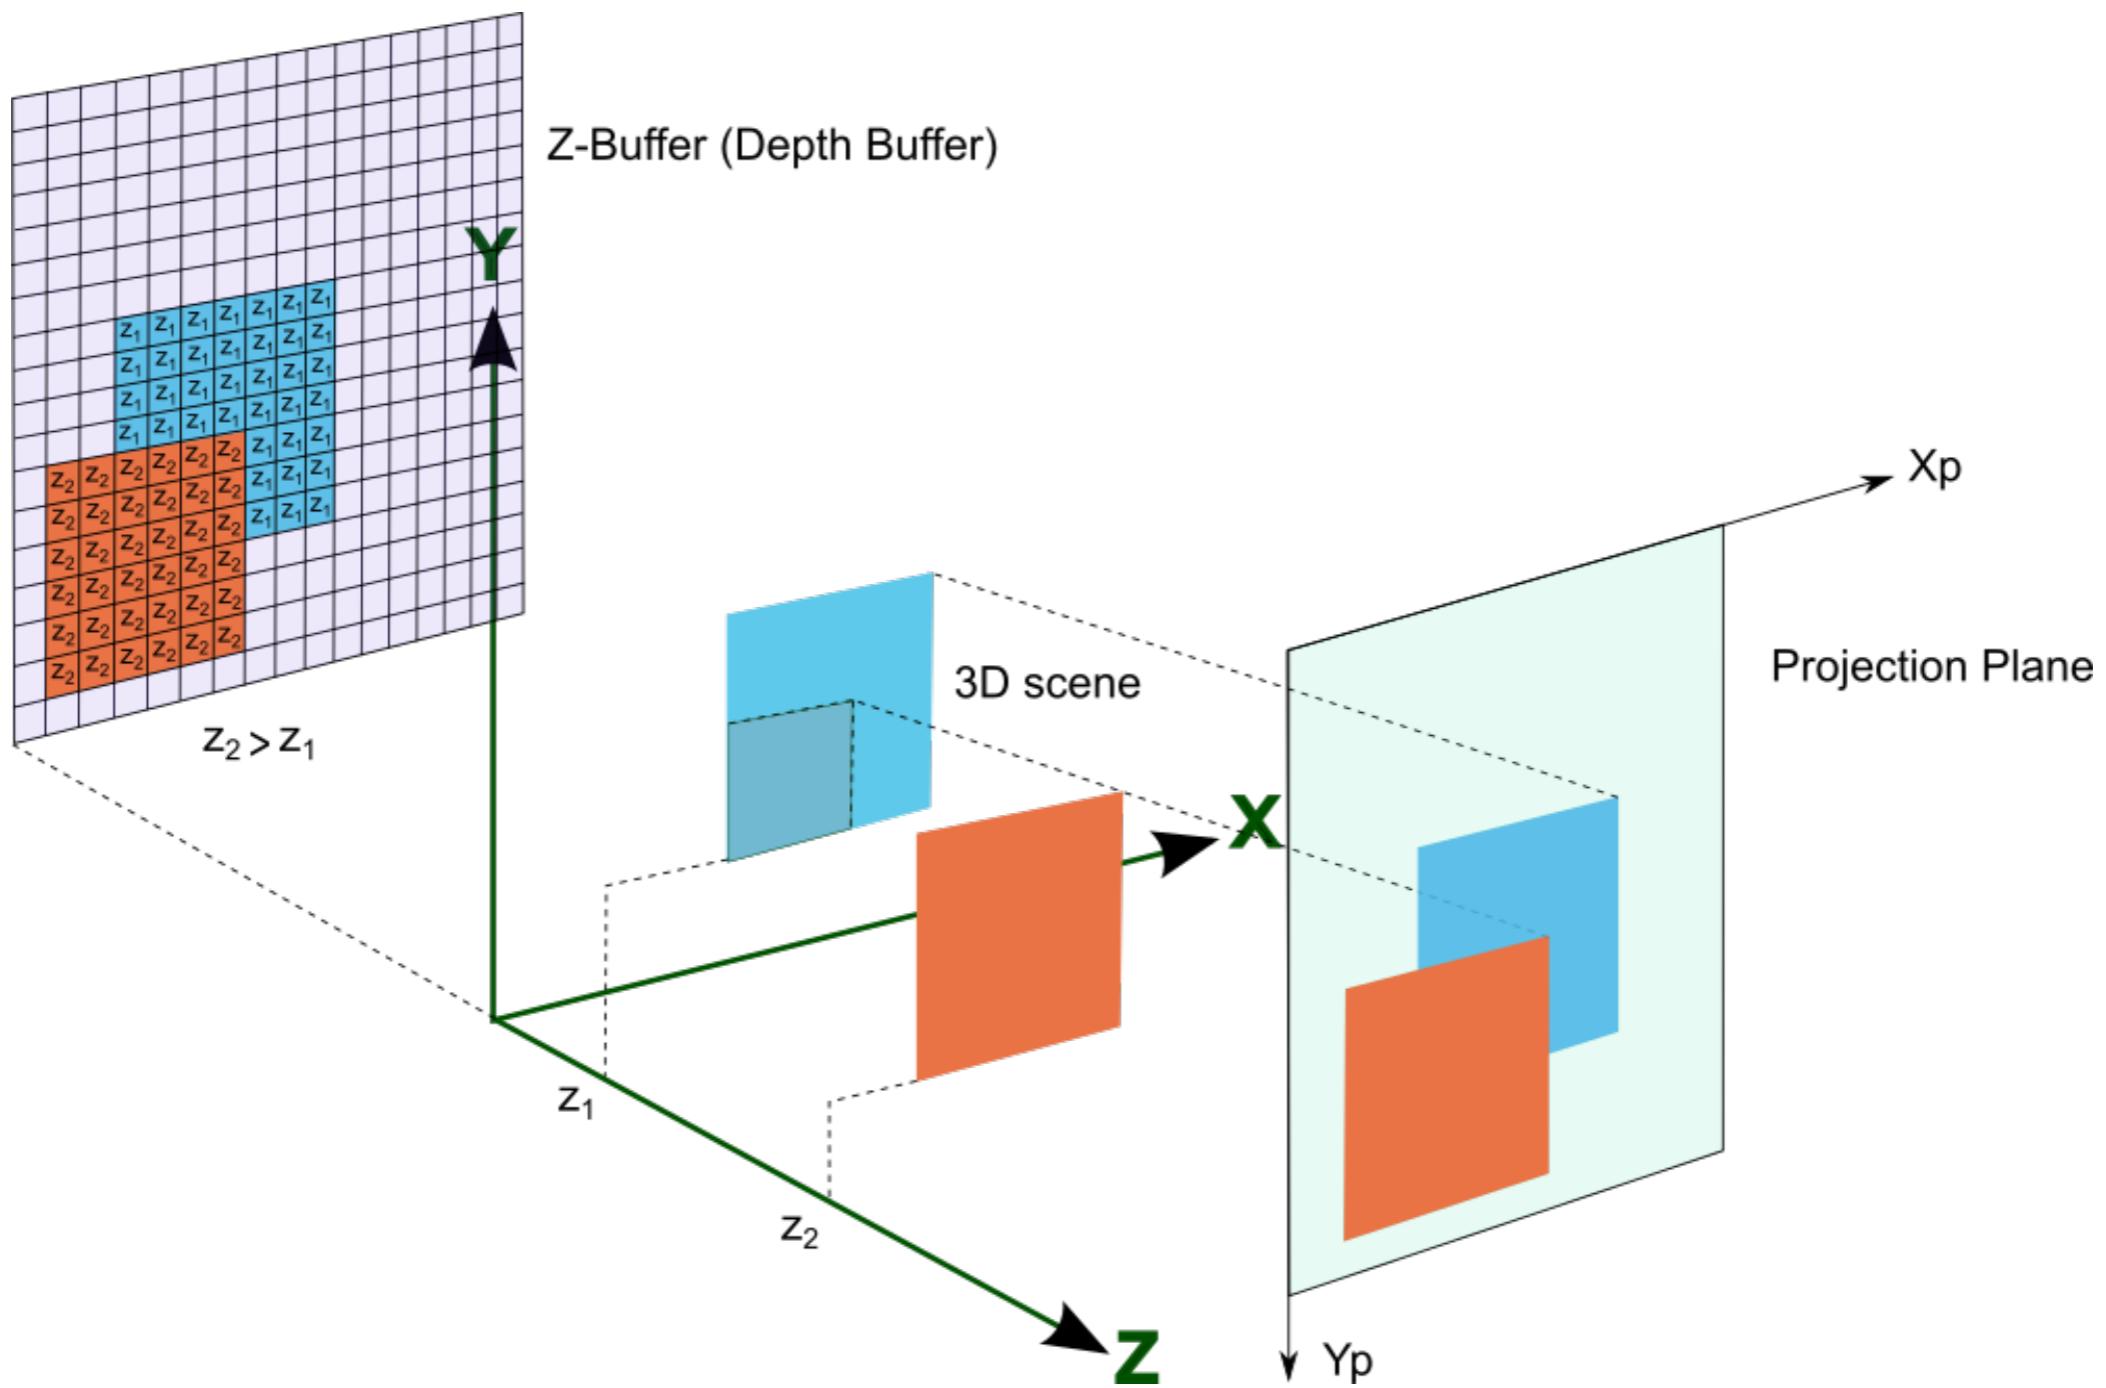
Fig.3.4. Painter's algorithm

inefficient for large numbers of objects (many won't be visible)



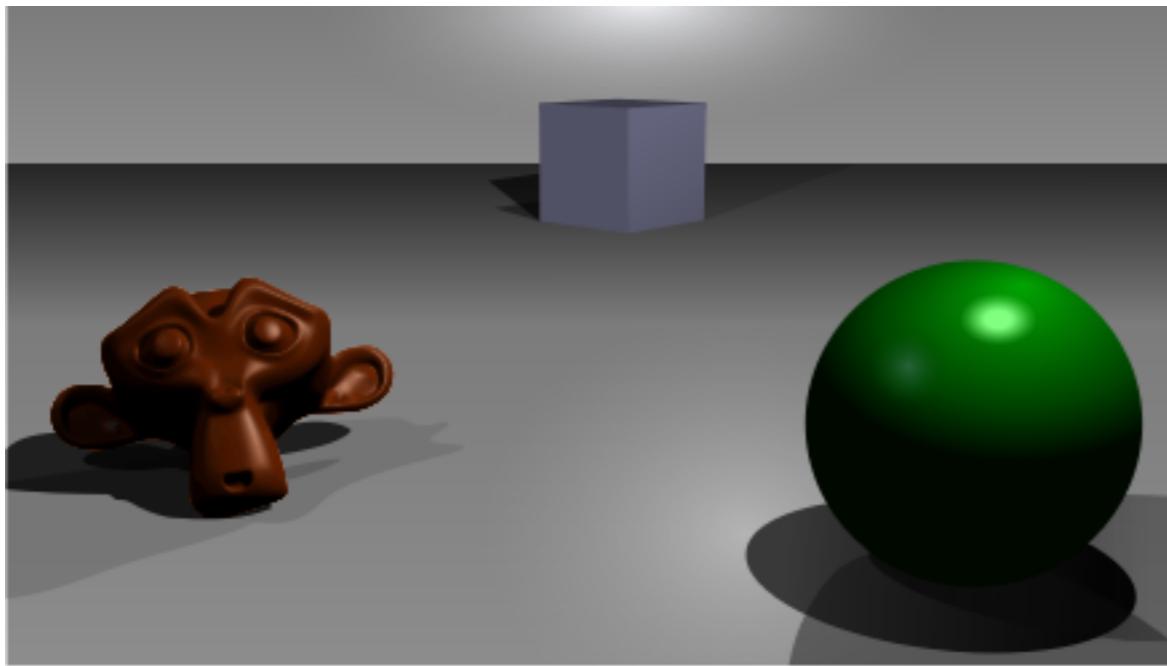
how to solve this?

Z-buffer

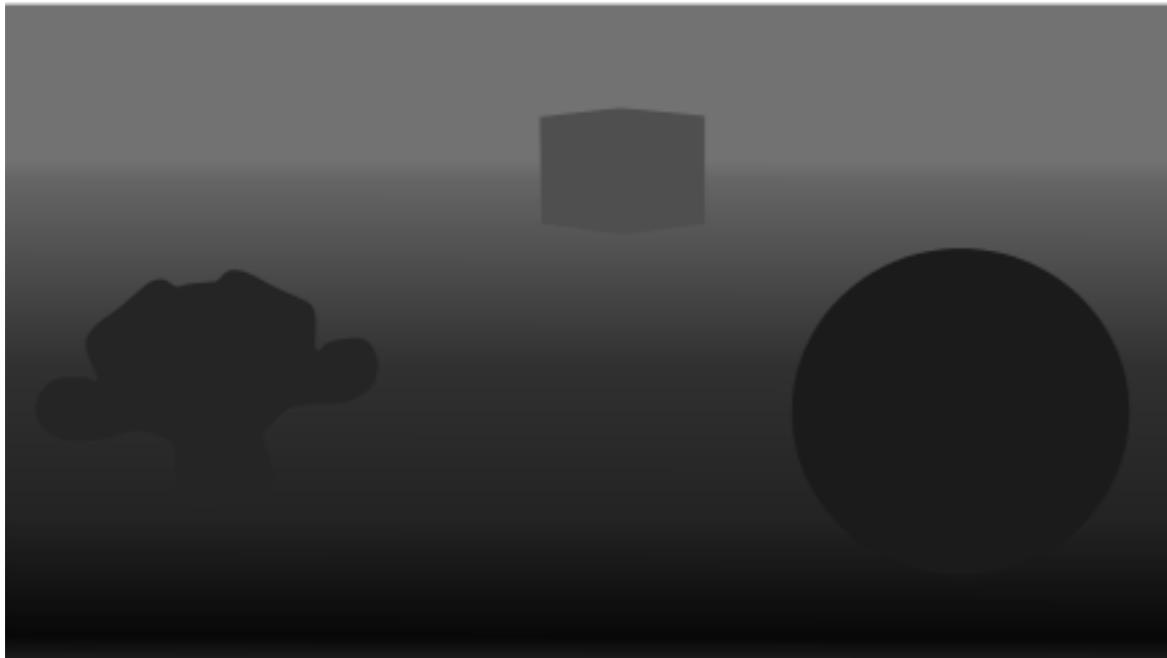


Keep track of the depth of every pixel that is visible

Z-buffer



A simple three-dimensional scene



Z-buffer representation

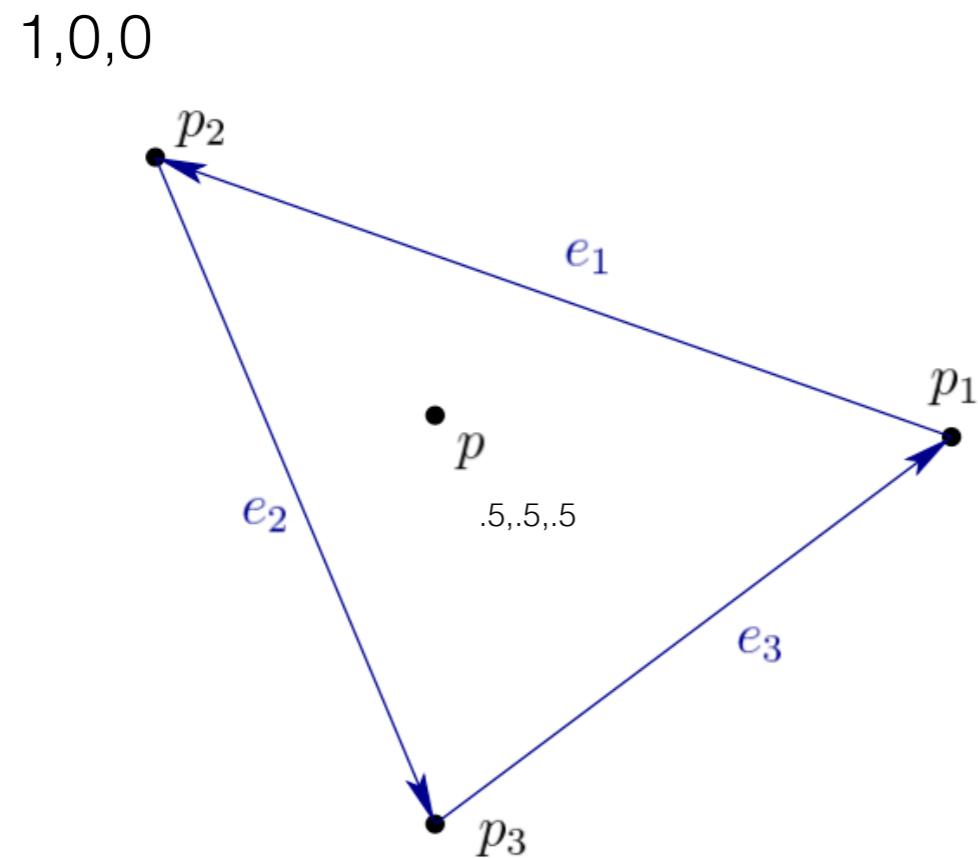
How It Works:

- Each pixel in the framebuffer has a corresponding depth value stored in the Z-buffer.
- When rendering a new pixel, its depth (Z-value) is compared to the existing depth in the buffer:
 - If the new pixel is **closer**, it replaces the old one.
 - If it's **farther**, it's discarded.

Benefits:

- ✓ Efficient hidden surface removal (only the closest surfaces are drawn).
- ✓ Works well with modern GPUs for real-time 3D rendering.
- ✓ Supports complex scenes with overlapping objects.

Barycentric coordinates



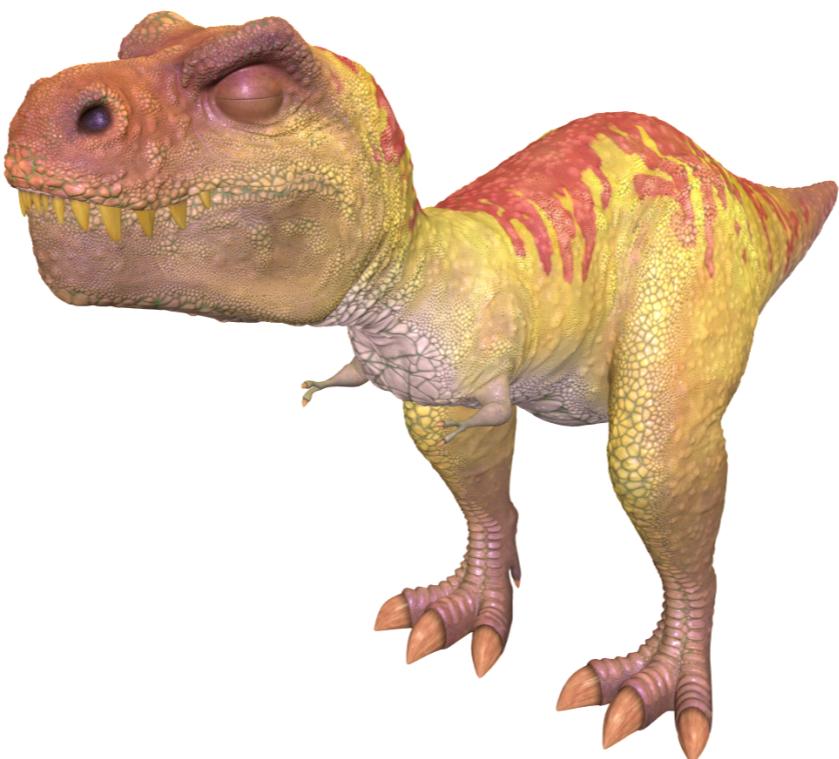
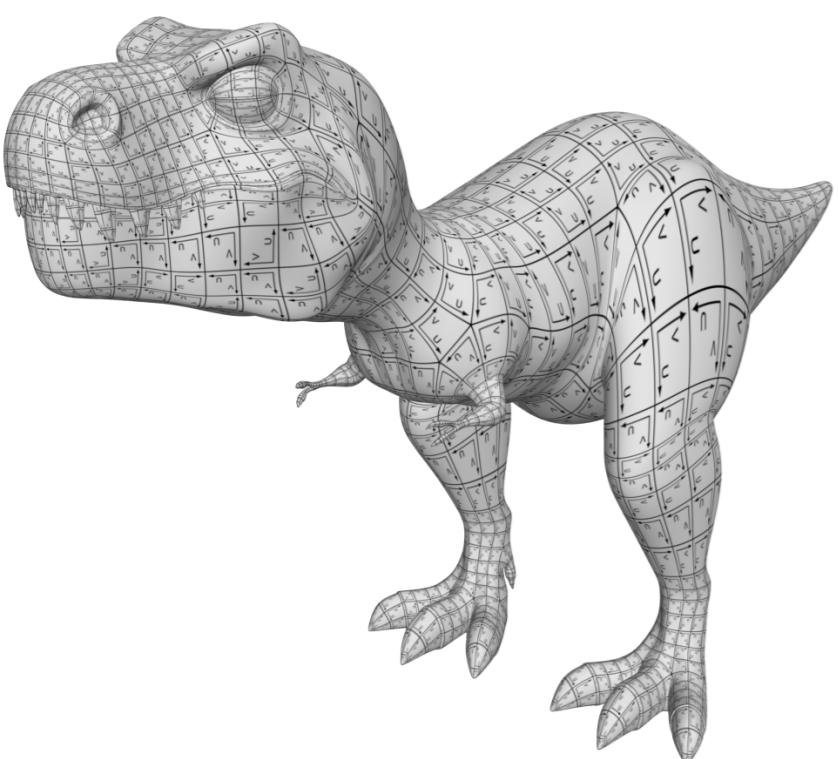
Barycentric coordinates specify the location of every point p in a triangle as a weighted average of its vertices p_1 , p_2 and p_3

Simplifies test if a point is inside or outside the triangle

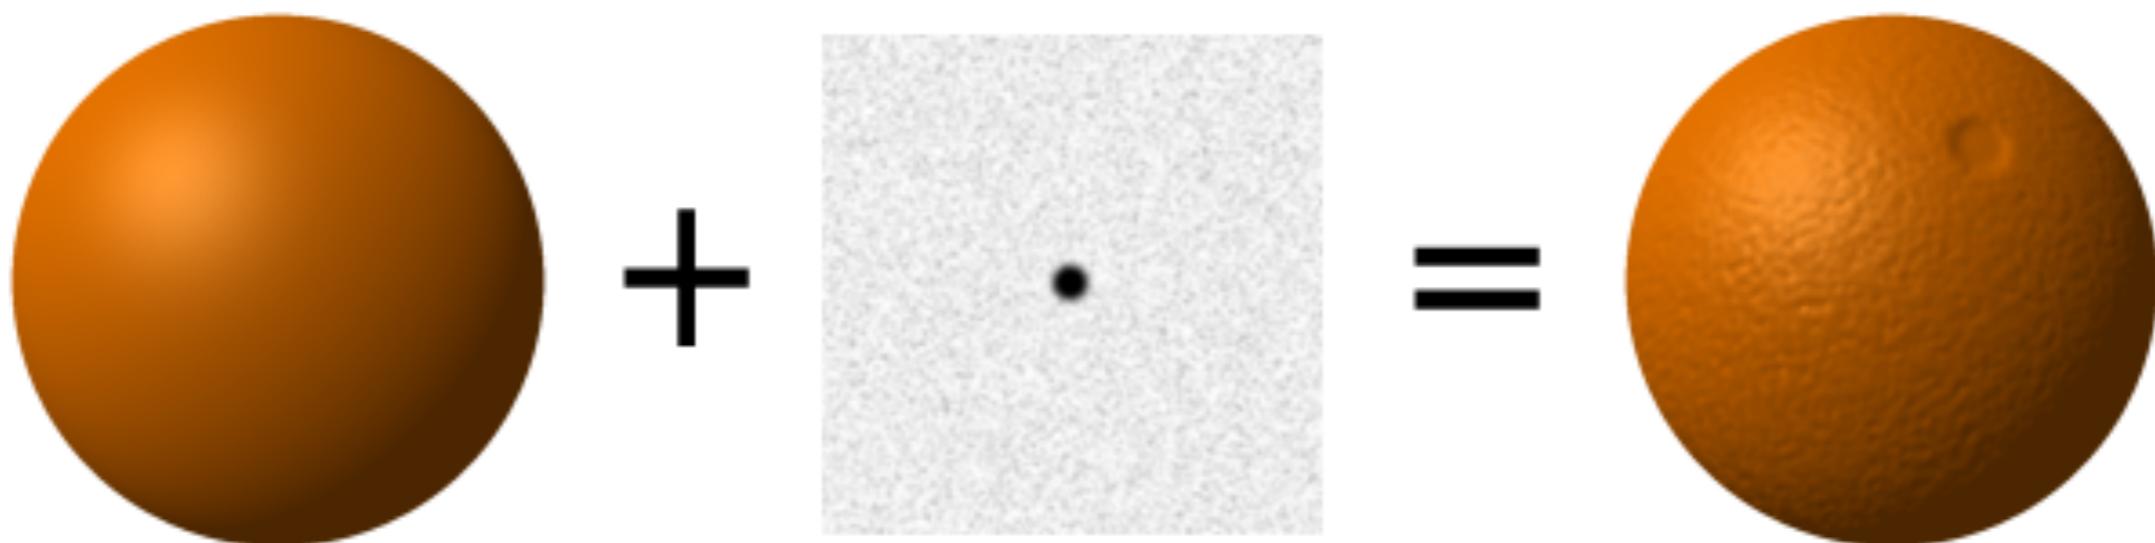
Comparison

Feature	Object-Order Rendering	Image-Order Rendering
Processing Approach	Iterates over objects	Iterates over pixels
Example Algorithms	Z-buffer, rasterization	Ray tracing, ray casting
Performance	Fast, hardware-accelerated	Slower, but high-quality lighting
Overdraw Issues	Yes (depth tests needed)	No (only visible pixels computed)
Real-time Suitability	Excellent (used in games)	Traditionally slow, but improving
Best for	Real-time graphics, games	High-quality offline rendering, movies

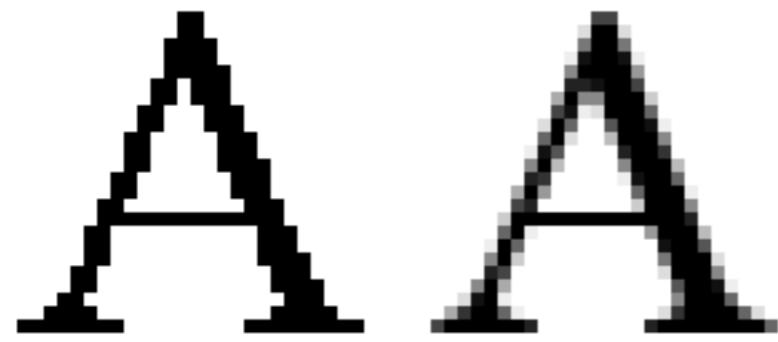
texture mapping



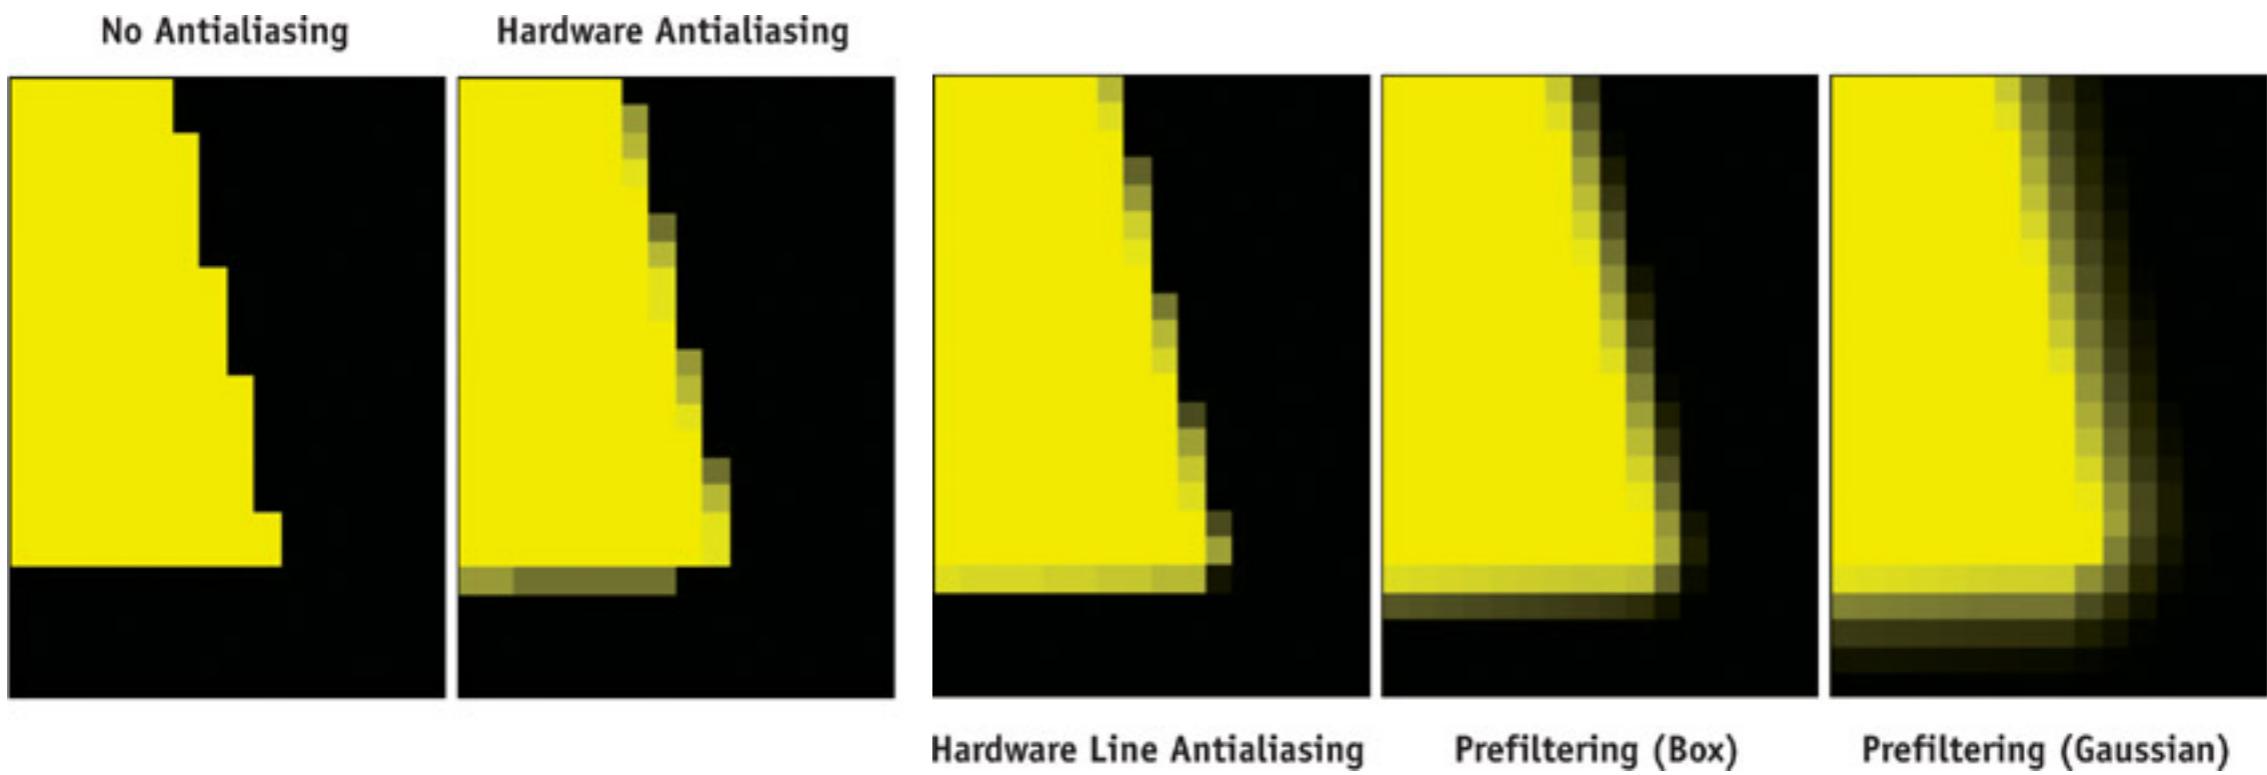
Bump mapping



alters the surface normal
Alters shading/appearance



Aliasing



Artifacts appear due to discretization
not a huge problem for modern displays (>300 ppi)
but it is for VR (eyes close to screen)

Mipmapping

A **ipmap** is a precomputed sequence of progressively smaller versions of a texture. When rendering, the system selects the most appropriate texture level based on the object's distance from the camera. This reduces aliasing (unwanted visual artifacts) and improves performance by using lower-resolution textures for distant objects.



256x256

LOD0



128x128

LOD1



64x64

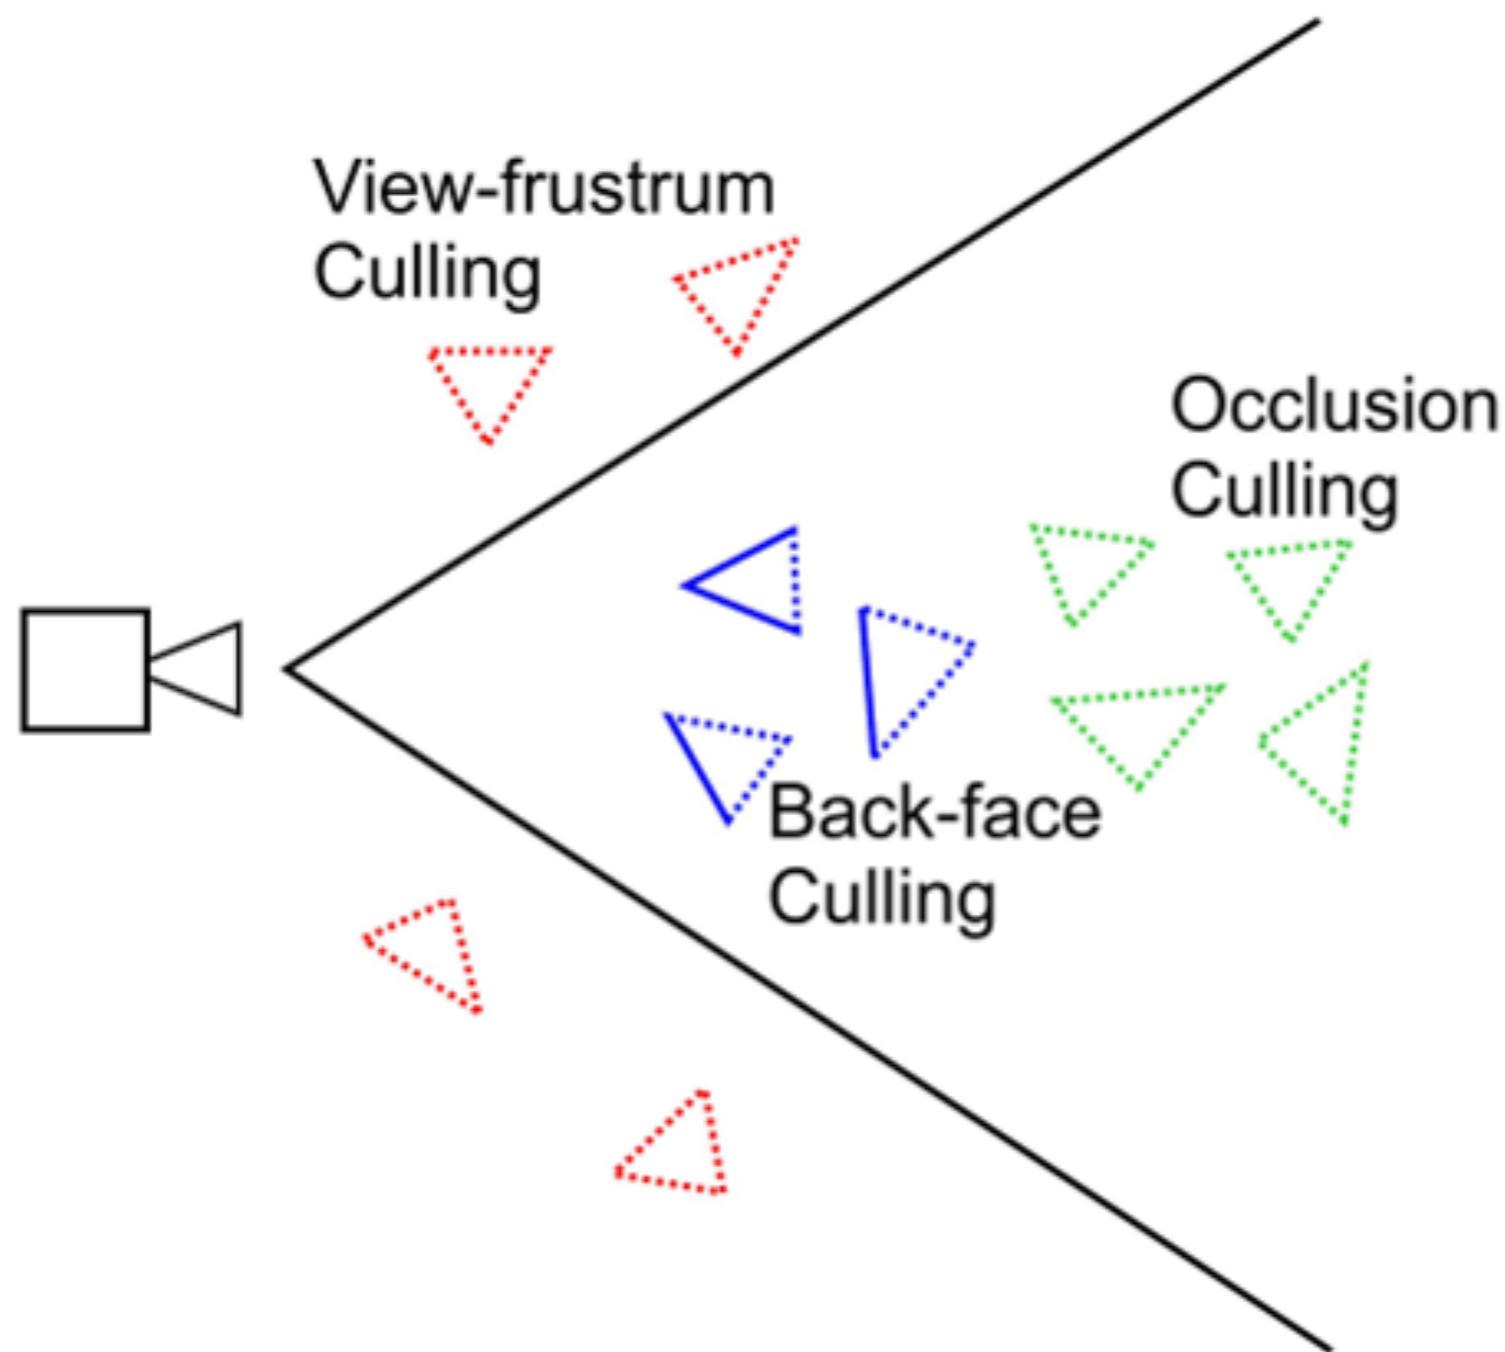
LOD2

32x32

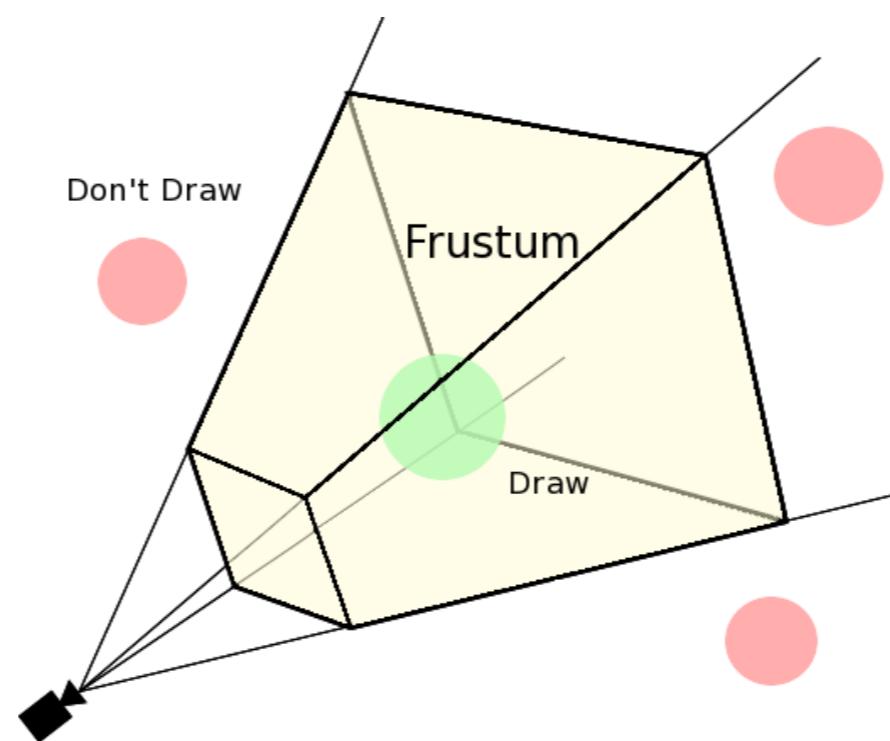
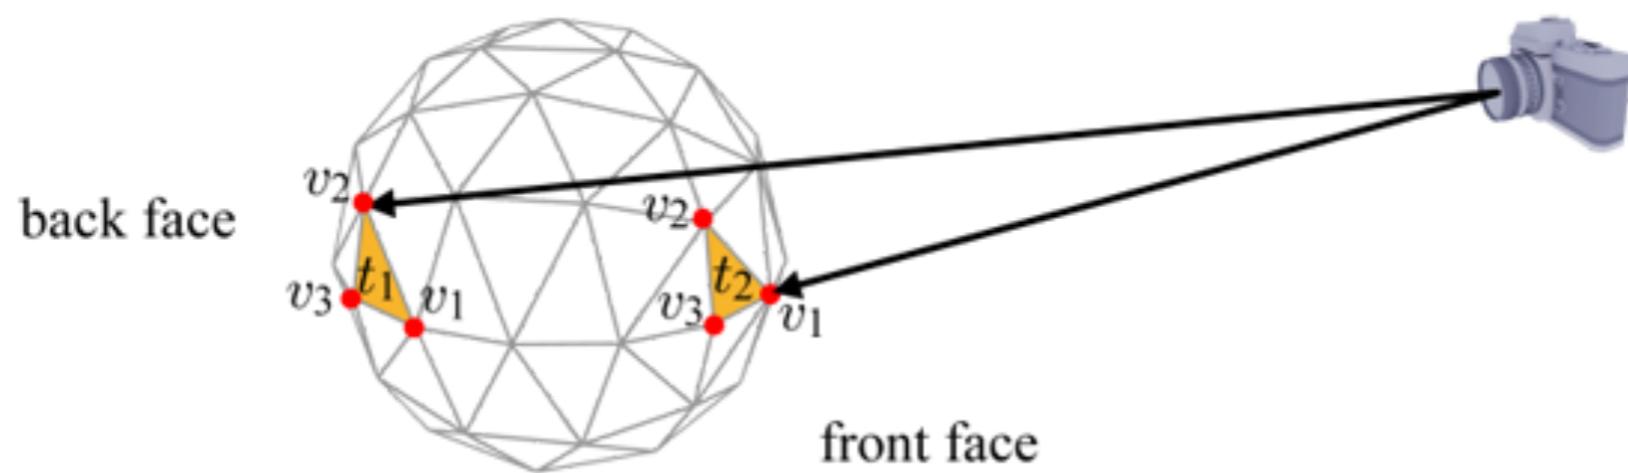
LOD3



culling



back-face / frustum culling



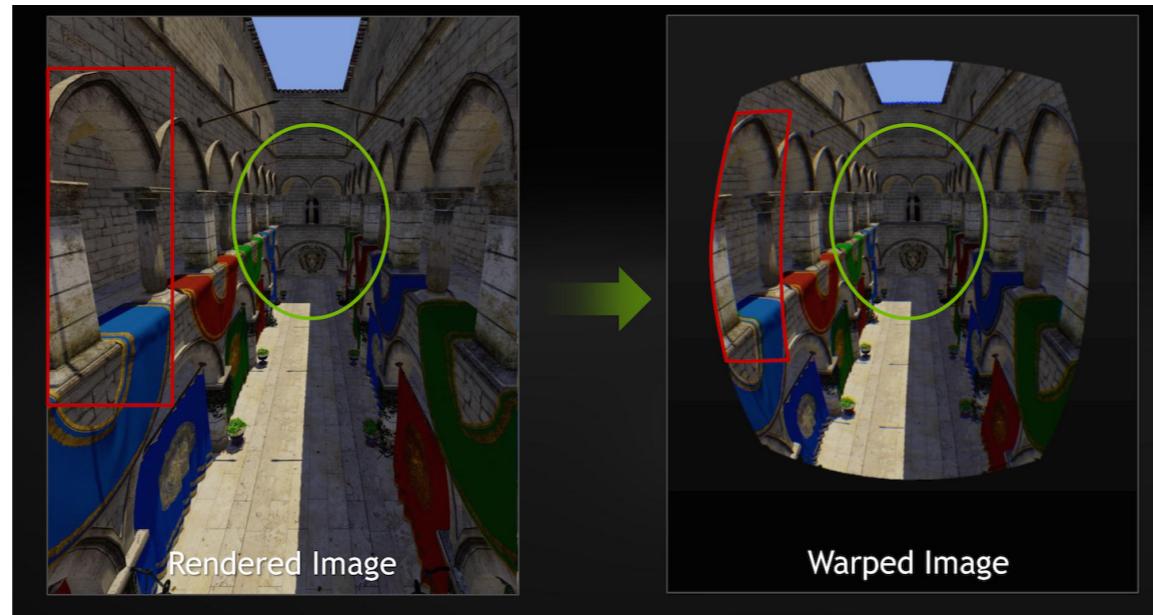
VR rasterization problems

- Aliasing problem due to VR displays not having a high pixel density (pixels vs refresh rate).
- Provide moving mismatch (tradeoff refresh rate versus pixel density)
- Magnified by stereo perception
- Stereo vision; texture mapping may look fake

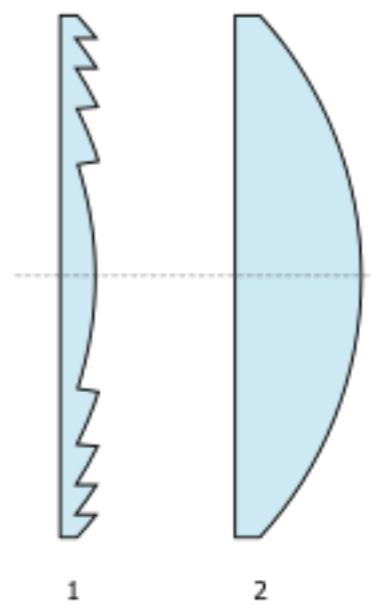
Screen door effect



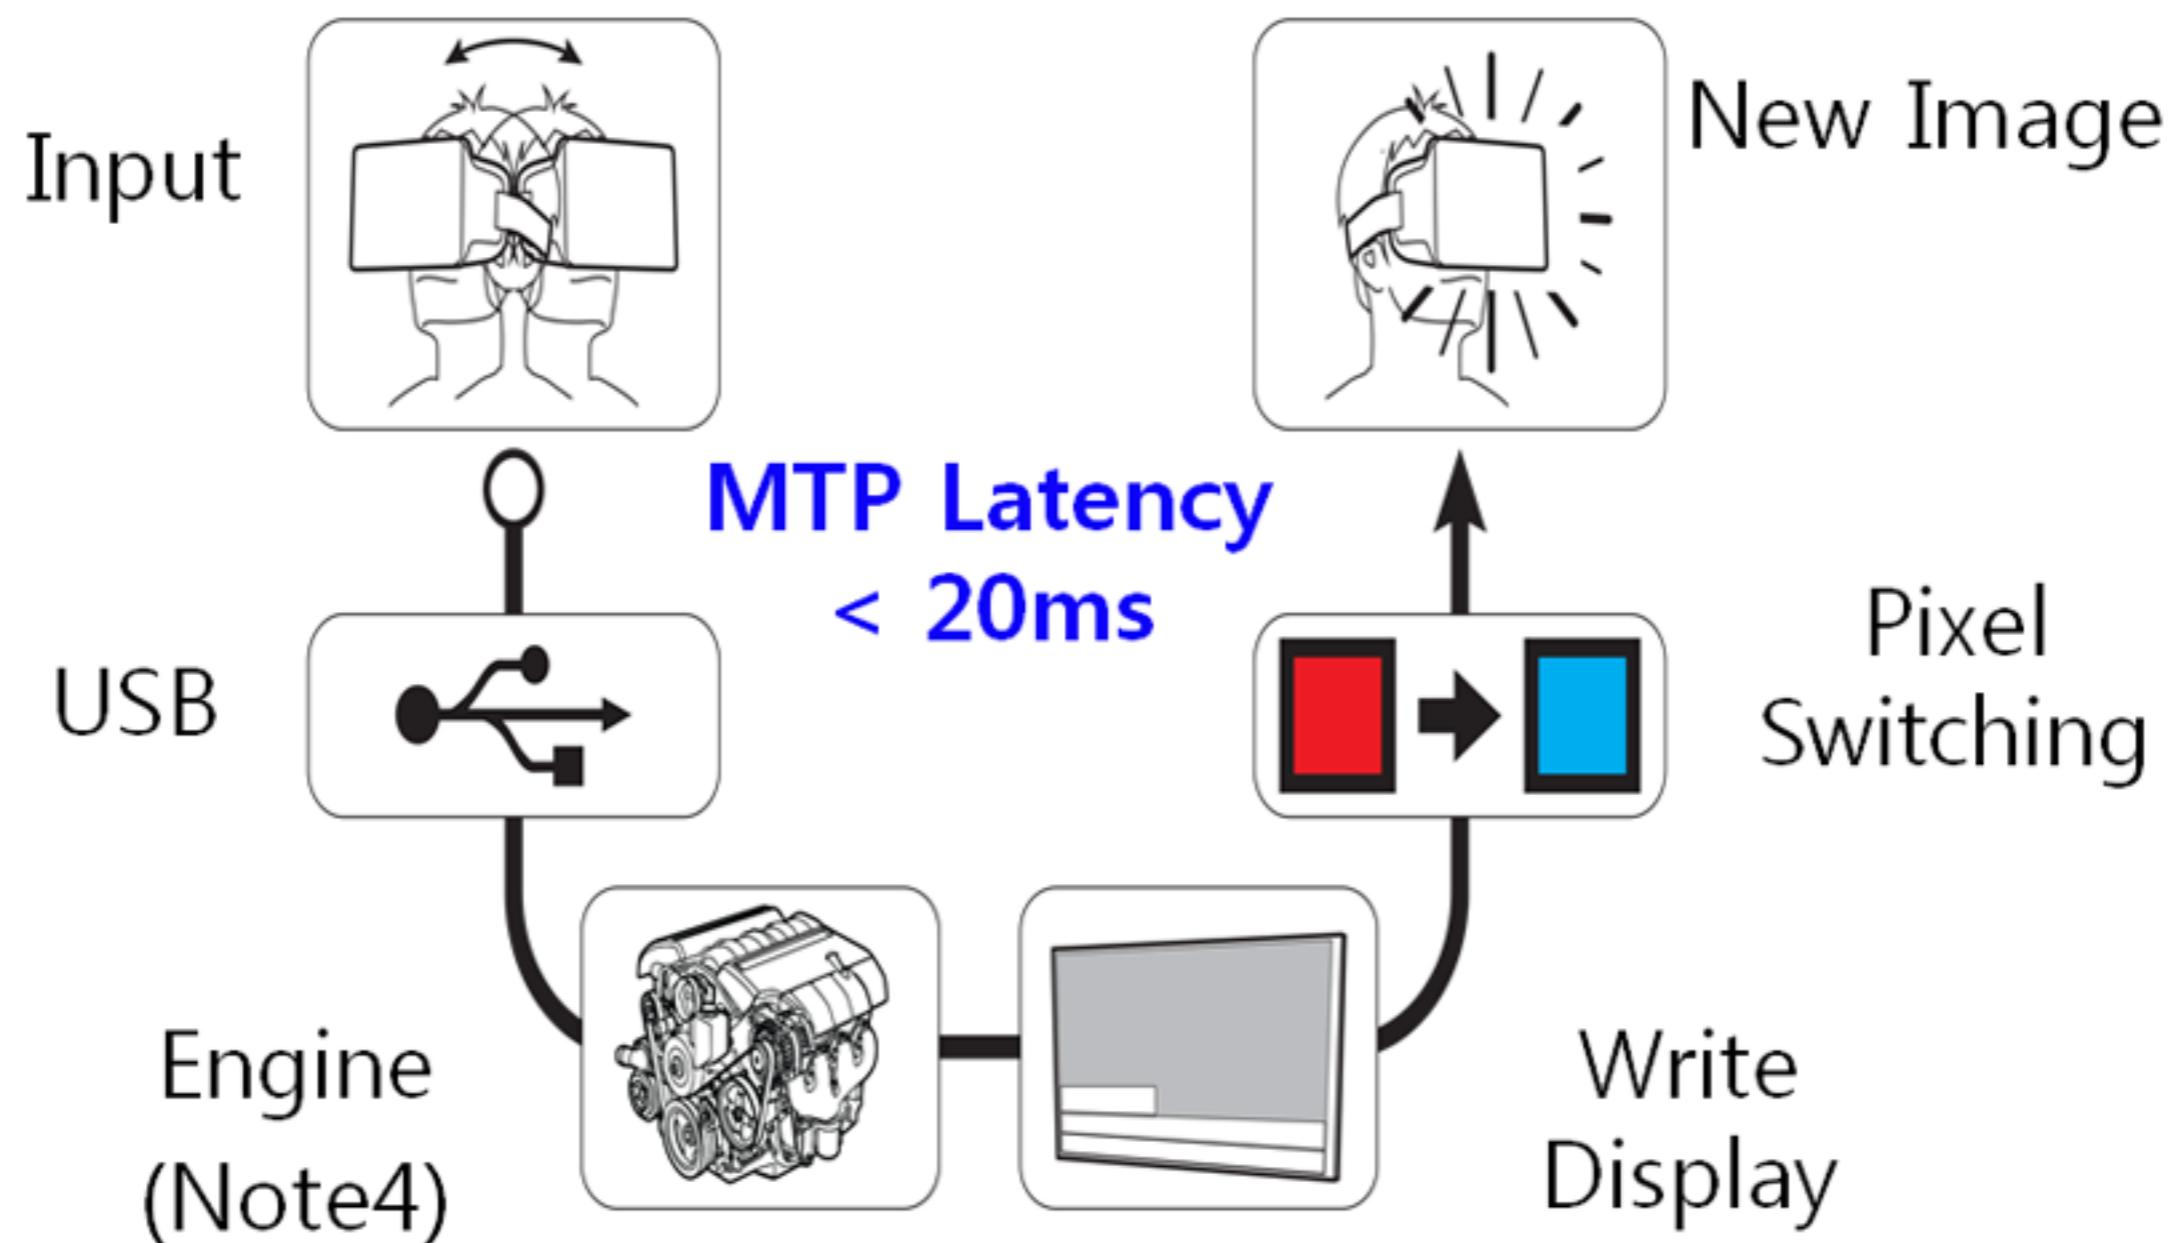
Correcting distortions



- barrel /pincushion distortion - because of a high FOV
- correct using fresnel lens
- More accurately bends light
- Also thinner
- But can have a bit of glare with light coming from behind



Latency



High latency

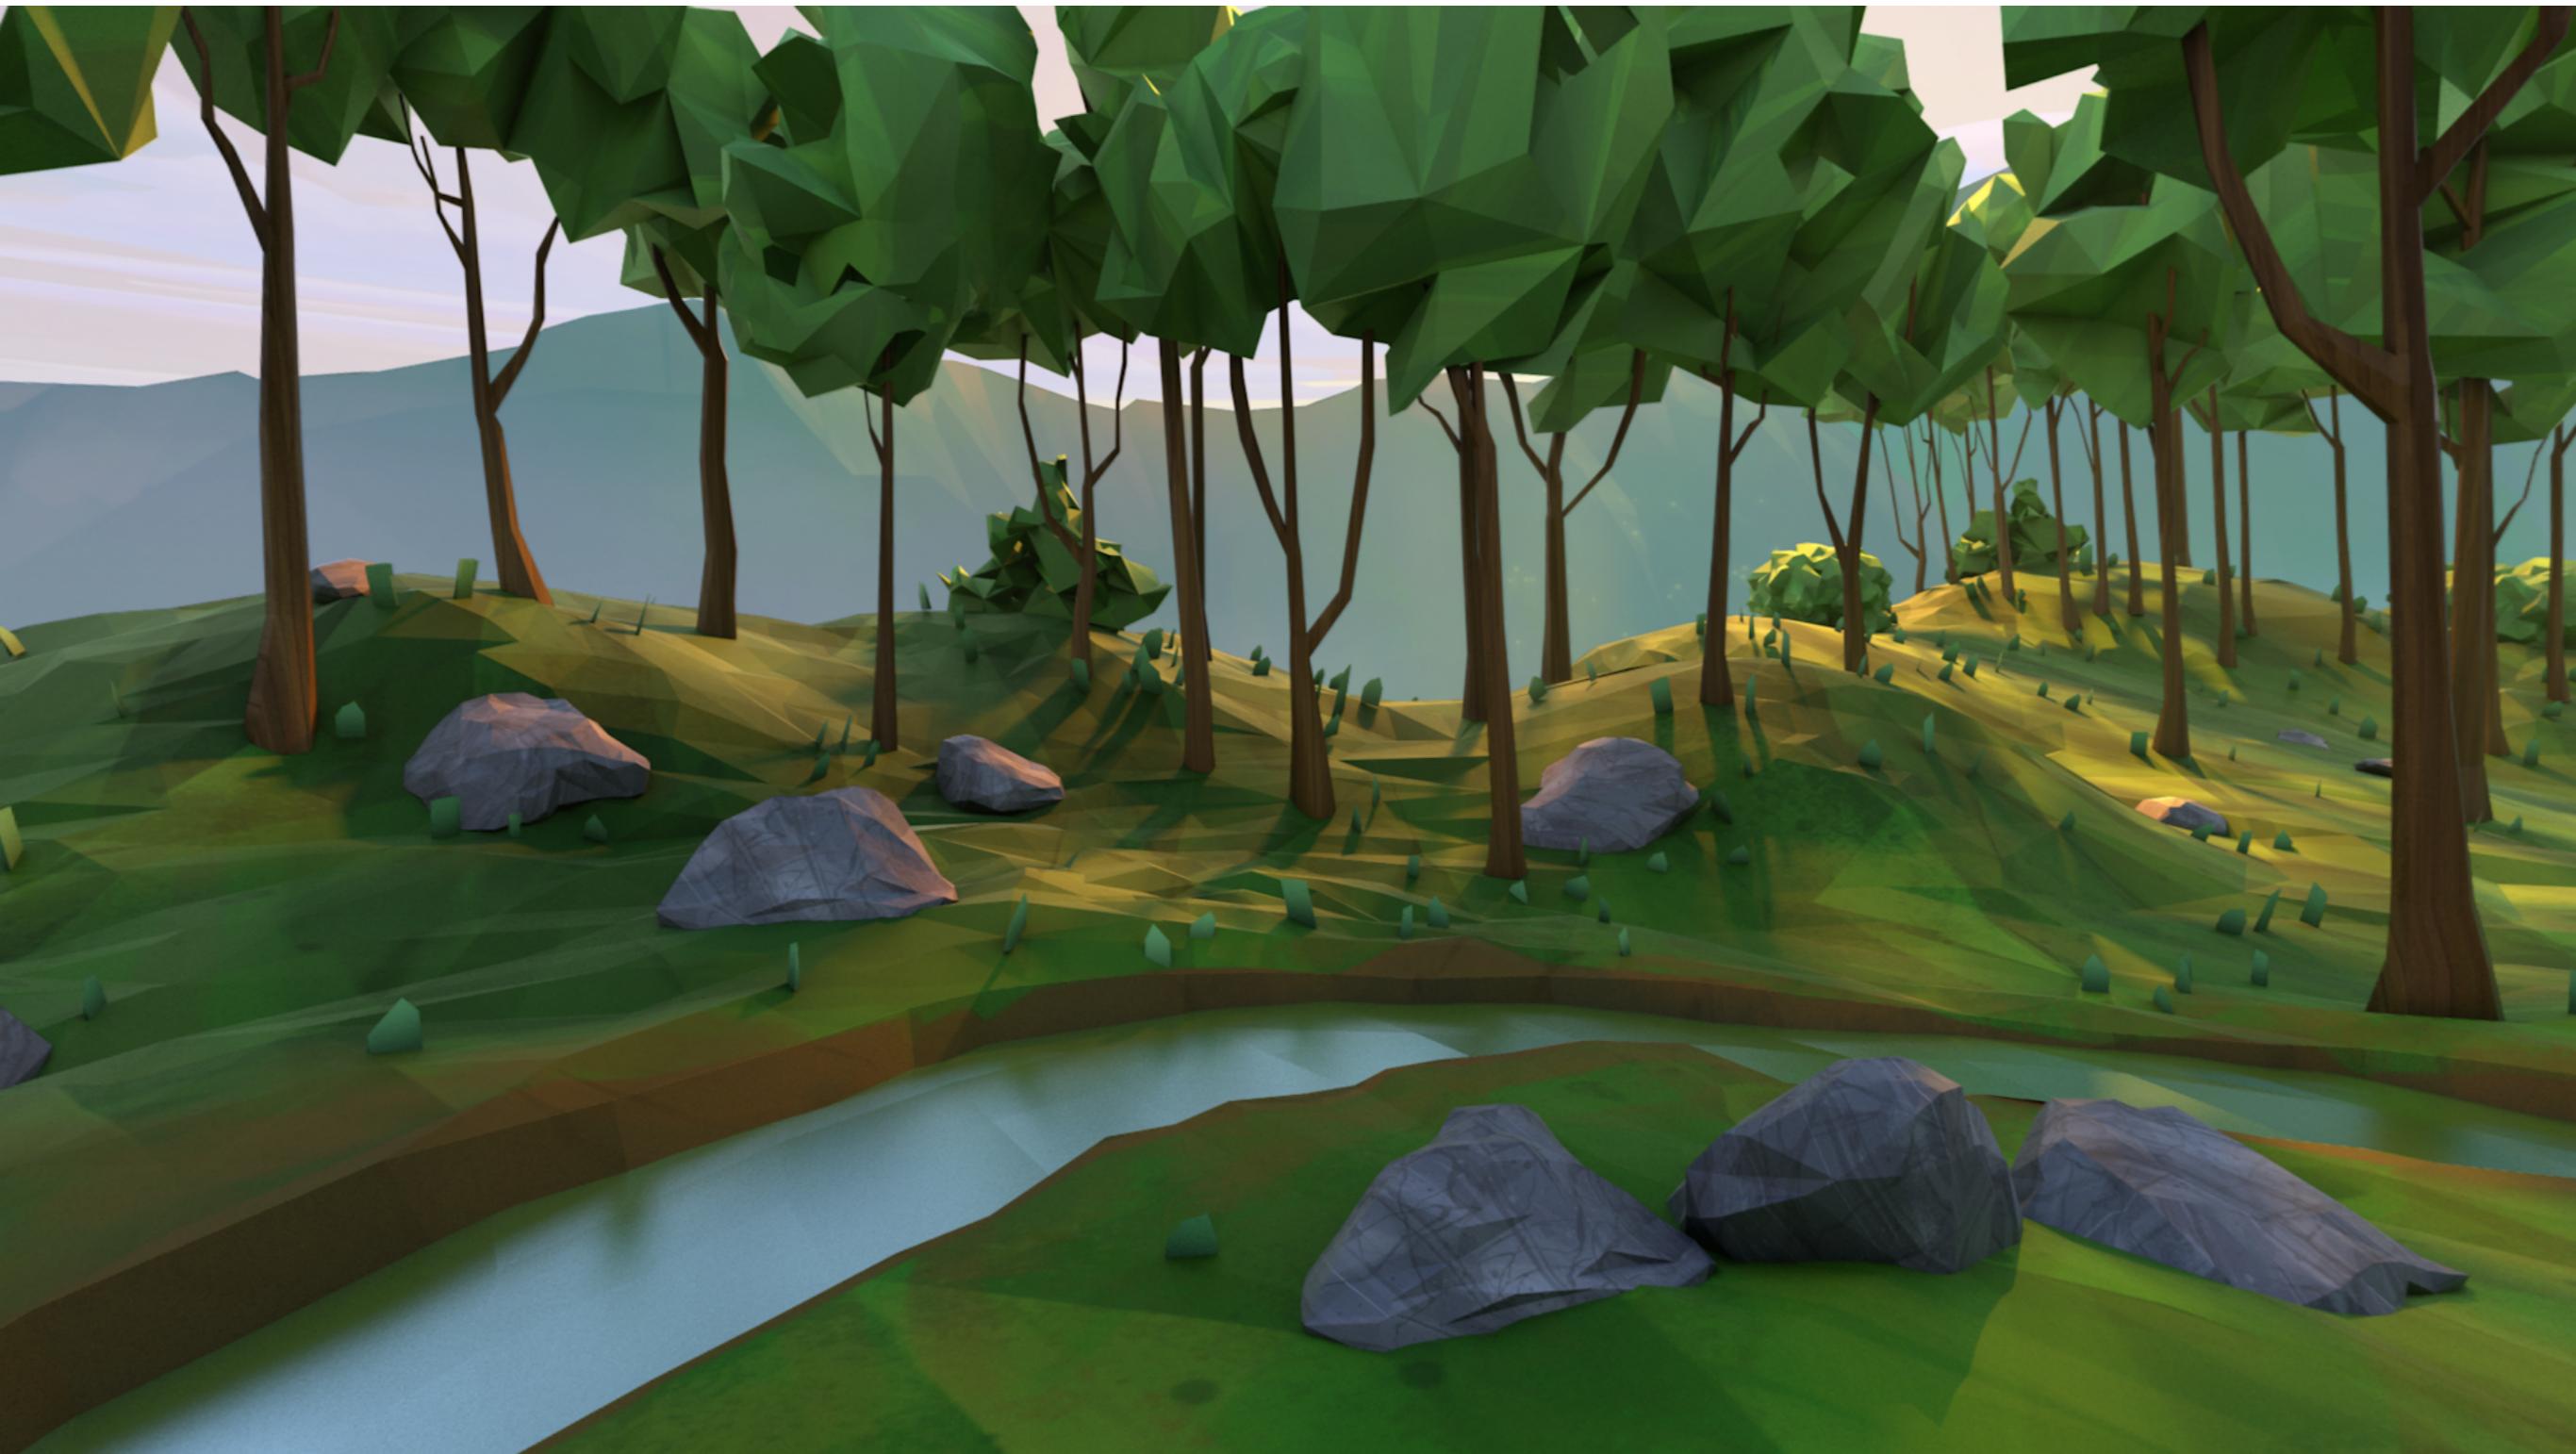
- Reduces presence
- causes VR sickness
- Inefficient interaction



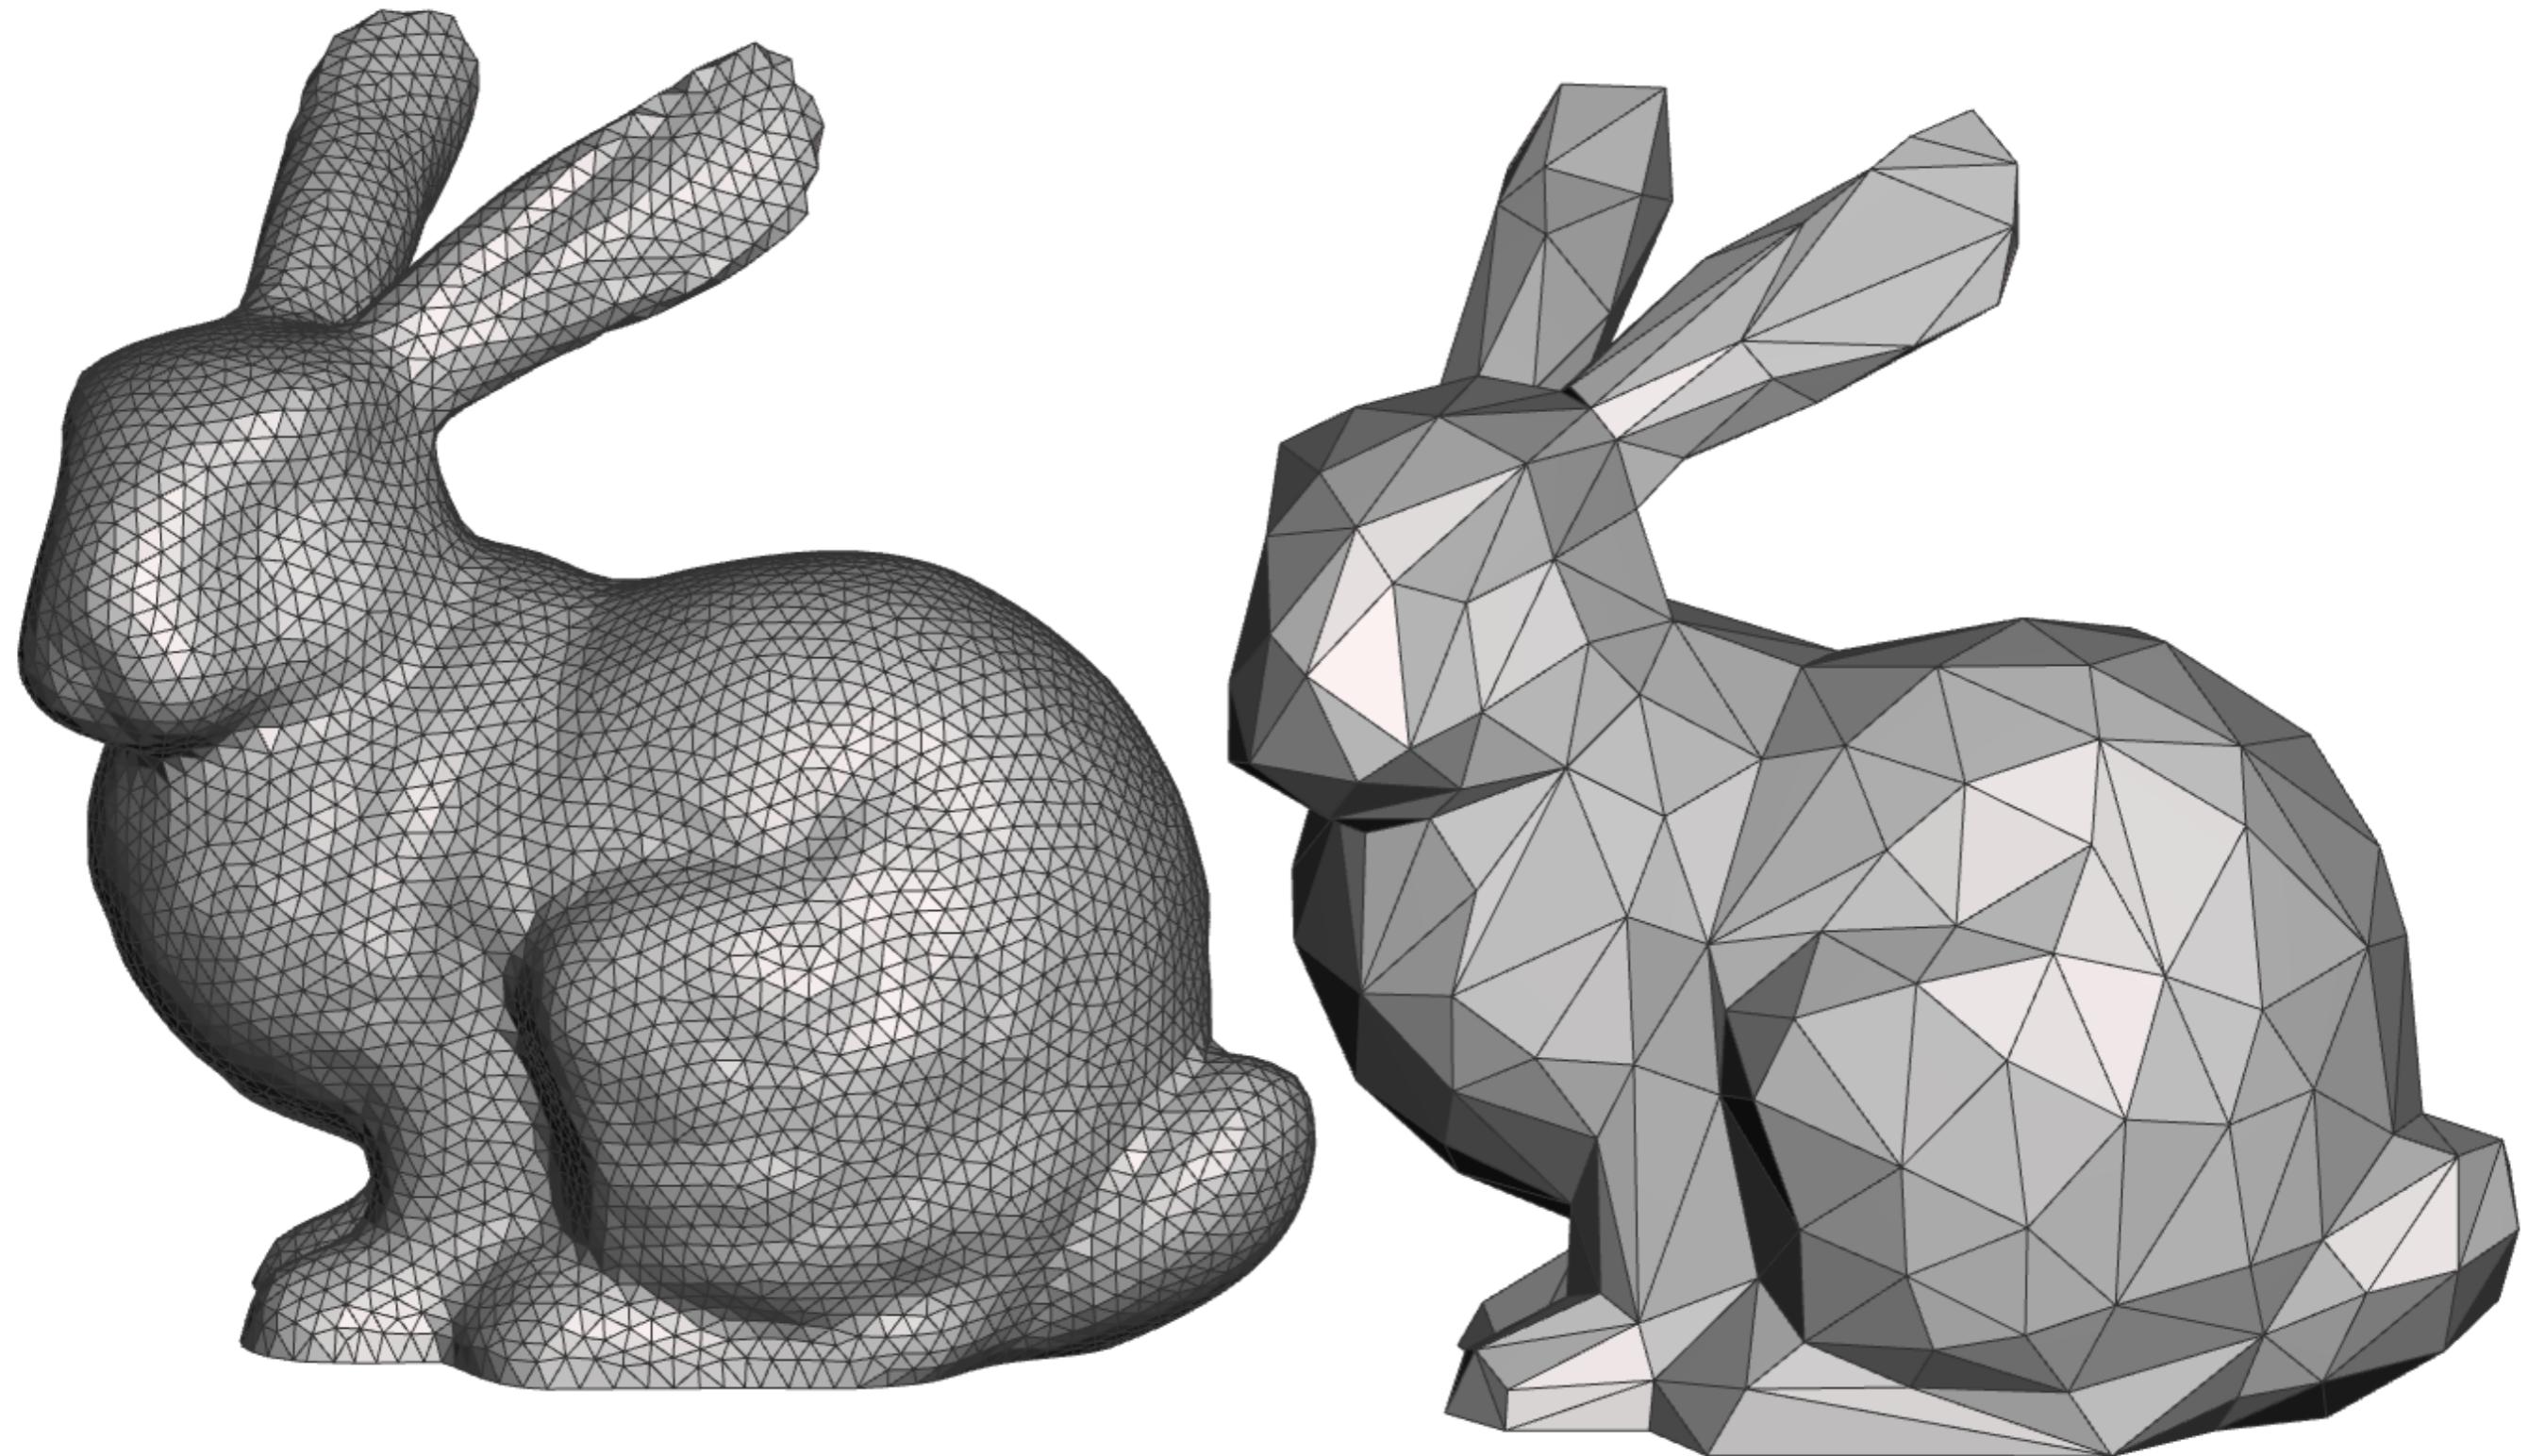
Remedies

- Lower complexity
- Improve rendering pipeline performance
- remove pipeline delays
- user prediction to estimate future viewpoints
- Shift/distort the rendered image to compensate for missing frames.

Simplifying



Simplifying



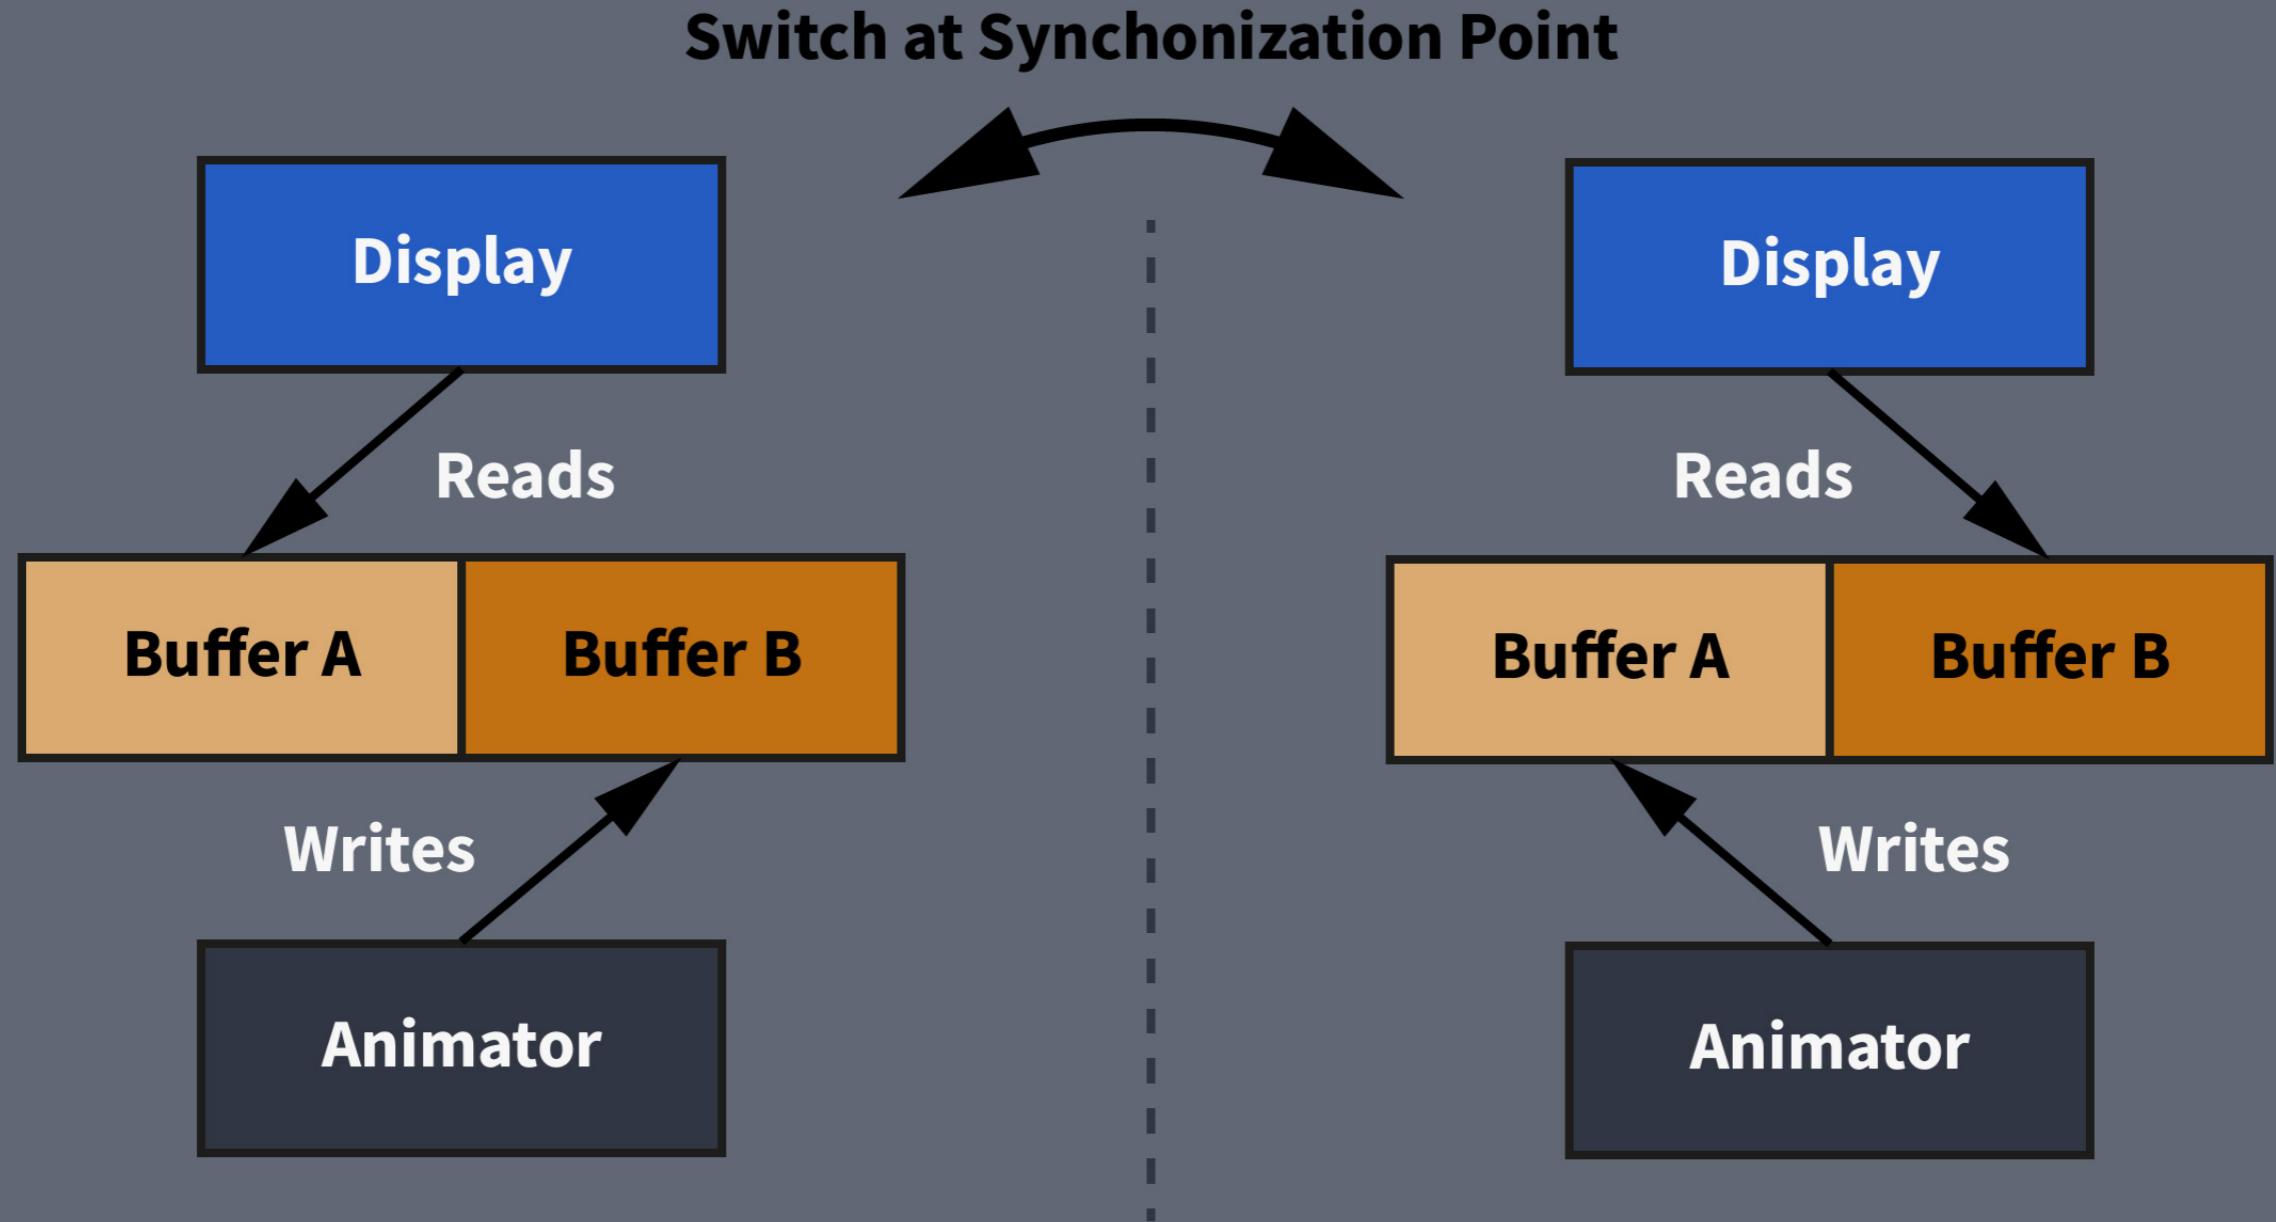
low poly



Image tearing



Double-Buffering



How It Works:

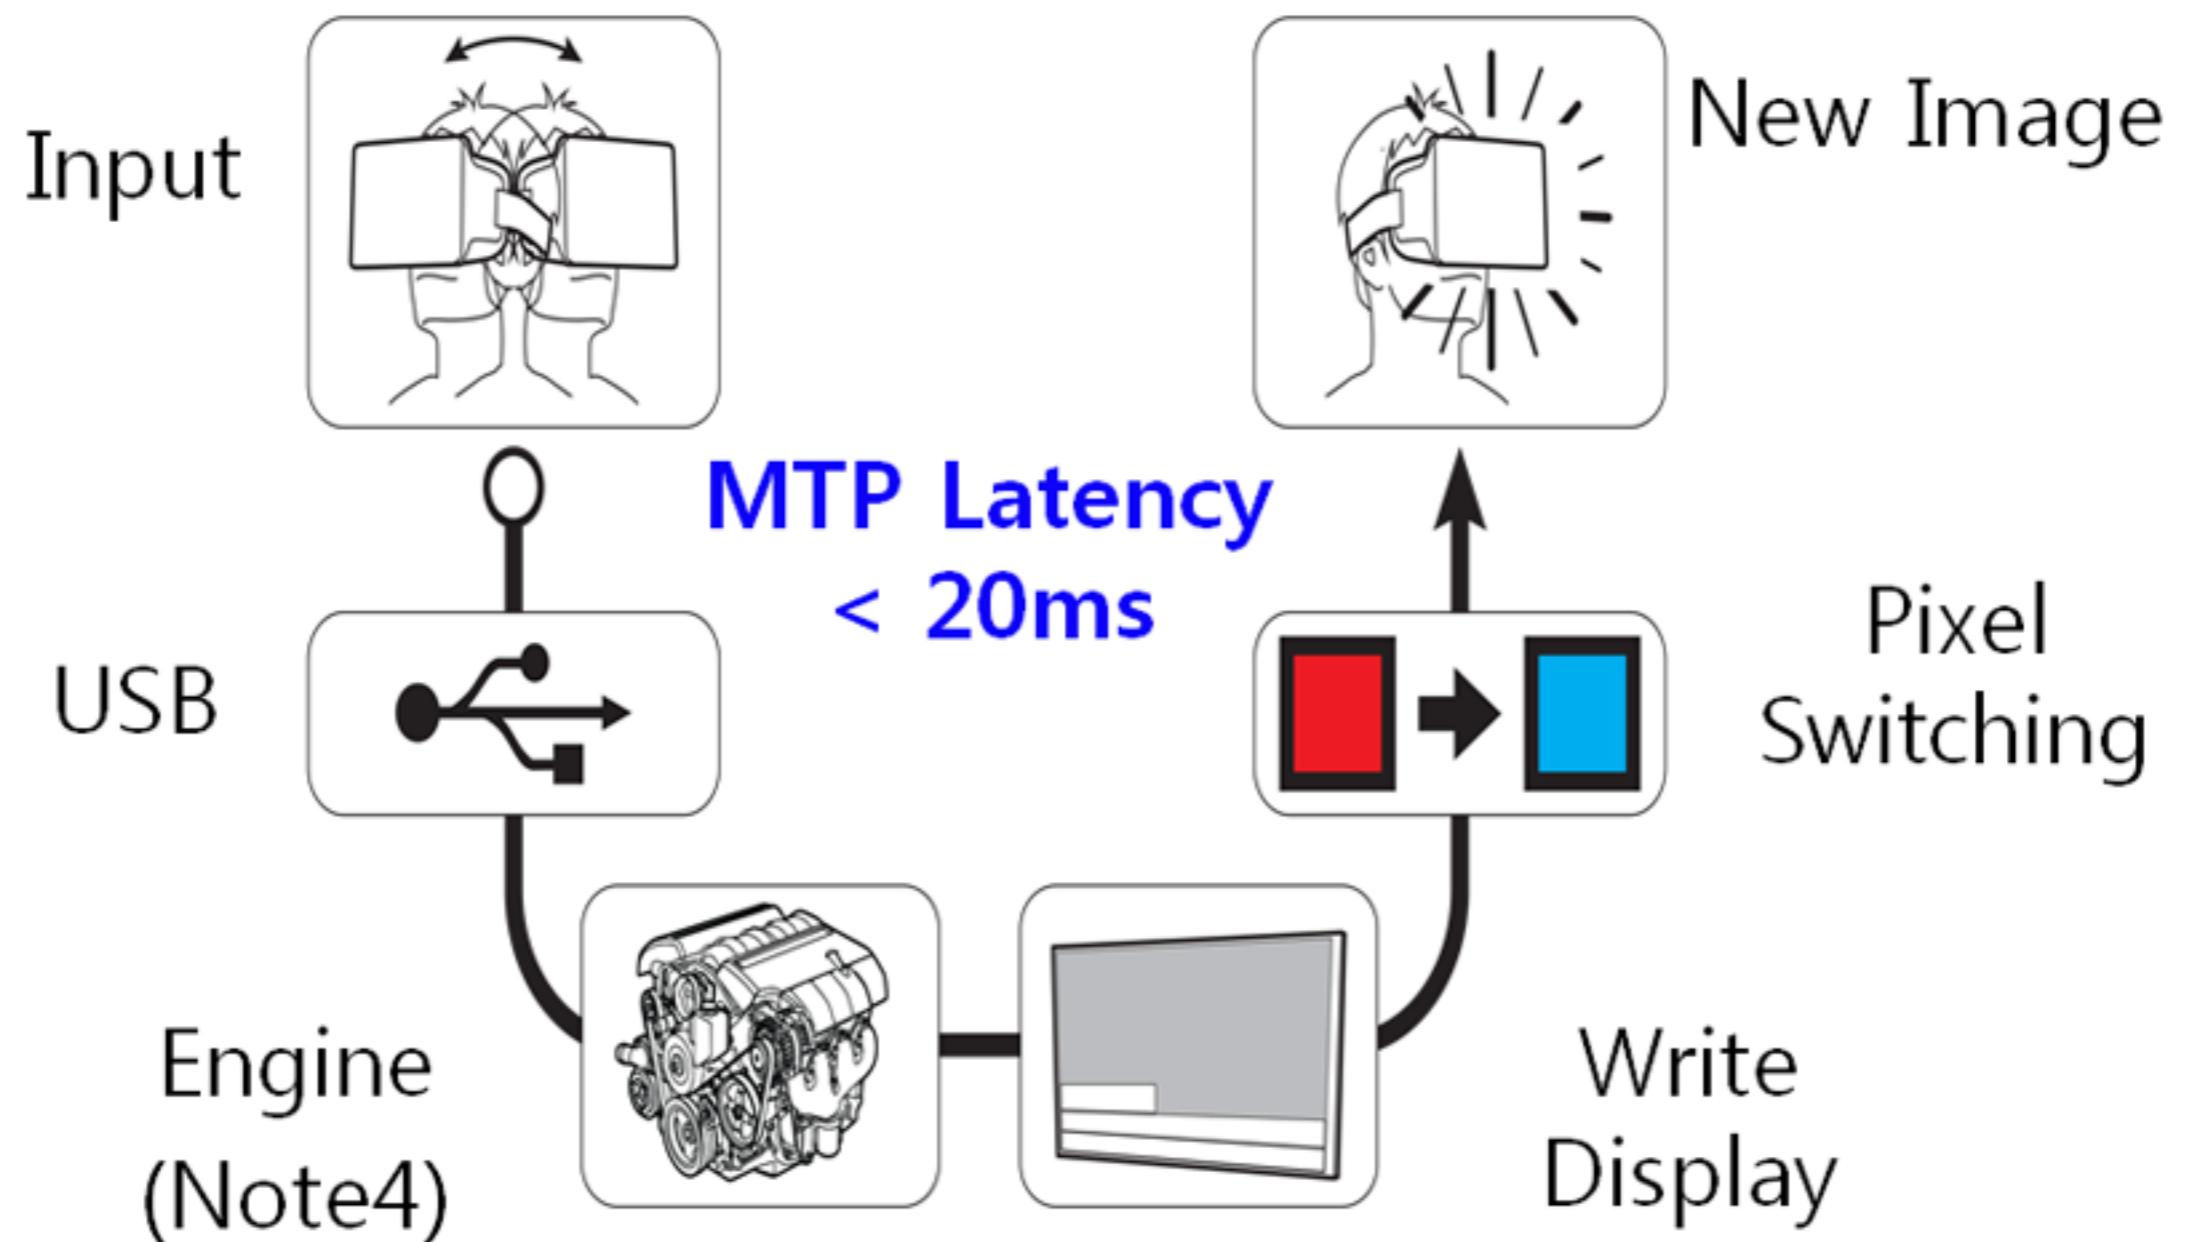
- There are two buffers:
 1. **Front buffer** – The buffer currently being displayed on the screen.
 2. **Back buffer** – The buffer where the next frame is drawn.
- Once the frame is fully rendered in the **back buffer**, the buffers **swap**, making the new frame visible.

Benefits:

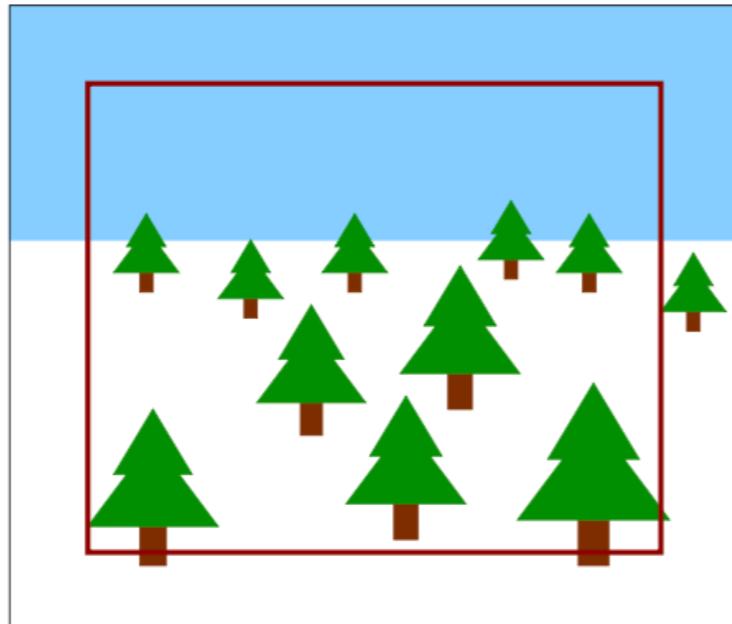
- ✓ Eliminates **flickering** by ensuring only fully rendered frames are displayed.
- ✓ Prevents **tearing** by synchronizing frame updates with the display refresh rate.
- ✓ Smooth animations and improved rendering ✗ increases latency

Prediction

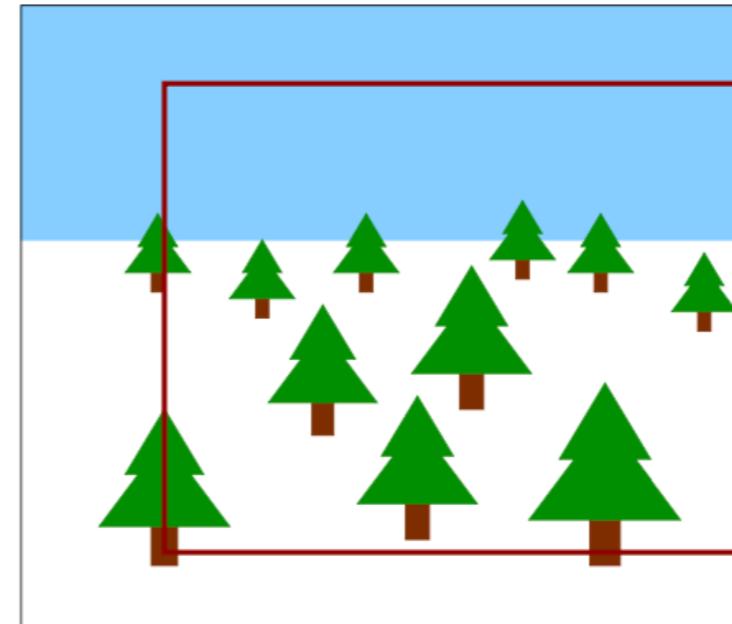
Head has a lot of inertia



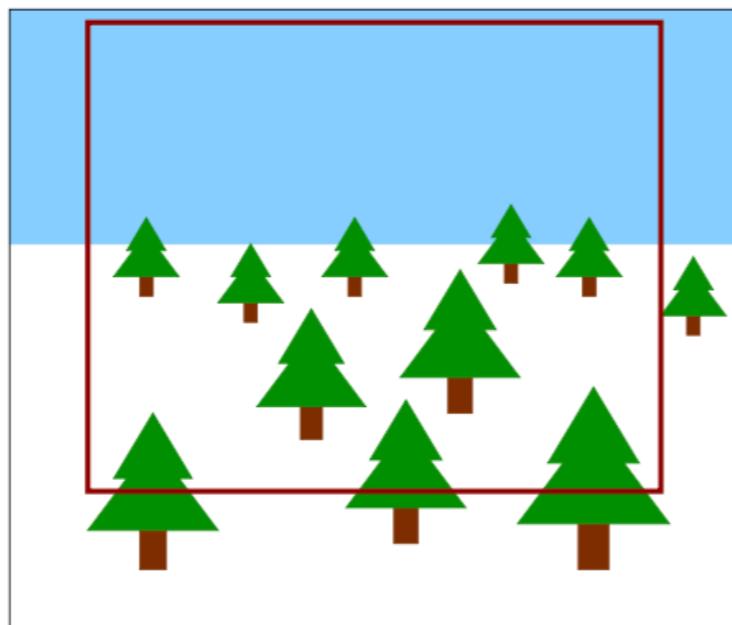
Post rendering Image warping



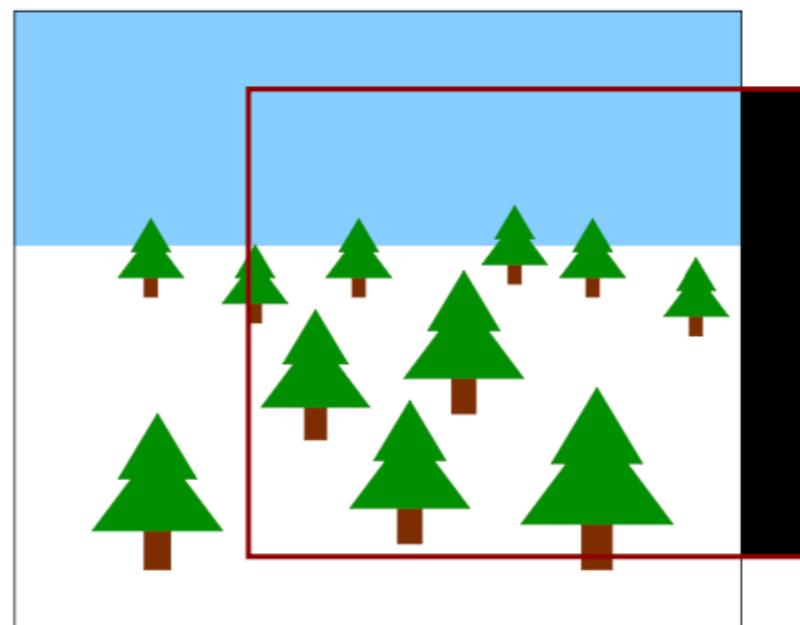
(a)



(b)



(c)

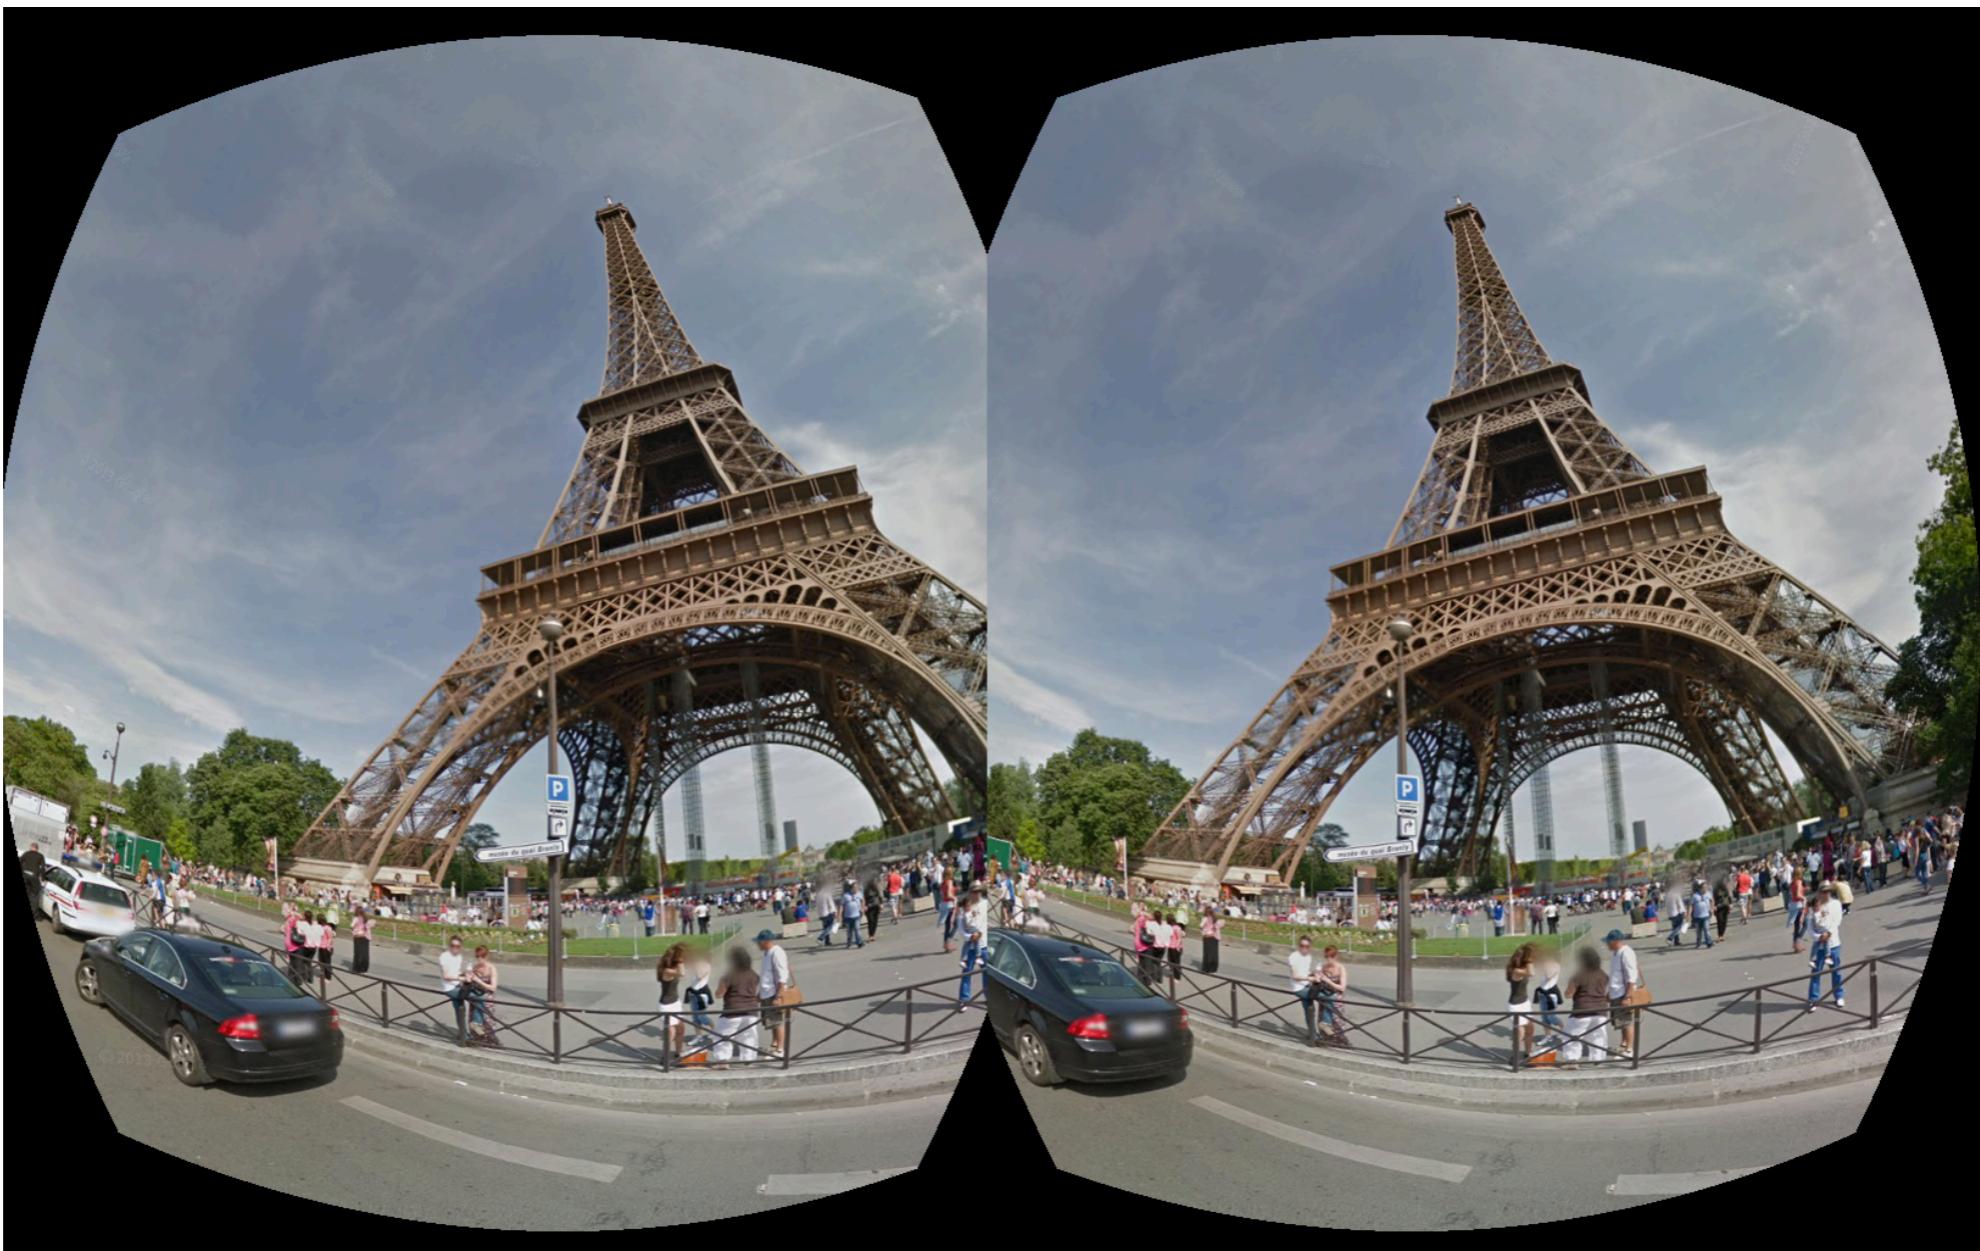


(d)

How It Works:

- The rendered image is **transformed** using geometric operations like **perspective correction**, **distortion removal**, or **viewpoint changes**.
- This is done **without re-rendering** the entire scene, making it computationally efficient.

Immersive photos/videos



VR streetview

Immersive videos



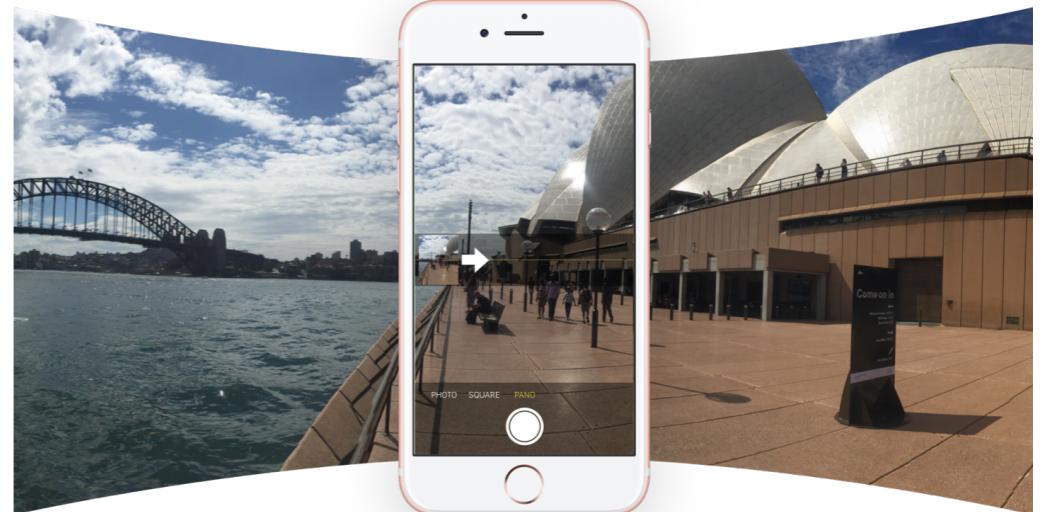
Hulu: [Virtually Mike and Norah](#)

Projection

- Video/photo on VR is form of texture mapping
- Mismatch frame rates (24fps versus 60-96hz)
- VR allows for simulating 3D movies
- Panoramic photo = photosphere
- Panoramic video = videosphere

capturing a wider FOV

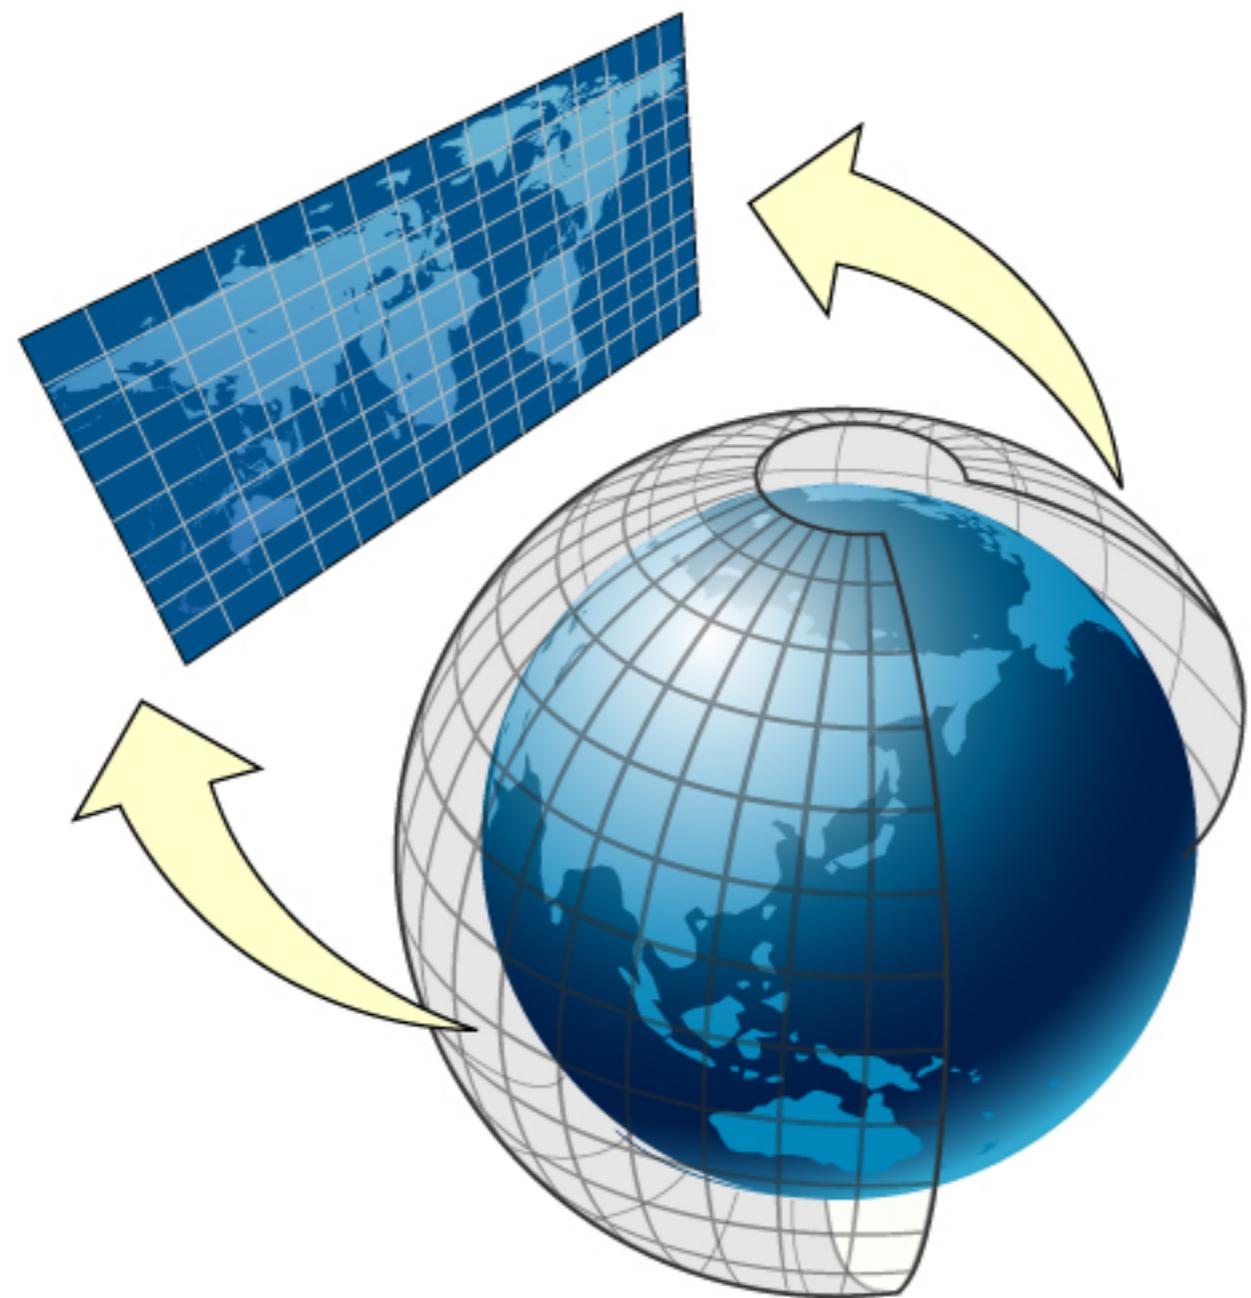
cameras have a limited FOV



Camera rigs



mapping



light field technology

