


IS475/675 Agenda for 04/14/2025

- Present the use of Common Table Expressions (CTE's).
- Compare and contrast Views and CTE's.
- Discuss applications for Views and CTE's – a “group of a group.”
- While waiting for class to start, if you didn't do SQL Lab Exercise 8 or attend class on Wednesday(04/10/2025) then execute SQL Server Management Studio and run this script file:
K:\cob\is475\labfiles\SQLLab8.sql

We are creating more complex queries



Simple queries usually:

Generate large result tables with relatively simple filtering operations.

Tend to require only one transaction table.

Do not require significant changes to the structure of the data; one row in the transaction table produces one filtered row in the result table.



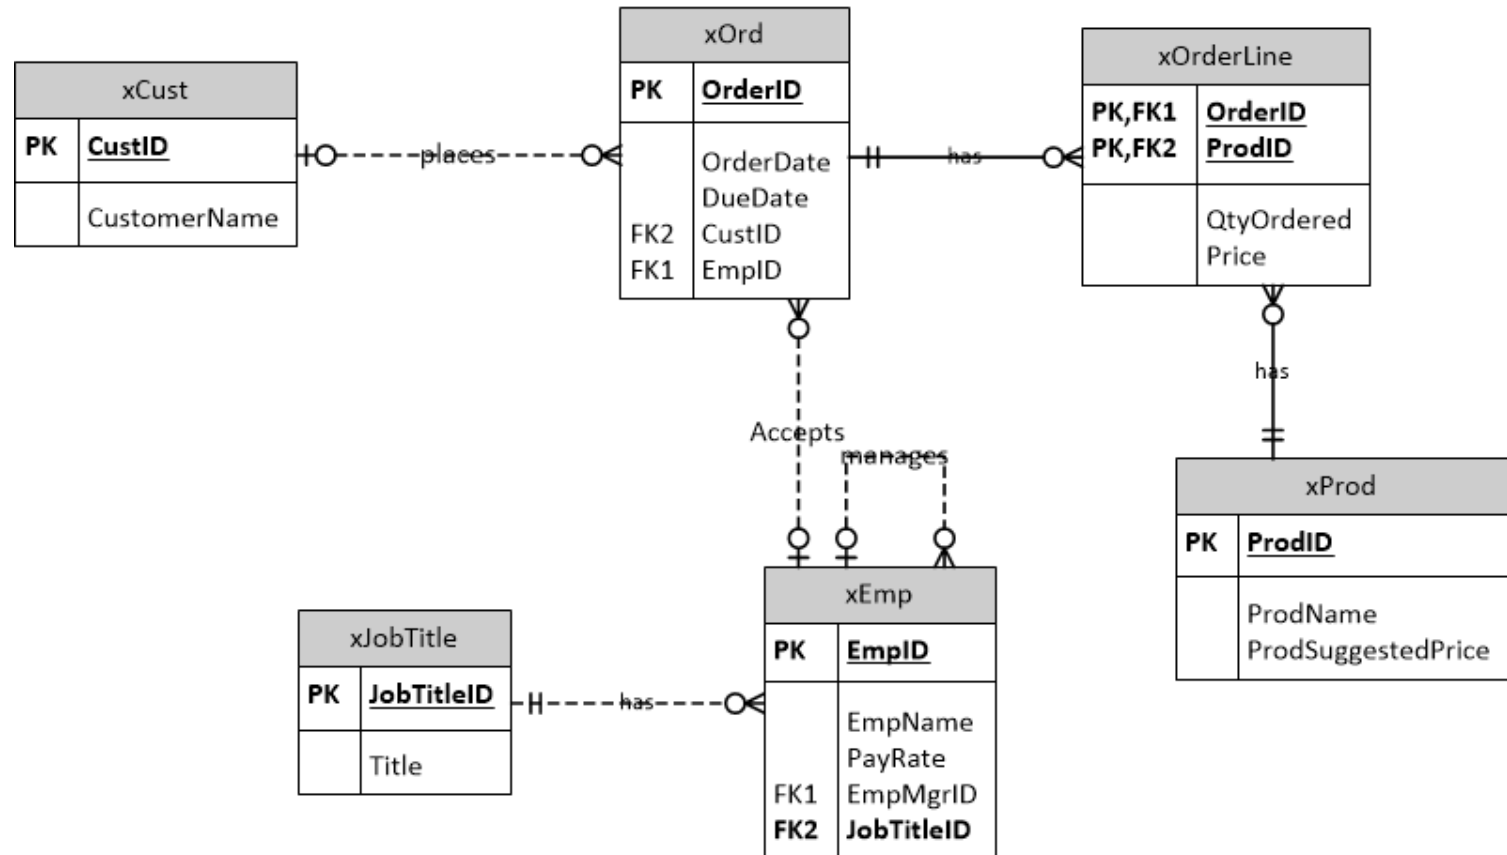
More complex queries can:

Require a combination of joins, group functions, and sub-queries.

Return one row per group because they make greater use of group functions.

Provide direct information for decision makers.

Remember the database design for exercises 7 & 8



Look at the content of the tables

- SELECT * FROM xEmp;
- SELECT * FROM xProd;
- SELECT * FROM xOrd;
- SELECT * FROM xOrderLine;
- SELECT * FROM xCust
- SELECT * FROM xJobTitle;

Recap: Which employees have a payrate than is higher than the average payrate for their job title?

	EmpID	EmpName	PayRate	Title	AveragePayRate
1	2	Polanski	45.00	Database Designer	35.00
2	3	Torquez	85.00	Manager	80.00
3	4	Ling	65.00	Interface Programmer	55.00
4	6	Martinez	35.00	Web Programmer	33.00
5	9	Fukamota	40.00	Web Programmer	33.00
6	11	Nguyen	35.00	Web Programmer	33.00
7	12	Duong	28.00	Business Analyst	25.60
8	13	Patel	30.00	Business Analyst	25.60

Recap - we can simplify queries with the use of SQL Views.

What is a SQL view?

- A “virtual” table.
 - A set of SQL statements that creates a result table which can be accessed by other SQL statements.
- A database object.
 - The code for a view is stored in the database.
 - A view contains no data of its own.
 - A view relies on the data in the base tables used to create the view.
- A set of stored SQL code.
 - Stores code; not data.

We created a SQL View to solve the problem
(from SQL Lab Exercise 8, Task 4, pg. 8)

```
CREATE VIEW vAvg AS
SELECT      jobtitleID,
            AVG(payrate) AS AveragePayRate,
            MAX(payrate) AS MaxPayRate,
            MIN(payrate) AS MinPayRate
FROM        xemp
GROUP BY    jobtitleID;
```

Join the view to the xemp table

Manages

xEmp	
PK	<u>EmpID</u>
FK2	EmpName PayRate EmpMgrID JobTitleID

vAvg	
	JobTitleID AveragePayRate MaxPayRate MinPayRate

is related to

The view does not have a primary key because it does not contain data. As long as you include a field that can be used to join (JobTitleID in this situation), then the view can be joined to other tables or other views.

```
SELECT *
FROM xemp
INNER JOIN vAvg
ON xemp.jobtitleid =
vAvg.jobtitleid
```

	empid	empname	payrate	empmgrid	jobtitleid	jobtitleid	AveragePayRate	MaxPayRate	MinPayRate
1	1	Martinson	75.00	NULL	10	10	80.00	85.00	75.00
2	3	Torquez	85.00	1	10	10	80.00	85.00	75.00
3	12	Duong	28.00	2	20	20	25.60	30.00	22.50
4	13	Patel	30.00	2	20	20	25.60	30.00	22.50
5	14	Agarwal	25.00	2	20	20	25.60	30.00	22.50
6	15	Anand	22.50	2	20	20	25.60	30.00	22.50
7	16	Smith	22.50	3	20	20	25.60	30.00	22.50
8	2	Polanski	45.00	1	40	40	35.00	45.00	25.00
9	7	Johnson	25.00	3	40	40	35.00	45.00	25.00
10	5	Bassett	25.00	1	45	45	33.00	40.00	25.00
11	6	Martinez	35.00	1	45	45	33.00	40.00	25.00
12	9	Fukamota	40.00	3	45	45	33.00	40.00	25.00
13	10	Stein	30.00	1	45	45	33.00	40.00	25.00
14	11	Nguyen	35.00	3	45	45	33.00	40.00	25.00
15	8	Cheng	45.00	1	50	50	55.00	65.00	45.00
16	4	Ling	65.00	3	50	50	55.00	65.00	45.00

Filter the rows and sort the result table

```
SELECT *  
FROM xemp  
INNER JOIN vAvg  
ON xemp.jobtitleid = vAvg.jobtitleid  
WHERE Payrate > AveragePayrate  
ORDER BY empid
```

	empid	empname	payrate	empmgrid	jobtitleid	jobtitleid	AveragePayRate	MaximumPayRate	MinimumPayRate
1	2	Polanski	45.00	1	40	40	35.00	45.00	25.00
2	3	Torquez	85.00	1	10	10	80.00	85.00	75.00
3	4	Ling	65.00	3	50	50	55.00	65.00	45.00
4	6	Martinez	35.00	1	45	45	33.00	40.00	25.00
5	9	Fukamota	40.00	3	45	45	33.00	40.00	25.00
6	11	Nguyen	35.00	3	45	45	33.00	40.00	25.00
7	12	Duong	28.00	2	20	20	25.60	30.00	22.50
8	13	Patel	30.00	2	20	20	25.60	30.00	22.50

Add the job title and SELECT the columns

```
SELECT      xemp.empid, xemp.EmpName, xemp.PayRate,  
            jt.Title,  
            vAvg.AveragePayrate  
FROM        xemp  
INNER JOIN  vAvg  
ON xemp.jobtitleid = vAvg.jobtitleid  
INNER JOIN  xJobTitle jt  
on xemp.jobtitleid = jt.jobtitleid  
WHERE Payrate > AveragePayrate  
ORDER BY empid
```

	empid	EmpName	PayRate	Title	AveragePayrate
1	2	Polanski	45.00	Database Designer	35.00
2	3	Torquez	85.00	Manager	80.00
3	4	Ling	65.00	Interface Programmer	55.00
4	6	Martinez	35.00	Web Programmer	33.00
5	9	Fukamota	40.00	Web Programmer	33.00
6	11	Nguyen	35.00	Web Programmer	33.00
7	12	Duong	28.00	Business Analyst	25.60
8	13	Patel	30.00	Business Analyst	25.60

What is a Common Table Expression (CTE)?

- A CTE is much like a view.
- A CTE creates a named virtual result table, just like a view.
- A CTE, however, is not a database object – it is only available in the session that is actively using the code.
- It is a temporary virtual result table, while a view is a more permanent virtual result table.
- A CTE is not ANSI-standard. It is available in MS SQL Server T-SQL.

Do the same thing with a CTE

```
WITH cteAvgRate AS
(SELECT jobtitleid,
      avg(payrate) AveragePayRate
 FROM   xemp
 GROUP BY jobtitleid)

SELECT emp.empid,
       emp.empname,
       emp.payrate,
       Title,
       cteAvgRate.AveragepayRate
FROM   xemp emp
inner join      cteAvgRate
ON      emp.jobtitleid = cteAvgRate.jobtitleid
inner join      xJobTitle jt
ON      emp.jobtitleid = jt.jobtitleid
WHERE   payrate > averagepayrate
ORDER BY 1
```

View vs. CTE

	View	Common Table Expression
Create/Store	Stored as a database object.	Not stored as a database object. Local to a single query.
Extent of Use	Use when the result table will be used in more than one query.	Use when the result table is local to single query.
Portability	Can be accessed by programs other than SQL.	Can only be used by SQL.

Moving on! Which customer(s) placed the most orders with our company based on a count of the orders?

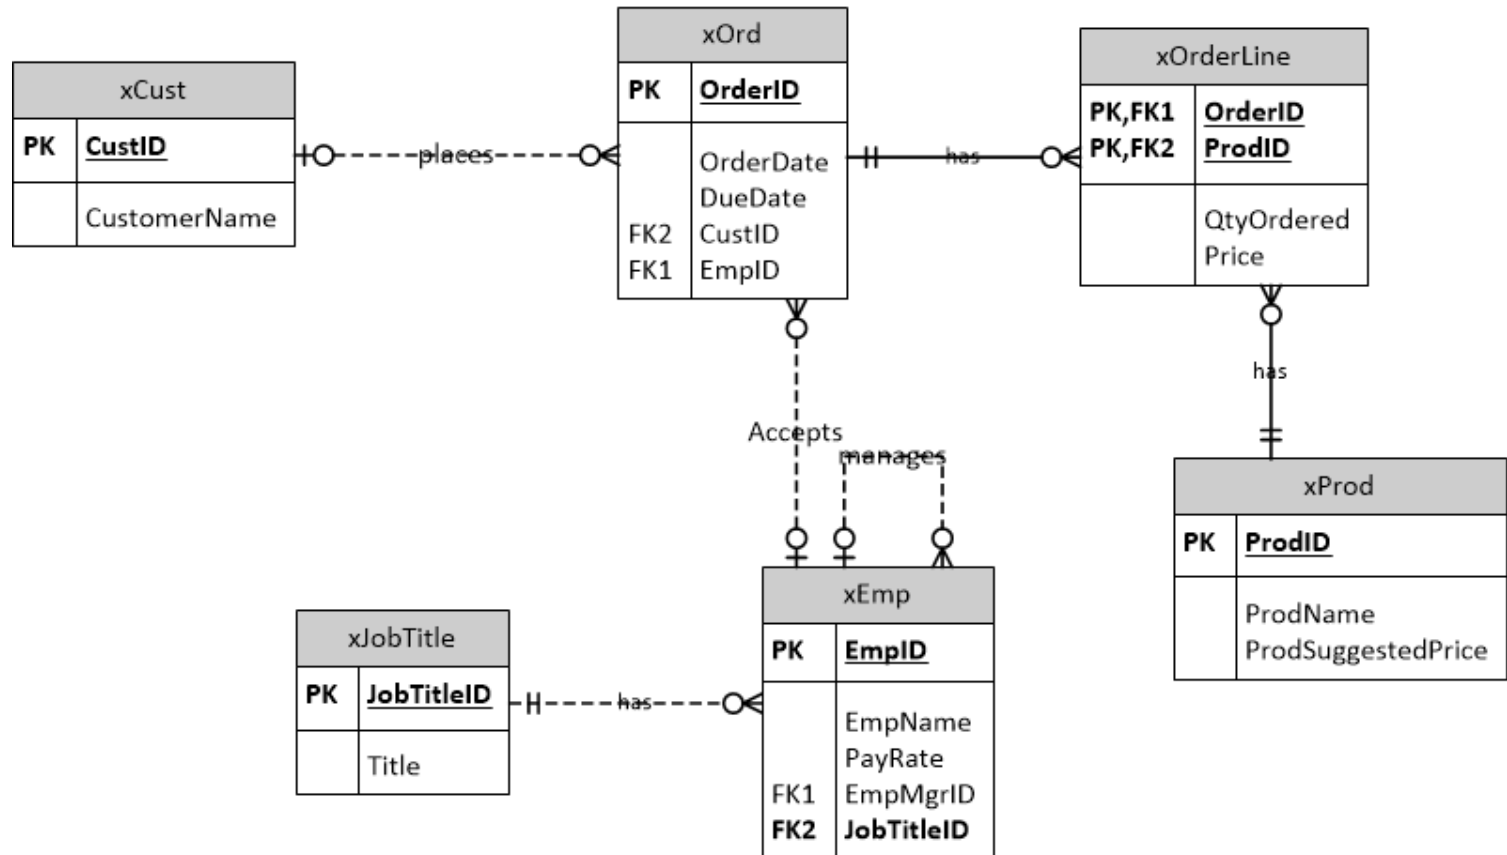
	CustID	CustomerName	CountOfOrders
1	2555	Mountain Design	3
2	6899	Opaka Sporting Goods	3

Pseudocode:

```
SELECT customer data
FROM cust, ord
WHERE COUNT(orderID) = MAX(COUNT(orderID))
```

This is a “group of a group” – in this example it is a MAX of a COUNT

Please note that this is not possible. First, a group function cannot be included in the WHERE clause. Second, it is not possible to nest GROUP functions.



Let's explore the problem – what are we counting?

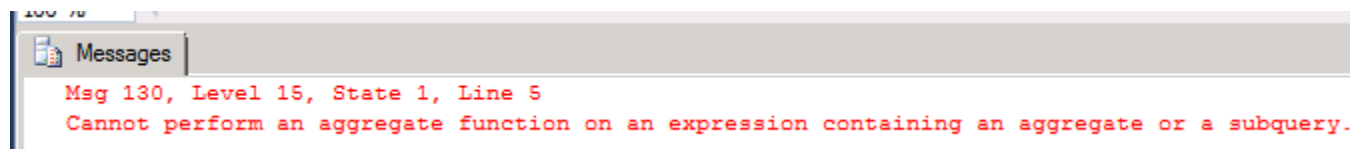
Example of “playing with code” to get an idea of the logic:

```
SELECT custID,  
       count(*) CountOfOrders  
FROM   xOrd  
GROUP BY custID
```

	custID	CountOfOrders
1	1234	1
2	2555	3
3	6773	2
4	6899	3
5	8372	2

How do you see just the customers with the most orders? Especially if you don't know how many there are? Not this way!

```
SELECT custID,  
       count(*) CountOfOrders  
FROM   xOrd  
GROUP BY custID  
HAVING count(*) = max(count(*))
```



Let's use a View and a sub-query for the basic logic

```
CREATE VIEW vCountOrders AS
SELECT      custID,
            count(*)    CountOfOrders


FROM        xOrd
GROUP BY    custID;
```

	custID	CountOfOrders
1	1234	1
2	2555	3
3	6773	2
4	6899	3
5	8372	2

```
SELECT      *
FROM        vCountOrders vCount
WHERE       countoforders =
            (SELECT MAX(CountOfOrders)
             FROM    vCountOrders);
```

	CustID	CountOfOrders
1	2555	3
2	6899	3

Join in the customer table to see the customer name



```
SELECT *  
FROM xCust cust  
INNER JOIN vCountOrders vCount  
ON cust.custid = vCount.custID  
WHERE countoforders =  
(SELECT MAX(CountOfOrders)  
FROM vCountOrders);
```

	CustID	CustomerName	custID	CountOfOrders
1	2555	Mountain Design	2555	3
2	6899	Opaka Sporting Goods	6899	3

Choose the columns you want to display

```
SELECT      vCount.CustID,  
            cust.CustomerName,  
            vCount.CountOfOrders  
FROM        xCust cust  
INNER JOIN  vCountOrders vCount  
ON          cust.custid = vCount.custID  
WHERE       countoforders =  
            (SELECT MAX(CountOfOrders)  
             FROM   vCountOrders);
```

	CustID	CustomerName	CountOfOrders
1	2555	Mountain Design	3
2	6899	Opaka Sporting Goods	3

Must separate the group function of a COUNT from the group function of a MAX. Let's use a CTE and a sub-query to accomplish the same goal as the VIEW in the prior slide.

```
WITH cteCountOrders AS
(SELECT      custID,
             count(*)    CountOfOrders
FROM        xOrd
GROUP BY    custID)

SELECT      cust.CustID,
             CustomerName,
             CountOfOrders
FROM        xCust cust
INNER JOIN  cteCountOrders cteCount
ON          cust.custid = ctecount.custID
WHERE       countoforders =
             (SELECT MAX(CountOfOrders)
              FROM    cteCountOrders);
```

	CustID	CustomerName	CountOfOrders
1	2555	Mountain Design	3
2	6899	Opaka Sporting Goods	3

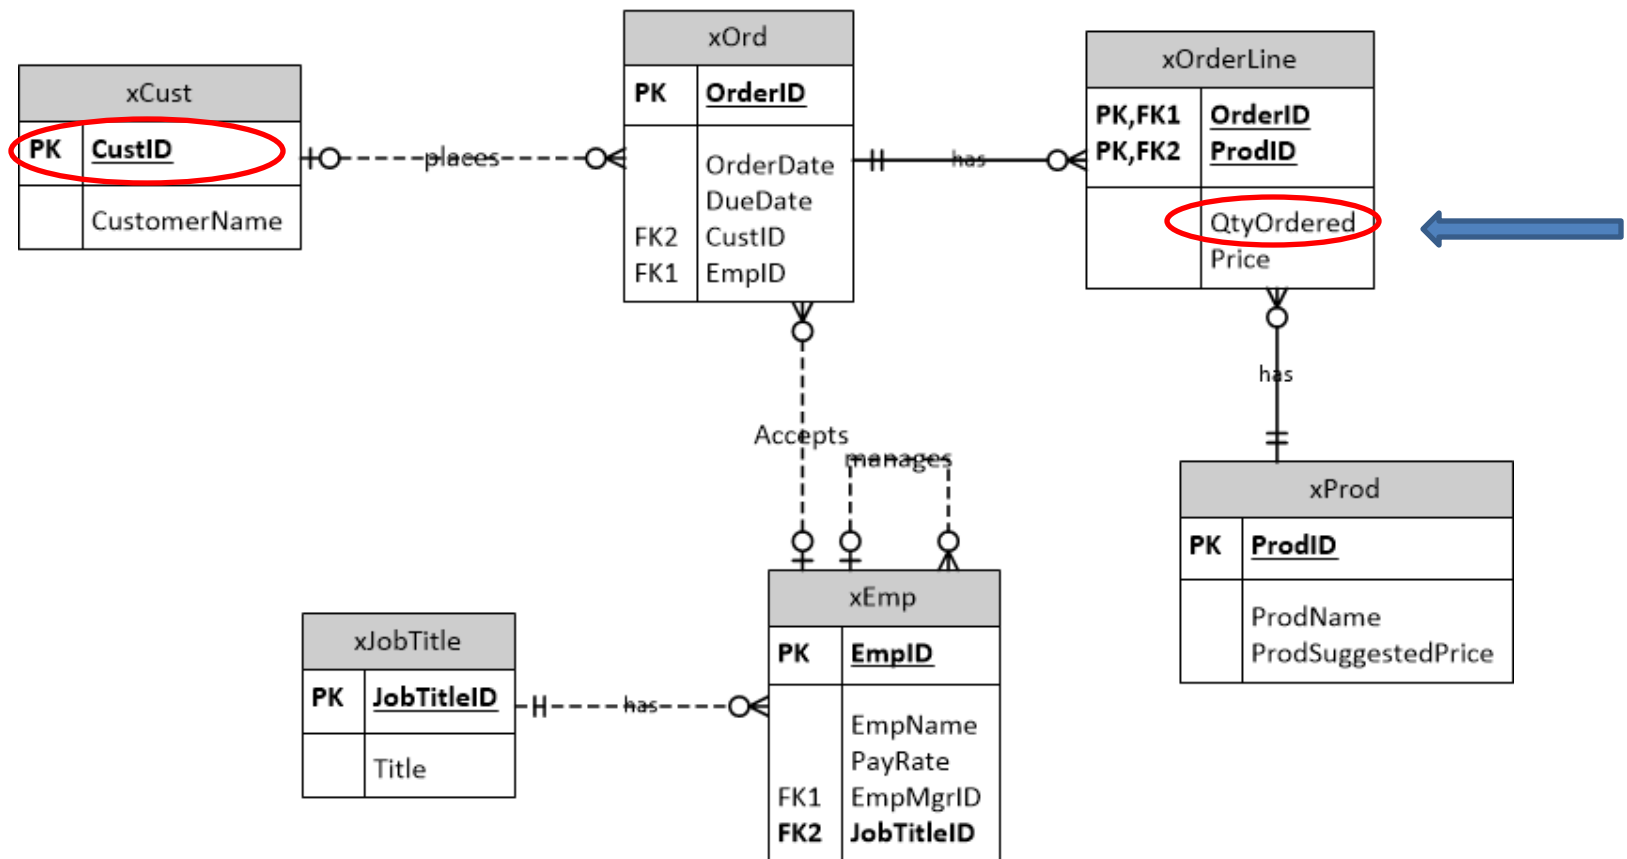
New query: Which customer bought the most items from us based on the **quantity** of items purchased?

	CustID	CustomerName	TotalQtyOrdered
1	6899	Opaka Sporting Goods	645.250

Where do the columns come from (which tables)?

What is the basic logic of the query (which table or tables are necessary to find the required rows in the result table)?

What are the simplest requirements necessary to accomplish the basic logic?



Let's explore the problem – what are we adding up?

Example of “playing with code” to get an idea of the logic:

```
SELECT *  
FROM xOrderline  
INNER JOIN xOrd  
ON xOrd.orderid = xOrderline.orderid
```

This requires the data stored in two different tables – the Orderline table for the qtyOrdered by product, and then the ord table to access the customer who purchased the product.

How do you sum up the qtyOrdered by customer?

```
SELECT custID,  
       sum(qtyOrdered) TotalQtyOrdered  
FROM xOrderline  
INNER JOIN xOrd  
ON xOrd.orderid = xOrderline.orderid  
GROUP BY custid
```

	custID	TotalQtyOrdered
1	1234	37.560
2	2555	88.550
3	6773	13.000
4	6899	645.250
5	8372	31.000

“Replace” the xOrderLine and xOrd tables with a VIEW


```
CREATE VIEW vCustQty AS
SELECT custID,
       sum(qtyOrdered) TotalQtyOrdered
FROM xOrderline
INNER JOIN xOrd
ON xOrd.orderid = xOrderline.orderid
GROUP BY custid
```

	custID	TotalQtyOrdered
1	1234	37.560
2	2555	88.550
3	6773	13.000
4	6899	645.250
5	8372	31.000

Test out the basic logic

```
SELECT      *
FROM        vCustQty
WHERE       TotalqtyOrdered =
           (SELECT MAX(TotalqtyOrdered)
            FROM   vCustQty);
```

	CustID	TotalQtyOrdered
1	6899	645.250



SELECT

vCustQty.CustID,
cust.CustomerName,
vCustQty.TotalQtyOrdered

FROM

xCust Cust

INNER JOIN

vCustQty

ON

cust.custid = vCustQty.custID

WHERE

TotalQtyOrdered =
(SELECT MAX(TotalQtyOrdered)
FROM vCustQty);

Join the view with the
Cust table to access the
customer name

Do the same thing, except with a CTE instead of a VIEW

```
WITH cteSumqtyOrdered AS
(SELECT custID,
       sum(qtyOrdered) TotalqtyOrdered
FROM xOrderline
INNER JOIN xOrd
ON xOrd.orderid = xOrderline.orderid
GROUP BY custid
)

SELECT      ctesq.CustID,
            cust.CustomerName,
            ctesq.TotalqtyOrdered
FROM        xCust Cust
INNER JOIN  cteSumqtyOrdered as ctesq
ON          cust.custid = ctesq.custID
WHERE      TotalqtyOrdered =
            (SELECT MAX(TotalqtyOrdered)
             FROM   ctesumqtyOrdered);
```

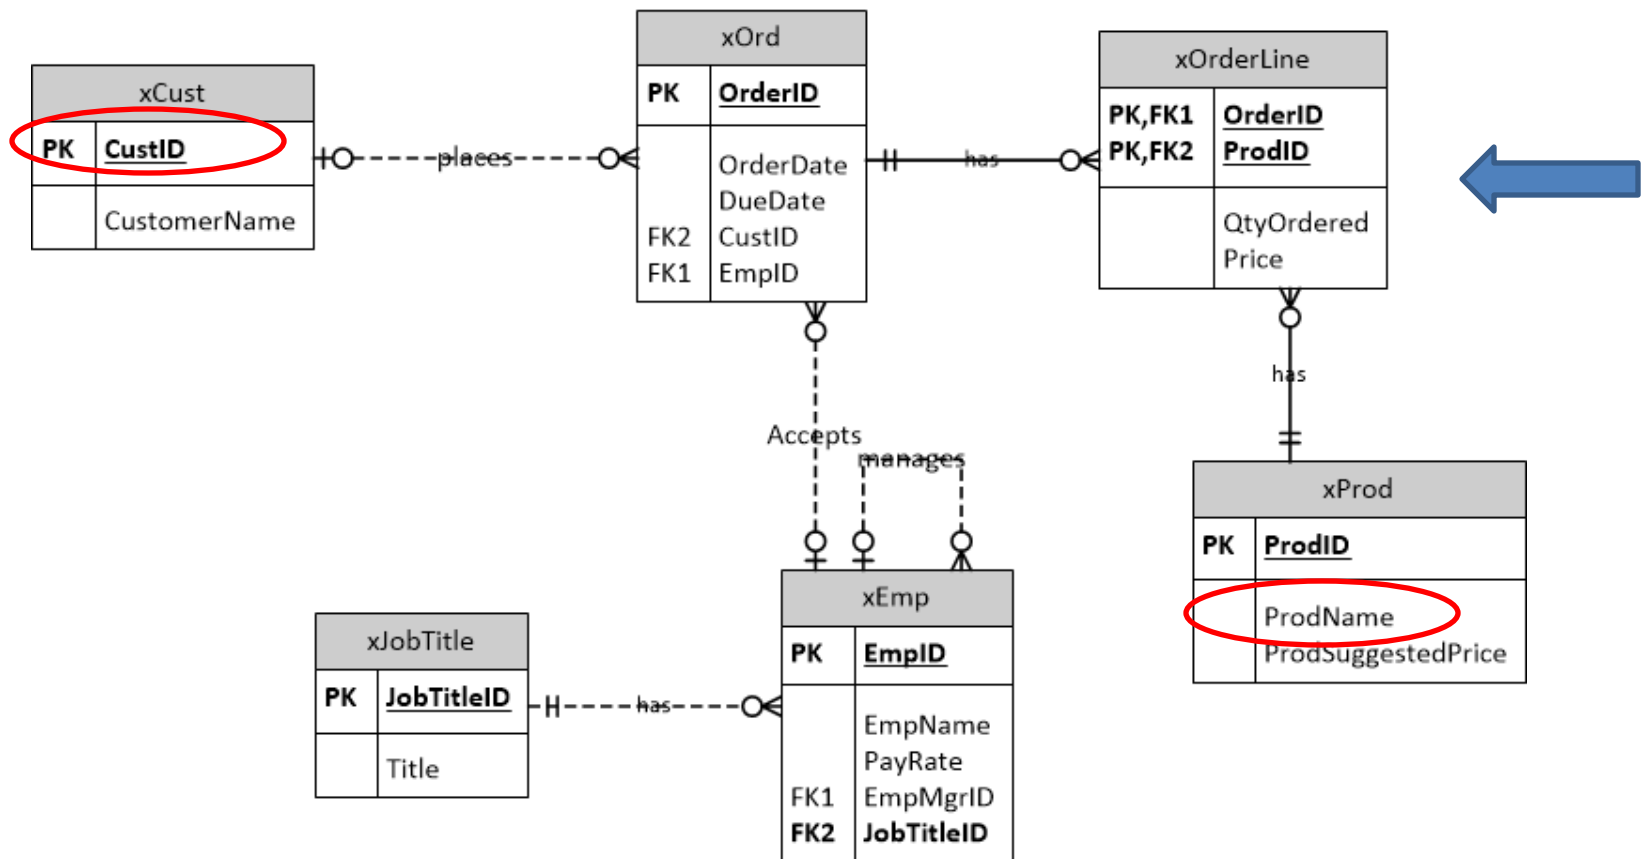
New query: Which customer spent the most for desks?

	CustID	CustomerName	ProdName	TotalExtendedPrice
1	2555	Mountain Design	Desk	4346.9400000

Where do the columns come from (which tables)?

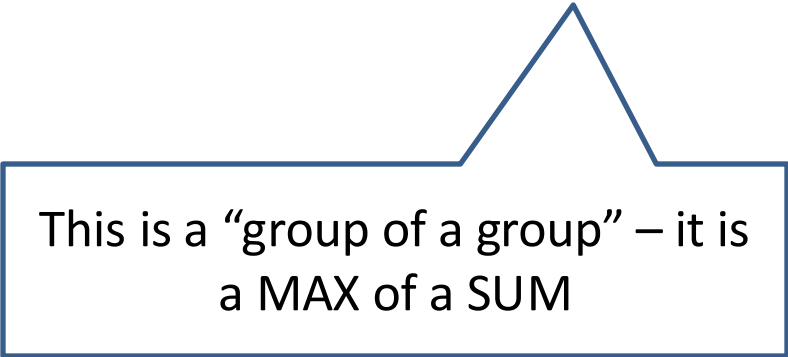
What is the basic logic of the query (which table or tables are necessary to find the required rows in the result table)?

What are the simplest requirements necessary to accomplish the basic logic?



Write the basic logic in pseudocode

```
SELECT  customer data
FROM    cust, ord, orderline and prod
WHERE   SUM(qtyOrdered*price for desks by customer) =
        MAX(SUM(qtyOrdered*price for desks by customer))
```



This is a “group of a group” – it is
a MAX of a SUM

This code isn't designed to actually work as a SQL query. It is just written to get an understanding of the basic logic necessary to accomplish the query.

Separate the two group functions – focus on the first GROUP function
 – the SUM of qtyOrdered*price for a product for a customer

```
SELECT *
FROM xOrderline
INNER JOIN xOrd
ON xOrderline.orderid = xOrd.orderid
ORDER BY custid
```

This requires the data stored in two different tables – the Orderline table for the qtyOrdered by product, and then the ord table to access the customer who purchased the product.

	OrderID	ProdID	QtyOrdered	Price	OrderID	OrderDate	CustID	DueDate	empid
1	100	10	3.000	135.95	100	2025-03-15 00:00:00.000	1234	2025-03-19 00:00:00.000	4
2	100	45	1.000	450.00	100	2025-03-15 00:00:00.000	1234	2025-03-19 00:00:00.000	4
3	100	67	30.560	35.87	100	2025-03-15 00:00:00.000	1234	2025-03-19 00:00:00.000	4
4	100	81	3.000	1925.99	100	2025-03-15 00:00:00.000	1234	2025-03-19 00:00:00.000	4
5	400	10	10.000	120.99	400	2025-03-27 00:00:00.000	2555	2025-04-16 00:00:00.000	7
6	400	12	2.000	678.99	400	2025-03-27 00:00:00.000	2555	2025-04-16 00:00:00.000	7
7	400	25	8.000	425.99	400	2025-03-27 00:00:00.000	2555	2025-04-16 00:00:00.000	7
8	400	64	3.000	381.00	400	2025-03-27 00:00:00.000	2555	2025-04-16 00:00:00.000	7
9	400	67	20.550	40.99	400	2025-03-27 00:00:00.000	2555	2025-04-16 00:00:00.000	7
10	600	12	5.000	455.99	600	2025-04-15 00:00:00.000	2555	2025-04-27 00:00:00.000	7
11	600	64	4.000	312.00	600	2025-04-15 00:00:00.000	2555	2025-04-27 00:00:00.000	7

First 11 rows of the result table

Let's add a calculation to the SELECT list.

```
SELECT    *,
          qtyOrdered * price ExtendedPrice
FROM xOrderline
INNER JOIN xOrd
ON xOrderline.orderid = xOrd.orderid
ORDER BY custid
```

First 14
rows of the
result table

	OrderID	ProdID	QtyOrdered	Price	OrderID	OrderDate	CustID	DueDate	empid	ExtendedPrice
1	100	10	3.000	135.95	100	2024-03-15 00:00:00.000	1234	2024-03-19 00:00:00.000	4	407.8500000
2	100	45	1.000	450.00	100	2024-03-15 00:00:00.000	1234	2024-03-19 00:00:00.000	4	450.0000000
3	100	67	30.560	35.87	100	2024-03-15 00:00:00.000	1234	2024-03-19 00:00:00.000	4	1096.1872000
4	100	81	3.000	1925.99	100	2024-03-15 00:00:00.000	1234	2024-03-19 00:00:00.000	4	5777.9700000
5	400	10	10.000	120.99	400	2024-03-27 00:00:00.000	2555	2024-04-16 00:00:00.000	7	1209.9000000
6	400	12	2.000	678.99	400	2024-03-27 00:00:00.000	2555	2024-04-16 00:00:00.000	7	1357.9800000
7	400	25	8.000	425.99	400	2024-03-27 00:00:00.000	2555	2024-04-16 00:00:00.000	7	3407.9200000
8	400	64	3.000	381.00	400	2024-03-27 00:00:00.000	2555	2024-04-16 00:00:00.000	7	1143.0000000
9	400	67	20.550	40.99	400	2024-03-27 00:00:00.000	2555	2024-04-16 00:00:00.000	7	842.3445000
10	600	12	5.000	455.99	600	2024-04-15 00:00:00.000	2555	2024-04-27 00:00:00.000	7	2279.9500000
11	600	64	4.000	312.00	600	2024-04-15 00:00:00.000	2555	2024-04-27 00:00:00.000	7	1248.0000000
12	700	10	25.000	99.99	700	2024-04-11 00:00:00.000	2555	2024-06-04 00:00:00.000	10	2499.7500000
13	700	45	5.000	410.99	700	2024-04-11 00:00:00.000	2555	2024-06-04 00:00:00.000	10	2054.9500000
14	700	64	6.000	325.99	700	2024-04-11 00:00:00.000	2555	2024-06-04 00:00:00.000	10	1955.9400000

Must now decide what to group on to create the
SUM of the ExtendedPrice

Now group it!

```
SELECT  custid,  
        prodid,  
        SUM(qtyOrdered*Price) SumQtyPrice  
FROM    xOrderline  
INNER JOIN xOrd  
ON      xOrderline.orderid = xOrd.orderid  
GROUP BY custid, prodid  
ORDER BY custid, prodID
```

We don't know which prodID represents a desk yet, but that is OK. We now know which customer bought what product and the extended price for that product. This is assuming that a customer is capable of buying the same product more than once – which is likely quite true!

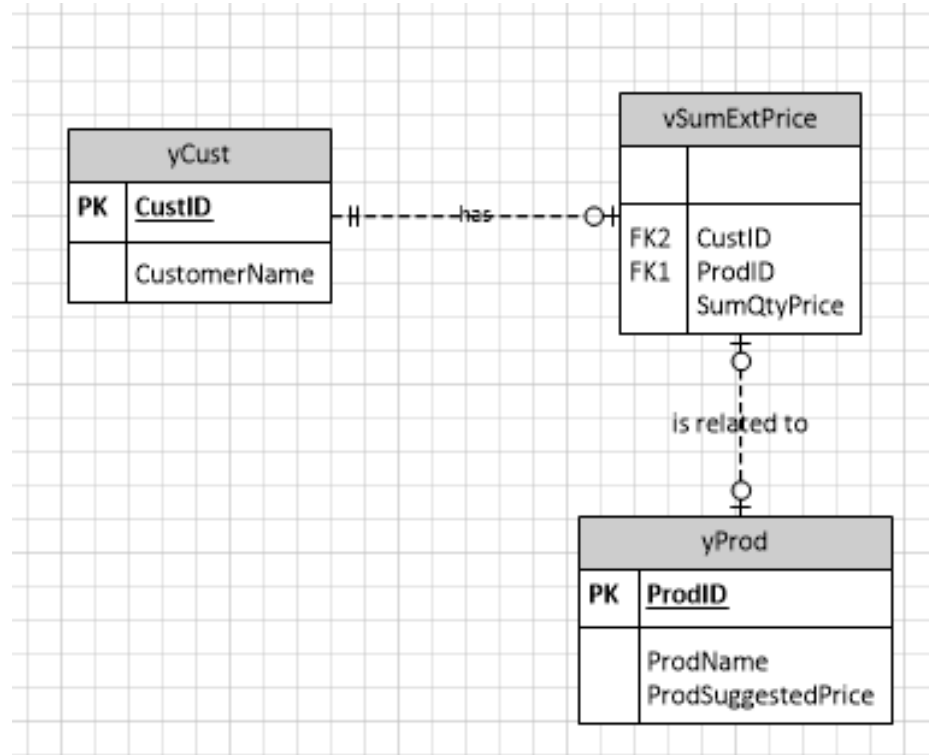
	custid	prodid	SumQtyPrice
1	1234	10	407.8500000
2	1234	45	450.0000000
3	1234	67	1096.1872000
4	1234	81	5777.9700000
5	2555	10	3709.6500000
6	2555	12	3637.9300000
7	2555	25	3407.9200000
8	2555	45	2054.9500000
9	2555	64	4346.9400000
10	2555	67	842.3445000
11	6773	12	1191.9800000
12	6773	45	2079.9500000
13	6773	64	731.9800000
14	6773	81	8179.5000000
15	6899	10	34437.5500000
16	6899	64	625.9800000
17	6899	67	10505.7475000
18	6899	77	1351.9800000
19	8372	10	2590.0000000
20	8372	25	5100.0000000
21	8372	64	975.3600000
22	8372	81	2598.9900000

Let's make a view out of that code so that we don't have to deal with the GROUP function SUM in other queries that need to use the total data.

```
CREATE VIEW vSumExtPrice AS
SELECT  custid,
        prodid,
        SUM(qtyOrdered*Price) SumQtyPrice
FROM    xOrderline
INNER JOIN xOrd
ON      xOrderline.orderid = xOrd.orderid
GROUP BY custid, prodid
```

Notice that the ORDER BY statement is gone.
A view cannot include an ORDER BY statement.

The view relates to both the Prod and Cust tables because it contains two potential FK's – A CustID and a ProdID



Look at the orders placed for a product name of a desk. The product name is stored in the xProd table, so let's join those two tables to get an idea of what data would be available.

```
SELECT *
FROM vsumextprice
INNER JOIN xProd
ON vsumextprice.prodid = xProd.prodid
ORDER BY 5
```

	custid	prodid	SumQtyPrice	ProdID	ProdName	ProdSuggestedPrice
1	1234	45	450.0000000	45	Bed	400.00
2	2555	45	2054.9500000	45	Bed	400.00
3	6773	45	2079.9500000	45	Bed	400.00
4	1234	10	407.8500000	10	Bookcase	135.99
5	2555	10	3709.6500000	10	Bookcase	135.99
6	6899	10	34437.5500000	10	Bookcase	135.99
7	8372	10	2590.0000000	10	Bookcase	135.99
8	2555	64	4346.9400000	64	Desk	330.00
9	6773	64	731.9800000	64	Desk	330.00
10	6899	64	625.9800000	64	Desk	330.00
11	8372	64	975.3600000	64	Desk	330.00
12	6899	77	1351.9800000	77	Platform Storage Bed	680.99
13	1234	81	5777.9700000	81	Sofa	1799.99
14	6773	81	8179.5000000	81	Sofa	1799.99
15	8372	81	2598.9900000	81	Sofa	1799.99
16	2555	25	3407.9200000	25	Table	460.99
17	8372	25	5100.0000000	25	Table	460.99
18	1234	67	1096.1872000	67	Walnut Finishing Wood	35.99
19	2555	67	842.3445000	67	Walnut Finishing Wood	35.99
20	6899	67	10505.7475000	67	Walnut Finishing Wood	35.99

```
SELECT *
FROM vsumextprice
INNER JOIN xProd
ON vsumextprice.prodid = xProd.prodid
WHERE prodname = 'desk'
```

	custid	prodid	SumQtyPrice	ProdID	ProdName	ProdSuggestedPrice
1	2555	64	4346.9400000	64	Desk	330.00
2	6773	64	731.9800000	64	Desk	330.00
3	6899	64	625.9800000	64	Desk	330.00
4	8372	64	975.3600000	64	Desk	330.00

Turn the query into a view. We are creating a view of a table joined with a view.

Will need to specify the columns for the View because it isn't possible to create a view when the columns have the same name (like ProdID in the query on the previous page) – so can't use the asterisk to declare the columns for a view.

```
create view vDesk as
SELECT  custid,
        yprod.prodid,
        prodname,
        sumqtyprice,
        ProdSuggestedPrice
FROM    vsumextprice
INNER JOIN xProd
ON vsumextprice.prodid = xProd.prodid
WHERE prodname = 'desk'
```

Find the correct row or rows for the goal of the query

```
SELECT *  
FROM vDesk  
WHERE sumqtyprice =  
      (SELECT MAX(sumqtyprice)  
        FROM vDesk)
```

	custid	prodid	prodname	sumqtyprice	ProdSuggestedPrice
1	2555	64	Desk	4346.9400000	330.00

Join the table(s) and additional columns
desired for the result table.

```
SELECT  vDesk.CustID,  
        CustomerName,  
        ProdName,  
        SumQtyPrice TotalExtendedPrice  
FROM vDesk  
INNER JOIN xCust  
ON vDesk.custID = xCust.custid  
WHERE sumqtyprice =  
      (SELECT MAX(sumqtyprice)  
        FROM vDesk)
```

	CustID	CustomerName	ProdName	TotalExtendedPrice
1	2555	Mountain Design	Desk	4346.9400000

Can accomplish the
same goal with CTE

```
WITH cteSumExtPrice AS
(SELECT  custid,
        prodid,
        SUM(qtyOrdered*Price) SumQtyPrice
FROM    xOrderline
INNER JOIN xOrd
ON      xOrderline.orderid = xOrd.orderid
GROUP BY custid, prodid),

cteDesk as
(SELECT  custid,
        yprod.prodid,
        prodname,
        sumqtyprice,
        ProdSuggestedPrice
FROM      cteSumExtPrice
INNER JOIN xProd
ON      ctesumextprice.prodid = xProd.prodid
WHERE   prodname = 'desk')

SELECT  cteDesk.CustID,
        CustomerName,
        ProdName,
        SumQtyPrice TotalExtendedPrice
FROM    cteDesk
INNER JOIN xCust
ON      cteDesk.custID = xCust.custid
WHERE   sumqtyprice =
        (SELECT MAX(sumqtyprice)
         FROM cteDesk)
```

```

WITH cteSumExtPrice AS
(SELECT  custid,
        xOrderline.prodid,
        prodname,
        SUM(qtyOrdered*price) AS SumQtyPrice
FROM xOrderline
INNER JOIN xOrd
ON xOrderline.orderid = xOrd.orderid
INNER JOIN xProd
ON xOrderline.prodid = xProd.prodid
WHERE prodname = 'desk'
GROUP BY custid, xOrderline.prodid, prodname)

```

```

SELECT  cteSumExtPrice.CustID,
        CustomerName,
        ProdName,
        SumQtyPrice TotalExtendedPrice
FROM cteSumExtPrice
INNER JOIN xCust
ON cteSumExtPrice.custID = xCust.custid
WHERE  sumqtyprice =
        (SELECT MAX(sumqtyprice)
         FROM cteSumExtPrice)

```

Can accomplish the
same goal with a
single CTE
Or a single VIEW...

Why bother using Views or CTEs?

- Group functions and joins are complex.
- GROUP BY should only be used with group functions (AVG, MAX, MIN, SUM, COUNT).
- Should not use a GROUP BY just to suppress rows in the result set. Use the GROUP BY only with a group function!
- Must have all non-group attributes that are in the SELECT list also in the GROUP BY statement.
- Difficult to do a group function of a group function. Examples:
 - The maximum of the sum of hours.
 - The minimum of a count of products.
- Joining multiple tables can yield full or partial cross joins making it difficult to trouble-shoot the SQL code.

IS475/675 Agenda: April 16, 2025



Complete one more example of a “group of a group.”



Do a bulk load of data to copy a table.




Connect SQL Server to Excel (if we have time).



Answer any questions.

Tables used in this example – from SQL Lab

Exercise 8



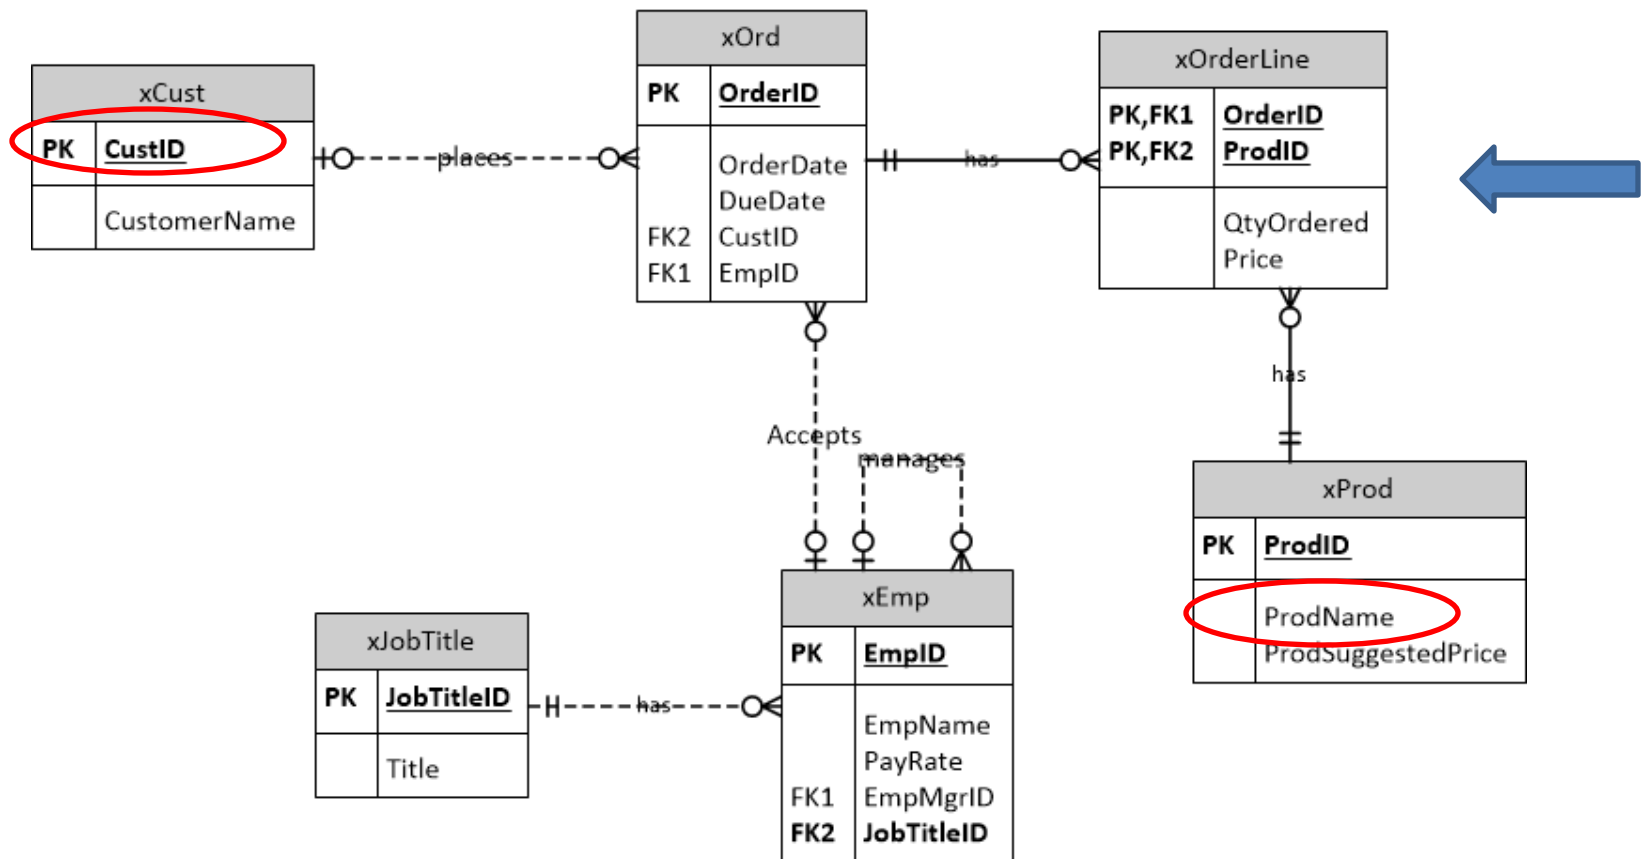
New query: Which customer spent the most for desks?

	CustID	CustomerName	ProdName	TotalExtendedPrice
1	2555	Mountain Design	Desk	4346.9400000

Where do the columns come from (which tables)?


What is the basic logic of the query (which table or tables are necessary to find the required rows in the result table)?

What are the simplest requirements necessary to accomplish the basic logic?



Write the basic logic in pseudocode

```
SELECT  customer data
FROM    cust, ord, orderline and prod
WHERE   SUM(qtyOrdered*price for desks by customer) =
        MAX(SUM(qtyOrdered*price for desks by customer))
```



This is a “group of a group” – it is
a MAX of a SUM

This code isn't designed to actually work as a SQL query. It is just written to get an understanding of the basic logic necessary to accomplish the query.

Let's add a calculation to the SELECT list.

```
SELECT    *,
          qtyOrdered * price ExtendedPrice
FROM xOrderline
INNER JOIN xOrd
ON xOrderline.orderid = xOrd.orderid
ORDER BY custid
```

First 14
rows of the
result table

	OrderID	ProdID	QtyOrdered	Price	OrderID	OrderDate	CustID	DueDate	empid	ExtendedPrice
1	100	10	3.000	135.95	100	2025-03-15 00:00:00.000	1234	2025-03-19 00:00:00.000	4	407.8500000
2	100	45	1.000	450.00	100	2025-03-15 00:00:00.000	1234	2025-03-19 00:00:00.000	4	450.0000000
3	100	67	30.560	35.87	100	2025-03-15 00:00:00.000	1234	2025-03-19 00:00:00.000	4	1096.1872000
4	100	81	3.000	1925.99	100	2025-03-15 00:00:00.000	1234	2025-03-19 00:00:00.000	4	5777.9700000
5	400	10	10.000	120.99	400	2025-03-27 00:00:00.000	2555	2025-04-16 00:00:00.000	7	1209.9000000
6	400	12	2.000	678.99	400	2025-03-27 00:00:00.000	2555	2025-04-16 00:00:00.000	7	1357.9800000
7	400	25	8.000	425.99	400	2025-03-27 00:00:00.000	2555	2025-04-16 00:00:00.000	7	3407.9200000
8	400	64	3.000	381.00	400	2025-03-27 00:00:00.000	2555	2025-04-16 00:00:00.000	7	1143.0000000
9	400	67	20.550	40.99	400	2025-03-27 00:00:00.000	2555	2025-04-16 00:00:00.000	7	842.3445000
10	600	12	5.000	455.99	600	2025-04-15 00:00:00.000	2555	2025-04-27 00:00:00.000	7	2279.9500000
11	600	64	4.000	312.00	600	2025-04-15 00:00:00.000	2555	2025-04-27 00:00:00.000	7	1248.0000000
12	700	10	25.000	99.99	700	2025-04-11 00:00:00.000	2555	2025-06-04 00:00:00.000	10	2499.7500000
13	700	45	5.000	410.99	700	2025-04-11 00:00:00.000	2555	2025-06-04 00:00:00.000	10	2054.9500000
14	700	64	6.000	325.99	700	2025-04-11 00:00:00.000	2555	2025-06-04 00:00:00.000	10	1955.9400000

Must now decide what to group
on to create the SUM of the
ExtendedPrice

Now group it!

```
SELECT  custid,  
        prodid,  
        SUM(qtyOrdered*Price) SumQtyPrice  
FROM    xOrderline  
INNER JOIN xOrd  
ON      xOrderline.orderid = xOrd.orderid  
GROUP BY custid, prodid  
ORDER BY custid, prodID
```

We don't know which prodID represents a desk yet, but that is OK. We now know which customer bought what product and the extended price for that product. This is assuming that a customer is capable of buying the same product more than once – which is likely quite true!


	custid	prodid	SumQtyPrice
1	1234	10	407.8500000
2	1234	45	450.0000000
3	1234	67	1096.1872000
4	1234	81	5777.9700000
5	2555	10	3709.6500000
6	2555	12	3637.9300000
7	2555	25	3407.9200000
8	2555	45	2054.9500000
9	2555	64	4346.9400000
10	2555	67	842.3445000
11	6773	12	1191.9800000
12	6773	45	2079.9500000
13	6773	64	731.9800000
14	6773	81	8179.5000000
15	6899	10	34437.5500000
16	6899	64	625.9800000
17	6899	67	10505.7475000
18	6899	77	1351.9800000
19	8372	10	2590.0000000
20	8372	25	5100.0000000
21	8372	64	975.3600000
22	8372	81	2598.9900000

Let's make a view out of that code so that we don't have to deal with the GROUP function SUM in other queries that need to use the total data.

```
CREATE VIEW vSumExtPrice AS
SELECT  custid,
        prodid,
        SUM(qtyOrdered*Price) SumQtyPrice
FROM    xOrderline
INNER JOIN xOrd
ON      xOrderline.orderid = xOrd.orderid
GROUP BY custid, prodid
```

Notice that the ORDER BY statement is gone.
A view cannot include an ORDER BY statement.

The view relates to both the Prod and Cust tables because it contains two potential FK's – A CustID and a ProdID



Look at the orders placed for a product name of a desk. The product name is stored in the xProd table, so let's join those two tables to get an idea of what data would be available.

```
SELECT *
FROM vsumextprice
INNER JOIN xProd
ON vsumextprice.prodid = xProd.prodid
ORDER BY 5
```

	custid	prodid	SumQtyPrice	ProdID	ProdName	ProdSuggestedPrice
1	1234	45	450.0000000	45	Bed	400.00
2	2555	45	2054.9500000	45	Bed	400.00
3	6773	45	2079.9500000	45	Bed	400.00
4	1234	10	407.8500000	10	Bookcase	135.99
5	2555	10	3709.6500000	10	Bookcase	135.99
6	6899	10	34437.5500000	10	Bookcase	135.99
7	8372	10	2590.0000000	10	Bookcase	135.99
8	2555	64	4346.9400000	64	Desk	330.00
9	6773	64	731.9800000	64	Desk	330.00
10	6899	64	625.9800000	64	Desk	330.00
11	8372	64	975.3600000	64	Desk	330.00
12	6899	77	1351.9800000	77	Platform Storage Bed	680.99
13	1234	81	5777.9700000	81	Sofa	1799.99
14	6773	81	8179.5000000	81	Sofa	1799.99
15	8372	81	2598.9900000	81	Sofa	1799.99
16	2555	25	3407.9200000	25	Table	460.99
17	8372	25	5100.0000000	25	Table	460.99
18	1234	67	1096.1872000	67	Walnut Finishing Wood	35.99
19	2555	67	842.3445000	67	Walnut Finishing Wood	35.99
20	6899	67	10505.7475000	67	Walnut Finishing Wood	35.99

```
SELECT *
FROM vsumextprice
INNER JOIN xProd
ON vsumextprice.prodid = xProd.prodid
WHERE prodname = 'desk'
```

	custid	prodid	SumQtyPrice	ProdID	ProdName	ProdSuggestedPrice
1	2555	64	4346.9400000	64	Desk	330.00
2	6773	64	731.9800000	64	Desk	330.00
3	6899	64	625.9800000	64	Desk	330.00
4	8372	64	975.3600000	64	Desk	330.00

Turn the query into a view. We are creating a view of a table joined with a view.

Will need to specify the columns for the View because it isn't possible to create a view when the columns have the same name (like ProdID in the query on the previous page) – so can't use the asterisk to declare the columns for a view.

```
create view vDesk as
SELECT  custid,
        xprod.prodid,
        prodname,
        sumqtyprice,
        ProdSuggestedPrice
FROM    vsumextprice
INNER JOIN xProd
ON vsumextprice.prodid = xProd.prodid
WHERE prodname = 'desk'
```

Find the correct row or rows for the goal of the query

```
SELECT *  
FROM vDesk  
WHERE sumqtyprice =  
      (SELECT MAX(sumqtyprice)  
        FROM vDesk)
```

	custid	prodid	prodname	sumqtyprice	ProdSuggestedPrice
1	2555	64	Desk	4346.9400000	330.00

Join the table(s) and additional columns
desired for the result table.

```
SELECT  vDesk.CustID,  
        CustomerName,  
        ProdName,  
        SumQtyPrice TotalExtendedPrice  
FROM vDesk  
INNER JOIN xCust  
ON vDesk.custID = xCust.custid  
WHERE sumqtyprice =  
      (SELECT MAX(sumqtyprice)  
        FROM vDesk)
```

	CustID	CustomerName	ProdName	TotalExtendedPrice
1	2555	Mountain Design	Desk	4346.9400000

Can accomplish the
same goal with CTE

```
WITH cteSumExtPrice AS
(SELECT  custid,
        prodid,
        SUM(qtyOrdered*Price) SumQtyPrice
FROM    xOrderline
INNER JOIN xOrd
ON      xOrderline.orderid = xOrd.orderid
GROUP BY custid, prodid),

cteDesk as
(SELECT  custid,
        xprod.prodid,
        prodname,
        sumqtyprice,
        ProdSuggestedPrice
FROM      cteSumExtPrice
INNER JOIN xProd
ON      ctesumextprice.prodid = xProd.prodid
WHERE   prodname = 'desk')

SELECT  cteDesk.CustID,
        CustomerName,
        ProdName,
        SumQtyPrice TotalExtendedPrice
FROM    cteDesk
INNER JOIN xCust
ON      cteDesk.custID = xCust.custid
WHERE   sumqtyprice =
        (SELECT MAX(sumqtyprice)
         FROM cteDesk)
```

```

WITH cteSumExtPrice AS
(SELECT  custid,
        xOrderline.prodid,
        prodname,
        SUM(qtyOrdered*price) AS SumQtyPrice
FROM xOrderline
INNER JOIN xOrd
ON xOrderline.orderid = xOrd.orderid
INNER JOIN xProd
ON xOrderline.prodid = xProd.prodid
WHERE prodname = 'desk'
GROUP BY custid, xOrderline.prodid, prodname)

```

```

SELECT  cteSumExtPrice.CustID,
        CustomerName,
        ProdName,
        SumQtyPrice TotalExtendedPrice
FROM cteSumExtPrice
INNER JOIN xCust
ON cteSumExtPrice.custID = xCust.custid
WHERE  sumqtyprice =
        (SELECT MAX(sumqtyprice)
         FROM cteSumExtPrice)

```

Can accomplish the
same goal with a
single CTE
Or a single VIEW...

Why bother using Views or CTEs?

- Group functions and joins are complex. Helps to split up a problem into smaller pieces of code.
- Difficult to do a group function of a group function. Examples:
 - The maximum of the sum of hours.
 - The minimum of a count of products.
- Joining multiple tables can yield full or partial cross joins making it difficult to trouble-shoot your SQL code.
- Good way to secure the data – users only see a view and don't know the actual structure of the tables.

Inserting Data vs. Bulk Load of Data

- A DBMS keeps an audit trail of each individual INSERT, UPDATE, and DELETE statement executed against the database. The log of the transactions can be extensive.
- A “bulk load” of data allows the database programmer to insert much data without also creating a huge listing of all the data that is being INSERTed.

Two general methods of bulk loading data

- One method requires that a table be created prior to loading the data.
- Another method creates a new table and bulk loads the data in a single statement.

Two general methods of “bulk loading” data

Examples of SQL Code used to “bulk load” data

```
INSERT INTO xCust (CustID, CustomerName)
SELECT      CustomerID,
            Name
FROM        xAnotherTable
WHERE       CustomerID NOT IN
            (SELECT CustID
             from xCust)
```

This syntax requires that a table be created prior to executing the INSERT statement. The SELECT can include a WHERE clause to determine which rows to insert.

```
SELECT      CustomerID,
            Name
INTO        xCustNew
FROM        xAnotherTable
```

This syntax creates a table and inserts the data in a single statement. The SELECT can also include a WHERE clause to determine which rows to insert.

Bulk load example

Create a new table called "xCustPast" by executing this script file: k:\CoB\IS475\LabFiles\SQLCustPast. This is the data that we want to add to the xCust table. But we only want to add those customers who don't already exist in the xCust table.

```
SELECT * FROM xCustPast
```

	customerID	Name	DateArchived
1	1234	Reston Supplies	2025-01-15 00:00:00.000
2	1284	Taran Singh	2024-10-29 00:00:00.000
3	2839	Marissa Wong	2023-12-29 00:00:00.000
4	2927	Chancey Motors Corp	2025-03-29 00:00:00.000
5	3408	Great Cupcakes Co.	2024-10-29 00:00:00.000
6	4500	Accessible Wheels, LLC	2024-04-04 00:00:00.000
7	5711	Rodriguez Markets	2025-03-31 00:00:00.000
8	6899	Opaka Sporting Goods	2025-03-15 00:00:00.000
9	8119	Chen Antiques	2024-02-15 00:00:00.000
10	8233	Right Way Massage	2024-04-07 00:00:00.000
11	8372	CutGlass Tile Co.	2025-04-02 00:00:00.000
12	9029	Bodega Bay Florist	2019-10-30 00:00:00.000
13	9886	Emma Wilson	2024-03-30 00:00:00.000

Use bulk load to create a new table called xCustHold

```
SELECT      CustID,  
            CustomerName  
INTO        xCustHold  
FROM        xCust
```

We are going to modify xCust by adding new data, so it is a good idea to make a backup copy of the xCust table. xCustHold is our backup copy of xCust.

Use bulk load to add data to the existing customer table

```
INSERT INTO xCust (CustID, CustomerName)
SELECT CustomerID,
       Name
FROM   xCustPast
WHERE  CustomerID NOT IN
      (SELECT CustID
       FROM xCust)
```

We are adding rows to xCust that do not already exist in the table.

The two tables have different field names, but that is okay – the data is entered on the position of the field name rather than the actual name

Bulk load will be used for HW#10



SQL Lab exercise 9 also shows how to do a bulk load of data.



Creating a backup copy of a table is optional, but a good idea when you intend to modify the contents of a table.